

The Regina Rexx Interpreter

Version 3.1

Anders Christensen

<Anders.Christensen@idi.ntnu.no>
Norwegian Institute of Technology
University of Trondheim

April 5, 2003

The logo for Regina Rexx. The word "Regina" is in a serif font with a 3D effect. The word "Rexx" is in a stylized, outlined font with a 3D effect.

Additions and corrections by Mark Hessling <M.Hessling@qut.edu.au>

Copyright©1992-1998AndersChristensen

Copyright©1998-2003MarkHessling

Permissionisgrantedtocopy,distributeand/ormodifythisdocumentunderthetermsoftheGNUFree DocumentationLicense,Version1.1oranylaterversionpublishedbytheFreeSoftwareFoundation; withnoInvariantSections,withnoFront-CoverTexts,andwithnoBack-CoverTexts.

Acopyofthelicenseisincludedinthesectionentitled"GNUFreeDocumentationLicense".

Introduction to Regina

This chapter provides an introduction to Regina, an Open Source Rexx Interpreter distributed under the GNU General Library License.

1 Purpose of this document

The purpose of this document is to provide an overview of the Rexx language and the Regina implementation of the Rexx language. It is not intended as a definitive reference to Rexx; you should really have a copy of the Rexx "bible"; *The Rexx Language*, by Mike Cowlishaw [TRL2].

2 Implementation

The Regina Rexx Interpreter is implemented as a library suitable for linking into third-party applications. Access to Regina from third-party applications is via the Regina API, which is consistent with the IBM's REXX SAA API. This API is implemented on most other Rexx interpreters.

The library containing Regina is available either as a static library or as a dynamically loadable library. The only functional difference between the two libraries is that the ability to dynamically load Rexx external function packages via the built-in function; `RxFuncAdd`, is only available with the dynamically loadable library.

The Regina distribution also includes a frontend to the Regina library, to enable the execution of Rexx programs directly from the command line. The *commandline* referred to here relates to the Unix shell, an OS/2 or DOS command window or a Windows NT/9x command prompt.

On platforms where both a static and a dynamic executable exist, it should be noted that the ability to load and execute external functions via the `RxFuncAdd` function, is only available by running the dynamic executable.

3 PortsofRegina

Reginahasbeenportedtomanyoperatingsystems.Thefollowingtableprovidesimplementation detailsofeachoftheportsof Regina.

OperatingSystem	Dynamic Library	Static Library	ThreadSafe Library	Dynamic Executable	Static Executable
Linux	libregina.so	libregina.a	libregina_ts.so	regina	rexx
HP-UX	libregina.sl	libregina.a	N/A	regina	rexx
AIX	libregina.a	libregina.a	libregina_ts.a	regina	rexx
OtherUnix	libregina.so	libregina.a	libregina_ts.so	regina	rexx
32-bitDOS(DJGPP) (UsesDPMI memorymanager)	N/A	libregina.a	N/A	N/A	rexx.exe
32-bitDOS(EMX) (UsesVCPImemory manager)	N/A	regina.a	N/A	N/A	rexx.exe
OS/2(EMX)	regina.dll (regina.lib)	regina.a	reginats.dll (reginats.lib)	regina.exe	rexx.exe
Windows 9x/Me/NT/2k/XP	regina.dll (regina.lib)	rexx.lib	reginats.dll (reginats.lib)	regina.exe	rexx.exe
BeOS	libregina.so	libregina.a	N/A	regina	rexx
AmigaOS	N/A	libregina.a	N/A	N/A	rexx
EPOC32	N/A	N/A	N/A	N/A	rexx.exe
AtheOS	libregina.so	libregina.a	N/A	regina	rexx
QNX4.2x	N/A	regina.lib	N/A	N/A	rexx
QNX6.x	libregina.so	libregina.a	libregina_ts.so	regina	rexx

4 ExecutingRexxprogramswithRegina

Rexxprogramsaregenerallyexecutedby Reginafromthe *commandline* inthefollowingmanner:

regina[*switches*][program][*programparameters*]

where:

regina is the name of the Regina executable (see table above)
switches are optional switches. See the section below for an explanation of the switches currently supported by Regina

program the name of the Rexx program to be executed. See the section **External Rexx Programs**, below, for detail on how Regina interprets this argument. If no program name is specified, Regina waits for Rexx command to be typed in and will execute those commands when the appropriate end-of-file character (^D on Unix and ^Z on DOS, OS/2 and Windows NT/95) is typed.

program parameters any optional parameters to be passed to the Rexx program.

Rexx programs to be executed by Regina can take advantage of a feature of Unix shell programs called *magic numbers*. By having the first line of a Rexx program consist of the special sequence of **#!** followed by the full filename of the Regina executable, you can invoke this program simply by typing the name of the Rexx program on the *command line* followed by any parameters you wish to pass to the Rexx program. The filename must also have the appropriate execute bit set for this to work. As an example, suppose your Rexx program, **myprog**, contained:

```
#!/usr/local/bin/regina
Parse Version ver
Say ver
```

When executing this program from the *command line* by typing **myprog**, the Unix shell program would execute the program `/usr/local/bin/regina` and pass the remainder of the lines in the file to this program via *stdin*.

The special processing done by Regina to find the filename in **REGINA_MACROS** and the file extension searching is not able to be carried out when using the magic number method of invocation.

4.1 Switches

The following switches allow the user to control how Regina executes the supplied Rexx program. Switches are recognised by a leading hyphen character; `-`, followed immediately by a single alphabetic character. Some switches allow for optional parameters. These, too, must follow the switch without any intervening spaces. All switches and their optional parameters are case-sensitive.

- `-t[traceparameter]` Turn on the specified tracing level. The optional *traceparameter* indicates the tracing level to be used. See the TRACE command later in this document for an explanation of each trace level. Use of this switch will result in any TRACE commands in the program to be ignored.
- `-a` Without this switch, all command line parameters are passed to Regina as a single argument. Specifying `-a`, ensures that the Rexx program invoked has access to the command line parameters as separate arguments, as passed from the command line interpreter. ie. The BIF ARG() can return a value of other than 1 or 0. Also PARSE SOURCE will return SUBROUTINE instead of the normal COMMAND value.
- `-r` Run Regina in restricted mode. See the section on Regina Restricted Mode for more details.

4.2 External Rexx programs

Regina searches for Rexx programs, using a combination of the **REGINA_MACROS** environment variable and the addition of filename extensions. This rule applies to both external function calls and the **programs** specified on the *commandline*.

Assume you have a call to an external function, and it is coded as follows:

```
Call myextfunc arg1, arg2
```

First, Regina looks for a file called **myextfunc** in the current directory. If it can't find that file, it looks in each directory specified in the **REGINA_MACROS** environment variable for a file called **myextfunc**. If the file is not found, Regina then attempts to find a file called **myextfunc.rexx** in the current directory, then in each directory in **REGINA_MACROS**. Regina continues, next by appending **.rex** to the supplied external function name, followed by **.cmd** and **.rx**.

Only if a file does not exist in either the current directory, or any directory in **REGINA_MACROS**, either with the supplied filename or with that filename appended with **.rexx**, **.rex**, **.rx** or **.cmd** does Regina complain that the external function is unknown.

RexxLanguageConstructs

In this chapter, the concept and syntax of REXX clauses are explained. At the end of the chapter there is a section describing how Regin differs from standard REXX as described in the first part of the chapter.

5 Definitions

A program in the REXX language consists of clauses, which are divided into four groups: null clauses, commands, assignments, and instructions. The three latter groups (commands, assignments, and instructions) are collectively referred to as statements. This does not match the terminology in [TRL2], where "instruction" is equivalent to what is known here as "statement", and "keyword instruction" is equivalent to what is known here as "instruction". However, I find the terminology used here simpler and less confusing.

Incidentally, the terminology used here matches [DANEY].

A clause is defined as all non-clause-delimiters (i.e. blanks and tokens) up to and including a clause delimiter. A token delimiter can be:

- An end-of-line, unless it lies within a comment. An end-of-line within a constant string is considered a syntax error { 6 }.
- A semicolon character that is not within a comment or constant string.
- A colon character, provided that the sequence of tokens leading up to it consists of a single symbol and whitespace. If a sequence of two symbol tokens is followed by a colon, then this implies SYNTAX condition { 13 }.

Some systems have the ability to store a text file having a last line unterminated by an end-of-line character sequence. In general, this applies to systems that use an explicit end-of-line character sequence to denote end-of-lines, e.g. Unix and MS-DOS systems. Under these systems, if the last line is unterminated, it will strictly speaking not be a clause, since a clause must include its terminating clause delimiter. However, some interpreters are likely to regard the end-of-file as a clause delimiter too. The functionality of INTERPRET gives some weight to this interpretation. But other systems may ignore that last, unterminated line, or may be issue a syntax error. (However, there is no SYNTAX condition number adequately covering this situation.)

Example: Binary transferring files

Suppose a REXX program is stored on an MS-DOS machine. Then, an end-of-line sequence is marked in the file as the two characters carriage return and newline. If this file is transferred to a Unix system, then only the newline marks the end-of-line. For this to work, the file must be transferred as a text file. If it is (incorrectly) transferred as a binary file, the result is that on the Unix system, each line seems to contain a trailing carriage return character. In an editor, it might look like this:

```
say 'hello world'^M
say 'that"s it'^M
```

This will probably raise `SYNTAXcondition{13}`.

6 Null clauses

Null clauses are clauses that consist of only whitespace, or comments, or both; in addition to the terminating clause delimiter. These clauses are ignored when interpreting the code, except for one situation: null clauses containing at least one comment are traced when appropriate. Null clauses not containing any comments are ignored in every respect.

Example: Tracing comments

The tracing of comments may be a major problem, depending on the context. There are basically two strategies for large comments: either box multiple lines as a single comment, or make the text on each line an independent comment, as shown below:

```
trace all

/*
  This is a single, large comment, which spans multiple lines.
  Such comments are often used at the start of a subroutine or
  similar, in order to describe both the interface to and the
  functionality of the function.
*/

/* This is also a large comment, but it is written as multiple
*/
/* comments, each on its own line. Thus, this is several clauses
*/
/* while the comment above is a single comment.
*/
```

During tracing, the first of these will be displayed as one large comment, and during interactive tracing, it will only pause once. The second will be displayed as multiple lines, and will make several pauses during interactive tracing. An interpreter may solve this situation in several ways, the main objective must be to display the comments nicely to the programmer debugging the code. Preferably, the code is shown in a fashion that resembles how it is entered in the file.

If a label is multiple defined, the first definition is used and the rest are ignored. Multiple defined labels is not a `SYNTAXcondition`.

A null clause is not a statement. In some situations, like after the `THEN` subclause, only a statement

come. If a null clause is provided, then it is ignored, and the next statement is used instead.

Consider the following code:

```
parse pull foo

if foo=2 then
    say 'foo is not 2'
else
    /* do nothing */
    say 'that"s it'
```

This will not work the way indentation indicates, since the comment in this example is not a statement. Thus, the ELSE reads beyond the comment, and connects to the SAY instruction which becomes the ELSE part. (That what probably not what the programmer intended.) This code will say that's it, only when foo is different from 2. A separate instruction, NOP has been provided in order to fill the need that was inadequately attempted filled by the comment in the code fragment above.

Example: Trailing comments

The effect that comments are not statements can be exploited when documenting the program, and simultaneously making the program faster. Consider the following two loops:

```
sum = 0
do i=1 to 10
/* sum 1 2 3 ... 8 9 10 */
    sum = sum + i
end

sum = 0
do i=1 to 10
    sum = sum + i /* sum 1 2 3 ... 8 9 10 */
end
```

In the first loop, there are two clauses, while the second loop contains only one clause, because the comment is appended to an already existing clause. During execution, the interpreter has to spend time ignoring the null clause in the first loop, while the second loop avoids this problem (assuming tracing is unenabled). Thus, the second loop is faster; although only insignificantly faster for small loops. Of course, the comment could have been taken out of the loop, which would be equally fast to the second version above.

7 Commands

7.1 Assignments

Assignments are clauses where the first token is a symbol and the second token is the equals sign (=). This definition opens for some curious effects, consider the following clauses:

a == b

This is not a command, but an assignment of the expression `b` to the variable `a`. Of course, the expression is illegal (`=b`) and will trigger a `SYNTAX` condition for syntax error {35}. TRL2 defines the operator `==` as consisting of two tokens. Thus, in the first of these examples, the second token is `=`, the third token is also `=`, while the fourth token is `b`.

3 = 5

This is an assignment of the value `5` to the symbol `3`, but since this is not a variable symbol, this is an illegal assignment, and will trigger the `SYNTAX` condition for syntax error {31}.

"hello" = foo

This is not an invalid assignment, since the first token in the clause is not a symbol. Instead, this becomes a command.

arg =(foo) bar

The fourth statement is a valid assignment, which will space-concatenate the two variable symbols `foo` and `bar`, and assign the result to the variable symbol `arg`. It is specifically not an `ARG` instruction, even though it might look like one. If you need an `ARG` instruction which template starts with an absolute indirect positional pattern, use the `PARSE UPPER ARG` instruction instead, or prepend a dot in front of the template.

An assignment can assign a value to a simple variable, a stem variable or a compound variable. When assigning to a stem variable, all possible variable symbols having that stem are assigned the value. Note specifically that this is not like setting a default, it is a one time multiple assignment.

Example: Multiple assignment

The difference between REXX's multiple assignment and a default value can be seen from the following code:

```
foo. = 'bar'
foo.1 = 'baz'
drop foo.1
say foo.1          /* says "FOO.1" */
```

Here, the `SAY` instruction writes out `FOO.1`, not `bar`. During the `DROP` instruction, the variable `FOO.1` regains its original, uninitialized value `FOO.1`, not the value of its stem variable `FOO.`, i.e. `bar`, because stem assignments does not setup a default.

Example:Emulatingadefaultvalue

If you want to set the compound variable to the value of its stem variable, if the stem is initialized, then you may use the following code:

```
if (symbol('foo.')) then
    foo.1 = foo.
else
    drop foo.1
```

In this example, the `foo.1` variable is set to the value of its stem if the stem currently is assigned a value. Else, the `foo.1` variable is dropped.

However, this is probably not exactly the same, since the internal storage of the computer is likely to store variables like `foo.2` and `foo.3` only implicitly (after all, it cannot explicitly store every compound having `foo.` as stem). After the assignment of the value of `foo.` to `foo.1`, the `foo.1` compound variable is likely to be explicitly stored in the interpreter.

There is now a way you can discover the difference, but the effects are often that more memory is used, and some functionality that dumps all variables may dump `foo.1` but not `foo.2` (which is inconsistent). See section `RexxVariablePool`.

Example:Spaceconsiderations

Even more strange are the effects of the following short example:

```
foo. = 'bar'
drop foo.1
```

Although apparently very simple, there is now a way that an interpreter can release all memory referring to `foo.1`. After all, `foo.1` has a different value than `foo.2`, `foo.3`, etc., so the interpreter must store information that tells it that `foo.1` has the uninitialized value.

These considerations may seem like nit-picking, but they will matter if you drop lots of compound variables for a stem which has previously received a value. Some programming idioms do this, so be aware. If you can do without assigning to the stem variable, then it is possible for the interpreter to regain all memory used for that stem's compound variables.

8 Instructions

In this section, all instructions in standard REXX are described.

Extensions are listed later in this chapter.

First, some notes on the terminology. What is called an instruction in this document is equivalent to a "unit" of clauses. That is, each instruction can consist of one or more clauses. For instance, the `SAY`

instruction is always a single instruction, but the IF instruction is a multi-clause instruction. Consider the following script, where each clause has been boxed:

```
if a=b then
    say 'hello'
else
    say 'bye'
```

Further, the THEN or ELSE part of this instruction might consist of a DO/END pair, in which case the IF instruction might consist of an virtually unlimited number of clauses.

Then, some notes on the syntax diagrams used in the following description of the instructions. The rules applying to these diagrams can be listed as:

- Anything written in *courier* font in the syntax diagrams indicates that it should occur as-is in the REXX program. Whenever something is written in *italic* font, it means that the term should be substituted for another value, expression, or terms.
- Anything contained within matching pair of square brackets ([...]) are optional, and may be left out.
- Whenever a pair of curly braces is used, it contains two or more subclauses that are separated by the vertical bar (|). It means that the curly braces will be substituted for one of the subclauses it contains.
- Whenever the ellipsis (...) is used, it indicates that the immediately following subclauses may be repeated zero or more times. The scope of the ellipsis is limited to the contents of a set of square brackets or curly braces, if it occurs there.
- Whenever the vertical bar | is used in any of the syntax diagrams, it means that either the term to the left, or the term to the right can be used, but not both, and at least one of them must be used. This "operator" is associative (can be used in sequence), and it has lower priority than the square brackets (the scope of the vertical bar located within a pair of square brackets or curly braces is limited to the text within those square brackets or curly braces).
- Whenever a semicolon (;) is used in the syntax diagram, it indicates that a clause separator must be present at the point. It may either be a semicolon character, or an end-of-line.
- Whenever the syntax diagram is spread out over more lines, it means that any of the lines can be used, but that the individual lines are mutually exclusive. Consider the syntax:

```
SAY = symbol
      string
```

This is equivalent to the syntax:

```
SAY [ symbol | string ]
```

Because in the first of these two syntaxes, the SAY part may be continued at either line.

- Sometimes the syntax of an instruction is so complex that part of the syntax has been extracted, and is shown below in its expanded state. The following is an example of how this looks:

SAY something TO someone

something : = HI
 HELLO
 BYE

someone : = THE BOSS
 YOUR NEIGHBOR

You can generally identify these situations by the fact that they come a bit below the real syntax diagram, and that they contain a colon character after the name of the term to be expanded.

In the syntax diagrams, some generic names have been used for the various parts, in order to indicate common attributes for the term. For instance, whenever a term in the syntax diagrams is called *expr*, it means that any valid REXX expression may occur instead of that term. The most common such names are:

condition

Indicates that the subclause can be any of the names of the conditions, e.g. SYNTAX, NOVALUE, HALT, etc.

expr

Indicates that the subclause can be any valid REXX expression, and will in general be evaluated as normal during execution.

statement

Indicates that extra clauses may be inserted into the instruction, and that exactly one of them must be a true statement.

string

Indicates that the subclause is a constant string, i.e. either enclosed by single quotes ('...') or double quotes ("...").

symbol

Indicates that the subclause is a single symbol. In general, whenever *symbol* is used as the name for a subclause, it means that the symbol will not automatically be expanded to the value of the symbol. But instead, some operation is performed on the name of the symbol.

template

Indicates that the subclause is a parsing template. The exact syntax of this is explained in a chapter on tracing, to be written later.

In addition to this, variants may also exist. These variants will have an extra letter or number appended to the name of the subclause, and is used for differing between two or more subclauses having the same "type" in one syntax diagram. In the case of other names for the subclauses, these are explained in the description of the instruction.

8.1 The ADDRESS Instruction

```
ADDRESS [ environment [ command ] [ redirection ] ] ;  
        [ [ VALUE ] expression [ redirection ] ] ;
```

and *redirection* has one of the forms:

```
WITH INPUT standard_redir [ OUTPUT out_redir ] [ ERROR out_redir  
] ;  
WITH INPUT standard_redir [ ERROR out_redir ] [ OUTPUT out_redir  
] ;  
WITH OUTPUT out_redir [ INPUT standard_redir ] [ ERROR out_redir  
] ;  
WITH OUTPUT out_redir [ ERROR out_redir ] [ INPUT standard_redir  
] ;  
WITH ERROR out_redir [ INPUT standard_redir ] [ OUTPUT out_redir  
] ;  
WITH ERROR out_redir [ OUTPUT out_redir ] [ INPUT standard_redir  
] ;
```

standard_redir is defined as:

```
NORMAL ;  
[ STREAM | STEM | LIFO | FIFO ] symbol ;
```

and *out_redir* is defined as:

```
NORMAL ;  
[ APPEND | REPLACE ] [ STREAM | STEM | LIFO | FIFO ] symbol ;
```

We will discuss redirection later.

The ADDRESS instruction controls where commands to an external environment are sent. If both *environment* and *command* are specified, the given command will be executed in the given environment. The effect is the same as issuing an expression to be executed as a command (see section **Commands**), except that the environment in which it is to be executed can be explicitly specified in the ADDRESS clause. In this case, the special variable RC will be set as usual, and the ERROR or FAILURE conditions might be raised, as for normal commands.

In other words: All *normal* commands are ADDRESS statements with a suppressed keyword and an environment.

The *environment* term must be a symbol or a literal string. If it is a symbol, its "name" is used, i.e. it is not tail substituted or swapped for a variable value. The *command* and *expression* terms can be any REXX expression. eg.

```
SYSTEM= 'PATH'  
ADDRESS SYSTEM "echo Hello"
```

is equivalent to

ADDRESS SYSTEM "echo Hello"

or

ADDRESS "SYSTEM" "echo Hello"

for the external *echo* command.

As symbols specified as an environment name is not case-sensitive, whereas a string must match the case. Built-in environments are always uppercased.

REXX maintains a list of environments, the size of this list is at least two. If you select a new environment, it will be put in the front of this list. Note that if *command* is specified, the contents of the environment stack is not changed. If you omit *command*, *environment* will always be put in the front of the list of environments. Regina has an infinite list and never pushes out any entry. Possible values are listed below. If you supply a *command* with the ADDRESS statement, the *environment* is interpreted as a temporary change for just this command.

What happens if you specify an environment that is already in the list, is not completely defined. Strictly speaking, you should end up with both entries in the list pointing to the same environment, but some implementations will probably handle this by reordering the list, leaving the selected environment in the front. This is Regina's behaviour. Every environment exists only once. The redirection command below always changes the behaviour of one--the given--environment. You can imagine a set of playing cards in your hand. The operation is to draw one card by name and put it to the front.

If you do not specify any subkeywords or parameters to ADDRESS, the effect is to swap the two first entries in the list of environments. Consequently, executing ADDRESS multiple times will toggle between two environments.

The second syntax form of ADDRESS is a special case of the first form with *command* omitted. If the first token after ADDRESS is VALUE, then the rest of the clause is taken to be an expression, naming the environment which is to be made the current environment. Using VALUE makes it possible to circumvent the restriction that the name of the new environment must be a symbol or literal string. However, you cannot combine both VALUE and *command* in a single clause.

Example: Examples of the ADDRESS instruction

Let's look at some examples, they can sometimes be combined with a redirection:

```
ADDRESS COMMAND
ADDRESS SYSTEM 'copy' fromfile tofile
ADDRESS system
ADDRESS VALUE newenv
ADDRESS
ADDRESS (oldenv)
```

The first of these sets the environment `COMMAND` as the current environment.

The second performs the command " `copy`" in the environment `SYSTEM`, using the values of the symbols `fromfile` and `tofile` as parameters. Note that this will not set `SYSTEM` as the current environment.

The third example sets `SYSTEM` as the current environment (it will be automatically converted to upper case).

The fourth example sets as the current environment the contents of the symbol `newenv`, pushing `SYSTEM` down one level in the stack.

The fifth clause swaps the two uppermost entries on the stack; and `SYSTEM` ends up at the top pushing the environments specified in `newenv` below it.

The sixth clause is equivalent to the fourth example, but is not allowed by ANSI. Since Regina 3.0 this style is deprecated and can't be used if `OPTIONS STRICT_ANSI` is in effect. Again, avoid this kind of `ADDRESS` statement style, and use the `VALUE` version instead.

Example: The VALUE subkeyword

Let us look a bit closer at the last example. Note the differences between the two clauses:

```
ADDRESS ENV
ADDRESS VALUE ENV
```

The first of these sets the current default environment to `ENV`, while the second sets it to the value of the symbol `ENV`.

If you are still confused, Don't Panic; the syntax of `ADDRESS` is somewhat bizarre, and you should not

put too much effort into learning all aspects of it. Just make sure that you understand how to use it in simple situations. Chances are that you will not have use for its more complicated variants for quite some time.

Then, what names are legal as environment names? Well, that is implementation-specific, but some names seem to be in common use. The name `COMMAND` is sometimes used to refer to an environment that sends the command to the operating system. Likewise, the name of the operating system is often used for this (`CMS`, `UNIX`, etc.). You have to consult the implementation's specific documentation for more information about this. Actually, there is not really any restriction on what constitutes a legal environment name (even the null string is legal). Some interpreters will allow you to select anything as the current environment; and if it is an illegal name, the interpreter will complain only when the environment is actually used. Other implementations may not allow you to select an invalid environment name at all.

`Regina` allows every name as an environment name. `Regina` gives an error message about wrong names only when the name is used. The error string looks somewhat strange if `Regina` is used as a separate program, since the extension of the environment namespace is only useful when running as part of a program which extends the standard names.

`Regina` uses three kinds of environments. Some have alias names. The environment names are:

`SYSTEM`

```
alias OS2ENVIRONMENT
alias ENVIRONMENT
```

This is the default environment which is selected at startup. The standard operating system command line interpreter will be loaded to execute the commands. You can use the builtin commands of the command line interpreter, often called shell, or any other program which the command line interpreter can find and load.

`COMMAND`

```
alias CMD
```

This environment loads the named program directly. You have to supply a path if this is needed for the current operating system to load the program. You can't use builtin shell functionality like system redirections like you can with `SYSTEM`. `Regina`'s redirections are more powerful and work in either environment.

`PATH`

This works like the environment `COMMAND` but `Regina` uses the standard operating system search rules for programs. This is done by searching through the items of the `PATH` system-environment variable in most operation systems.

The definition of `REXX` says nothing about which environment is preselected when you invoke the interpreter, although `TRL` defines that one environment is automatically preselected when starting up a `REXX` script. Note that there is no `NONE` environment in standard `REXX`, i.e. an environment that ignores commands, but some interpreters implement the `TRACE` setting ??? to accomplish this. `Regina` uses the environment `SYSTEM` as the preselected environment as mentioned above. More implementation specific details can be found in the section implementation specific documentation for `Regina`.

The list of environments will be saved across subroutine calls; so the effect of any `ADDRESS` clauses in a subroutine will cease upon return from the subroutine.

ADDRESS Redirections

ANSI defines redirections for the ADDRESS statement. This feature has been missing from **Regina** until version 3.0; although you have had the chance to redirect input and output by using **LIFO** and **FIFO** modifiers on command strings.

These command modifiers still exist and have a higher precedence than the ANSI defined redirections. Note, that **LIFO** and **FIFO** can be used by the newer redirection system. But, first of all, some examples show the usage of ADDRESS redirections.

```
ADDRESS SYSTEM "sort" WITH INPUT STEM names. OUTPUT STEM names.  
ADDRESS SYSTEM "myprog" WITH INPUT STEM somefood. OUTPUT STREAM  
    prg.out ERROR STEM oops.  
ADDRESS PATH WITH INPUT FIFO '' OUTPUT NORMAL  
ADDRESS SYSTEM WITH INPUT FIFO '' OUTPUT FIFO '' ERROR NORMAL  
ADDRESS SYSTEM "fgrep 'bongo'" WITH INPUT STREAM feeder
```

The first command instructs the default command line interpreter to call the program called *sort*. The input for the command is read from the stem *names*. (note the trailing period) and the output is sent back to the same stem variable after the command terminates. Thus, bothering about the implementation of a fast sort algorithm for a stem is as simple as calling a program which can actually do the sort.

A program called *myprog* is called in the second case. The input is fetched from the stem *somefood*. (again note the trailing period), and the standard output of the program is redirected to the stream called PRG.OUT (note it is upper case using standard REXX rules). Any generated error messages via the standard error stream are redirected to the stem called *oops*. Note the problematic PRG.OUT. You have to use a symbol and can't use strings.

In the third example, the redirection behaviour of the environment PATH is changed for all future uses. The input for all commands addressed to this environment is fetched from the standard stack in FIFO order. After each call the stack will be flushed. The output is sent to the default output stream, which is the current console in most cases. The behaviour for error messages is not changed.

The fourth example allows pipes between commands in the environment; SYSTEM for all future uses. The input is fetched from the default stack and sent to the default stack after each command. The stack itself is flushed in between. Each executed program will write to something which is the input to the next called command. The error redirection is set or set back to the initial behaviour of writing to the standard error stream.

The fifth example relates to the fourth. The default stack has to be filled with something initially. This is done by the redirection to the stream FEEDER while writing the output of the *fgrep* command to the default FIFO as declared in example four. After this, a single line with a simple *sort* command will sort the output of *fgrep* and place it in the default stack. You can fetch the final output of your pipe cascade by reading the stack contents. This statement overwrites some of the rules of the fourth example temporarily.

You can see the powerful possibilities of the redirection command. The disadvantage is the loss of a direct overview of what happens after a permanent redirection command has executed.

It's now the time to show you all rules and semantics of the redirection.

Rules for the redirection by the keyword WITH of the ADDRESS statement:

- Every environment has its own default *redirectionset*.
- Every *redirectionset* consists of three independent *redirectionstreams*; standard input (INPUT), standard output (OUTPUT) and standard error (ERROR). Users with some experiences with Unix, DOS & Windows or OS/2 may remember the redirection commands of the command line interpreter which can redirect each of the streams, too. This is nearly the same.

- Each *redirectionstream* starts with the program-startup streams given to REXX when invoking the interpreter. These can be reset to the startup default by specifying the argument NORMAL for each *redirectionstream*.
- The sequence of the *redirectionstreams* is irrelevant.
- You can specify each stream only once per statement.
- Redirections can be intermixed. This means you can let both the OUTPUT and the ERROR redirection point to the same "thing". The data from the different channels will be put to the assigned "thing" as they arrive. ANSI's point of view isn't very clear at this point. The state to keep the output different for files and put them together after the called program finished while the data shall be mixed at once when using stems. Regina always mixes the fetched data at once.
- Redirections from and to the same source/destination try to keep the data consistent. If the INPUT/OUTPUT pair or the INPUT/ERROR pair point to the same destination, the content of the input or output channel is buffered so that writing to the output won't overwrite the input.
- All *redirectionstreams* are entered by its name (e.g. INPUT), a redirection processor (e.g. STREAM) and a destination symbol (e.g. OUT_FN) following the rule to the redirection processor. This means that you have to enter a dot after a symbol name for a stem, or any symbol for the rest of the processors, in which case the content of the symbol is used as a normal variable.
- Both OUTPUT and ERROR streams can replace or append the data to the destination. Simply append either APPEND or REPLACE immediately after the OUTPUT or ERROR keywords. REPLACE is the default.
- The destination is checked or cleared prior to the execution of the command.
- ANSI defines two redirection processors: STEM and STREAM. The processors LIFO and FIFO are allowed extensions to the standard.
- The processor STEM uses the content of the symbol destination. 0 to access the count of the currently accessible lines. destination is the given destination name, of course. destination.0 must be filled with a whole, non-negative number in terms of the DATATYPE builtin function. Each of *n* lines can be addressed by appending the whole number on to the stem. Example: STEMfoo.is given, FOO.0 contains 3. This indicates three content lines. They are the contents of the symbols FOO.1 and FOO.2 and FOO.3. nto
- The processor STREAM uses the content of the symbol destination to use a stream as known in the STREAM builtin function. The usage is nearly equivalent to the commands LINEIN destination or LINEOUT destination for accessing the contents of the file. An empty variable (content set to the empty string) as the content of the destination is allowed and indicates the default input, output or error streams given to the REXX program. This is equivalent to the NORMAL keyword.
- The processor LIFO uses the content of the symbol destination as a queue name. New lines are pushed in last-in, first-out order to the queue. An empty destination string is allowed and describes the default queue. Lines are fetched from the queue if this processor is used for the INPUT stream.
- The processor FIFO uses the content of the symbol destination as a queue name. New lines are pushed in first-in, first-out order to the queue. An empty destination string is allowed and describes the default queue. Lines are fetched from the queue if this processor is used for the INPUT stream.
- On INPUT, all the data in the input stream is read up to either the end of the input data or until the called process terminates. The latter one may be determined after feeding up the input stream of the called process with unused data. Thus, there is no way to say if data is used or not. This isn't a problem with STEMs. But all file related sequential access objects including LIFOs and FIFOs may have lost data between two calls. Imagine an input file (STREAM) with three lines:

```
One line
DELIMITER
Second line
```

and furthermore two processes **p1** and **p2** called WITHINPSTREAM **f** with **f** containing the three lines above. **p1** reads lines up until a line containing DELIMITER and **p2** processes the rest. It is very likely that the second process won't fetch any line because the stream may be processed by REXX, and REXX may have put one or more lines ahead into the feeder pipe to the process. This might or might not happen. It is implementation dependent and Regina shows this behaviour. The input object is checked for existence and if it is properly setup before the command is started.

In short: INPUT may or may not use the entire input.

- Both OUTPUT and ERROR objects are checked for being properly setup just before the command starts. REPLACE is implemented as a deletion just before the command starts. Note that ANSI doesn't force STEM lines to be dropped in case of a replacement. A big stem with thousands of lines will still exist after a replacement operation if the called

command doesn't produce any output. Just destination.0 is set to 0.

The redirection of commands is a mystery to many people and it will continue to be. You can thank all the people who designed stacks, queues, pipelines and all the little helper utilities of a witch's kitchen of process management.

8.2 The ARG Instruction

```
ARG [ template ] ;
```

The ARG instruction will parse the argument strings at the current procedural level into the template. Parsing will be performed in uppercase mode. This clause is equivalent to:

```
PARSE UPPER ARG [ template ] ;
```

For more information, see the PARSE instruction. Note that this is the only situation where a multi-string template is relevant.

Example: Beware assignments

The similarity between ARG and PARSE UPPER ARG has one exception. Suppose the PARSE UPPER ARG has an absolute positional pattern as the first element in the template, like:

```
parse upper arg =(foo) bar
```

This is not equivalent to an ARG instruction, because ARG instruction would become an assignment. A simple trick to avoid this problem is just to prepend a placeholder period (`.`) to the pattern, thus the equals sign (`=`) is no longer the second token in the new ARG instruction. Also, unless the absolute positional pattern is indirect, the equals sign can be removed without changing the meaning of the statement.

8.3 The CALL Instruction

```
CALL = routine [ parameter ]  
      [ , [ parameter ] ... ] ;  
      { ON | OFF } condition [ NAME label ] ;
```

The CALL instruction invokes a subroutine, named by *routine*, which can be internal, built-in, or external; and the three repositories of functions are searched in that order. *routine* is searched for *routine* in that order. The token *routine* must be either a literal string or a symbol (which is taken literally). However, if *routine* is a literal string, the pool of internal subroutines is not searched. Note that some interpreters may have additional repositories of labels to search.

In a CALL instruction, each *parameter* is evaluated, strictly in order from left to right, and passed as an argument to the subroutine. A *parameter* might be left out (i.e. an empty argument), which is not the same as passing the null string as an argument.

Users often confuse a parameter which is the null string with leaving out the parameter. However, this is two very different situations. Consider the following call to the built-in function TRANSLATE () :

```

say translate('abcDEF' ) /* says ABCDEF */

say translate('abcDEF', "") /* says abcDEF */

say translate('abcDEF', , "") /* says ' ' */

```

The `TRANSLATE()` function is able to differ between receiving the null string (i.e. a defined string having zero length), from the situation where a parameter was not specified (i.e. the undefined string). Since `TRANSLATE()` is one of the few functions where the parameters' default values are very different from the null string, the distinction becomes very visible.

BreakageAlert!!

Prior to Version 3.1 of Regina, the following syntactical use of the `CALL` instruction was valid:

```
call myfunc('abcDEF', , "")
```

This syntax is not allowed by ANSI and use of this syntax will now result in Error 37.1

BreakageAlert!!

For the `CALL` instruction, watch out for interference with line continuation. If there are trailing commas, it might be interpreted as line continuation. If a `CALL` instruction uses line continuation between two parameters, two commas are needed: one to separate the parameters, and one to denote line continuation.

A number of settings are stored across internal subroutine calls. An internal subroutine will inherit the values in effect when the call is made, and the settings are restored on exit from the subroutine. These settings are:

- Condition traps, see chapter `Conditions`.
- Current trapped condition, see section `CTS`.
- `NUMERIC` settings, see section `Numeric`.
- `ADDRESS` environments, see section `Address`.
- `TRACE` mode, see section `Trace` and chapter [not yet written].
- The elapsed time clock, see section `Time`.

Also, the `OPTIONS` settings may or may not be restored, depending on the implementation. Further, a number of other things may be saved across internal subroutines. The effect on variables are controlled by the `PROCEDURE` instruction in the subroutine itself. The state of all `DO`-loops will be preserved during subroutine calls.

Example: Subroutines and trace settings

Subroutines cannot be used to set various settings like trace settings, `NUMERIC` settings, etc. Thus, the following code will not work as intended:

```

say digits() /* says 9, maybe */
call inc_digits
say digits() /* still says 9 */
exit

```

```

inc_digits:
    numeric digits digits() + 1
    return

```

The programmer probably wanted to call a routine which incremented the precision of arithmetic operations. However, since the setting of `NUMERIC DIGITS` is saved across subroutine calls, the new value set in `inc_digits` is lost at return from that routine. Thus, in order to work correctly, the `NUMERIC` instruction must be located in the main routine itself.

Built-in subroutines will have no effect on the settings, except for explicitly defined side effects. Nor will external subroutines change the settings. For all practical purposes, an external subroutine is conceptually equivalent to reinvoking the interpreter in a totally separated process.

If the name of the subroutine is specified by a literal string, then the name will be used as-is; it will not be converted to uppercase. This is important because a routine which contains lowercase letters can only be invoked by using a literal string as the routine name in the `CALL` instruction.

Example: Labels are literals

Labels are literal, which means that they are neither tail-substituted nor substituted for the value of the variable. Further, this also means that the setting of `NUMERIC DIGITS` has no influence on the section of labels, even when the labels are numeric symbols. Consider the following code:

```

call 654.32
exit

654.321:
    say here
    return

654.32:
    say there
    return

```

In this example, the second of the two subroutines are always chosen, independent of the setting of `NUMERIC DIGITS`. Assuming that `NUMERIC DIGITS` are set to 5, then the number 654.321 is converted to 654.32, but that does not affect labels. Nor would a statement `CALL 6.5432E2` call the second label, even though the numeric value of that symbol is equal to that of one of the labels.

The called subroutines may or may not return data to the caller. In the calling routine, the special variable `RESULT` will be set to the return value or dropped, depending on whether any data was returned or not. Thus, the `CALL` instruction is equivalent to calling the routine as a function, and

assigning the return value to `RESULT`, except when the *routine* does not return data.

In REXX, recursive routines are allowed. A minimum number of 100 nested internal and external subroutine invocations, and support for a minimum of 10 parameters for each call are required by REXX. See chapter `Limits` for more information concerning implementation limits.

When the token following `CALL` is either `ON` or `OFF`, the `CALL` instruction is not used for calling a subroutine, but for setting up condition traps. In this case, the third token of the clause must be the name of a condition, which is set up to be changed.

If the second token was `ON`, then there can be either three or five tokens. If the five token version is used, then the fourth token must be `NAME` and the fifth token is taken to be the symbolic name of a label, which is the condition handler. This name can be either a constant string, or a symbol, which is taken literally. When `OFF` is used, then a named condition trap is turned off.

Note that the `ON` and `OFF` forms of the `CALL` instruction were introduced in TRL2. Thus, they are not likely to be present on older interpreters. More information about conditions and condition traps are given in a chapter `Conditions`.

8.4 The DO/END Instruction

```
DO [ repetitor ] [ conditional ] ;
    [ clauses ]
END [ symbol ] ;

repetitor : = symbol = expr1 [ TO expr2 ]
    [ BY expr3 ] [ FOR expr4 ]
    expr5
    FOREVER

conditional : = WHILE expr1
    UNTIL expr2
```

The `DO/END` instruction is the instruction used for looping and grouping several statements into one block. This is a multi-clause instruction.

The most simple case is when there is no *repetitor* or *conditional*, in which case it works like `BEGIN/END` in Pascal or `{...}` in C. I.e. it groups zero or more REXX clauses into one conceptual statement.

The *repetitor* subclause controls the control variable of the loop, or the number of repetitions. The *expr*₅ subclause may specify a certain number of repetitions, or you may use `FOREVER` to go on looping forever.

If you specify the control variable *symbol*, it must be a variable symbol, and it will get the initial value *expr*₁ at the start of the loop. At the start of each iteration, including the first, it will be checked whether it has reached the value specified by *expr*₂. At the end of each iteration the value *expr*₃ is added to the control variable. The loop will terminate after at most *expr*₄ iterations. Note that all these expressions

are evaluated only once, before the loop is entered for the first iteration.

You may also specify `UNTIL` or `WHILE`, which take a boolean expression. `WHILE` is checked before each iteration, immediately after the maximum number of iteration has been performed. `UNTIL` is checked after each iteration, immediately before the control variable is incremented. It is not possible to specify both `UNTIL` and `WHILE` in the same `DO` instruction.

The `FOREVER` keyword is only needed when there is no *conditional*, and the *repetitor* would also be empty if `FOREVER` was not specified. Actually, you could rewrite this as `DO WHILE 1`. The two forms are equivalent, except for tracing output.

The subclauses `TO`, `BY`, and `FOR` may come in any order, and their expressions are evaluated in the order in which they occur. However, the initial assignment must always come first. Their order may affect your program if these expressions have any side effects. However, this is seldom a problem, since it is quite intuitive. Note that the counting of iterations, if the `FOR` subclause has been specified, is never affected by the setting of `NUMERIC DIGITS`.

Example: Evaluation order

What may prove a real trap, is that although the value to which the control variable is set is evaluated before any other expressions in the *repetitor*, it is assigned to the control variable after all expressions in the *repetitor* have been evaluated.

The following code illustrates this problem:

```
ctrl = 1
do ctrl=f(2) by f(3) to f(5)
    call func(6)
end
call func(7)
exit

f:
    say 'ctrl='ctrl 'arg='arg(1)
    return arg(1)
```

This code produces the following output:

```
ctrl=1 arg=2
ctrl=1 arg=3
ctrl=1 arg=5
ctrl=2 arg=6
ctrl=5 arg=6
ctrl=8 arg=7
```

Makesure you understand why the program produces this output. Failure to understand this may give you a surprise later, when you happen to write a complex `DO`-instruction, and do not get the expected

result.

If the `TO` expression is omitted, there is no checking for an upper bound of the expression. If the `BY` subclause is omitted, then the default increment of 1 is used. If the `FOR` subclause is omitted, then there is no checking for a maximum number of iterations.

Example: Loop convergence For the reasons just explained, the instruction:

```
do ctrl=1
    nop /* and other statements */
end
```

will start with `CTRL` being 1, and then iterate through 2, 3, 4, ..., and never terminate except by `LEAVE`, `RETURN`, `SIGNAL`, or `EXIT`.

Although similar constructs in other languages typically provoke an overflow at some point, something "strange" happens in `REXX`. Whenever the value of `ctrl` becomes too large, the incrementation of that variable produces a result that is identical to the old value of `ctrl`. For `NUMERIC DIGITS` set to 9, this happens when `ctrl` becomes 1.00000000E+9. When adding 1 to this number, the result is still 1.00000000E+9. Thus, the loop "converges" at that value.

If the value of `NUMERIC DIGITS` is 1, then it will "converge" at 10, or 1E+1 which is the "correct" way of writing that number under `NUMERIC DIGITS 1`. You can in general disregard loop "convergence", because it will only occur in very rare situations.

Example: Difference between `UNTIL` and `WHILE`

One frequent misunderstanding is that the `WHILE` and `UNTIL` subclauses of the `DO/END` instruction are equivalent, except that `WHILE` is checked before the first iteration, while `UNTIL` is first checked before the second iteration.

This may be so in other languages, but in `REXX`. Because of the order in which the parts of the loop are performed, there are other differences. Consider the following code:

```
count = 1
do i=1 while count \= 5
    count = count + 1
end
say i count

count = 1
do i=1 until count=5
    count = count + 1
end
say i count
```

After the first loop, the numbers 6 and 5, while in the second loop, the numbers 5 and 5 are written out. The reason is that a `WHILE` clause is checked after the control variable of the loop has been incremented, but an `UNTIL` expression is checked before the incrementation.

A loop can be terminated in several ways. A `RETURN` or `EXIT` instruction terminates all active loops in the procedure level terminated. Further, a `SIGNAL` instruction transferring control (i.e. neither a `SIGNAL ON` nor `SIGNAL OFF`) terminates all loops at the current procedural level. This applies even to "implicit" `SIGNAL` instructions, i.e. when triggering a condition handler by the method of `SIGNAL`. A `LEAVE` instruction terminates one or more loops. Last but not least, a loop can terminate itself, when it has reached its specified stop conditions.

Note that the `SIGNAL` instruction terminates also non-repetitive loops (or rather: `DO/END` pairs), thus after an `SIGNAL` instruction, you must not execute an `END` instruction without having executed its corresponding `DO` first (and after the `SIGNAL` instruction). However, as long as you stay away from the `END`s, it is all right according to TRM to execute code within a loop without having properly activated the loop itself.

Note that on exit from a loop, the value of the control variable has been incremented once after the last iteration of the loop, if the loop was terminated by the `WHILE` expression, by exceeding the number of max iterations, or if the control variable exceeded the stop value. However, the control variable has the value of the last iteration if the loop was terminated by the `UNTIL` expression, or by an instruction inside the loop (e.g. `LEAVE`, `SIGNAL`, etc.).

The following algorithm in `REXX` codes shows the execution of a `DO` instruction, assuming that *expri*, *exprt*, *exprb*, *exprf*, *exprw*, *expru*, and *symbol* have been taken from the syntax diagram of `DO`.

```
@expri = expri
@exprt = exprt
@exprb = exprb
@exprf = exprf
@iters = 0

symbol = @expri

start_of_loop:
    if symbol > @exprt then signal after_loop
    if @iters > @exprf then signal after_loop
    if exprw then signal after_loop
        instructions
end_of_loop:
    if \expru then signal after_loop
    symbol = symbol + @exprb
    signal start_of_loop

after_loop:
```

Some notes are in order for this algorithm. First, it uses the `SIGNAL` instruction, which is defined to

terminate all active loops. This aspect of the `SIGNAL` instruction has been ignored for the purpose of illustrating the `DO`, and consequently, the codes shown above is not suitable for nested loops. Further, the order of the first four statements should be identical to the order in the corresponding subclauses in the *repetitor*. The code has also ignored that the `WHILE` and the `UNTIL` subclauses cannot be used in the same `DO` instruction. And in addition, all variables starting with the `@` sign, are assumed to be internal variables, private to this particular loop. Within *instructions*, a `LEAVE` instruction is equivalent to `signal after_loop`, while a `ITERATE` instruction is equivalent to `signal end_of_loop`.

8.5 The DROP Instruction

```
DROP symbol [ symbol ... ] ;
```

The `DROP` instruction makes the named *variables* uninitialized, i.e. the same state that they had at the startup of the program. The list of variable names are processed strictly from left to right and dropped in that order. Consequently, if one of the variables to be dropped is used in a tail of another, then the order might be significant. E.g. the following two `DROP` instructions are not equivalent:

```
bar = 'a'
drop bar foo.bar /* drops 'BAR' and 'FOO.BAR' */
bar = 'a'
drop foo.bar bar /* drops 'FOO.a' and 'BAR' */
```

The *variable* terms can be either a variable symbol or a symbol enclosed in parentheses. The former form is first tail-substituted, and then taken as the literal name of the symbol to be dropped. The result name is the variable to drop. In the latter form, the value of the variable symbol inside the parentheses is retrieved and taken as a space-separated list of symbols. Each of these symbols is tail-substituted (if relevant); and the result is taken as the literal name of a variable to be dropped. However, this process is not recursive, so that the list of names referred to indirectly cannot itself contain parentheses. Note that the second form was introduced in TRL2, mainly in order to make `INTERPRET` unnecessary.

In general, things contained in parentheses can be any valid `REXX` expression, but this does not apply to the `DROP`, `PARSE`, and `PROCEDURE` instructions.

Example: Dropping compound variables

Note a potential problem for compound variables: when a stem variable is set, it will not set a default value, rather it will assign "all possible variables" in that stem collection at once. So dropping a compound variable in a stem collection for which the stem variable has been set, will set that compound variable to the original uninitialized value; not the value of the stem variable. See section *Assign* for further notes on assignments. To illustrate consider the code:

```
foo. = 'default'
drop baz bar foo.bar
say foo.bar foo.baz /* says 'FOO.BAR default' */
```

In this example, the `SAY` instruction writes out the value of the two compound variables `FOO.BAR` and `FOO.BAZ`. When performing tail-substitution for these, the interpreter finds that both `BAR` and `BAZ` are

uninitialized. Further, `FOO.BAR` has also been made uninitialized, while `FOO.BAZ` has the value assigned to it in the assignment to the stem variable.

Example: Tail-substitution in `DROP`

For instance, suppose that the variable `FOO` has the value `bar`. After being dropped, `FOO` will have its uninitialized value, which is the same as its name: `FOO`. If the variable to be dropped is a stem variable, then both the stem variable and all compound variables of that stem become uninitialized.

```
bar = 123
drop foo.bar /* drops 'FOO.123' */
```

Technically, it should be noted that some operations involving dropping of compound variables can be very space consuming. Even though the standard does not operate with the term "default value" for the value assigned to a stem variable, that is the way in which it is most likely to be implemented. When a stem is assigned a value, and some of its compound variables are dropped afterwards, then the interpreter must use memory to store references to the variables dropped. This might seem counterintuitive at first, since dropping ought to release memory, not allocate more.

There is a parallel between `DROP` and `PROCEDURE EXPOSE`. However, there is one important difference, although `PROCEDURE EXPOSE` will expose the name of a variable enclosed in parentheses before starting to expose the symbol that variable refers to, this is not so for `DROP`. If `DROP` had mimicked the behavior of `PROCEDURE EXPOSE` in this matter, then the whole purpose of indirect specifying of variables in `DROP` would have been defeated.

Dropping a variable which does not have a value is not an error. There is no upper limit on the number of variables that can be dropped in one `DROP` clause, other than restrictions on the clause length. If an exposed variable is dropped, the variable in the caller is dropped, but the variable remains exposed. If it is reassigned a value, the value is assigned to a variable in the caller routine.

8.6 The `EXIT` Instruction

```
EXIT [ expr ] ;
```

Terminates the `REXX` program, and optionally returns the expression *expr* to the caller. If specified, *expr* can be any string. In some systems, there are restrictions on the range of valid values for the *expr*. Often the return expression must be an integer, or even a non-negative integer. This is not really a restriction on the `REXX` language itself, but a restriction in the environment in which the interpreter operates, check the system dependent documentation for more information.

If *expr* is omitted, nothing will be returned to the caller. Under some circumstances that is not legal, and might be handled as an error or a default value might be used. The `EXIT` instruction behaves differently in a "program" than in an external subroutine. In a "program", it returns control to the caller e.g. the operating system command interpreter. While for an external routine, it returns control to the calling `REXX` script, independent of the level of nesting inside the external routine being terminated.

	RETURN	EXIT
At the main level of the program	Exits program	Exits program
At an internal subroutine level of the program	Exit subroutine, and returns to caller	Exits program
At the main level of an external subroutine	Exit the external subroutine	Exit the external subroutine
At a subroutine level within an external subroutine	Exit the subroutine, returning to calling routine within external subroutine script	Exit the external subroutine

Actionsof **RETURN**and **EXIT**Instructions

If terminating an external routine (i.e. returning to the calling REXX script) any legal REXX string value is allowed as a return value. Also, no return value can be returned, and in both cases, this information is successfully transmitted back to the calling routine. In the case of a function call (as opposed to a subroutine call), returning no value will raise SYNTAX condition { 44 }. The table above describes the action taken by the EXIT and RETURN instruction in various situations.

8.7 The IF/THEN/ELSE Instruction

```
IF expr [ ; ] THEN [ ; ] statement
                        [ ELSE [ ; ] statement ]
```

This is a normal if-construct. First the boolean expression *expr* is evaluated, and its value must be either 0 or 1 (everything else is a syntax error which raises SYNTAX condition number { 34 }). Then, the statement following either THEN or ELSE is executed, depending on whether *expr* was 1 or 0, respectively.

Note that there must come a statement after THEN and ELSE. It is not allowed to put just a null-clause (i.e. a comment or a label) there. If you want the THEN or ELSE part to be empty, use the NOP instruction. Also note that you cannot directly put more than one statement after THEN or ELSE; you have to package them in a DO-END pair to make them a single, conceptual statement.

After THEN, after ELSE, and before THEN, you might put one or more clause delimiters (new lines or semicolons), but these are not required. Also, the ELSE part is not required either, in which case no code is executed if *expr* is false (evaluates to 0). Note that there must also be a statement separator before ELSE, since the statement must be terminated. This also applies to the statement after ELSE. However, since *statement* includes a trailing clause delimiter itself, this is not explicitly shown in the syntax diagram.

Example: Dangling ELSE

Note the case of the "dangling" ELSE. If an ELSE part can correctly be thought of as belonging to more than one IF/THEN instruction pair, it will be parsed as belonging to the closest (i.e. innermost) IF instruction:

```

parse pull foo bar
if foo then
    if bar then
        say 'foo and bar are true'
    else
        say 'one or both are false'

```

In this code, the `ELSE` instruction is nested to the innermost `IF`, i.e. to `IF BAR THEN`.

8.8 The `INTERPRET` Instruction

```
INTERPRET expr ;
```

The `INTERPRET` instruction is used to dynamically build and execute `REXX` instructions during run-time. First, it evaluates the expression *expr*, and then parses and interprets the result as a (possibly empty) list of `REXX` instructions to be executed. For instance:

```

foo = 'hello, world'
interpret 'say "'foo'!"'

```

executes the statement `SAY "hello, world!"` after having evaluated the expression following `INTERPRET`. This example shows several important aspects of `INTERPRET`. Firstly, it's very easy to get confused by the levels of quotes, and a bit of caution should be taken to nest the quotes correctly. Secondly, the use of `INTERPRET` does not exactly improve readability.

Also, `INTERPRET` will probably increase execution time considerably if put inside loops, since the interpreter may be forced to reparse the source code for each iteration. Many optimizing `REXX` interpreters (and in particular `REXX` compilers) have little or no support for `INTERPRET`. Since virtually anything can happen inside it, it is hard to optimize, and it often invalidates assumptions in other parts of the script, forcing it to ignore other possible optimizations. Thus, you should avoid `INTERPRET` when speed is a premium.

There are some restrictions on which statements can be inside an `INTERPRET` statement. Firstly, labels cannot occur there. `TRL` states that they are not allowed, but you may find that in some implementations labels occurring there will not affect the label symbol table of the program being run. Consider the statement:

```

interpret 'signal there; there: say hallo'
there:

```

This statement transfers control to the label `THERE` in the program, never to the `THERE` label inside the expression of the `INTERPRET` instruction. Equivalently, any `SIGNAL` to a label `THERE` elsewhere in the program never transfers control to the label inside the `INTERPRET` instruction. However, labels are strictly speaking not allowed inside `INTERPRET` strings.

Example: Self-modifying Program

There is an idea for a self-modifying program in REXX which is basically like this:

```
string = ''
do i=1 to sourceline()
    string = string ';' sourceline(i)
end

string = transform( string )
interpret string
exit

transform: procedure
    parse arg string
    /* do some transformation on the argument */
    return string
```

Unfortunately, there are several reasons why this program will not work in REXX, and it may be instructive to investigate why. Firstly, it uses the label TRANSFORM, which is not allowed in the argument to INTERPRET. The interpreter will thus refer to the TRANSFORM routine of the "outermost" invocation, not the one "in" the INTERPRET string.

Secondly, the program does not take line continuations into mind. Worse, the SOURCELINE () built-in function refers to the data of the main program, even inside the code executed by the INTERPRET instruction. Thirdly, the program will never end, as it will nest itself up till an implementation-dependent limit for the maximum number of nested INTERPRET instructions.

In order to make this idea work better, temporary files should be used.

On the other hand, loops and other multi-clause instructions, like IF and SELECT occur inside an INTERPRET expression, but only if the whole instruction is there; you cannot start a structured instruction inside an INTERPRET instruction and end it outside, or vice-versa. However, the instruction SIGNAL is allowed even if the label is not in the interpreted string. Also, the instructions ITERATE and LEAVE are allowed in an INTERPRET, even when they refer to a loop that is external to the interpreted string.

Most of the time, INTERPRET is not needed, although it can yield compact and interesting code. If you do not strictly need INTERPRET, you should consider not using it, for reasons of compatibility, speed, and readability. Many of the traditional uses of INTERPRET have been replaced by other mechanisms in order to decrease the necessity of INTERPRET; e.g. indirect specification of variables in EXPOSE and DROP, the improved VALUE () built-in function, and indirect specification of patterns in templates.

Only semicolon (;) is allowed as a clause delimiter in the string interpreted by an INTERPRET instruction. The colon of labels cannot be used, since labels are not allowed. Nor does specific end-of-line character sequences have any defined meaning there. However, most interpreters probably allow

the end-of-line character sequence of the host operating system as an alternative clause delimiter. It is interesting to note that in the context of the `INTERPRET` instruction, an implicit, trailing clause delimiter is always appended to the string to be interpreted.

8.9 The `ITERATE` Instruction

```
ITERATE [ symbol ] ;
```

The `ITERATE` instruction will iterate the innermost, active loop in which the `ITERATE` instruction is located. If *symbol* is specified, it will iterate the innermost, active loop having *symbol* as a control variable. The simple `DO/END` statement without a *repetitor* and *conditional* is not affected by `ITERATE`. All active multi-clause structures (`DO`, `SELECT`, and `IF`) within the loop being iterated are terminated.

The effect of an `ITERATE` is to immediately transfer control to the `END` statement of the affected loop, so that the next (if any) iteration of the loop can be started. It only affects loops on the current procedural level. All actions normally associated with the end of an iteration are performed.

Note that *symbol* must be specified literally; i.e. tail substitution is not performed for compound variables. So if the control variable in the `DO` instruction is `FOO.BAR`, then *symbol* must use `FOO.BAR` if it is to refer to the control variable, no matter the value of the `BAR` variable.

Also note that `ITERATE` (and `LEAVE`) are means of transferring control in the program, and therefore they are related to `SIGNAL`, but they do have the effect of automatically terminating all active loops on the current procedural level, which `SIGNAL` has.

Two types of errors can occur. Either *symbol* does not refer to any loop active at the current procedural level; or (if *symbol* is not specified) there does not exist any active loops at the current procedural level. Both errors are reported as `SYNTAXcondition{28}`.

8.10 The `LEAVE` Instruction

```
LEAVE [ symbol ] ;
```

This statement terminates the innermost, active loop. If *symbol* is specified, it terminates the innermost, active loop having *symbol* as a control variable. As for scope, syntax, errors, and functionality, it is identical to `ITERATE`, except that `LEAVE` terminates the loop, while `ITERATE` lets the loop start on the next iteration normal iteration. No actions normally associated with the normal end of an iteration of a loop are performed for a `LEAVE` instruction.

Example: Iterating a simple `DO/END`

In order to circumvent this, a simple `DO/END` can be rewritten as this:


```

if foo then do until 1
    say 'This is a simple DO/END group'
    say 'but it can be terminated by'
    leave
    say 'iterate or leave'
end

```

This shows how `ITERATE` has been used to terminate what for all practical purposes is a simple `DO/END` group. Either `ITERATE` or `LEAVE` can be used for this purpose, although `LEAVE` is perhaps marginally faster.

8.11 The `NOP` Instruction

```
NOP ;
```

The `NOP` instruction is the "no operation" statement; it does nothing. Actually, that is not totally true, since the `NOP` instruction is a "real" statement (and a placeholder), as opposed to null clauses. I've only seen this used in two circumstances.

- After any `THEN` or `ELSE` keyword, where a statement is required, when the programmer wants an empty `THEN` or `ELSE` part. By the way, this is the intended use of `NOP`. Note that you cannot use a null clause there (label, comment, or empty lines), since these are not parsed as "independent" statements.
- I have seen it used as "trace-bait". That is, when you start interactive trace, the statement immediately after the `TRACE` instruction will be executed before you receive interactive control. If you don't want that to happen (or maybe the `TRACE` instruction was the last in the program), you need to add an extra dummy statement. However, in this context, labels and comments can be used, too.

8.12 The `NUMERIC` Instruction

```

NUMERIC DIGITS [ expr ] ;
          FORM [ SCIENTIFIC | ENGINEERING | [ VALUE ] expr ] ;
          FUZZ [ expr ] ;

```

`REXX` has an unusual form of arithmetic. Most programming languages use integer and floating point arithmetic, where numbers are coded as bits in the computer's native memory words. However, `REXX` uses floating point arithmetic of arbitrary precision, that operates on strings representing the numbers. Although much slower, this approach gives lots of interesting functionality. Unless number-crunching is your task, the extra time spent by the interpreter is generally quite acceptable and often almost unnoticeable.

The `NUMERIC` statement is used to control most aspects of arithmetic operations. It has three distinct forms: `DIGITS`, `FORM` and `FUZZ`; which to choose is given by the second token in the instruction:

DIGITS

Is used to set the number of significant digits in arithmetic operations. The initial value is 9, which is also the default value if `expr` is not specified. Large values for `DIGITS` tend to slow

downsome arithmetic operations considerably. If specified, *expr* must be a positive integer.

FUZZ

Is used in numeric comparisons, and its initial and default value is 0. Normally, two numbers must have identical numeric values for a number of their most significant digits in order to be considered equal. How many digits are considered is determined by `DIGITS`. If `DIGITS` is 4, then 12345 and 12346 are equal, but not 12345 and 12356. However, when `FUZZ` is non-zero, then only the `DIGITS` minus `FUZZ` most significant digits are checked. E.g. if `DIGITS` is 4 and `FUZZ` is 2, then 1234 and 1245 are equal, but not 1234 and 1345.

The value for `FUZZ` must be a non-negative integer, and less than the value of `DIGITS`. `FUZZ` is seldom used, but is useful when you want to make comparisons less influenced by inaccuracies. Note that using values of `FUZZ` that is close to `DIGITS` may give highly surprising results.

FORM

Is used to set the form in which exponential numbers are written. It can be set to either `SCIENTIFIC` or `ENGINEERING`. The former uses a mantissa in the range 1.000... to 9.999..., and an exponent which can be any integer; while the latter uses a mantissa in the range 1.000... to 999.999..., and an exponent which is divisible by 3. The initial and default setting is `SCIENTIFIC`. Following the subkeyword `FORM` may be the subkeywords `SCIENTIFIC` and `ENGINEERING`, or the subkeyword `VALUE`. In the latter case, the rest of the statement is considered an expression, which will evaluate to either `SCIENTIFIC` or `ENGINEERING`. However, if the first token of the expression following `VALUE` is neither a symbol nor a literal string, then the `VALUE` subkeyword can be omitted.

The setting of `FORM` never affects the decision about whether to choose exponential form or normal floating point form; it only affects the appearance of the exponential form once that form has been selected.

Many things can be said about the usefulness of `FUZZ`. My impression is that it is seldom used in `REXX` programs. One problem is that it only addresses relative inaccuracy: i.e. that the small value must be within a certain range, that is determined by a percentage of the large value. Often one needs absolute inaccuracy, e.g. two measurements are equal if their difference are less than a certain absolute threshold.

Example: Simulating relative accuracy with absolute accuracy

As explained above, `REXX` arithmetic has only relative accuracy, in order to obtain absolute accuracy, one can use the following trick:

```
numeric fuzz 3
if a=b then
    say 'relative accuracy'
if abs(a-b)<=500 then
    say 'absolute accuracy'
```

In the first `IF` instruction, if `A` is 100,000, then the range of values for `B` which makes the expression

true is 99,500-100,499, i.e. an inaccuracy of about ± 500 . If `A` has the value 10,000,000, then `B` must be within the range 9,950,000-10,049,999; i.e. an inaccuracy of about $\pm 50,000$.

However, in the second `IF` instruction, assuming `A` is 100,000, the expression becomes true for values of `B` in the range 99,500-100,500. Assuming that `A` is 10,000,000, the expression becomes true for values of `B` in the range 9,999,500-10,000,500.

The effect is largely to force an absolute accuracy for these second example, no matter what the values of `A` and `B` are. This transformation has taken place since an arithmetic subtraction is not affected by the `NUMERIC FUZZ`, only numeric comparison operations. Thus, the effect of `NUMERIC FUZZ` on the implicit subtraction in the operation `=` in the first `IF` has been removed by making the subtraction explicit.

Note that there are some minor differences in how numbers are rounded, but this can be fixed by transforming the expression into something more complex.

To retrieve the values set for `NUMERIC`, you can use the built-in functions `DIGITS ()`, `FORM ()`, and `FUZZ ()`. These values are saved across subroutine calls and restored upon return.

8.13 The `OPTIONS` Instruction

```
OPTIONS expr ;
```

The `OPTIONS` instruction is used to set various interpreter-specific options. Its typical uses are to select certain `REXX` dialects, enable optimizations (e.g. time versus memory considerations), etc. No standard dictates what may follow the `OPTIONS` keyword, except that it should be a valid `REXX` expression, which is evaluated. Currently, no specific options are required by any standard.

The content of *expr* is supposed to be word-based, and it is the intention that more than one option can be specified in one `OPTIONS` instruction. `REXX` interpreters are specifically instructed to ignore `OPTIONS` words which they do not recognize. That way, a program can use run-time options for one interpreter, without making other interpreters strip when they see those options. An example of `OPTION` may be:

```
OPTIONS 4.00 NATIVE_FLOAT
```

The instruction might instruct the interpreter to start enforcing language level 4.00, and to use native floating point numbers instead of the `REXX` arbitrary precision arithmetic. On the other hand, it might also be completely ignored by the interpreter.

It is uncertain whether the mode selected by `OPTIONS` will be saved across subroutine calls. Refer to implementation-specific documentation for information about this.

Example: Drawback of `OPTIONS`

Unfortunately, the processing of the `OPTIONS` instruction has a drawback. Since an interpreter is instructed to ignore option-settings that it does not understand, it may ignore options which are

essential for further processing of the program. Continuing might cause a fatal error later, although the behavior that would most precisely point out the problem is a complaint about the non-supported `OPTION` setting. Consider:

```
options 'cms_bifs'  
pos = find( haystack, needle )
```

If this code fragment is run on an interpreter that does not support the `cms_bifs` option setting, then the `OPTIONS` instruction may still seem to have been executed correctly. However, the second clause will generally crash, since the `FIND()` function is still not available. Even though the real problem is in the first line, the error message is reported for the second line.

8.14 The `PARSE` Instruction

```
PARSE [ UPPER ] type [ template ] ;  
      type = { ARG | LINEIN | PULL | SOURCE | VERSION }  
      VALUE [ expr ] WITH  
      VAR symbol
```

The `PARSE` instruction takes one or more source strings, and then parses them using the `template` for directions. The process of parsing is one where parts of a source string are extracted and stored in variables. Exactly which parts, is determined by the patterns. A complete description of parsing is given in chapter [not yet written].

Which strings are to be the source of the parsing is defined by the `type` subclause, which can be any of:

ARG.

The data to use as the source during the parsing is the argument strings given at the invocation of this procedure level. Note that this is the only case where the source may consist of multiple strings.

LINEIN.

Make the `PARSE` instruction read a line from the standard input stream, as if the `LINEIN()` built-in function had been called. It uses the contents of that line (after stripping off end-of-line characters, if necessary) as the source for the parsing. This may raise the `NOTREADY` condition if problems occurred during the read.

PULL.

Retrieves as the source string for the parsing the top most line from the stack. If the stack is empty, the default action for reading an empty stack is taken. That is, it will read a whole line from the standard input stream, strip off any end-of-line characters (if necessary), and use that string as the source.

SOURCE.

The source string for the parsing is a string containing information about how this invocation of the REXX interpreter was started. This information will not change during the execution of a REXX script. The format of the string is:

system invocation filename

Here, the first space-separated word (*system*) is a single word describing the platform on which the system is running. Often, this is the name of the operating system. The second word describes how the script was invoked. TRL2 suggests that *invocation* could be `COMMAND`, `FUNCTION`, or `SUBROUTINE`, but notes that this may be specific to VM/CMS.

Everything after the second word is implementation-dependent. It is indicated that it should refer to the name of the `REXX` script, but the format is not specified. In practice, the format will differ because the format of filenames differs between various operating systems. Also, the part after the second word might contain other types of information. Refer to the implementation-specific notes for exact information.

VALUE *expr* WITH.

This form will evaluate *expr* and use the result of that evaluation as the source string to be parsed. The token `WITH` may not occur inside *expr*, since it is a reserved subkeyword in this context.

VAR *symbol*.

This form uses the current value of the named variable *symbol* (after tail-substitution) as the source string to be parsed. The variable may be any variable symbol. If the variable is uninitialized, then a `NOTREADY` condition will be raised.

VERSION.

This format resembles `SOURCE`, but it contains information about the version of `REXX` that the interpreter supports. The string contains five words, and has the following format:

language level date month year

Where *language* is the name of the language supported by the `REXX` interpreter. This may seem like overkill, since the language is `REXX`, but there may be various different dialects of `REXX`. The word can be just about anything, except for two restrictions, the first four letters should be `REXX` (in uppercase), and the word should not contain any periods. [TRL2] indicates that the remainder of the word (after the fourth character) can be used to identify the implementation.

The second word is the `REXX` language level supported by the interpreter. Note that this is not the same as the version of the interpreter, although several implementations make this mistake. Strictly speaking, neither [TRL1] nor [TRL2] define the format of this word, but a numeric format is strongly suggested.

The last three words (*date*, *month*, and *year*) make up the date part of the string. This is the release date of the interpreter, in the default format of the `DATE ()` built-in function.

Much confusion seems to be related to the second word of `PARSE VERSION`. It describes the language level, which is not the same as the version number of the interpreter. In fact, most interpreters have a version numbering which is independent of the `REXX` language level. Unfortunately, several interpreters make the mistake of using this field as for their own version number. This is very unfortunate for two reasons; first, it is incorrect, and second, it makes it difficult to determine which `REXX` language level the interpreter is supposed to support.

Chances are that you can find the interpreter version number in `PARSE SOURCE` or the first word of `PARSE VERSION`.

The format of the `REXX` language level is not rigidly defined, but `TRL1` corresponds to the language level 3.50, while `TRL2` corresponds to the language level 4.00. Both implicitly indicate that the language level description is a number, and states that an implementation less than a certain number "may be assumed to indicate a subset" of that language level. However, this must not be taken too literally, since language level 3.50 has at least two features which are missing in language level 4.00 (the `Scan` trace setting, and the `PROCEDURE` instruction that is not forced to be the first instruction in a subroutine). [TRH:PRICE] gives a very good overview over the varying functionality of different language levels of `REXX` up to level 4.00.

With the release of the `ANSI REXX Standard` [ANSI] in 1996, the `REXX` language is now rigidly defined. The language level of `ANSI REXX` is 5.00. `Regina` is attempting to keep pace with the `ANSI Standard`. It includes some features of language level 5.00 such as date and time conversions in the `DATE()` and `TIME()` BIFs plus the new BIFs `COUNTSTR()` and `CHANGESTR()`. `Regina` does not supply a complete set of multiple-level error messages as defined in the `ANSI Standard`, nor the extensions to `ADDRESS`, so does not comply to language level 5.00, but currently is a hybrid between 4.00 and 5.00. Thus `PARSE VERSION` will return 4.xx:-)

Note that even though the information of the `PARSE SOURCE` is constant throughout the execution of a `REXX` script, this is not necessarily correct for the `PARSE VERSION`. If your interpreter supports multiple language levels (e.g. through the `OPTIONS` instruction), then it will have to change the contents of the `PARSE VERSION` string in order to comply with different language levels. To some extent, this may also apply to `PARSE SOURCE`, since it may have to comply with several implementation-specific standards.

After the source string has been selected by the `types` subclause in the `PARSE` instruction, this string is parsed into the template. The functionality of templates is common for the `PARSE`, `ARG` and `PULL` instructions, and is further explained in chapter [not yet written].

8.15 The `PROCEDURE` Instruction

```
PROCEDURE [ EXPOSE [ varref [ varref ... ] ] ] ;  
          varref = { symbol | ( symbol ) }
```

The `PROCEDURE` instruction is used by `REXX` subroutines in order to control how variables are shared among routines. The simplest use is without any parameters; then all future references to variables in that subroutine refer to local variables. If there is no `PROCEDURE` instruction in a subroutine, then all variable references in that subroutine refer to variables in the calling routine's namespace.

If the `EXPOSE` subkeyword is specified too, then any references to the variables in the list following `EXPOSE` refer to local variables, but to variables in the namespace of the calling routine.

Example: Dynamic execution of `PROCEDURE`

The definition opens for some strange effects, consider the following code:

```
call testing

testing:
    say foo
    procedure expose bar
    say foo
```

Here, the first reference to `FOO` is to the variable `FOO` in the caller routine's namespace, while the second reference to `FOO` is to a local variable in the called routine's namespace. This is difficult to parse statically, since the name to expose (and even when to expose them) is determined dynamically during run-time. Note that this use of `PROCEDURE` is allowed in [TRL1], but not in [TRL2].

Several restrictions have been imposed on `PROCEDURE` in [TRL2] in order to simplify the execution of `PROCEDURE` (and in particular, to ease the implementation of optimizing interpreters and compilers).

- The first restriction, to which all `REXX` interpreters adhere as far as I know, is that each invocation of a subroutine (i.e. not the main program) may execute `PROCEDURE` at most once. Both TRL1 and TRL2 contain this restriction. However, more than one `PROCEDURE` instruction may exist "in" each routine, as long as at most one is executed at each invocation of the subroutine.
- The second restriction is that the `PROCEDURE` instruction must be the first statement in the subroutine. This restriction was introduced between `REXX` language level 3.50 and 4.00, but several level 4.00 interpreters may not enforce it, since there is no breakage when allowing it.

There are several important consequences of this second restriction:

(1) it implicitly includes the first restriction listed above, since only one instruction can be the first; (2) it prohibits selecting one of several possible `PROCEDURE` instructions; (3) it prohibits using the same variable name twice; first as an exposed and then as a local variable, as indicated in the example above; (4) it prohibits the customary use of `PROCEDURE` and `INTERPRET`, where the latter is used to create a level of indirectness for the `PROCEDURE` instruction. This particular use can be exemplified by:

```
testing:
    interpret 'procedure expose' bar
```

where `BAR` holds a list of variable names which are to be exposed. However, in order to make this functionality available without having to resort to `INTERPRET`, which is generally considered "bad" programming style, new functionality has been added to `PROCEDURE` between language levels 3.50 and 4.00. If one of the variables in the list of variables is enclosed in parentheses, that means indirectness. Then, the variables exposed are: (1) the variable enclosed in parentheses; (2) the value of that variable is read, and its contents is taken to be a space-separated list of variable names; and (3) all the variable names are exposed strictly in order from left to right.

Example: Indirect exposing

Consider the following example:

```
testing:
    procedure expose foo (bar) baz
```

Assuming that the variable `BAR` holds the value `one two`, then variables exposed are the following: `FOO`, `BAR`, `ONE`, `TWO`, `BAZ`, in that order. In particular, note that the variable `FOO` is exposed immediately before the variables which it names are exposed.

Example: Order of exposing

Then there is another fine point about exposing, the variables are hidden immediately after the subkeyword, so they are not initially available when the variable list is processed. Consider the following code:

```
testing:
    procedure expose bar foo.bar foo.baz baz
```

which exposes variables in the orders specified. If the variable `BAR` holds the value `123`, then `FOO.123` is exposed as the second item, since `BAR` is visible after having already been exposed as the first item. On the other hand, the third item will always expose the variable `FOO.BAZ`, no matter what the value of `BAZ` is in the caller, since the `BAZ` variable is visible only after it has been used in the third item. Therefore, the order in which variables are exposed is important. So, if a compound variable is used inside parentheses in a `PROCEDURE` instruction, then any simple symbols needed for tail substitution must previously have been explicitly exposed. Compare this to the `DROP` instruction.

What exactly is exposing? Well, the best description is to say that it makes all future uses (within that procedural level) to a particular variable name refer to the variable in the calling routine rather than in the local subroutine. The implication of this is that even if it is dropped or it has never been set, an exposed variable will still refer to the variable in the calling routine. Another important thing is that it is the tail-substituted variable name that is exposed. So if you expose `FOO.BAR`, and `BAR` has the value `123`, then only `FOO.123` is exposed, and continue to be so, even if `BAR` later changes its value to e.g. `234`.

Example: Global variables

A problem lurking on new `REXX` users, is the fact that exposing a variable only exposes it to the calling routine. Therefore, it is incorrect to speak of global variables, since the variable might be local to the calling routine. To illustrate, consider the following code:


```

foo = 'bar'
call sub1
call sub2
exit

sub1: procedure expose foo
      say foo /* first says 'bar', then 'FOO' */
      return

sub2: procedure
      say foo /* says 'FOO' */
      call sub1
      return

```

Here, the first subroutine call in the "main" program writes out `bar`, since the variable `FOO` in `SUB1` refers to the `FOO` variable in the main program's (i.e. its caller routine's) namespace. During the second call from the main program, `SUB2` writes out `FOO`, since the variable is not exposed. However, `SUB2` calls `SUB1`, which exposes `FOO`, but that subroutine also writes out `FOO`. The reason for this is that `EXPOSE` works on the run-time nesting of routines, not on the typographical structure of the code. So the `PROCEDURE` in `SUB1` (on its second invocation) exposes `FOO` to `SUB2`, not to the main program as typography might falsely indicate.

The often confusing consequence of the run-time binding of variable names is that an exposed variable of `SUB1` can be bound to different global variables, depending on from where it was called. This differs from most compiled languages, which bind their variables independently of from where a subroutine is called. In turn, the consequence of this is that `REXX` has severe problems storing a persistent, static variable which is needed by one subroutine only. A subroutine needing such a variable (e.g. a count variable which is incremented each time the subroutine is called), must either use an operating system command, or all subroutines calling that subroutine (and their calling routines, etc.) must expose the variable. The first of these solutions is very inelegant and non-standard, while the second is at best troublesome and at worst seriously limits the maximum practical size of a `REXX` program. There are hopes that the `VALUE ()` built-in function will fix this in future standards of `REXX`.

Another important drawback with `PROCEDURE` is that it only works for internal subroutines; for external subroutines it either does not work, or `PROCEDURE` may not even be allowed on the main level of the external subroutine. However, in internal subroutines inside the external subroutines, `PROCEDURE` is allowed, and works like usual.

8.16 The PULL Instruction

```
PULL [ template ] ;
```

This statement takes a line from the top of the stack and parses it into the variables in the *template*. It will also translate the contents of the line to upper case.

This statement is equivalent to `PARSE UPPER PULL [template]` with the same exceptions as explained for the `ARG` instruction. See chapter [not yet written] for a description of parsing and chapter

Stack for a discussion of the stack.

8.17 The PUSH Instruction

```
PUSH [ expr ] ;
```

The PUSH instruction will add a string to the stack. The string added will either be the result of the *expr*, or the null string if *expr* is not specified.

The string will be added to the top of the stack (LIFO), i.e. it will be the first line normally extracted from the stack. For a thorough discussion of the stack and the methods of manipulating it, see chapter [Stack](#) for a discussion of the stack.

8.18 The QUEUE Instruction

```
QUEUE [ expr ] ;
```

The QUEUE instruction is identical to the PUSH instruction, except for the position in the stack where the new line is inserted. While the PUSH puts the line on the "top" of the stack, the QUEUE instruction inserts it at the bottom of the stack (FIFO), or in the bottom of the topmost buffer, if buffers are used.

For further information, refer to documentation for the PUSH instruction, and see chapter [Stack](#) for general information about the stack.

8.19 The RETURN Instruction

```
RETURN [ expr ] ;
```

The RETURN instruction is used to terminate the current procedure level, and return control to a level above. When RETURN is executed inside one or more nesting constructs, i.e. DO, IF, WHEN, or OTHERWISE, then the nesting constructs (in the procedural levels being terminated) are terminated too.

Optionally, an expression can be specified as an argument to the RETURN instruction, and the string resulting from evaluating this expression will be the return value from the procedure level terminated to the caller procedure level. Only a single value can be returned. When RETURN is executed with no argument, no return value is returned to the caller, and then a SYNTAX condition {44} is raised if the subroutine was invoked as a function.

Example: Multiple entry points

A routine can have multiple exit points, i.e. a procedure can be terminated by any of several RETURN instructions. A routine can also have multiple entry points, i.e. several routine entry points can be terminated by the same RETURN instruction. However, this is rarer than having multiple exit points, because it is generally perceived that it creates less structured and readable code. Consider the following code:

```

call foo
call bar
call baz
exit

foo:
    if datatype(name, 'w') then
        drop name
    signal baz
bar:
    name = 'foo'
baz:
    if symbol('name')== 'VAR' then
        say 'NAME currently has the value' name
    else
        say 'NAME is currently an unset variable'
    return

```

Although this is hardly a very practical example, it shows how the main bulk of a routine can be used together with three different entry points. The main part of the routine is the `IF` statement having two `SAY` statements. It can be invoked by calling `FOO`, `BAR`, or `BAZ`.

There are several restrictions to this approach. For instance, the `PROCEDURE` statement becomes cumbersome, but not impossible, to use.

Also note that when a routine has multiple exit points, it may choose to return a return value only at some of those exit points.

When a routine is located at the very end of a source file, there is an implicit `RETURN` instruction after the last explicit clause. However, according to good programming practice, you should avoid taking advantage of this feature, because it can create problems later if you append new routines to the source file and forget to change the implied `RETURN` to an explicit one.

If the current procedure level is the main level of either the program or an external subroutine, then a `RETURN` instruction is equivalent to an `EXIT` instruction, i.e. it will terminate the `REXX` program or the external routine. The table in the `Exit` section shows the actions of both the `RETURN` and `EXIT` instructions depending on the context in which they occur.

The `SAY` Instruction

```
SAY [ expr ] ;
```

Evaluate the expression *expr*, and print the resulting string on the standard output stream. If *expr* is not specified, then the null string is used instead. After the string has been written, an implementation-specific action is taken in order to produce an end-of-line.

The `SAY` instruction is roughly equivalent to

```
call lineout , expr
```

The differences are that there is no way of determining whether the printing was successfully completed if SAY is used, and the special variable `RESULT` is never set when executing a SAY instruction. Besides, the effect of omitting `expr` is different. In SAA API, the `RXSIO SAY` subfunction of the `RXSIO` exit handler is able to trap a SAY instruction, but not a call to the `LINEOUT ()` built-in function. Further, the `NOTREADY` condition is never raised for a SAY instruction.

8.20 The SELECT/WHEN/OTHERWISE Instruction

```
SELECT ; whenpart [ whenpart ... ] [ OTHERWISE [ ; ]  
      [ statement ... ] ] END ;
```

```
whenpart : WHEN expr [ ; ] THEN [ ; ] statement
```

This instruction is used for general purpose, nested IF structures. Although it has certain similarities with CASE in Pascal and switch in C, it is in some respects very different from these. An example of the general use of the SELECT instruction is:

```
select  
    when expr1 then statement1  
    when expr2 then do  
        statement2a  
        statement2b  
    end  
    when expr3 then statement3  
    otherwise  
        ostatement1  
        ostatement2  
end
```

When the SELECT instruction is executed, the next statement after the SELECT statement must be a WHEN statement. The expression immediately following the WHEN token is evaluated, and must result in a valid boolean value. If it is true (i.e. 1), the statement following the THEN token matching the WHEN is executed, and afterwards, control is transferred to the instruction following the END token matching the SELECT instruction. This is not completely true, since an instruction may transfer control elsewhere, and thus implicitly terminate the SELECT instruction; e.g. LEAVE, EXIT, ITERATE, SIGNAL, or RETURN or a condition trapped by method SIGNAL.

If the expression of the first WHEN is not true (i.e. 0), then the next statement must be either another WHEN or an OTHERWISE statement. In the former case, the process explained above is iterated. In the latter case, the clauses following the OTHERWISE up to the END statement are interpreted.

It is considered a SYNTAX condition, {7} if no OTHERWISE statement when none of the WHEN-expression evaluate to true. In general this can only be detected during runtime. However, if one of the WHENs is selected, the absence of an OTHERWISE is not considered an error.

By the nature of the `SELECT` instruction, the `WHEN`s are tested in the sequence they occur in the source. If more than one `WHEN` has an expression that evaluates to true, the first one encountered is selected.

If the programmer wants to associate more than one statement with a `WHEN` statement, a `DO/END` pair must be used to enclose the statements, to make them one statement conceptually. However, zero, one, or more statements may be put after the `OTHERWISE` without having to enclose them in a `DO/END` pair. The clause delimiter is optional after `OTHERWISE`, and before and after `THEN`.

Example: Writing `SWITCH` as `IF`

Although `CASE` in Pascal and `switch` in C are in general table-driven (they check an integer constant and jump directly to the correct case, based on the value of the constant), `SELECT` in REXX is not so. It is just a shorthand notation for nested `IF` instructions. Thus a `SWITCH` instruction can always be written as a set of nested `IF` statements; but for very large `SWITCH` statements, the corresponding nested `IF` structure may be too deeply nested for the interpreter to handle.

The following code shows how the `SWITCH` statements shown above can be written as a nested `IF` structure:

```
if expr1 then statement1
else if expr2 then do
    statement2a
    statement2b
end else if expr3 then statement3
else
    ostatement1
    ostatement2
end
```

8.21 The `SIGNAL` Instruction

```
SIGNAL = { string | symbol } ;
          [ VALUE ] expr ;
          { ON | OFF } condition [ NAME
          { string | symbol } ] ;
```

The `SIGNAL` instruction is used for two purposes: (a) to transfer control to a named label in the program, and (b) to set up a named condition trap.

The first form in the syntax definition transfers control to the named label, which must exist somewhere in the program; if it does not exist, a `SYNTAX` condition { 16 } is raised. If the label is multiple defined, the first definition is used. The parameter can be either a symbol (which is taken literally) or a string. If it is a string, then be sure that the case of the string matches the case of the label where it is defined. In practice, labels are in uppercase, so the strings should contain only uppercase letters too, and no space characters.

The second form of the syntax is used if the second token of the instruction is `VALUE`. Then, the rest of the instruction is taken as a general `REXX` expression, which result after reevaluation is taken to be the name of the label to transfer control to. This form is really just a special case of the first form, where the programmer is allowed to specify the label as an expression. Note that if the start of `expr` is such that it cannot be misinterpreted as the first form (i.e. the first token of `expr` is neither a string nor a symbol), then the `VALUE` subkeyword can be omitted.

Example: Transferring control to inside a loop

When the control of execution is transferred by a `SIGNAL` instruction, all active loops at the current procedural level are terminated, i.e. they cannot continue later, although they can of course be reentered from the normal start. The consequence of this is that the following code is illegal:

```
do forever
    signal there
there:
nop
end
```

The fact that the jump is altogether within the loop does not prevent the loop from being terminated. Thus, after the jump to the loop, the `END` instruction is attempted executed, which will result in a SYNTAX condition { 10 }. However, if control is transferred out of the loop after the label, but before the `END`, then it would be legal, i.e. the following is legal:

```
do forever
    signal there
there:
nop
signal after
end

after:
```

This is legal, simply because the `END` instruction is never seen during this script. Although both TRL1 and TRL2 allow this construct, it will probably be disallowed in ANSI.

Just as loops are terminated by a `SIGNAL` instruction, `SELECT` and `IF` instructions are also terminated. Thus, it is illegal to jump to a location within a block of statements contained in a `WHEN`, `OTHERWISE`, or `IF` instruction, unless the control is transferred out of the block before the execution reaches the end of the block.

Whenever execution is transferred during a `SIGNAL` instruction, the special variable `SIGL` is set to the line number of the line containing the `SIGNAL` instruction, before the control is transferred. If this instruction extends over several lines, it refers to the first of this. Note that even blanks are part of a clause, so if the instruction starts with a line continuation, then all line of the instruction is different from that line where the instruction keyword is located.

The third form of syntax is used when the second token in the instruction is either `ON` or `OFF`. In both cases must the third token in the instruction be the name of a condition (as a constant string or a symbol, which is taken literally), and the setup of that condition trap is changed. If the second token is `OFF`, then the trap of the named condition is disabled.

If the second token is `ON`, then the trap of the named condition is enabled. Further, in this situation two more tokens may be allowed in the instruction: the first must be `NAME` and the second must be the name of a label (either as a constant string or as a symbol, which is taken literally). If the five token form is used, then the label of the condition handler is set to the named label, else the name of the condition handler is set to the default, which is identical to the name of the condition itself.

Note that the `NAME` subclause of the `SIGNAL` instruction was a new construct in TRL2, and is not a part of TRL1. Thus, older interpreters may not support it.

Example: Naming condition traps

Note that the default value for the condition handler (if the `NAME` subclause is not specified) is the name of the condition, not the condition handler from the previous time the condition was enabled. Thus, after the following code, the name of the condition handler for the condition `SYNTAX` is `SYNTAX`, not `FOOBAR`:

```
signal on syntax name foobar
signal on syntax
```

Example: Named condition traps in TRL1

A common problem when trying to port `REXX` code from a TRL2 interpreter to a TRL1 interpreter, is that explicitly named condition traps are not supported. There exist ways to circumvent this, like:

```
syntax_name = 'SYNTAX_HANDLER'
signal on syntax
if 1 + 2 then /* will generate SYNTAX condition */
    nop
syntax:
oldsigl = sigl
signal value translate(syntax_name)

syntax_handler:
say 'condition at line' oldsigl 'is being handled...'
exit
```

Here, a "global" variable is used to store the name of the real condition handler, in the absence of a field for this in the interpreter. This works fine, but there are some problems: the variable `SYNTAX_NAME` must be exposed to everywhere, in order to be available at all times. It would be far better if this value could be stored somewhere from which it could be retrieved from any part of the script, no matter the current state of the call-stack. This can be fixed with programs like `GLOBALV` under VM/CMS and

putenv under Unix.

Another problem is that this destroys the possibility of setting up the condition handler with the default handler name. However, to circumvent this, add a new `DEFAULT_SYNTAX_HANDLER` label which becomes the new name for the old `SYNTAX` label.

Further information about conditions and condition traps are given in chapter [Conditions](#).

8.22 The TRACE Instruction

```
TRACE [ number | setting | [ VALUE ] expr ] ;  
      setting = A | S | C | E | F | I | L | N | O | R | S
```

The `TRACE` instruction is used to set a tracing mode. Depending on the current mode, various levels of debugging information is displayed for the programmer. Also interactive tracing is allowed, where the user can re-execute clauses, change values of variables, or in general, execute `REXX` code interactively between the statements of the `REXX` script.

If *setting* is not specified, then the default value `N` is assumed. If the second token after `TRACE` is `VALUE`, then the remaining part of the clause is interpreted as an expression, which value is used as the trace setting. Else, if the second token is either a string of a symbol, then it is taken as the trace setting; and a symbol is taken literally. In all other circumstances, whatever follows the token `TRACE` is taken to be an expression, which value is the trace setting.

If a parameter is given to the `TRACE` instruction, and the second token in the instruction is not `VALUE`, then there must only be one token after `TRACE`, and it must be either a constant string or a symbol (which is taken literally). The value of this token can be either a whole number or a trace setting.

If it is a whole number and the number is positive, then the number specifies how many of interactive pauses to skip. This assumes interactive tracing; if interactive tracing is not enabled, this `TRACE` instruction is ignored. If the parameter is a whole, negative number, then tracing is turned off temporarily for a number of clauses determined by the absolute value of *number*.

If the second token is a symbol or string, but not a whole number, then it is taken to be one of the settings below. It may optionally be preceded by one or more question mark (`?`) characters. Of the rest of the token, only the first letter matters; this letter is translated to uppercase, and must be one of the following:

[A]

(All) Traces all clauses before execution.

[C]

(Commands) Traces all command clauses before execution.

[E]

(Errors) Traces any command that would raise the `ERROR` condition (whether enabled or not) after execution. Both the command clause and the return value is traced.

[F]

(Failures) Traces any command that would raise the `FAILURE` condition (whether enabled or not) after execution. Both the command clause and the return value is traced.

[I]

(Intermediate) Traces not only all clauses, but also traces all evaluation of expressions; even intermediate results. This is the most detailed level of tracing.

[L]

(Labels) Traces the name of any label clause executed; whether the label was jumped to or not.

[N]

(Normal or Negative) This is the same as the `Failure` setting.

[O]


(Off) Turns off all tracing.

[R]

(Results) Traces all clauses and the results of evaluating expressions. However, intermediate expressions are not traced.

The `Errors` and `Failure` settings are not influenced by whether the `ERROR` or `FAILURE` conditions are enabled or not. These `TRACE` settings will trace the command and return value after the command has been executed, but before the respective condition is raised.

The level of tracing might be set up graphically, as in the figure below. An arrow indicates that the setting pointed to is a superset of the setting pointed from.



Hierarchy of `TRACE` settings

According to this figure, `Intermediate` is a superset of `Result`, which is a superset of `All`. Further, `All` is a superset of both `Commands` and `Labels`. `Commands` is a superset of `Errors`, which is a superset of `Failures`. Both `Failure` and `Labels` are a superset of `Off`. Actually, `Command` is strictly speaking not a superset of `Errors`, since `Error` traces after the command, while `Command` traces before the command.

`Scan` is not part of this diagram, since it provides a completely different tracing functionality. Note that `Scan` is part of TRL1, but was removed in TRL2. It is not likely to be part of newer `REXX` interpreters.

8.23 The UPPER Instruction

```
UPPER symbol [ symbol [ symbol [...] ] ] ;
```

The UPPER instruction is used to translate the contents of one or more variables to uppercase. The variables are translated in sequence from left to right.

Each symbol is separated by one or more blanks.

While it is more convenient and probably faster than individual calls to TRANSLATE, UPPER is not part of the ANSI standard and is not common in other interpreters so should be avoided. It is provided to ease porting of programs from CMS.

Only simple and compound symbols can be specified. Specification of a stem variable results in an error.

9 Operators

An operator represents an operation to be carried out between two terms, such as division. There are 5 types of operators in the Rexx Language: *Arithmetic*, *Assignment*, *Comparative*, *Concatenation*, and *Logical* Operators. Each is described in further details below.

9.1 Arithmetic Operators

Arithmetic operators can be applied to numeric constants and Rexx variables that evaluate to valid Rexx numbers. The following operators are listed in decreasing order of precedence:

-	Unary prefix. Same as 0-number .
+	Unary prefix. Same as 0+number .
**	Power
*	Multiply
/	Divide
%	Integer divide. Divide and return the integer part of the division.
//	Remainder divide. Divide and return the remainder of the division.
+	Add
-	Subtract.

9.2 Assignment Operators

Assignment operators are means to change the value of a variable. Rexx only has one assignment operator.

=	Assign the value on the right side of the "=" to the variable on the left.
---	--

9.3 Comparative Operators

The Rexx comparative operators compare two terms and return the logical value **1** if the result of the comparison is true, or **0** if the result of the comparison is false. Thenon-strict comparative operators

will ignore leading or trailing blanks for string comparisons, and leading zeros for numeric comparisons. Numeric comparisons are made if both terms to be compared are valid REXX numbers, otherwise string comparison is done. String comparisons are case sensitive, and the shorter of the two strings is padded with blanks.

The following list shows the non-strict comparative operators.

=	Equal
\=,^=	Not equal
>	Greater than.
<	Less than.
>=	Greater than or equal.
<=	Less than or equal
<>, ><	Greater than or less than. Same as Not equal.

The following list shows the strict comparative operators. For two strings to be considered equal when using the strict equal comparative operator, both strings must be the same length.

==	Strictly equal
\==,^==	Strictly not equal.
>>	Strictly greater than.
<<	Strictly less than.
>>=	Strictly greater than or equal.
<<=	Strictly less than or equal.

9.4 Concatenation Operators

The concatenation operators combine two strings to form one, by appending the second string to the right side of the first. The REXX concatenation operators are:

(blank)	Concatenation of strings with one space between them.
(abuttal)	Concatenation of strings with no intervening space.
	Concatenation of strings with no intervening space.

Examples:

```

a = abc; b = 'def'
Say a b           -> results in 'abc def'
Say a || b        -> results in 'abcdef'
Say a'xyz'        -> results in 'abcxyz'
```

9.5 Logical Operators

Logical operators work with the REXX strings 1 and 0, usually as a result of a comparative operator. These operators also only result in logical TRUE; 1 or logical FALSE; 0.

&	And	Returns 1 if both terms are 1.
	Inclusive or	Returns 1 if either term is 1.
&&	Exclusive or	Returns 1 if either term is 1 but NOT both terms.

\ Logicalnot Reversestheresult;0becomes1and1becomes0.

10 Implementation-SpecificInformation

10.1 Miscellaneous

OPTIONS settings

Aresavedacrosssubroutines,justlikeotherpiecesofinformation,likeconditionssettings,
NUMERICsettings,etc.Seechapter Optionsformoreinformationabout OPTIONSsettings.

Return value

Totheprogramthatcalled Reginaislimitedtobeinganinteger,whenthisisrequiredbythe
operatingsystems.Allcurrentimplementationsareforoperatingsystemsthatrequirethis.

Default return value

Froma REXXprogramis 0undermostsystems,specificallyUnix,OS/2,MS-DOS.Here,
VMSdeviates,sinceituses 1asthedefaultreturnvalue.Using 0underVMStendstomake
VMSissueawarningsayingthatnoerroroccurred.

Transferring control into a loop

Worksfinein Regina,aslongasno END, THEN, ELSE, WHEN,or OTHERWISEinstructions
areexecutedafterwards;unlessthenormalentrypointfortheconstructhasbeenexecutedafter
thetransferofcontrol.

PARSE SOURCE information

PARSE VERSION information

Last line of source code

Isimplicitlytakentobeterminatedbyanend-of-linesequencein Regina,evenifsucha
sequenceisnotpresentinthesourcecodeofthe REXXscript.Thisappliesonlytosourcecode.
Also,theend-of-stringin INTERPRETstringsistakentobeimplicitlyterminatedbyanend-of-
linecharactersequence.

Moving code MS-DOS to Unix

Issimplifiedby Regina,sinceitwillaccepttheMS-DOSypeendoflinesequencesasvalid.
I.e.anyCtrl-MinfrontofaCtrl-JinthesourcefileisignoredonUnixsystemsby Regina.
Thisappliesonlytosourcecode.

Labels in INTERPRET

Ishandledby Reginainthefollowingway:Alabelcanoccurinsidean INTERPRETstring,
butitisignored,andcanneverbejumpedtoina SIGNALor CALLinstruction.

10.2 ImplementationoftheADDRESSenvironment

Mostpeoplehaveproblemsinvokingexternalprograms.Thissectionshowsthebasicrules,andsome
trickstoletyouuse Reginaandother Rexxinterpreterssucessfully.

Everycalltoanexternalprogramisexecutedbyanimplicit ADDRESSstatement.

```
'echo Hello planet'
           is equivalent to
ADDRESS currentenvironment 'echo Hello planet'
```

The default environment is `SYSTEM` in Regina and many other Rexx interpreters.

Every ADDRESS environment has its own purpose and advantages. It is a good idea to use ADDRESS in front of each command. Everybody knows what happens in this case. And you can choose the *best* environment for the command.

10.2.1 SYSTEM aka ENVIRONMENT aka OS2 ENVIRONMENT

This is the all-purpose solution for every command. The command is passed to the current command interpreter. It is generally the best option for most commands, but it has some disadvantages:

- You don't have control over the different interpreters. You can get ugly errors in Windows NT, 2000, XP or in Unixes if you don't know how the interpreter interprets your command.
- You have some trouble passing special characters to the command. Have you ever tried to pass a ">" sign to a command? You won't get what you expect if you don't know how to quote it to bypass the interpreter.
- You invoke a separate program just to invoke another program. It costs time and memory usage. Choosing another environment may lead to a quicker and safer execution.

Use `SYSTEM` if you want to use pipelines and redirections of the interpreter or if you want to use a builtin command of the shell. "echo" is a builtin command in command interpreters. Also, the Unix pipeline of commands like "prog1|prog2|prog3" cannot be represented shorter in Regina.

10.2.2 PATH

This is the right ADDRESS environment if you know the called program's name but not where it is on disk. One example is "sort" in many systems.

Since Regina has ANSI's extremely useful ADDRESS WITH technique, you can very effectively sort queue contents or stem leaves by:

```
ADDRESS PATH 'sort' WITH INPUT STEM unsort. OUTPUT STEM sort.
```

You let Regina find the program 'sort' (or SORT.EXE if you use Windows) and get the fastest way to do it. You don't have to bother about the current interpreter. Regina acts as a one. You can pass every character you want and Regina does its best to let it appear in the called program.

10.2.3 CMD aka COMMAND

This is a special variant of `PATH`. It acts like `PATH`, but you have to pass the fully qualified filename of the program to execute. You usually use this if you want to use a distinct version of a program, e.g. if you can choose the version.

Another purpose is to call a program which isn't listed in the path and you don't want to change the

path's environment variable by the builtin function `VALUE`.
The operating system and runtime system decides what program is looked for if you omit the path component; the current directory is used in most cases.

10.2.4 REXX or REGINA

Use this if you want to execute a REXX program in a separate instance of the interpreter. Whereas a normal `CALL` on an external program will run the external REXX program in the current instance of Regina, this allows the external REXX program to run in a new, independent instance of Regina.

Use it if:

1. The called interpreter is unstable and a crash in it will not affect the current execution. A common situation where you want it is an external program library you can bind with `RxFuncAdd`. Such a library can crash or terminate the interpreter. The calling interpreter won't be affected by this termination.
2. You want to take advantage of the powerful `ADDRESS WITH` redirection. The general mechanism to communicate with external scripts is a queue, but you don't have this in cases where you want to pass error messages in a different way or if you use a script which wasn't designed to use queues originally.
3. The current interpreter shall be reused and you want to take advantage of the second point. You may have different Regina interpreters and you want to use just the current interpreter even if it isn't in your path. Regina tries to load the current interpreter a second time if you use this `ADDRESS` environment. You should get "rex" if you use `ADDRESS REXX`, "regina" if you use `ADDRESS REGINA`. Regina also attempts to load the same executable that the current instance was started from, but not every system passes enough information to Regina to find its own executable in all cases.

10.3 ADDRESS WITH on Windows

Redirection of program's input and/or output in general is relatively predictable on most operating systems, however mention must be made of behaviour specific to the Windows platform.

Windows and to a lesser degree OS/2, have techniques to hide windows, to start programs in separate windows and other *cool* features. Florian did some significant testing of this on all different Windows platforms and there is bad news. There is no consistent mechanism to start external programs without error and full control. Sounds strange, is strange. We have the options to:

- use the interpreter(`ADDRESS SYSTEM`) or not(`ADDRESS PATH` or `CMD`)
- start GUI or text mode programs
- choose the interpreter(`CMD.EXE` or `COMMAND.COM`)

The main goal was to start GUI separately and text mode programs under the control of the caller (GUI or text). Regina can be part of a GUI program and must be treated as GUI in this case. Most people get upset with console windows popping up showing nothing.

Some combinations of the interpreter, the target programs, and the options we can pass along to the system lead to nonstarting, nonstopping, crashing programs. Or we may lose control by means of

brokencommunicationstothesubprocess(ADDRESS WITH...).

So we had to choose either to let program run safely OR to let program run pretty. Blame the guys who designed Windows, not the Regina crew!

So, if you have a DOS graphical extension known as Windows 95, Windows 98 or Windows Millennium you will get console windows popping up if run from a GUI program. We are sorry for this, we can't change it.

Those Systems with a 32 bit startup kernel known as Windows NT, Windows 2000, Windows XP will hide the console windows when starting a text mode program from a GUI program.

ATTENTION: Your programs might crash or you may lose control either of the called program or of Regina if you change the interpreter inside your Rexx program. Never use
CALL VALUE 'COMSPEC', something, 'SYSTEM'
in your program if you don't know the consequences! Unpredictable behaviour is likely to occur; use at your own peril!

10.4 Regina Restricted Mode

Many language interpreters provide a mechanism where code executed within that interpreter is limited to affecting the environment of the interpreter and cannot change the external environment in which the interpreter runs.

Restricted mode is used in situations where you need to guarantee that the author of a Rexx program is unable to affect the user's environment.

Situations where restricted mode is applicable include, using Regina as a database procedural language, or as a language plugin for a Web browser.

Features of Regina that are disabled in restricted mode are:

- LINEOUT, CHAROUT, POPEN, RXFUNCADDRBIFs
- "OPENWRITE", "OPENBOTH" subcommands of STREAMBIF
- The "built-in" environment segments SYSTEM, CMD or PATH of ADDRESS command
- Setting the value of a variable in the external environment with VALUEBIF.
- Calling external functions

To run Regina in restricted mode, you can start the Regina interpreter from the command line with the '-r' switch, or when using the Rexx SAA API, ORing, RXRESTRICTED to the CallType parameter of RexxStart() function.

10.5 Native Language Support

Regina provides native language support in the following ways:

- Error messages can be displayed in a user-selectable native language.

10.5.1 Error Messages

All native language error messages are contained in binary files (*.mtb) that are built with the Regina

executables from source files (*.mts).

The mechanism Regina uses to determine what native language to use to display error messages depends on the operating system.

On EPOC32, the language is supplied when installing; the selected language is contained in default.mtb. On all other platforms, Regina uses environment variables if you want to use a language other than English.

The English language messages are built into the interpreter for two reasons:

1. to satisfy the ANSI requirement that error messages can be obtained in English using the `ERRORTEXTBIF` and specifying a value of 'S' for argument 2.
2. used as a fallback position when non-native language support is available

10.5.2 Implementation

To specify a native language, up to 2 environment variables are used.

REGINA_LANG environment variable is set to an ISO 639-2 character language abbreviation as defined in the following table.

REGINA_LANG	Language	Translation By
de	German	Florian Grosse-Coosmann
es	Spanish	Pablo Garcia-Abia
no	Norwegian	Vidar Tysse
pt	Portuguese	Susana and Brian Carpenter, Josie Medeiros

(to get your name in this table, contact the maintainer with the language you wish to support)

If **REGINA_LANG** is not set, the default is **en**. The case of the value is irrelevant; **EN** is the same as **en**.

REGINA_LANG_DIR is required if Regina does not know where the language files will be at runtime.

Any binary distribution that includes an installation routine; RPM, Windows InstallShield or EPOC32, will set the location of the .mtb files automatically. Similarly building and installing Regina on Unix-like platforms using configure; make install combination will also set the location automatically. All other platforms will require this environment variable to be set manually.

REXX Built-in Functions

This chapter describes the *REXX* library of built-in functions. It is divided into three parts:

- First a general introduction to built-in functions, pointing out concepts, pitfalls, parameter conventions, peculiarities, and possible system dependencies.
- Then there is the reference section, which describes in detail each function in the built-in library.
- At the end, there is documentation that describes where and how *Regina* differs from standard *REXX*, as described in the two other sections. It also lists *Regina*'s extensions to the built-in library.

It is recommended that you read the first part on first reading of this documentation, and that you use the second part as reference. The third part is only relevant if you are going to use *Regina*.

11 General Information

This section is an introduction to the built-in functions. It describes common behavior, parameter conventions, concepts and lists possible system-dependent parts.

11.1 The Syntax Format

In the description of the built-in functions, the syntax of each one is listed. For each of the syntax diagrams, the parts written in *italic* font name the parameters. Terms enclosed in [square brackets] denote optional elements. And the *courier* font is used to denote that something should be written as is, and it is also used to mark output from the computer. At the right of each function's syntax is an indication of where the function is defined.

(ANSI)	ANSI Standard for <i>REXX</i> 1996
(EXT-ANSI)	Extended <i>REXX</i>
(SAA)	System Application Architecture-IBM
(OS/2)	IBM OS/2 <i>REXX</i>
(CMS)	<i>REXX</i> on CMS
(AREXX)	<i>AREXX</i> on Amiga
(REGINA)	Additional function provided by <i>Regina</i>

Definitions of the *AREXX* built-in functions have been taken verbatim from
<http://dfduck.homeip.net/dfd/ados/arexx/main.shtml>

Note that in standard *REXX* it is not really allowed to let the last possible parameter be empty if all commas are included, although some implementations allow it. In the following calls:

```
say D2X( 61 )
say D2X( 61, 1 )
say D2X( 61, )
```

The two first return the string consisting of a single character `A`, while the last should return error. If the last argument of a function call is omitted, you cannot safely include the immediately preceding comma.

11.2 Precision and Normalization

The built-in library uses its own internal precision for whole numbers, which may be the range from `-999999999` to `+999999999`. That is probably far more than you will ever need in the built-in functions. For most functions, neither parameters nor return values will be effected by any setting of `NUMERIC`. In the few cases where this does not hold, it is explicitly stated in the description of the function.

In general, only parameters that are required to be whole numbers are used in the internal precision, while numbers not required to be whole numbers are normalized according to the setting of `NUMERIC` before use. But of course, if a parameter is a numeric expression, that expression will be calculated and normalized under the setting of `NUMERIC` before it is given to the function as a parameter.

11.3 Standard Parameter Names

In the descriptions of the built-in functions, several generic names are used for parameters, to indicate something about the type and use of that parameter, e.g. `validrange`. To avoid repeating the same information for the majority of the functions, some common "rules" for the standard parameter names are stated here. These rules implicitly apply for the rest of this chapter.

Note that the following list does not try to classify any general `REXX` "datatypes", but provides a binding between the sub-datatypes of strings and the methodology used when naming parameters.

- *Length* is an non-negative whole number within the internal precision of the built-in functions. Whether it denotes a length in characters or in words, depends on the context.
- *String* can be any normal character string, including the null string. There are no further requirements for this parameter. Sometimes a string is called a "packed string" to explicitly show that it usually contains more than the normal printable characters.
- *Option* is used in some of the functions to choose a particular action, e.g. in `DATE ()` to set the format in which the date is returned. Everything except the first character will be ignored, and case does not matter. Note that the strings should consequently not have any leading space.
- *Start* is a positive whole number, and denotes a start position in e.g. a string. Whether it refers to characters or words depends on the context. The first position is always numbered `1`, unless explicitly stated otherwise in the documentation. Note that when return values denote positions, the number `0` is generally used to denote an nonexistent position.
- *Padchar* must be a string, exactly one character long. That character is used for padding.
- *Streamid* is a string that identifies a `REXX` stream. The actual contents and format of such a string is implementation dependent.
- *Number* is any valid `REXX` number, and will be normalized according to the settings of `NUMERIC`.

before it is used by the function.

If you see one of these names having a number appended, that is only to separate several parameters of the same type, e.g. *string1*, *string2* etc. They still follow the rules listed above. There are several parameters in the built-in functions that do not easily fall into the categories above. These are given other names, and their type and functionality will be described together with the functions in which they occur.

11.4 Error Messages

There are several errors that might occur in the built-in functions. Just one error message is only relevant for all the built-in functions, that is number 40 (*Incorrect call to routine*). In fact, an implementation of REXX can choose to use that for any problem it encounters in the built-in functions. Regina also provides further information in errors in built-in functions, as defined by the ANSI standard. This additional information is provided as sub-error messages and usually provides a more detailed explanation of the error.

Depending on the implementation, other error messages might be used as well. Error message number 26 (*Invalid whole number*) might be used for any case where a parameter should have been a whole number, or where a whole number is out of range. It is implied that this error message can be used in these situations, and it is not explicitly mentioned in the description of the functions.

Other general error messages that might be used in the built-in functions are error number 41 (*Bad arithmetic conversion*) for any parameter that should have been a valid REXX number. The error message 15 (*Invalid binary or hexadecimal string*) might occur in any of the conversion routines that convert from binary or hexadecimal format (*B2X* (), *X2B* (), *X2C* (), *X2D* ()). And of course the more general error messages like error message 5 (*Machineresources exhausted*) can occur.

Generally, it is taken as granted that these error messages might occur for any relevant built-in function, and this will not be restated for each function. When other error messages than these are relevant, it will be mentioned in the text.

In REXX, it is in general not an error to specify a start position that is larger than the length of the string, or a length that refers to parts of a string that is beyond the end of that string. The meaning of such instances will depend on the context, and are described for each function.

11.5 Possible System Dependencies

Some of the functions in the built-in library are more or less system or implementation dependent. The functionality of these may vary, so you should use defensive programming and be prepared for any side-effects that they might have. These functions include:

- *ADDRESS* () is dependent on your operating system and the implementation of REXX, since there is not a standard for naming environments.
- *ARG* () at the main level (not in subroutines and functions) is dependent on how your implementation handles and parses the parameters it got from the operating system. It is also dependent on whether the user specifies the -a command line switch.

- `BITAND()`, `BITOR()` and `BITXOR()` are dependent on the character set of your machine. Seemingly identical parameters will in general return very different results on ASCII and EBCDIC machines. Results will be identical if the parameter was given to these functions as a binary or hexadecimal literal.
- `C2X()`, `C2D()`, `D2C()` and `X2C()` will be effected by the character set of your computers since they convert to or from characters. Note that if `C2X()` and `C2D()` get their first parameter as a binary or hexadecimal literal, the result will be unaffected by the machine type. Also note that the functions `B2X()`, `X2B()`, `X2D()` and `D2X()` are not effected by the character set, since they do not use character representation.
- `CHARIN()`, `CHAROUT()`, `CHARS()`, `LINEIN()`, `LINEOUT()`, `LINES()` and `STREAM()` are the interface to the file system. They might have system dependent peculiarities in several ways. Firstly, the naming of streams is very dependent on the operating system. Secondly, the operation of stream is very dependent on both the operating system and the implementation. You can safely assume very little about how streams behave, so carefully read the documentation for your particular implementation.
- `CONDITION()` is dependent on the conditions system, which in turn depends on such implementation dependent things as file I/O and execution of commands. Although the general operation of this function will be fairly equal among systems, the details may differ.
- `DATATYPE()` and `TRANSLATE()` know how to recognize upper and lowercase letters, and how to transform letters to uppercase. If your `REXX` implementation supports national character sets, the operation of these two functions will depend on the language chosen.
- `DATE()` has the options `Month`, `Weekday` and `Normal`, which produce the name of the day or month in text. Depending on how your implementation handles national character sets, the result from these functions might use the correct spelling of the currently chosen language.
- `DELWORD()`, `SUBWORD()`, `WORD()`, `WORDINDEX()`, `WORDLENGTH()`, `WORDPOS()` and `WORDS()` require the concept of a "word", which is defined as a non-blank character separated by blanks. However, the interpretation of what is a blank character depends upon the implementation.
- `ERRORTEXT()` might have slightly different wordings, depending on the implementation, but the meaning and numbering should be the same. However, note that some implementations may have additional error messages, and some might not follow the standard numbering. Error messages may also be returned in the user's native language.
- `QUEUED()` refers to the system specific concept of a "stack", which is either internal or external to the implementation. The result of this function may therefore be dependent on how the stack is implemented on your system.
- `RANDOM()` will differ from machine to machine, since the algorithm is implementation dependent. If you set the seed, you can safely assume that the same interpreter under the same operating system

and on the same hardware platform will return a reproducible sequence. But if you change to another interpreter, another machine or even just another version of the operating system, the same seed might not give the same pseudo-random sequence.

- `SOURCELINE ()` has been changed between REXX language level 3.50 and 4.00. In 4.00 it can return 0 if the REXX implementation finds it necessary, and any request for a particular line may get a null string as result. Before assuming that this function will return anything useful, consult the documentation.
- `TIME ()` will differ somewhat on different machines, since it is dependent on the underlying operating system to produce the timing information. In particular, the granularity and accuracy of this information may vary.
- `VALUE ()` will be dependent on implementation and operating system if it is called with its third parameter specified. Consult the implementation specific documentation for more information about how each implementation handles this situation.
- `XRANGE ()` will return a string, which contents will be dependent on the character set used by your computer. You can safely make very few assumptions about the visual representation, the length, or the character order of the string returned by this function.

As you can see, even REXX interpreters that are within the standard can differ quite a lot in the built-in functions. Although the points listed above seldom are any problem, you should never assume anything about them before you have read the implementation specific documentation. Failure to do so will give you surprises sooner or later.

And, by the way, many implementations (probably the majority) do not follow the standard completely. So, in fact, you should never assume anything at all. Sorry...

11.6 Blanks vs. Spaces

Note that the description differs between "blanks" and the `<space>` character. A blank is any character that might be used as "whitespace" to separate text into groups of characters. The `<space>` character is only one of several possible blanks. When this text says "blank" it means any one from a set of characters that are used to separate visual characters into words. When this text says `<space>`, it means one particular blank, that which is generally bound to the space bar on a normal computer keyboard.

All implementation can be trusted to treat the `<space>` character as blank. Additional characters that might be interpreted as blanks are `<tab>` (horizontal tabulator), `<ff>` (form feed), `<vt>` (vertical tabulator), `<nl>` (new line) and `<cr>` (carriage return). The interpretation of what is blank will vary between machines, operating systems and interpreters. If you are using support for national character sets, it will even depend on the language selected. So be sure to check the documentation before you assume anything about blank characters.

Some implementations use only one blank character, and perceive this set of blank characters as equivalent to the `<space>` character. This will depend on the implementation, the character set, the customs of the operating system and various other reasons.

12 ReginaBuilt-inFunctions

Below follows an in-depth description of all the functions in the library of built-in functions. Note that all functions in this section are available on all ports of **Regina**. Each function is designated as being part of the ANSI standard, or from other implementations. Following sections describe those built-in functions that are available on specific ports of Regina, or when Regina is built with certain switches.

ABBREV(long , short [, length]) **(ANSI)**

Returns 1 if the string *short* is strictly equal to the initial first part of the string *long*, and returns 0 otherwise. The minimum length which *short* must have, can be specified as *length*. If *length* is unspecified, no minimum restrictions for the length of *short* applies, and thus the null string is an abbreviation of any string.

Note that this function is case sensitive, and that leading and trailing spaces are not stripped off before the two strings are compared.

ABBREV('Foobar' , 'Foo')	1
ABBREV('Foobar' , 'Foo' , 4)	0 /*Too short */
ABBREV('Foobar' , 'foo')	0 /*Different case */

ABS(number) **(ANSI)**

Returns the absolute value of the *number*, which can be any valid **REXX** number. Note that the result will be normalized according to the current setting of **NUMERIC**.

ABS(-42)	42
ABS(100)	100

ADDRESS() **(ANSI)**

Returns the current default environment to which commands are sent. The value is set with the **ADDRESS** clause, for more information, see documentation on that clause.

ADDRESS()	UNIX /* Maybe */
------------	------------------

ARG([argno [, option]]) **(ANSI)**

Returns information about the arguments of the current procedure level. For subroutines and functions it will refer to the arguments with which they were called. For the "main" program it will refer to the arguments used when the **REXX** interpreter was called.

Note that under some operating systems, **REXX** scripts are run by starting the **REXX** interpreter as a program, giving it the name of the script to be executed as a parameter. Then the **REXX** interpreter might process the command line and "eat" some or all of the arguments and options. Therefore, the result of

this function at the main level is implementation dependent. The part of the command line which are not available to the REXX script might for instance be the options and arguments meaningful only to the interpreter itself.

Also note that how the interpreter on the main level divides the parameter line into individual arguments, is implementation dependent. The standard seems to define that the main procedure level can only get one parameter string, but don't count on it.

For more information on how the interpreter processes arguments when called from the operating system, see the documentation on how to run a REXX script.

When called without any parameters, ARG() will return the number of comma-delimited arguments. Unspecified (omitted) arguments at the end of the call are not counted. Note the difference between using comma and using space to separate strings. Only comma-separated arguments will be interpreted by REXX as different arguments. Space-separated strings are interpreted as different parts of the same argument.

ARGNO must be a positive whole number. If only ARGNO is specified, the arguments specified will be returned. The first argument is numbered 1. If ARGNO refers to an unspecified argument (either omitted or ARGNO is greater than the number of arguments), a null string is returned.

If OPTION is also specified, the return value will be 1 or 0, depending on the value of OPTION and on whether the numbered parameter was specified or not. Option can be:

[O]

(Omitted) Returns 1 if the numbered argument was omitted or unspecified. Otherwise, 0 is returned.

[E]

(Existing) Returns 1 if the numbered argument was specified, and 0 otherwise.

If called as:

```
CALL FUNCTION 'This' 'is', 'a',, 'test',,
```

ARG()	4 /*Last parameter omitted */	
ARG(1)	'This is'	
ARG(2)	'a'	
ARG(3)	' '	
ARG(9)	' ' /*Ninth parameter doesn't exist*/	
ARG(2, 'E')		1
ARG(2, 'O')		0
ARG(3, 'E')	0 /*Third parameter omitted */	
ARG(9, 'O')		1

B2C(binstring) (AREXX)

Converts a string of binary digits (0,1) into the corresponding (packed) character representation. The conversion is the same as though the argument string had been specified as a literal binary string (e.g. '1010'B). Blanks are permitted in the string, but only at byte boundaries. This function is particularly useful for creating strings that are to be used as bit masks.

B2C('00110011')	'3'
B2C('01100001')	'A'

B2X(binstring) (ANSI)

Takes a parameter which is interpreted as a binary string, and returns a hexadecimal string which represent the same information. *Binstring* can only contain the binary digits 0 and 1. To increase readability, blanks may be included in *binstring* to group the digits into groups. Each such group must have a multiple of four binary digits, except from the first group. If the number of binary digits in the first group is not a multiple of four, that group is padded at the left with up to three leading zeros, to make it a multiple of four. Blanks can only occur between binary digits, not as leading or trailing characters.

Each group of four binary digits is translated into one hexadecimal digit in the output string. There will be no extra blanks in the result, and the upper six hexadecimal digits are in uppercase.

B2X('0010 01011100 0011')	'26C3'
B2X('10 0101 11111111')	'26FF'
B2X('0100100 0011')	'243'

BEEP(frequency[,duration]) (OS/2)

Sounds the machine's bell. The *frequency* and *duration* (in milliseconds) of the tone are specified. If no *duration* value is specified, it defaults to 1. Not all operating systems can sound their bells with the given specifications.

BEEP(50, 1000)	
------------------	--

BITAND(string1[,string2][,padchar]) (ANSI)

Returns the result from bitwise applying the operator AND to the characters in the two strings *string1* and *string2*. Note that this is not the logical AND operation, but the bitwise AND operation. default to an null string. The two strings are left-justified; the first characters in both strings will be AND'ed, then the second characters and so forth. *String2*

The behavior of this function when the two strings do not have equal length is defined by the *padchar* character. If it is undefined, the remaining part of the longer string is appended to the result after all characters in the shorter string have been processed. If *padchar* is defined, each character in the remaining

part of the longer string is logically AND'ed with the *padchar* (or rather, the shorter string is padded on the right length, using *padchar*).

When using this function on character strings, e.g. to uppercase or lowercase a string, the result will be dependent on the character set used. To lowercase a string in EBCDIC, use `BITAND()` with a *padchar* value of `'bf'x`. To do the same in ASCII, use `BITOR()` with a *padchar* value of `'20'x`.

<code>BITAND('123456'x, '3456'x)</code>	<code>'101456'x</code>
<code>BITAND('foobar' , , 'df'x)</code>	<code>'FOOBAR' /*For ASCII*/</code>
<code>BITAND('123456'x, '3456'x, 'f0'x)</code>	<code>'101450'x</code>

BITCHG(string,bit) (AREXX)

Changes the state of the specified *bit* in the argument *string*. Bit numbers are defined such that bit 0 is the low-order bit of the rightmost byte of the string.

<code>BITCHG('0313'x, 4)</code>	<code>'0303'x</code>
-----------------------------------	----------------------

BITCLR(string,bit) (AREXX)

Clears (sets to zero) the specified *bit* in the argument *string*. Bit numbers are defined such that bit 0 is the low-order bit of the rightmost byte of the string.

<code>BITCLR('0313'x, 4)</code>	<code>'0303'x</code>
-----------------------------------	----------------------

BITCOMP(string1,string2,bit[,pad]) (AREXX)

Compares the argument strings bit-by-bit, starting at bit number 0. The returned value is the bit number of the first bit in which the strings differ, or -1 if the strings are identical.

<code>BITCOMP('7F'x, 'FF'x)</code>	<code>'7'</code>
<code>BITCOMP('FF'x, 'FF'x)</code>	<code>'-1'</code>

BITOR(string1[,string2][,padchar]) (ANSI)

Returns the result from byte-wise applying the operator OR to the characters in the two strings *string1* and *string2*. Note that this is not the logical OR operation, but the bitwise OR operation. *String2* defaults to an null string. The two strings are left-justified; the first characters in both strings will be OR'ed, then the second characters and so forth.

The behavior of this function when the two strings do not have equal length is defined by the *padchar* character. If it is undefined, the remaining part of the longer string is appended to the result after all characters in the shorter string have been processed. If *padchar* is defined, each character in the remaining part of the longer string is logically OR'ed with the *padchar* (or rather, the shorter string is padded on the right length, using *padchar*).

When using this function on character strings, e.g. to uppercase or lowercase a string, the result will be dependent on the character set used.

BITOR('12x')	'12'x
BITOR('15'x, '24'x)	'35'x
BITOR('15'x, '2456'x)	'3556'x
BITOR('15'x, '2456'x, 'F0'x)	'35F6'x
BITOR('1111'x, , '4D'x)	'5D5D'x
BITOR('pQrS' , , '20'x)	'pqrS' /* ASCII */

BITSET(string,bit) (AREXX)

Sets the specified *bit* in the argument *string* to 1. Bit numbers are defined such that bit 0 is the low-order bit of the rightmost byte of the string.

BITSET('0313'x, 2)	'0317'x
----------------------	---------

BITTST(string,bit) (AREXX)

The boolean return indicates the state of the specified bit in the argument string. Bit numbers are defined such that bit 0 is the low-order bit of the rightmost byte of the string.

BITTST('0313'x, 4)	'1'
----------------------	-----

BITXOR(string1[, [string2][, padchar]]) (ANSI)

Works like BITAND(), except that the logical function XOR (exclusive OR) is used instead of AND. For more information see BITAND().

BITXOR('123456'x, '3456'x)	'266256'x
BITXOR('FooBar' , , '20'x)	'f00bAR' /* For ASCII */
BITXOR('123456'x, '3456'x, 'f0'x)	'2662A6'x

BUFTYPE() (CMS)

This function is used for displaying the contents of the stack. It will display both the string and notify where the buffers are displayed. It is meant for debugging, especially interactive, when you need to obtain information about the contents of the stack. It always returns the null string, and takes no parameters.

Here is an example of the output from calling BUFTYPE (note that the second and fourth buffers are empty):

```

==> Lines: 4
==> Buffer: 3
"fourth line pushed, in third buffer"
==> Buffer: 2
==> Buffer: 1
"third line pushed, in first buffer"
==> Buffer: 0
"second line pushed, in 'zeroth' buffer"
"first line pushed, in 'zeroth' buffer"
==> End of Stack

```

C2B(string)

(AREXX)

Converts the supplied *string* into the equivalent string of binary digits.

C2B('abc')	'011000010110001001100011'
------------	----------------------------

C2D(string[,length])

(ANSI)

Returns a whole number, which is the decimal representation of the packed string *string*, interpreted as a binary number. If *length* (which must be a non-negative whole number) is specified, it denotes the number of characters in *string* to be converted, and *string* is interpreted as a two's complement representation of a binary number, consisting of the length rightmost characters in *string*. If *length* is not specified, *string* is interpreted as an unsigned number.

If *length* is larger than the length of *string*, *string* is sign-extended on the left. I.e. if the most significant bit of the leftmost character of *string* is set, *string* is padded with 'ff' x chars at the left side. If the bit is not set, '00' x chars are used for padding.

If *length* is too short, only the *length* rightmost characters in *string* are considered. Note that this will not only in general change the value of the number, but it might even change the sign.

Note that this function is very dependent on the character set that your computer is using.

If it is not possible to express the final result as a whole number under the current settings of **NUMERIC DIGITS**, an error is reported. The number to be returned will not be stored in the internal representation of the built-in library, so size restrictions on whole numbers that generally apply for built-in functions, do not apply in this case.

C2D('foo')	'6713199' /*For ASCII machines */
C2D('103'x)	'259'
C2D('103'x,1)	'3'
C2D('103'x,2)	'259'
C2D('0103'x,3)	'259'
C2D('ffff'x,2)	'-1'
C2D('ffff'x)	'65535'
C2D('ffff'x,3)	'65535'
C2D('fff9'x,2)	'-6'
C2D('ff80'x,2)	'-128'

C2X(string) (ANSI)

Returns a string of hexadecimal digits that represents the character string *string*. Converting is done bytewise, the six highest hexadecimal digits are in uppercase, and there are no blank characters in the result. Leading zeros are not stripped off in the result. Note that the behavior of this function is dependent on the character set that your computer is running (e.g. ASCII or EBCDIC).

C2X('ffff'x)	'FFFF'
C2X('Abc')	'416263' /*For ASCII Machines */
C2X('1234'x)	'1234'
C2X('011 0011 1101'b)	'033D'

CD(directory) (REGINA)

CHDIR(directory) (REGINA)

Changes the current process's directory to the *directory* specified. A more portable, though non-standard alternative is to use the DIRECTORYBIF.

CHDIR('/tmp/aa')	/* new directory now /tmp/aa */
--------------------	---------------------------------

CENTER(string, length [, padchar]) (ANSI)

CENTRE(string, length [, padchar]) (ANSI)

This function has two names, to support both American and British spelling. It will center *string* in a string total of length *length* characters. If *length* (which must be a non-negative whole number) is greater than the length of *string*, *string* is padded with *padchar* or <space> if *padchar* is unspecified. If *length* is smaller than the length of *string*, *string* character will be removed.

If possible, both ends of *string* receives (or loses) the same number of characters. If an odd number of

characters are to be added (or removed), one character more is added to (or removed from) the right end than the left end of *string*.

CENTER(' Foobar ' ,10)	' Foobar '
CENTER(' Foobar ' ,11)	' Foobar '
CENTRE(' Foobar ' ,3)	' oob '
CENTER(' Foobar ' ,4)	' ooba '
CENTER(' Foobar ' ,10, '*')	' **Foobar** '

CHANGESTR(needle, haystack, newneedle) (ANSI)

The purpose of this function is to replace all occurrences of *needle* in the string *haystack* with *newneedle*. The function returns the changed string.

If *haystack* does not contain *needle*, then the original *haystack* is returned.

CHANGESTR('a' , 'fred' , 'c')	' fred '
CHANGESTR(' ' , ' ' , 'x')	' '
CHANGESTR('a' , 'abcdef' , 'x')	' xbcdef '
CHANGESTR('0' , '0' , '1')	' 1 '
CHANGESTR('a' , 'def' , 'xyz')	' def '
CHANGESTR('a' , ' ' , 'x')	' '
CHANGESTR(' ' , 'def' , 'xyz')	' def '
CHANGESTR('abc' , 'abcdef' , 'xyz')	' xyzdef '
CHANGESTR('abcdefg' , 'abcdef' , 'xyz')	' abcdef '
CHANGESTR('abc' , 'abcdefabccdabcd' , 'z')	' zdefzcdzd '

CHARIN([streamid][,[start][,length]]) (ANSI)

This function will in general read characters from a stream, and return a string containing the characters read. The *streamid* parameter names a particular stream to read from. If it is unspecified, the default input stream is used.

The *start* parameter specifies a character in the stream, on which to start reading. Before anything is read, the current read position is set to that character, and it will be the first character read. If *start* is unspecified, no repositioning will be done. Independent of any convention of the operating system, the first character in a stream is always numbered 1. Note that transient streams do not allow repositioning, and an error is reported if the *start* parameter is specified for a transient stream.

The *length* parameter specifies the number of characters to read. If the reading did work, the return string will be of length *length*. There are no other way to know how many characters were read than checking the length of the return value. After the read, the current read position is moved forward as

many characters as was read. If *length* is unspecified, it defaults to 1. If *length* is 0, nothing is read, but the file might still be repositioned if *start* was specified.

Note that this function reads the stream raw. Some operating systems use special characters to differ between separate lines in text files. On these systems these special characters will be returned as well. Therefore, never assume that this function will behave identically for text streams on different systems.

What happens when an error occurs or the End-Of-File (EOF) is seen during reading, is implementation dependent. The implementation may choose to set the NOTREADY condition (does not exist in REXX language level 3.50). For more information, see chapter on **Stream Input and Output**.

(Assuming that the file "/tmp/file" contains the first line: " This is the first line"):

CHARIN()	'F' /*Maybe*/
CHARIN(, , 6)	'Foobar' /*Maybe*/
CHARIN('/tmp/file' , , 6)	'This i'
CHARIN('/tmp/file' , 4, 6)	's is t'

CHAROUT([*streamid*][,*string*][,*start*])) (ANSI)

In general this function will write *string* to a *streamid*. If *streamid* is not specified the default output stream will be used.

If *start* is specified, the current write position will be set to the *start*th character in *streamid*, before any writing is done. Note that the current write position cannot be set for transient streams, and attempt to do so will report an error. Independent of any convention that the operating system might have, the first character in the stream is numbered 1. If *start* is not specified, the current write position will not be changed before writing.

If *string* is omitted, nothing is written, and the effect is to set the current write position if *start* is specified. If neither *string* nor *start* is specified, the implementation can really do whatever it likes, and many implementations use this operation to close the file, or flush any changes. Check implementation specific documentation for more information.

The return value is the number of characters in *string* that was not successfully written, so 0 denotes a successful write. Note that in many REXX implementations there is no need to open a stream; it will be implicitly opened when it is first used in a read or write operation.

(Assuming the file referred to by *outdata* was empty, it will contain the string FoobWow afterwards. Note that there might will not be an End-Of-Line marker after this string, it depends on the implementation.)

CHAROUT(, 'Foobar')	'0'
CHAROUT(outdata, 'Foobar')	'0'
CHAROUT(outdata, 'Wow', 5)	'0'

CHARS([streamid]) (ANSI)

Returnsthenumberofcharactersleftinthenamed *streamid*,orthedefaultinputstreamif *streamid*is unspecified.Fortransientstreamsthiswillalwaysbeeither 1ifmorecharactersareavailable,or 0if theEnd-Of-Fileconditionhasbeenmet.Forpersistentstreamsthenumberofremainingbytesinthe filewillbepossibleto calculateandthetrue numberofremainingbyteswillbereturned.

However,onsomesystems,itisdifficultto calculatethenumberofcharactersleftinapersistent stream;therequirementsto CHARS() hasthereforebeenrelaxed,soitcanreturn 1insteadofany numberotherthan 0.If itreturns 1,youcanthereforenotassumeanythingmorethanthatthereisat leastonemorecharacterleftintheinputstream.

CHARS()	'1' /* more data on def. input stream */
CHARS()	'0' /* EOF for def. input stream */
CHARS('outdata')	'94' /* maybe */

CLOSE(file) (AREXX)

Closethe *file*specifiedbythegivenlogicalname.The returnedvalueisabooleansuccessflag,and willbe 1unless thespecifiedfilewasnotopen.

CLOSE('input')	'1'
-----------------	-----

COMPARE(string1,string2[,padchar]) (ANSI)

Thisfunctionwillcompare *string1*to *string2*,andreturnawholenumberwhichwillbe 0iftheyare equal,otherwisethepositionofthefirstcharacteratwhichthetwostringsdifferisreturned.The comparisoniscase-sensitive,andleadingandtrailingspacedomatter.

Ifthestringsareofunequallength,theshorterstringwillbepaddedattherighthandendwiththe *padchar*character tothelengthofthelongerstringbeforethecomparison.If a *padchar*isnotspecified, <space>isused.

COMPARE('FooBar' , 'Foobar')	'4'
COMPARE('Foobar' , 'Foobar')	'0'
COMPARE('Foobarr' , 'Fooba')	'6'
COMPARE('Foobarr' , 'Fooba' , 'r')	'0'

COMPRESS(string[,list]) (AREXX)

If the *list* argument is omitted, the function removes leading, trailing, or embedded blank characters from the *string* argument. If the optional *list* is supplied, it specifies the characters to be removed from the *string*.

COMPRESS(' why not ')	'whynot'
COMPRESS('++12-34-+', '+-')	'1234'

CONDITION([option]) (ANSI)

Returns information about the current trapped condition. A condition becomes the current trapped condition when a condition handler is called (by `CALL` or `SIGNAL`) to handle the condition. The parameter *options* specifies what sort of information to return:

- [C]** (Condition) The name of the current trapped condition is returned, this will be one of the condition names legal to `SIGNAL ON`, like `SYNTAX`, `HALT`, `NOVALUE`, `NOTREADY`, `ERROR` or `FAILURE`.
- [D]** (Description) A text describing the reason for the condition. What to put into this variable is implementation and system dependent.
- [I]** (Instruction) Return either `CALL` or `SIGNAL`, depending on which method was current when the condition was trapped.
- [S]** (State) The current state of the current trapped condition. This can be one of `ON`, `OFF` or `DELAY`. Note that this option reflects the current state, which may change, not the state at the time when the condition was trapped.

For more information on conditions, consult the chapter **Conditions**. Note that condition may in several ways be dependent on the implementation and system, so read system and implementation dependent information too.

COPIES(string,copies) (ANSI)

Returns a string with *copies* concatenated copies of *string*. *Copies* must be a non-negative whole number. No extra space is added between the copies.

COPIES('Foo',3)	'FooFooFoo'
COPIES('*',16)	'*****'
COPIES('Bar ',2)'Bar Bar '	
COPIES('',10000)	' '

COUNTSTR(needle, haystack) (ANSI)

Returns a count of the number of occurrences of *needle* in *haystack* that do not overlap.

COUNTSTR(' ', ' ')	0
COUNTSTR('a', 'abcdef ')	1
COUNTSTR(0, 0)	1
COUNTSTR('a', 'def ')	0
COUNTSTR('a', ' ')	0
COUNTSTR(' ', 'def ')	0
COUNTSTR('abc', 'abcdef ')	1
COUNTSTR('abcdefg', 'abcdef ')	0
COUNTSTR('abc', 'abcdefabccdabcd ')	3

CRYPT(string, salt) (REGINA)

Encrypts the given *string* using the supplied *salt* and returns the encrypted string. Only the first two characters of *salt* are used. Not all operating systems support encryption, and on these platforms, the string is returned unchanged. It is also important to note that the encrypted string is not portable between platforms.

CRYPT('a string', '1x')	'1xYwPPW11zRJs' /* maybe */
-------------------------	-----------------------------

DATATYPE(string[, option]) (ANSI)

With only one parameter, this function identifies the "datatype" of *string*. The value returned will be "NUM" if *string* is a valid REXX number. Otherwise, "CHAR" is returned. Note that the interpretation of whether *string* is a valid number will depend on the current setting of **NUMERIC**.

If *option* is specified too, it will check if *string* is of a particular datatype, and return either "1" or "0" depending on whether *string* is or is not, respectively, of the specified datatype. The possible values of *option* are:

[A]

(Alphanumeric) Consisting of only alphabetic characters (in upper, lower or mixed case) and decimal digits.

[B]

(Binary) Consisting of only the two binary digits 0 and 1. Note that blanks are not allowed within *string*, as would have allowed been within a binary string.

[L]

(Lower) Consisting of only alphabetic characters in lowercase.

[M]

(Mixed) Consisting of only alphabetic characters, but the case does not matter (i.e. upper, lower

ormixed.)

[N]

(Numeric) If *string* is a valid REXX number, i.e. `DATATYPE(string)` would return `NUM`.

[S]

(Symbolic) Consists of characters that are legal in REXX symbols. Note that this test will pass several strings that are not legal symbols. The characters include plus, minus and the decimal point.

[U]

(Upper) Consists of only uppercase alphabetic characters.

[W]

(Whole) If *string* is a valid REXX whole number under the current setting of `NUMERIC`. Note that `13.0` is a whole number since the decimal part is zero, while `13E+1` is not a whole number, since it must be interpreted as `130` plus/minus 5.

[X]

(Hexadecimal) Consists of only hexadecimal digits, i.e. the decimal digits 0-9 and the alphabetic characters A-F in either case (ormixed.) Note that blanks are not allowed within *string*, as it would have been within a hexadecimal string.

If you want to check whether a string is suitable as a variable name, you should consider using the `SYMBOL()` function instead, since the `Symbolic` option only verifies which characters *string* contains, not the order. You should also take care to watch out for lowercase alphabetic characters, which are allowed in the tail of a compound symbol, but not in a simple or stem symbol or in the head of a compound symbol.

Also note that the behavior of the options `A`, `L`, `M` and `U` might depend on the setting of language, if you are using an interpreter that supports national character sets.

DATATYPE(' - 1.35E-5 ')	' NUM '
DATATYPE(' 1E999999999 ')	' CHAR '
DATATYPE(' 1E9999999999 ')	' CHAR '
DATATYPE(' !@#&#\$(&*%` ')	' CHAR '
DATATYPE(' FooBar ', 'A')	' 1 '
DATATYPE(' Foo Bar ', 'A')	' 0 '
DATATYPE(' 010010111101 ', 'B')	' 1 '
DATATYPE(' 0100 1011 1101 ', 'B')	' 0 '
DATATYPE(' foobar ', 'L')	' 1 '
DATATYPE(' FooBar ', 'M')	' 1 '
DATATYPE(' -34E3 ', 'N')	' 1 '
DATATYPE(' A_SYMBOL!?! ', 'S')	' 1 '
DATATYPE(' 1.23.39E+4.5 ', 'S')	' 1 '
DATATYPE(' Foo bar ', 'S')	' 0 '
DATATYPE(' FOOBAR ', 'U')	' 1 '
DATATYPE(' 123deadbeef ', 'X')	' 1 '

DATE([option_out [,date [,option_in]]) **(ANSI)**

This function returns information relating to the current date. If the *option_out* character is specified, it will set the format of the return string. The default value for *option_out* is " N".

Possible options are:

- [B]** (Base) The number of completed days from January 1st 0001 until yesterday inclusive, as a whole number. This function uses the Gregorian calendar extended backwards. Therefore Date('B')//7 will equal the day of the week where 0 corresponds to Monday and 6 Sunday.
- [C]** (Century) The number of days in this century from January 1st -00 until today, inclusive. The return value will be a positive integer.
- [D]** (Days) The number of days in this year from January 1st until today, inclusive. The return value will be a positive integer.
- [E]** (European) The date in European format, i.e. " dd/mm/yy". If any of the numbers is single digit, it will have a leading zero.
- [M]** (Month) The unabbreviated name of the current month, in English.
- [N]** (Normal) Return the date with the name of the month abbreviated to three letters, with only the

first letter in uppercase. The format will be " dd Mmm yyyy", where Mmm is the month abbreviation (in English) and dd is the day of the month, without leading zeros.

[O]

(Ordered) Return the date in the ordered format, which is " yy/mm/dd".

[S]

(Standard) Return the date according to the formats specified by International Standards Organization Recommendation ISO/R2014-1971(E). The format will be " yyyyymmdd", and each part is padded with leading zero where appropriate.

[U]

(USA) Return the date in the format that is normally used in USA, i.e. " mm/dd/yy", and each part is padded with leading zero where appropriate.

[W]

(Weekday) Return the English unabbreviated name of the current weekday for today. The first letter of the result is in uppercase, the rest is in lowercase.

[T]

(time_t) Return the current date/time in UNIX time_t format. time_t is the number of seconds since January 1st 1970.

Note that the " C" option is present in REXX language level 3.50, but was removed in level 4.00. The new " B" option should be used instead. When porting code that uses the " C" option to an interpreter that only has the " B" option, you will need to use the conversion that January 1st 1900 is day 693595 in the Gregorian calendar.

Note that none of the formats in which DATE () returns its answer are effected by the settings of NUMERIC. Also note that if there are more than one call to DATE () (and TIME ()) in a single clause of REXX code, all of them will use the same basis data for calculating the date (and time).

If the REXX interpreter contains national support, some of these options may return different output for the names of months and weekdays.

Assuming that today is January 6th 1992:

DATE ('B')	' 727203 '
DATE ('C')	' 33609 '
DATE ('D')	' 6 '
DATE ('E')	' 06/01/92 '
DATE ('M')	' January '
DATE ('N')	' 6 Jan 1992 '
DATE ('O')	' 92/01/06 '
DATE ('S')	' 19920106 '
DATE ('U')	' 01/06/92 '
DATE ('W')	' Monday '
DATE ('T')	694620000

If the *date* option is specified, the function provides for date conversions. The optional *option_in* specifies the format in which *date* is supplied. The possible values for *option_in* are: **BDEOUNST**. The default value for *option_in* is **N**.

DATE('O','13 Feb 1923')	'23/02/13'
DATE('O','06/01/50','U')	'50/06/01'

If the *date* supplied does not include a century in its format, then the result is chosen to make the year within 50 years past or 49 years future of the current year.

The date conversion capability of the DATEBIF was introduced with the ANSI standard.

DELSTR(string,start[,length]) (ANSI)

Returns *string*, after the substring of length *length* starting at position *start* has been removed. The default value for *length* is the rest of the string. *Start* must be a positive whole number, while *length* must be a non-negative whole number. It is not an error if *start* or *length* (or a combination of them) refer to more characters than *string* holds.

DELSTR('Foobar',3)	'Foo'
DELSTR('Foobar',3,2)	'Foor'
DELSTR('Foobar',3,4)	'Foo'
DELSTR('Foobar',7)	'Foobar'

DELWORD(string,start[,length]) (ANSI)

Removes *length* words and all blanks between them, from *string*, starting at word number *start*. The default value for *length* is the rest of the string. All consecutive spaces immediately after the last deleted word, but no spaces before the first deleted word, are removed. Nothing is removed if *length* is zero.

The valid range of *start* is the positive whole numbers; the first word in *string* is numbered 1. The valid range of *length* is the non-negative integers. It is not an error if *start* or *length* (or a combination of them) refer to more words than *string* holds.

DELWORD('This is a test',3)	'This is '
DELWORD('This is a test',2,1)	'This a test'
DELWORD('This is a test',2,5)	'This'
DELWORD('This is a test',1,3)	'test' /*No leading space*/

DESBUF () (CMS)

This function removes all buffers on the stack, it is really just a way of clearing the whole stack for buffers as well as strings. Functionally, it is equivalent to executing `DROPBUF` with a parameter of `0`. (Actually, this is a lie, since `DROPBUF` is not able to take zero as a parameter. Rather, it is equivalent to executing `DROPBUF` with `1` as a parameter and then executing `DROPBUF` without a parameter, but this is a subtle point.) It will return the number of buffers left on the stack after the function has been executed. This should be `0` in all cases.

DESBUF ()	0
------------	---

DIGITS () (ANSI)

Returns the current precision of arithmetic operations. This value is set using the `NUMERIC` statement. For more information, refer to the documentation on `NUMERIC`.

DIGITS ()	'9' /* Maybe */
------------	-----------------

DIRECTORY([new directory]) (OS/2)

Returns the current directory for the running process, and optionally changes directory to the specified *newdirectory*. If the *newdirectory* exists, and the change to *newdirectory* succeeds, the *newdirectory* is returned. If the *newdirectory* does not exist or an error occurred changing to that *newdirectory*, the empty string is returned.

DIRECTORY ()	'/tmp' /* Maybe */
DIRECTORY ('c:\temp')	'c:\temp' /* Maybe */

D2C(integer[,length]) (ANSI)

Returns a (packed) string, that is the character representation of *integer*, which must be a whole number, and is governed by the settings of `NUMERIC`, not of the internal precision of the built-in functions. If *length* is specified the string returned will be *length* bytes long, with sign extension. If *length* (which must be a non-negative whole number) is not large enough to hold the result, an error is reported.

If *length* is not specified, *integer* will be interpreted as an unsigned number, and the result will have no leading `<nul>` characters. If *integer* is negative, it will be interpreted as a two's complement, and *length* must be specified.

D2C(0)	' '
D2C(127)	' 7F 'x
D2C(128)	' 80 'x
D2C(128 , 3)	' 000080 'x
D2C(-128)	' 80 'x
D2C(-10 , 3)	' fffff5 'x

D2X(integer[,length])

(ANSI)

Returns a hexadecimal number that is the hexadecimal representation of *integer*. *Integer* must be a whole number under the current settings of `NUMERIC`, it is not effected by the precision of the built-in functions.

If *length* is not specified, then *integer* must be non-negative, and the result will be stripped of any leading zeros.

If *length* is specified, then the resulting string will have that length. If necessary, it will be sign-extended on the left side to make it the right length. If *length* is not large enough to hold *integer*, an error is reported.

D2X(0)	' 0 '
D2X(127)	' 7F '
D2X(128)	' 80 '
D2X(128 , 5)	' 00080 'x
D2X(-128)	' 80 'x
D2X(-10 , 5)	' ffff5 'x

DROPBUF([number])

(CMS)

This function will remove zero or more buffers from the stack. Called without a parameter, it will remove the top most buffer from the stack, provided that there were at least one buffer in the stack. If there were no buffers in the stack, it will remove all strings in the stack, i.e. remove the zeroth buffer.

If the parameter *number* was specified, and the stack contains a buffer with an assigned number equal to *number*, then that buffer itself, and all strings and buffers above it on the stack will be removed; but no strings or buffers below the numbered buffer will be touched. If *number* refers to a buffer that does not exist in the stack, no strings or buffers in the stack is touched.

As an extra extension, in `Regina` the `DROPBUF()` built-in function can be given a non-positive integer as a parameter. If the name is negative then it will convert that number to its absolute value, and remove that many buffers, counted from the top. This is functionally equivalent to repeating `DROPBUF()` without parameters for so many times as the absolute value of the negative number specifies. Note that using `-0` as a parameter is equivalent to removing all strings and buffers in the stack, since `-0` is

equivalent to normal 0. The number is converted during evaluation of parameters prior to the call to the DROPBUF () routine, so the sign is lost.

The value returned from this function is the number of buffers left on the stack after the buffers to be deleted have been removed. Obviously, this will be a non-negative integer. This too, deviates from the behavior of the DROPBUF command under CMS, where zero is always returned.

DROPBUF(3)	2 /* remove buffer 3 and 4 */
DROPBUF(4)	0 /* no buffers on the stack */
DROPBUF()	4 /* if there were 5 buffers */

EOF(file) (AREXX)

Checks the specified logical *filename* and returns the boolean value 1 (True) if the end-of-file has been reached, and 0 (False) otherwise.

EOF('infile')	'1' /* maybe */
---------------	-----------------

ERRORTEXT(errno[, lang]) (ANSI)

Returns the REXX error message associated with error number *errno*. If the *lang* character is specified, it will determine the native language in which the error message is returned. The default value for *lang* is "N".

Possible options are:

- [N]**
(Normal) The error text is returned in the default native language.
- [S]**
(Standard English) The error text is returned in English.

For more information on how Regina supports different native languages, see **Native Language Support**.

If the error message is not defined, a null string is returned.

The error messages in REXX might be slightly different between the various implementations. The standard says that *errno* must be in the range 0-99, but in some implementations it might be within a less restricted range which gives room for system specific messages. You should in general not assume that the wordings and ordering of the error messages are constant between implementations and systems.

ERRORTEXT(20)	'Symbol expected'
ERRORTEXT(30)	'Name or string too long'
ERRORTEXT(40)	'Incorrect call to routine'

errno can also be specified as an *errno* followed by a sub-error number, with a period between. The resulting string will be the text of the sub-error number with placemarkers indicating where substitution values would normally be placed.

ERRORTEXT(40.24)	<bif> argument 1 must be a binary string; found "<value>"
------------------	--

Regina also supports messages in several native languages. See this section on **Native Language Support** for details on how this is configured. With **DE** as the native language in effect:

ERRORTEXT(40.24)	Routine <bif>, Argument 1 muß eine Binärzeichenkette sein; "<value>"
ERRORTEXT(40.24,'S')	<bif> argument 1 must be a binary string; found "<value>"

EXISTS(filename) (AREXX)

Tests whether the specified name of the given *filename* exists. The *filename* string may include any portion of a full file path specification. Note that the argument is not a logical filename used in other AREXX file functions. A more portable equivalent of this is to use the 'QUERY EXISTS' command of the STREAMBIF.

EXISTS('c:\temp\infile.txt')	'1' /* maybe */
------------------------------	-----------------

EXPORT(address,[string],[length],[pad]) (AREXX)

Copies data from the (optional) string into a previously-allocated memory area, which must be specified as a 4-byte *address*. The *length* parameter specifies the maximum number of characters to be copied; the default is the length of the string. If the specified *length* is longer than the string, the remaining area is filled with the *pad* character or nulls ('00'x). The returned value is the number of characters copied.

Caution is advised in using this function. Any area of memory can be overwritten, possibly causing a system crash.

See also STORAGE() and IMPORT().

Note that the *address* specified is subject to a machine's endianess.

EXPORT('0004 0000'x,'The answer')	'10'
-----------------------------------	------

FIND(string,phrase) (CMS)

Searches *string* for the first occurrence of the sequence of blank-delimited words *phrase*, and return the word number of the first word of *phrase* in *string*. Multiple blanks between words are treated as a single blank for the comparison. Returns 0 if *phrase* is not found. Deprecated: see WORDPOS().

FIND('now is the time','is the time')	2
FIND('now is the time','is the')	2
FIND('now is the time','is time')	0

FORK ()

(REGINA)

This function spawns a new process as a child of the current process at the current point in the program where FORK is called. The program then continues from this point as two separate processes; the parent and the child. FORK returns 0 to the child process, and the process id of the child process spawned to the parent (always non-zero). A negative return value indicates an error while attempting to create the new process. FORK is not available on all platforms. If FORK is not supported, it will always return '1'. It is safe to assume that a return value of '1' means that FORK is not supported. All platforms AFAIK, will never return '1' as a child process id; that number is usually reserved for the first process that starts on a machine.

FORK ()	'0' /* To child */
	'3456' /* maybe to parent */

FORM ()

(ANSI)

Returns the current "form", in which numbers are represented when exponential form is used. This might be either SCIENTIFIC (the default) or ENGINEERING. This value is set through the NUMERIC FORM clause. For more information, see the documentation on NUMERIC.

FORM ()	'SCIENTIFIC' /* Maybe */
----------	--------------------------

FORMAT (number [, [before] [, [after] [, [exp] [, [expt]]]]])

(ANSI)

This function is used to control the format of numbers, and you may request the size and format in which the number is written. The parameter *number* is the number to be formatted, and it must be a valid REXX number. Note that before any conversion or formatting is done, this number will be normalized according to the current setting of NUMERIC.

The *before* and *after* parameters determine how many characters that are used before and after the decimal point, respectively. Note that *before* does **not** specify the number of digits in the integer part, it specifies the size of the field in which the integer part of the number is written. Remember to allocate space in this field for a minus too, if that is relevant. If the field is not long enough to hold the integer part (including a minus if relevant), an error is reported.

The *after* parameter will dictate the size of the field in which the fractional part of the number is written. The decimal point itself is not a part of that field, but the decimal point will be omitted if the field holding the fractional part is empty. If there are less digits in the number than the size of the field, it is padded with zeros at the right. If there is more digits than it is possible to fit into the field, the number will be rounded (not truncated) to fit the field.

Before must at least be large enough to hold the integer part of *number*. Therefore it can never be less than 1, and never less than 2 for negative numbers. The integer field will have no leading zeros, except a single zero digit if the integer part of *number* is empty.

The parameter *exp* is the size of the field in which the exponent is written. This is the size of the numeric part of the exponent, so the "E" and the sign comes in addition, i.e. there all length if the exponent is two more than *exp* specifies. If *exp* is zero, it signals that exponential form should not be used. *Exp* must be a non-negative whole number. If *exp* is positive, but not large enough to hold the exponent, an error is reported.

Expt is the trigger value that decides when to switch from simple to exponential form. Normally, the default precision (`NUMERIC DIGITS`) is used, but if *expt* is set, it will override that. Note that if *expt* is set to zero, exponential form will always be used. However, if *expt* tries to force exponential form, simple form will still be used if *exp* is zero. Negative values for *expt* will give an error. Exponential form is used if more digits than *expt* is needed in the integer part, or more than twice *expt* digits are needed in the fractional part.

Note that the *after* number will mean different things in exponential and simple form. If *after* is set to e.g. 3, then in simple form it will force the precision to 0.001, no matter the magnitude of the number. If in exponential form, it will force the number to 4 digits precision.

<code>FORMAT(12.34, 3, 4)</code>	<code>' 12.3400 '</code>
<code>FORMAT(12.34, 3, , 3, 0)</code>	<code>' 1.234E+001 '</code>
<code>FORMAT(12.34, 3, 1)</code>	<code>' 12.3400 '</code>
<code>FORMAT(12.34, 3, 0)</code>	<code>' 12.3 '</code>
<code>FORMAT(12.34, 3, 4)</code>	<code>' 12 '</code>
<code>FORMAT(12.34, , , , 0)</code>	<code>' 1.234E+1 '</code>
<code>FORMAT(12.34, , , 0)</code>	<code>' 12.34 '</code>
<code>FORMAT(12.34, , , 0, 0)</code>	<code>' 12.34 '</code>

FREESPACE(address, length) (AREXX)

Returns a block of memory of the given *length* to the interpreter's internal pool. The address argument must be a 4-byte string obtained by a prior call to `GETSPACE()`, the internal allocator. It is not always necessary to release internally-allocated memory, since it will be released to the system when the program terminates. However, if a very large block has been allocated, returning it to the pool may avoid memory space problems. The return value is a boolean success flag. See also `GETSPACE()`

<code>FREESPACE('00042000'x, 32)</code>	<code>' 1 '</code>
---	--------------------

FUZZ () (ANSI)

Returns the current number of digits which are ignored when comparing numbers, during operations

like = and >. The default value for this is 0. This value is set using the `NUMERIC FUZZ` statement, for more information see that.

<code>FUZZ ()</code>	<code>'0' /* Maybe */</code>
-----------------------	------------------------------

GETENV(environmentvar) (REGINA)

Returns the named UNIX environment variable. If this variable is not defined, a null string is returned. It is not possible to use this function to determine whether the variable was unset, or just set to the null string.

This function is now obsolete, instead you should use:

```
VALUE( environmentvar, , 'SYSTEM' )
```

GETPID() (REGINA)

Returns the process id of the currently running process.

<code>GETPID ()</code>	<code>'234' /* Maybe */</code>
-------------------------	--------------------------------

GETSPACE(length) (AREXX)

Allocates a block of memory of the specified length from the interpreter's internal pool. The returned value is the 4-byte address of the allocated block, which is not cleared or otherwise initialized. Internal memory is automatically returned to the system when the Rexx program terminates, so this function should not be used to allocate memory for use by external programs. See also `FREESPACE()`

<code>GETSPACE (32)</code>	<code>'0003BF40' /* maybe */</code>
------------------------------	-------------------------------------

GETTID() (REGINA)

Returns the thread id of the currently running process.

<code>GETTID ()</code>	<code>'2' /* Maybe */</code>
-------------------------	------------------------------

HASH(string) (AREXX)

Returns the hash attribute of a string as a decimal number, and updates the internal hash value of the string.

<code>HASH ('1')</code>	<code>'49'</code>
---------------------------	-------------------

IMPORT(address[,length]) (AREXX)

Creates a string by copying data from the specified 4-byte *address*. If the *length* parameter is not supplied, the copy terminates when a null byte is found. See also EXPORT().

Note that the *address* specified is subject to a machine's endianness.

IMPORT('0004 0000'x,8)	'The answer' /* maybe */
------------------------	--------------------------

INDEX(haystack,needle[,start]) (CMS)

Returns the character position of the string *needle* in *haystack*. If *needle* is not found, 0 is returned. By default the search starts at the first character of *haystack* (*start* is 1). This can be overridden by giving a different *start*, which must be a positive, whole number. See POS function for an ANSI function that does the same thing.

INDEX('abcdef','cd')	'3'
INDEX('abcdef','xd')	'0'
INDEX('abcdef','bc',3)	'0'
INDEX('abcabc','bc',3)	'5'
INDEX('abcabc','bc',6)	'0'

INSERT(string1,string2[,position[,length[,padchar]]]) (ANSI)

Returns the result of inserting *string1* into a copy of *string2*. If *position* is specified, it marks the character in *string2* which *string1* is to be inserted after. *Position* must be a non-negative whole number, and it defaults to 0, which means that *string2* is put in front of the first character in *string1*.

If *length* is specified, *string1* is truncated or padded on the right side to make it exactly *length* characters long before it is inserted. If padding occurs, then *padchar* is used, or <space> if *padchar* is undefined.

INSERT('first','SECOND')	'SECONDfirst'
INSERT('first','SECOND',3)	'fiSECONDrst'
INSERT('first','SECOND',3,10)	'fiSECOND rst'
INSERT('first','SECOND',3,10,'*')	'fiSECOND****rst'
INSERT('first','SECOND',3,4)	'fiSECOrst'
INSERT('first','SECOND',8)	'first SECOND'

JUSTIFY(string,length[,pad]) (CMS)

Formats blank-delimited words in *string*, by adding *pad* characters between words to justify to both margins. That is, to width *length* (*length* must be non-negative). The default *pad* character is a blank.

string is first normalized as though `SPACE(string)` had been executed (that is, multiple blanks are converted to single blanks, and leading and trailing blanks are removed). If *length* is less than the width of the normalized string, the string is then truncated on the right and any trailing blank is removed. Extra *pad* characters are then added evenly from the left to right to provide the required length, and the blanks between words are replaced with the *pad* character.

<code>JUSTIFY('The blue sky',14)</code>	<code>'The blue sky'</code>
<code>JUSTIFY('The blue sky',8)</code>	<code>'The blue'</code>
<code>JUSTIFY('The blue sky',9)</code>	<code>'The blue'</code>
<code>JUSTIFY('The blue sky',9,'+'))</code>	<code>'The++blue'</code>

LASTPOS(*needle*,*haystack*[,*start*]) (ANSI)

Searches the string *haystack* for the string *needle*, and returns the position in *haystack* of the first character in the substring that matched *needle*. This search is started from the right side, so if *needle* occurs several times, the last occurrence is reported.

If *start* is specified, the search starts at character number *start* in *haystack*. Note that the standard only states that the search starts at the *start*th character. It is not stated whether a match can partly be to the right of the *start* position, so some implementations may differ on that point.

<code>LASTPOS('be','To be or not to be')</code>	17
<code>LASTPOS('to','to be or not to be',10)</code>	3
<code>LASTPOS('is','to be or not to be')</code>	0
<code>LASTPOS('to','to be or not to be',0)</code>	0

LEFT(*string*,*length*[,*padchar*]) (ANSI)

Returns the *length* leftmost characters in *string*. If *length* (which must be a non-negative whole number) is greater than the length of *string*, the result is padded on the right with `<space>` (or *padchar* if that is specified) to make it the correct length.

<code>LEFT('Foo bar',5)</code>	<code>'Foo b'</code>
<code>LEFT('Foo bar',3)</code>	<code>'Foo'</code>
<code>LEFT('Foo bar',10)</code>	<code>'Foo bar '</code>
<code>LEFT('Foo bar',10,'*')</code>	<code>'Foo bar***'</code>

LENGTH(*string*) (ANSI)

Returns the number of characters in *string*.

LENGTH(' ')	'0'
LENGTH('Foo')	'3'
LENGTH('Foo bar')	'7'
LENGTH(' foo bar ')	'10'

LINEIN([streamid][,[line][,count]]) **(ANSI)**

Returns a line read from a file. When only *streamid* is specified, the reading starts at the current read position and continues to the first End-Of-Line (EOL) mark. Afterwards, the current read position is set to the character after the EOL mark which terminated the read-operation. If the operating system uses special characters for EOL marks, these are not returned by a part of the string read..

The default value for *streamid* is default input stream. The format and range of the string *streamid* are implementation dependent.

The *line* parameter (which must be a positive whole number) might be specified to set the current position in the file to the beginning of line number *line* before the read operation starts. If *line* is unspecified, the current position will not be changed before the read operation. Note that *line* is only valid for persistent teams. For transient streams, an error is reported if *line* is specified. The first line in the stream is numbered 1.

Count specifies the number of lines to read. However, it can only take the values 0 and 1. When it is 1 (which is the default), it will read one line. When it is 0 it will not read any lines, and a null string is returned. This has the effect of setting the current read position of the file if *line* was specified.

What happens when the function finds a End-Of-File (EOF) condition is to some extent implementation dependent. The implementation may interpret the EOF as an implicit End-Of-Line (EOL) mark is none such was explicitly present. The implementation may also choose to raise the NOTREADY condition flag (this condition is new from REXX language level 4.00).

Whether or not *stream* must be explicitly opened before a read operation can be performed, is implementation dependent. In many implementations, a read or write operation will implicitly open the stream if not already open.

Assuming that the file /tmp/file contains the three lines: " Firstline ", Secondline " and " Thirdline ":

LINEIN('/tmp/file',1)	'First line'
LINEIN('/tmp/file')	'Second line'
LINEIN('/tmp/file',1,0)	' ' /* But sets read position */
LINEIN('/tmp/file')	'First line'
LINEIN()	'Hi, there!' /* maybe */

LINEOUT([streamid][,[string][,line]])

(ANSI)

Returnsthenumberoflinesremainingafterhavingpositionedthestream *streamid*tothestartofline *line*andwrittenout *string*asalineoftext.If *streamid*isomitted,thedefaultoutputstreamisused.If *line*(whichmustbeapositivewholenumber)isomitted,thestreamwillnotberepositionedbeforethewrite.If *string*isomitted,nothingsiswrittentothestream.If *string*isspecified,asystem-specificactionistakenafterithasbeenwrittentostream,tomarkanewline.

Theformatandcontentsofthefirstparameterwilldependupontheimplementationandhowitnames streams.Consultimplementation-specificdocumentationformoreinformation.

If *string*isspecified,butnot *line*,theeffectistowrite *string*tothestream,startingatthecurrentwrite position.If *line*isspecified,butnot *string*,theeffectisonlytopositionthestreamatthenewposition. Notethatthe *line*parameterisonlylegalifthestreamispersistent;youcannotpositionthecurrent writepositionfortransientstreams.

Ifneither *line*nor *string*isspecified,thestandardrequiresthatthecurrentwritepositionissettheend ofthestream,andimplementationspecificside-effectsmayoccur.Inpractice,thismeansthatan implementationcanusethissituationtodothingslikeclosingthestream,orflushingtheoutput. Consulttheimplementationspecificdocumentationformoreinformation.

Alsonotethatthereturnvalueofthisfunctionsmaybeoflittleornovalue.Ifjustahalflineiswritten, 1maystillbereturned,andtherearenowayoffindingouthowmuch(ifany)of *string*waswritten.If *string*isnotspecified,thereturnvaluewillalwaysbe 0,evenif LINEOUT()wasnotabletocorrectly positionthestream.

Ifitisimpossibletocorrectlywrite *string*tothestream,the NOTREADYflagwillberaised.Itisnot definedwhetherornotthe NOTREADYflagisraisedwhen LINEOUT()isusedforpositioning,and thisisnotpossible.

Notethatifyouwrite *string*toalineinthemiddleofthestream(i.e. *line*islessthanthetotalnumber oflinesinthestream),thenbehaviorissystemandimplementationspecific.Somesystemswill truncatethestreamafterthenewlywrittenline,otherwillonlytruncateifthenewlywrittenlinehasa differentlengththantheoldlinewhichitreplaced,andyetothersystemswilloverwriteandnever truncate.

Ingeneral,consultyoursystemandimplementationspecificdocumentationformoreinformationabout thisfunction.Youcansafelyassumeverylittleabouthowitbehaves.

LINEOUT(,'First line')	'1'
LINEOUT('/tmp/file','Second line',2)	'1'
LINEOUT('/tmp/file','Third line')	'1'
LINEOUT('/tmp/file','Fourth line',4)	'0'

LINES([streamid][,option]) **(ANSI)**

Returns 1 if there is at least one complete line remaining in the named file *stream* or 0 if no complete lines remain in the file. A complete line is not really as complete as the name might indicate; a complete line is zero or more characters, followed by an End-Of-Line (EOL) marker. So, if you have read half a line already, you still have a "complete" line left. Note that it is not defined what to do with a half-finished line at the end of a file. Some interpreters might interpret the End-Of-File as an implicit EOL mark too, while others might not.

The format and content of the stream *streamid* is system and implementation dependent. If omitted, the default input stream will be used.

The ANSI Standard has extended this function from TRL2. It allows an *option*:

[C]

(Count) Returns the actual number of complete lines remaining in the stream, irrespective of how expensive this operation is.

[N]

(Normal) Returns 1 if there is at least one complete line remaining in the file or 0 if no lines remain. This is the default. To maintain backwards compatibility with older releases of the `OPTION;NOFAST_LINES_BIF_DEFAULT` can be used to make the default option behave as though `LINES(streamid,'C')` was specified.

Regina,

`LINES` will only return 0 or 1 for all transient streams, as the interpreter cannot reposition in these files, and can therefore not count the number of remaining lines.

As a result, defensive programming indicates that you can safely only assume that this function will return either 0 or a non-zero result. If you want to use the non-zero result to more than just an indicator on whether more lines are available, you must check that it is larger than one. If so, you can safely assume that it holds the number of available lines left.

As with all the functions operating on streams, you can safely assume very little about this function, so consult the system and implementation specific documentation.

<code>LINES()</code>	<code>'1' /* Maybe */</code>
<code>LINES()</code>	<code>'0' /* Maybe */</code>
<code>LINES('/tmp/file','C')</code>	<code>'2' /* Maybe */</code>
<code>LINES('/tmp/file')</code>	<code>'1' /* Maybe */</code>

MAKEBUF() **(CMS)**

Creates a new buffer on the stack, at the current top of the stack. Each new buffer will be assigned a number; the first buffer being assigned the number 1. A new buffer will be assigned a number which is one higher than the currently highest number of any buffer on the stack. In practice, this means that the buffers are numbered, with the bottom-most having the number 1 and the top-most having a number

which value is identical to the number of buffers currently in the stack.

The value returned from this function is the number assigned to the newly created buffer. The assigned number will be one more than the number of buffers already in the stack, so the numbers will be "recycled". Thus, the assigned numbers will not necessarily be in sequence.

MAKEBUF ()	1 /* if no buffers existed */
MAKEBUF ()	6 /* if 5 buffers existed */

MAX(number1[,number2]...) (ANSI)

Takes any positive number of parameters, and will return the parameter that had the highest numerical value. The parameters may be any valid REXX number. The number that is returned, is normalized according to the current settings of NUMERIC, so the result need not be strictly equal to any of the parameters.

Actually, the standard says that the value returned is the first number in the parameter list which is equal to the result of adding a positive number or zero to any of the other parameters. Note that this definition opens for "strange" results if you are brave enough to play around with the settings of NUMERIC FUZZ.

MAX (1 , 2 , 3 , 5 , 4)	' 5 '
MAX (6)	' 6 '
MAX (-4 , .001E3 , 4)	' 4 '
MAX (1 , 2 , 05.0 , 4)	' 5.0 '

MIN(number[,number]...) (ANSI)

Like MAX () , except that the lowest numerical value is returned. For more information, see MAX () .

MAX (5 , 4 , 3 , 1 , 2)	' 1 '
MAX (6)	' 6 '
MAX (-4 , .001E3 , 4)	' -4 '
MAX (1 , 2 , 05.0E-1 , 4)	' 0.50 '

OPEN(file, filename, ['Append' | 'Read' | 'Write']) (AREXX)

Opens a file for the specified operation. The *file* argument defines the logical name by which the file will be referenced. The *filename* is the external name of the file, and may include any portion of a full file path.

The function returns a boolean value that indicates whether the operation was successful. There is no limit to the number of files that can be opened simultaneously, and all open files are closed automatically when the program exits.

See also CLOSE(), READ(), WRITE()

OPEN('myfile','c:\temp\aa.txt','R')	'1'
OPEN('infile','/tmp/fred.txt')	'1'

OVERLAY(string1,string2[, [start][, [length][, padchar]]) (ANSI)

Returns a copy of *string2*, totally or partially overwritten by *string1*. If these are the only arguments, the overwriting starts at the first character in *string2*.

If *start* is specified, the first character in *string1* overwrites character number *start* in *string2*. *Start* must be a positive whole number, and default to 1, i.e. the first character of *string1*. If the *start* position is to the right of the end of *string2*, then *string2* is padded at the right hand end to make it *start*-1 characters long, before *string1* is added.

If *length* is specified, then *string2* will be stripped or padded at the right hand end to match the specified length. For padding (of both strings) *padchar* will be used, or <space> if *padchar* is unspecified. *Length* must be non-negative, and default to the length of *string1*.

OVERLAY('NEW','old-value')	'NEW-value'
OVERLAY('NEW','old-value',4)	'oldNEWlue'
OVERLAY('NEW','old-value',4,5)	'oldNEW e'
OVERLAY('NEW','old-value',4,5,'*')	'oldNEW**e'
OVERLAY('NEW','old-value',4,2)	'oldNEalue'
OVERLAY('NEW','old-value',9)	'old-valuNEW'
OVERLAY('NEW','old-value',12)	'old-value NEW'
OVERLAY('NEW','old-value',12,, '*')	'old-value**NEW'
OVERLAY('NEW','old-value',12,5,'*')	'old-value**NEW**'

POPEN(command[, stem.]) (REGINA)

Run the operating system *command*. If the optional *stem.* is supplied all output from the *command* is placed in the specified stem variable as a REXX array. Note that only the command's stdout can be captured.

This command is now deprecated. ADDRESS WITH can do the same thing, and can also capture the command's stderr.

POPEN('ls -l', 'lists.')	/* LISTS. stem has list */
ADDRESS SYSTEM 'ls -l' WITH OUTPUT STEM LISTS.	/* same as above */

POS(*needle*,*haystack*[,*start*]) (ANSI)

Seeks for an occurrence of the string *needle* in the string *haystack*. If *needle* is not found, then 0 is returned. Else, the position in *haystack* of the first character in the part that matched is returned, which will be a positive whole number. If *start* (which must be a positive whole number) is specified, the search for *needle* will start at position *start* in *haystack*.

POS('be','to be or not to be')	4
POS('to','to be or not to be',10)	14
POS('is','to be or not to be')	0
POS('to','to be or not to be',18)	0

QUALIFY([*streamid*]) (ANSI)

Returns a name for the *streamid*. The two names are currently associated with the same resource and the result of this function may be more persistently associated with that resource.

QUALIFY(' ../mypath/fred.the')	/home/mark/mypath/fred.the '
--------------------------------	------------------------------

QUEUED() (ANSI)

Returns the number of lines currently in the external data queue (the "stack"). Note that the stack is a concept external to REXX, this function may depend on the implementation and system. Consult the system specific documentation for more information.

QUEUED()	'0' /* Maybe */
QUEUED()	'42' /* Maybe */

RANDOM(*max*) (ANSI)

RANDOM([*min*][,*max*][,*seed*]) (ANSI)

Returns a pseudo-random whole number. If called with only the first parameter, the first format will be used, and the number returned will be in the range 0 to the value of the first parameter, inclusive. Then the parameter *max* must be a non-negative whole number, not greater than 100000.

If called with more than one parameter, or with one parameter, which is not the first, the second format will be used. Then *min* and *max* must be positive whole numbers, and *max* cannot be less than *min*, and the difference *max*-*min* cannot be more than 100000. If one or both of them is unspecified, the default for *min* is 0, and the default for *max* is 999.

If *seed* is specified; (it must be a positive whole number) you may control which numbers the pseudo-random algorithm will generate. If you do not specify it, it will be set to some "random" value at the first call to RANDOM() (typically a function of the time). When specifying *seed*, it will effect the result

of the current call to `RANDOM()`.

The standard does not require that a specific method is to be used for generating the pseudo-random numbers, so there reproducibility can only be guaranteed as long as you use the same implementation on the same machine, using the same operating system. If any of these change, a given *seed* may produce a different sequence of pseudo-random numbers.

Note that depending on the implementation, some numbers might have a slightly increased chance of turning up than others. If the `REXX` implementation uses a 32-bit pseudo-random generator provided by the operating system and returns the remainder after integer dividing it by the difference of *min* and *max*, low numbers are favored if the 2^{32} is not a multiple of that difference. Supposing that the call is `RANDOM(100000)` and the pseudo-random generator generates any 32-bit number with equal chance, the change of getting a number in the range 0 – 67296 is about 0.000010000076, while the change of getting a number in the range 67297 – 100000 is about 0.0000099999843.

A much worse problem with pseudo-random numbers are that they sometimes do not tend to be random at all. Under one operating system (name withheld to protect the guilty), the system's pseudo-random routine returned numbers where the last binary digit alternated between 0 and 1. On that machine, `RANDOM(1)` would return the series 0, 1, 0, 1, 0, 1, etc., which is hardly random at all. You should therefore never trust the pseudo-random routine to give you random numbers.

Note that due to the special syntax, there is a big difference between using `RANDOM(10)` and `RANDOM(10 ,)`. The former will give a pseudo-random number in the range 0 – 10, while the latter will give a pseudo-random number in the range 10 – 999.

Also note that it is not clear whether the standard allows *min* to be equal to *max*, so to program compatible, makes sure that *max* is always larger than *min*.

<code>RANDOM()</code>	<code>'123' /*Between 0 and 999 */</code>
<code>RANDOM(10)</code>	<code>'5' /*Between 0 and 10 */</code>
<code>RANDOM(, 10)</code>	<code>'3' /*Between 0 and 10 */</code>
<code>RANDOM(20 , 30)</code>	<code>'27' /*Between 20 and 30 */</code>
<code>RANDOM(, , 12345)</code>	<code>'765' /*Between 0 and 999, and sets seed */</code>

RANDU([seed]) (AREXX)

Returns a uniformly-distributed pseudo-random number between 0 and 1. The number of digits of precision in the result is always equal to the current Numeric Digits setting. With the choice of suitable scaling and translation values, `RANDU()` can be used to generate pseudo-random numbers on an arbitrary interval.

The optional *seed* argument is used to initialize the internal state of the random number generator. See also `RANDOM()`

RANDU ()	'0.371902021'
RANDU (45)	'0.873' /*numeric digits 3*/

READCH(file,length) (AREXX)

Readsthespecifiednumberofcharactersfromthegivenlogicalfileandretrunsthem.Thelengthofthe returnedstringistheactualnumberofcharactersread,andmaybelessthantherequestedlengthif,for example,theend-of-filewasreached.

SeealsoREADLN()

READCH('infile',10)	'a string o'
---------------------	--------------

READLN(file) (AREXX)

Readscharactersfromthegivenlogicalfileintoastringuntila"newline"characterisfound.The returnedstringdoesnotincludethe"newline".

SeealsoREADCH()

READLN('infile')	'a string of chars'
------------------	---------------------

REVERSE(string) (ANSI)

Returnsastringofthesamelengthas *string*,buthavingtheorderofthecharactersreversed.

REVERSE('FooBar')	'raBooF'
REVERSE(' Foo Bar')	'raB ooF '
REVERSE('3.14159')	'95141.3'

RIGHT(string,length[,padchar]) (ANSI)

Returnsthe *length*rightmostcharactersin *string*.If *length*(whichmustbeanon-negativewhole number)isgreaterthanthelengthof *string*theresultispaddedontheleftwiththenecessarynumberof *padchar*stomakeitaslongas *length*specifies. *Padchar*defaultsto<space>.

RIGHT('Foo bar',5)	'o bar'
RIGHT('Foo bar',3)	'bar'
RIGHT('Foo bar',10)	' Foo bar'
RIGHT('Foo bar',10,'*')	'****Foo bar'

RXFUNCADD(externalname,library,internalname) (SAA)

Registersthe *internalname*in *library*asanexternalfunctioncallablefromwiththecurrentprogramby referencing *externalname*. *library*isaREXXexternalfunctionpackageintheformatofsharedlibrary

or dynamic link library (DLL). *library* and *internalname* are case-sensitive. *library* is the **basename** of the shared library or dynamic link library. On platforms that support DLLs, the full name of the external function package is *library.dll*. On Unix environments, the full name of the shared library is **liblibrary.a** (AIX), **liblibrary.sl** (HPUX) or **liblibrary.so** (most other Unixes). External function packages are searched for in the location where shared libraries or DLLs are normally found by the operating system. DLLs are normally located in directories specified in the **PATH** or **LIBPATH** environment variables. Shared libraries are normally searched for in **LD_LIBRARY_PATH** or **LIBPATH** environment variables.

This function returns 0 if the function is registered successfully.

RXFUNCADD('SQLLoadFuncs' , 'rexsql' , 'SQLLoadFuncs')	0
---	---

RXFUNCDROP(externalname) **(SAA)**

Removes the specified external name from the list of external functions available to be called. This function returns 0 if the function was successfully dropped.

RXFUNCDROP('SQLLoadFuncs')	0
------------------------------	---

RXFUNCERRMSG() **(REGINA)**

Returns the error message associated with the last call to RXFUNCADD. This function is generally used immediately after a failed call to RXFUNCADD to determine why it failed.

RXFUNCERRMSG()	'rexsql.dll not found' /* Maybe */
----------------	------------------------------------

RXFUNCQUERY(externalname) **(SAA)**

Returns 0 if the *externalname* is already registered, or 1 if the *externalname* is not registered.

RXFUNCQUERY('SQLLoadFuncs')	1 /* Maybe */
-------------------------------	---------------

RXQUEUE(command[,queue]) **(OS/2)**

This function interfaces to the Regina internal or external queue mechanism. If **OPTIONS INTERNAL_QUEUE** is set, all operations on queues are internal to the interpreter.

[C]

(Create) Request the interpreter or rxstack to create a new named *queue*. If the *queue* name already exists, a new unique queue name is generated. The name of the queue that was created (either the specified queue or the system-generated queue) is returned. All queue names are case-insensitive; i.e. the queue name **FRED** and **fred** are the same.

[D]

(Delete) Delete the specified *queue*. The default queue; **SESSION** becomes the current queue.

[G]

(Get)Returns the current *queue* name.

[S]

(Set)Sets the current queue name to that *queue* specified. The previously current queue is returned. It is valid to set a queue name to a queue that has not been created.

[T]

(Timeout)Sets the timeout period (in milliseconds) to wait for something to appear on the specified queue. By default, when a line is read from a queue with a PULL command, it either returns immediately with the top line in the stack, or it will wait for a line to be entered by the user via the process's stdin. If 0 is specified, Regina will wait forever for a line to be ready on the stack.

RXQUEUE('Create')	'S0738280'
RXQUEUE('Create','fred')	'FRED'
RXQUEUE('Create','fred')	'S88381'
RXQUEUE('Get')	'S88381'
RXQUEUE('Delete','fred')	'SESSION'
RXQUEUE('Set','fred')	'SESSION'
RXQUEUE('Timeout','fred',10)	' '

SEEK(file,offset,['Begin'|'Current'|'End']) (AREXX)

Moves to a new position in the given logical file, specified as an *offset* from an anchor position. The default anchor is Current. The returned value is the new position relative to the start of the file.

SEEK('infile',10,'B')	'10'
SEEK('infile',0,'E')	'356' /* file length */

SHOW(option,[name],[pad]) (AREXX)

Returns the names in the resource lists specified by the *option* argument, or tests to see whether an entry with the specified *name* is available. The currently implemented options keywords are Clip, Files, Libraries, and Ports, which are described below.

Clip.Examines the names in the ClipList.

Files.Examines the names of the currently open logical filenames.

Libraries.Examines the names in the LibraryList, which are either function libraries or function hosts.

Ports.Examines the names in the system PortsList.

If the *name* argument is omitted, the function returns a string with the resource names separated by a blank space or the *pad* character, if one was supplied. If the *name* argument is given, the returned boolean value indicates whether the *name* was found in the resource list. The *name* entries are case-sensitive.

Only the **Files** *option* is valid on all platforms. All other values for *option* are only applicable to the Amiga and AROS ports.

SIGN(number) (ANSI)

Return either -1, 0 or 1, depending on whether *number* is negative, zero, or positive, respectively. *Number* must be a valid REXX number, and are normalized according to the current settings of NUMERIC before comparison.

SIGN(-12)	'-1'
SIGN(42)	'1'
SIGN(-0.00000012)	'-1'
SIGN(0.000)	'0'
SIGN(-0.0)	'0'

SLEEP(seconds) (CMS)

Pauses for the supplied number of seconds.

SLEEP(5)	/* sleeps for 5 seconds */
----------	----------------------------

SOURCELINE([lineno]) (ANSI)

If *lineno* (which must be a positive whole number) is specified, this function will return a string containing a copy of the REXX script source code on that line. If *lineno* is greater than the number of lines in the REXX script source code, an error is reported.

If *lineno* is unspecified, the number of lines in the REXX script source code is returned.

Note that from REXX language level 3.50 to 4.00, the requirements of this function were relaxed to simplify execution when the source code is not available (compiled or pre-parsed REXX). An implementation might make two simplifications: to return 0 if called without a parameter. If so, any call to SOURCELINE() with a parameter will generate an error. The other simplification is to return a null string for any call to SOURCELINE() with a legal parameter.

Note that the code executed by the INTERPRET clause cannot be retrieved by SOURCELINE().

SOURCELINE()	'42' /*Maybe */
SOURCELINE(1)	'/* This REXX script will ... */'
SOURCELINE(23)	'var = 12' /*Maybe */'

SPACE(string[, [length][, padchar]]) (ANSI)

With only one parameter *string* is returned, stripped of any trailing or leading blanks, and any consecutive blanks inside *string* translated to a single <space> character (or *padchar* if specified).

Length must be a non-negative whole number. If specified, consecutive blanks within *string* is replaced

by exactly *length* instances of <space> (or *padchar* if specified). However, *padchar* will only be used in the output string, in the input string, blanks will still be the "magic" characters. As a consequence, if there exist any *padchars* in *string*, they will remain untouched and will not affect the spacing.

SPACE(' Foo bar ')	'Foo bar '
SPACE(' Foo bar ',2)	'Foo bar '
SPACE(' Foo bar ',,'*')	'Foo*bar '
SPACE('Foo bar ',3,'-')	'Foo--bar '
SPACE('Foo bar ',,'o')	'Fooobar '

STATE(*streamid*) (CMS)

Returns 0 if the *streamid* exists, or 1 if it does not. Use STREAM(*streamid*, 'C', 'QUERY EXISTS') for portability.

STORAGE([*address*],[*string*],[*length*],[*pad*]) (AREXX)

Calling STORAGE() with no arguments returns the available system memory. If the *address* argument is given, it must be a 4-byte string, and the function copies data from the optional *string* into the indicated memory area. The *length* parameter specifies the maximum number of bytes to be copied, and defaults to the length of the string. If the specified length is longer than the string, the remaining area is filled with the *pad* character or nulls ('00'x.)

The returned value is the previous contents of the memory area. This can be used in a subsequent call to restore the original contents.

Caution is advised in using this function. Any area of memory can be overwritten, possibly causing a system crash.

STORAGE()	'248400 '
STORAGE('0004 0000'x, 'The answer')	'question' /* maybe */

STREAM(*streamid*[,*option*],[*command*]) (ANSI)

This function was added to REXX in language level 4.00. It provides a general mechanism for doing operations on streams. However, very little is specified about how the internal of this function should work, so you should consult the implementation specific documentation for more information.

The *streamid* identifies a stream. The actual contents and format of this string is implementation dependent.

The *option* selects one of several operations which STREAM() is to perform. The possible operations are:

[C]

(Command) If this option is selected, a third parameter must be present, *command*, which is the

command to be performed on the stream. The contents of *command* is implementation dependent. For **Regina**, the valid commands follow. Commands consist of one or more space separated words.

[D]

(Description) Returns a description of the state of *streamid*. The return value is implementation dependent.

[S]

(Status) Returns a state which describes the state of *streamid*. The standard requires that it is one of the following: **ERROR**, **NOTREADY**, **READY** and **UNKNOWN**. The meaning of these are described in the chapter; **StreamInput and Output** .

Note that the options **Description** and **Status** really have the same function, but that **Status** in general is implementation independent, while **Description** is implementation dependent.

The *command* specifies the command to be performed on *streamid*. The possible operations are:

[READ]

Open for read access. The file pointer will be positioned at the start of the file, and only read operations are allowed. This command is **Regina-specific**; use **OPEN_READ** in its place.

[WRITE]

Open for write access and position the current write position at the end of the file. An error is returned if it was not possible to get appropriate access. This command is **Regina-specific**; use **OPEN_WRITE** in its place.

[APPEND]

Open for append access and position the current write position at the end of the file. An error is returned if it was not possible to get appropriate access. This command is **Regina-specific**; use **OPEN_WRITE_APPEND** in its place.

[UPDATE]

Open for append access and position the current write position at the end of the file. An error is returned if it was not possible to get appropriate access. This command is **Regina-specific**; use **OPEN_BOTH** in its place.

[CREATE]

Open for write access and position the current write position at the start of the file. An error is returned if it was not possible to get appropriate access. This command is **Regina-specific**; use **OPEN_WRITE_REPLACE** in its place.

[CLOSE]

Close the stream, flushing any pending writes. An error is returned if it was not possible to get appropriate access.

[FLUSH]

Flush any pending write to the stream. An error is returned if it was not possible to get appropriate access.

[STATUS]

Returns status information about the stream in human readable form that **Regina** stores about the stream.

[FSTAT]

Returns status information from the operating system about the stream. This consists of at least 8 words:

Device Number	UnderDOS,Win32,OS/2,thisrepresentsthedisknumber,with0 beingDriveA .
Inode Number	UnderDOS,Win32,OS/2,thisiszero.
Permissions	User/Group/Otherpermissionsmask.Consistsof3octalnumbers with4representingread,2representingwrite,and1representing execute.Thereforeavalueof750isread/write/executeforuser, read/executeforgroup,andnopermissionsforother.
Number Links	UnderDOS,Win32,OS/2,thiswillalwaysbe1.
User Name	Theownerofthestream.UnderDOS,Win32,OS/2,thiswill alwaysbe“USER”.
Group Name	Thegroupownerofthestream.UnderDOS,Win32,OS/2,this willalwaysbe“GROUP”.
Size	Sizeofstreaminbytes.
Stream Type	Oneormoreofthefollowing: RegularFile anormalfile. Directory adirectory. BlockSpecial ablockspecialfile. FIFO usuallyapipe. SymbolicLink asymboliclink Socket asocket SpecialName anamedspecialfile.

Ifthestreamisasymboliclink,thethedetailsreturnedaredetailsaboutthelink,notthe filethelinkpointsto. **CharacterSpecial** acharacterspecialfile.

[RESET]

Resets.thestreamafteranerror.Onlystreamsthatareresettablecanbereset.

[READABLE]

Returns1if.thestreamisreadableby.theuseror0otherwise.

[WRITABLE]

Returns1if.thestreamiswriteableby.theuseror0otherwise.

[EXECUTABLE]

Returns1if.thestreamisexecutableby.theuseror0otherwise.

[QUERY]

Returnsinformationaboutthenamedstream.Ifthenamedstreamdoesnotexists,thenthe emptystringisreturned.Thiscommandisfurtherbrokendownintothefollowingsub- commands:

DATETIME	returnsthe.dateand.time.of.last.modification.of.the.stream.in USDateformat;MM-DD-YYHH:MM:SS.	Rexx
EXISTS	returnsthefully-qualified.filename.of.the.specified.stream.	
HANDLE	returnstheinternal.file.handle.of.the.stream.This.will.only.return.a valid.value.if.the.stream.was.opened.explicitly.or.implicitly.by Regina.	
SEEK READ CHAR	returnsthe.current.read.position.of.the.open.stream.expressed.in characters.	

SEEK READ LINE	returnsthecurrentreadpositionoftheopenstreamexpressedin lines.
SEEK WRITE CHAR	returnsthecurrentwritepositionoftheopenstreamexpressedin characters.
SEEK WRITE LINE	returnsthecurrentwritepositionoftheopenstreamexpressedin lines.
SEEK SYS	returnsthecurrentreadpositionoftheopenstreamastheoperating reportsit. Thisisexpressedincharacters.
SIZE	returnsthesize,expressedincharacters,ofthepersistentstream.
STREAMTYPE	returnsthe typeofthestream. OneofTRANSIENT,PERSISTENT orUNKNOWNisreturned.
TIMESTAMP	returnsthedateandtimeoflastmodificationofthestream. The formatofthestringreturnedisYYYY-MM-DDHH:MM:SS.

You can use **POSITION** in place of **SEEK** in the above options.

[OPEN]

Open the stream in the optional modes specified. If no optional mode is specified, the default is **OPEN BOTH**.

READ	The file pointer will be positioned at the start of the file, and only read operations are allowed.
WRITE	Open for write access and position the current write pointer at the end of the file. On platforms where it is not possible to open a file for write without also allowing reads, the read pointer will be positioned at the start of the file. An error is returned if it was not possible to get appropriate access.
BOTH	Open for read and write access. Position the current read pointer at the start of the file, and the current write pointer at the end of the file. An error is returned if it was not possible to get appropriate access.
WRITE APPEND	Open for write access and position the write pointer at the end of the file. On platforms where it is not possible to open a file for write without also allowing reads, the read pointer will be positioned at the start of the file.
WRITE REPLACE	Open for write access and position the current write position at the start of the file. On platforms where it is not possible to open a file for write without also allowing reads, the read pointer will be positioned at the start of the file. This operation will clear the contents of the file. An error is returned if it was not possible to get appropriate access.
BOTH APPEND	Open for read and write access. Position the current read position at the start of the file, and the current write position at the end of the file. An error is returned if it was not possible to get appropriate access.
BOTH REPLACE	Open for read and write access. Position both the current read and write pointers at the start of the file. An error is returned if it was not possible to get appropriate access.

[**SEEK** *position* **READ**|**WRITE** [**CHAR**|**LINE**]]

Position the file's read or write pointer in the file to the specified *position*. **SEEK** is a synonym for **POSITION**.

position A position can be of the following forms. [*relative*] *offset*.

relative can be one of:

= The file pointer is moved to the specified *offset* relative to the start of the file. This is the default.

< The file pointer is moved to the specified *offset* relative to the end of the file.

- The file pointer is moved backwards relative to the current position.

+ The file pointer is moved forwards relative to the current position.

offset is a positive whole number.

READ The read file pointer will be positioned.

WRITE The write file pointer is positioned.

CHAR The *offset* specified in *position* above is in terms of characters.

LINE The *offset* specified in *position* above is in terms of lines.

Assume a file; '/home/mark/myfile' last changed March 30th 2002 at 15:07:56, with 100 lines, each line 10 characters long, and the following command executed in sequence.

STREAM('myfile','C','QUERY EXISTS')	'/home/mark/myfile'
STREAM('myfile','C','QUERY SIZE')	1100
STREAM('myfile','C','QUERY TIMESTAMP')	2002-03-30 15:07:56
STREAM('myfile','C','QUERY DATETIME')	03-30-02 15:07:56
STREAM('myfile','D')	
STREAM('myfile','S')	UNKNOWN
STREAM('myfile','C','QUERY SEEK READ')	
STREAM('myfile','C','OPEN READ')	READY:
STREAM('myfile','D')	
STREAM('myfile','S')	READY
STREAM('myfile','C','QUERY SEEK READ')	1
STREAM('myfile','C','CLOSE')	UNKNOWN
STREAM('myfile','C','STATUS')	
STREAM('myfile','C','FSTAT')	773 35006 064 1 mark
STREAM('myfile','C','READABLE')	1
STREAM('myfile','C','WRITABLE')	1
STREAM('myfile','C','EXECUTABLE')	0
STREAM('myfile','C','???')	

STRIP(string[, [option][, char]]) (ANSI)

Returns *string* after possibly stripping it of any number of leading and/or trailing characters. The default action is to strip off both leading and trailing blanks. If *char* (which must be a string containing exactly one character) is specified, that character will be stripped off instead of blanks. Inter-word blanks (or *chars* if defined, that are not leading or trailing) are untouched.

If *option* is specified, it will define what to strip. The possible values for *option* are:

[L]

(Leading) Only strip off leading blanks, or *chars* if specified.

[T]

(Trailing) Only strip off trailing blanks, or *chars* if specified.

[B]

(Both) Combine the effect of *L* and *T*, that is, strip off both leading and trailing blanks, or *chars* if it is specified. This is the default action.

STRIP(' Foo bar ')	'Foo bar'
STRIP(' Foo bar ','L')	'Foo bar '
STRIP(' Foo bar ','t')	' Foo bar'
STRIP(' Foo bar ','Both')	'Foo bar'
STRIP('0.1234500',,'0')	'.12345'
STRIP('0.1234500 ',' ','0')	'.1234500'

SUBSTR(string,start[,length][,padchar]) (ANSI)

Returnsthesubstringof *string*thatstartsat *start*,andhasthelength *length*. *Length*defaultsto therest of thestring. *Start*mustbeapositivewhole,while *length*canbeanynon-negativewholenumber.

Itisnotanerrorfor *start*to belargerthan thelengthof *string*.If *length*isspecifiedandthesumof *length*and *start*minus1isgreaterthan thelengthof *string*,thenthe resultwillbepaddedwith *padchar*tothespecifiedlength.The defaultvaluefor *padchar*isthe<space>character.

SUBSTR(' Foo bar ',3)	'o bar'
SUBSTR(' Foo bar ',3,3)	'o b'
SUBSTR(' Foo bar ',4,6)	' bar '
SUBSTR(' Foo bar ',4,6,'*')	' bar**'
SUBSTR(' Foo bar ',9,4,'*')	'****'

SUBWORD(string,start[,length]) (ANSI)

Returnsthepartof *string*thatstartsatblankdelimitedword *start*(whichmustbeapositivewhole number).If *length*(whichmustbeanon-negativewholenumber)isspecified,thatnumberof wordsare returned.The defaultvaluefor *length*istherestof thestring.

Itisnotanerrortospecify *length*to refertomorewordsthan *string*contains,orfor *start*and *length*togethertospecifymorewordsthan *string*holds.The resultstringwillbestrippedof anyleadingand trailingblanks,butinter-wordblanks willbepreservedasis.

SUBWORD('To be or not to be',4)	'not to be'
SUBWORD('To be or not to be',4,2)	'not to'
SUBWORD('To be or not to be',4,5)	'not to be'
SUBWORD('To be or not to be',1,3)	'To be or'

SYMBOL(name) (ANSI)

Checksif thestring *name*isavalidsymbol(a positivenumberorapossiblevariablename),andreturns athreeletterstringindicatingtheresultof thatcheck.If *name*is a symbol,andnamesacurrentlyset variable, VARisreturned;if *name*isalegalsymbolname,buthasnotabeengivenavalue(orisa

constant symbol, which cannot be used as a variable name), `LIT` is returned to signify that it is a literal. Else, if `name` is not a legal symbol name the string `BAD` is returned.

Watch out for the effect of "double expansion". `Name` is interpreted as an expression evaluating naming the symbol to be checked, so you might have to quote the parameter.

<code>SYMBOL(' Foobar ')</code>	<code>'VAR' /* Maybe */</code>
<code>SYMBOL(' Foo bar ')</code>	<code>'BAD'</code>
<code>SYMBOL(' Foo.Foo bar ')</code>	<code>'VAR' /* Maybe */</code>
<code>SYMBOL(' 3.14 ')</code>	<code>'LIT'</code>
<code>SYMBOL(' .Foo->bar ')</code>	<code>'BAD'</code>

TIME([option_out [,time [option_in]]) **(ANSI)**

Returns a string containing information about the time. To get the time in a particular format, an `option_out` can be specified. The default `option_out` is `Normal`. The meaning of the possible options are:

- [C]** (Civil) Return the time in civil format. The return value might be "`hh:mmXX`", where `XX` are either `am` or `pm`. The `hh` part will be stripped of any leading zeros, and will be in the range 1–12 inclusive.
- [E]** (Elapsed) Return the time elapsed in seconds since the internal stopwatch was started. The result will not have any leading zeros or blanks. The output will be a floating point number with six digits after the decimal point.
- [H]** (Hours) Return the number of complete hours that have passed since last midnight in the form "`hh`". The output will have no leading zeros, and will be in the range 0–23.
- [L]** (Long) Return the exact time, down to the microsecond. This is called the long format. The output might be "`hh:mm:ss.mmmmmm`". Be aware that most computers do not have a clock of that accuracy, so the actual granularity you can expect, will be about a few milliseconds. The `hh`, `mm` and `ss` parts will be identical to what is returned by the options `H`, `M` and `S` respectively, except that each part will have leading zeros as indicated by the format.
- [M]** (Minutes) Return the number of complete minutes since midnight, in a format having no leading zeros, and will be in the range 0–59.
- [N]** (Normal) The output format is "`hh:mm:ss`", and is padded with zeros if needed. The `hh`, `mm` and `ss` will contain the hours, minutes and seconds, respectively. Each part will be padded with leading zeros to make it double-digit.
- [R]** (Reset) Return the value of the internal stopwatch just like the `E` option, and using the same format. In addition, it will reset the stopwatch to zero after its contents has been read.

[S]

(Seconds) Return the number of complete seconds since midnight, in a format having no leading spaces, and will be in the range 0-59.

[T]

(*time_t*) Return the current date/time in UNIX *time_t* format. *time_t* is the number of seconds since January 1st 1970.

Note that the time is never rounded, only truncated. As shown in the examples below, these seconds do not get rounded upwards, even though the decimal part implies that they are close to 59th to 58. The same applies for the minutes, which are close to 33th to 32, but is truncated to 32. None of the formats will have leading or trailing spaces.

Assuming that the time is exactly 14:32:58.987654 on March 30th 2002, the following will be true:

TIME('C')	' 2 : 32pm '
TIME('E')	' 0.01200 ' /* Maybe */
TIME('H')	' 14 '
TIME('L')	' 14 : 32 : 58.987654 '
TIME('M')	' 32 '
TIME('N')	' 14 : 32 : 58 '
TIME('R')	' 0.430221 ' /* Maybe */
TIME('S')	' 58 '

If the *time* option is specified, the function provides for time conversions. The optional *option_in* specifies the format in which *time* is supplied. The possible values for *option_in* are: **CHLMNS**. The default value for *option_in* is **N**.

TIME('C' , ' 11 : 27 : 21 ')	' 11 : 27am '
TIME('N' , ' 11 : 27am ' , 'C')	' 11 : 27 : 00 '

The time conversion capability of the TIMEBIF was introduced with the ANSI standard.

TRACE([setting])

(ANSI)

Return the current value of the trace setting. If the string *setting* is specified, it will be used as the new setting for tracing, after the old value has been recorded for the return value. Note that the *setting* is not an option, but may be any of the trace settings that can be specified to the clause TRACE, except that the numeric variant is not allowed with TRACE(). In practice, this can be a word, of which only the first letter counts, optionally preceded by a question mark.

TRACE ()	'C' /* Maybe */
TRACE ('N') 'C'	
TRACE ('?') 'N'	

TRANSLATE(string[, [tableout][, [tablein][, padchar]]) (ANSI)

Performs a translation on the characters in *string*. As a special case, if neither *tablein* nor *tableout* is specified, it will translate *string* from lower case to upper case. Note that this operation may depend on the language chosen, if your interpreter supports national character sets.

Two translation tables might be specified as the strings *tablein* and *tableout*. If one or both of the tables are specified, each character in *string* that exists in *tablein* is translated to the character in *tableout* that occupies the same position as the character did in *tablein*. The *tablein* defaults to the whole character set (all 256) in numeric sequence, while *tableout* defaults to an empty set. Characters not in *tablein* are left unchanged.

If *tableout* is larger than *tablein*, the extra entries are ignored. If it is smaller than *tablein* it is padded with *padchar* to the correct length. *Padchar* defaults to <space>.

If a character occurs more than once in *tablein*, only the first occurrence will matter.

TRANSLATE ('FooBar')	'FOOBAR'
TRANSLATE ('FooBar' , 'ABFORabfor' , 'abforABFOR')	'fOObAR'
TRANSLATE ('FooBar' , 'abfor')	'F B '
TRANSLATE ('FooBar' , 'abfor' , , '#')	'F##B##'

TRIM(string) (AREXX)

Remove trailing blanks from the string argument. A more portable option is to use the Trailing option of the STRIPBIF.

TRIM (' abc ')	' abc '
------------------	---------

TRUNC(number[, length]) (ANSI)

Returns *number* truncated to the number of decimalsspecified by *length*. *Length* defaults to 0, that is return an whole number with no decimal part.

The decimal point will only be present if there is a non-empty decimal part, i.e. *length* is non-zero. The number will always be returned in simple form, never exponential form, no matter what the current settings of *NUMERIC* might be. If *length* specifies more decimal than *number* has, extra zeros are appended. If *length* specifies less decimal than *number* has, then the number is truncated. Note that *number* is never rounded, except for the rounding that might take place during normalization.

TRUNC (12.34)	' 12 '
TRUNC (12.99)	' 12 '
TRUNC (12.34, 4)	' 12.3400 '
TRUNC (12.3456, 2)	' 12.34 '

UNAME ([option]) (REGINA)

Returns details about the current platform. This function is basically a wrapper for the Unix command; `uname`. Valid values for *option* are:

[A]

(All) The default. Returns a string with the all following option values. Equivalent to: UNAME('S')UNAME('N')UNAME('R')UNAME('V')UNAME('M').

[S]

(System) The name of the operating system.

[N]

(Node name) The name of the machine.

[R]

(Release) The release of the operating system.

[V]

(Version) The version of the operating system.

[M]

(Machine) The machine's hardware type.

Example running Linux Redhat 6.1 on 'boojum', Athalon K7

UNAME (' S ')	Linux
UNAME (' N ')	boojum
UNAME (' R ')	2.2.12.-20
UNAME (' V ')	#1 Mon Sep 27 10:40:35 EDT 1999
UNAME (' M ')	i686

Example running Windows NT 4.0 on 'VM_NT', Intel Pentium

UNAME (' S ')	WINNT
UNAME (' N ')	VM_NT
UNAME (' R ')	0
UNAME (' V ')	4
UNAME (' M ')	i586

UNIXERROR (errno) (REGINA)

This function returns the string associated with the `errno` error number that *errno* specifies. When

some UNIX interface function returns an error, it really is a reference to an error message which can be obtained through `UNIXERROR`.

This function is just an interface to the `strerror()` function call in UNIX, and the actual error messages might differ with the operating system.

This function is now obsolete, instead you should use:

<code>ERRORTEXT(100 + <i>errno</i>)</code>	
--	--

UPPER(*string*)

(AREXX)

Translate the string to uppercase. The action of this function is equivalent to that of `TRANSLATE(string)`, but it is slightly faster for short strings.

<code>UPPER('One Fine Day')</code>	<code>'ONE FINE DAY'</code>
------------------------------------	-----------------------------

USERID()

(REGINA)

Return the name of the current user. A meaningful name will only be returned on those platforms that support multiple users, otherwise an empty string is returned.

<code>USERID()</code>	<code>'mark' /* Maybe */</code>
-----------------------	---------------------------------

VALUE(*symbol* [, [*value*] [, [*pool*]])

(ANSI)

This function expects as first parameter string *symbol*, which names an existing variable. The result returned from the function is the value of that variable. If *symbol* does not name an existing variable, the default value is returned, and the `NOVALUE` condition is not raised. If *symbol* is not a valid symbol name, and this function is used to access a normal REXX variable, an error occurs. Be aware of the "double-expansion" effect, and quote the first parameter if necessary.

If the optional second parameter is specified, the variable will be set to that value, after the old value has been extracted.

The optional parameter *pool* might be specified to select a particular pool of variables to search for *symbol*. The contents and format of *pool* is implementation dependent. The default is to search in the variables at the current procedural level in REXX. Which *pool*s that are available is implementation dependent, but typically one can set variables in application programs or in the operating system.

Note that if `VALUE()` is used to access variable in pools outside the REXX interpreter, the requirements to format (a valid symbol) will not in general hold. There may be other requirements instead, depending on the implementation and the system. Depending on the validity of the name, the value, or whether the variable can be set or read, the `VALUE()` function can give error messages when accessing variables in pools other than the normal. Consult the implementation and system specific documentation for more information.

If it is used to access compound variables inside the interpreter the tail part of this function can take any expression, even expression that are not normally legal in REXX script source code.

By using this function, it is possible to perform an extra level of interpretation of a variable.

VALUE('FOO')	'bar'
VALUE('FOO', 'new')	'bar'
VALUE('FOO')	'new'
VALUE('USER', 'root', 'SYSTEM')	'guest' /* If SYSTEM exists */
VALUE('USER', , 'SYSTEM')	'root'

VERIFY(string,ref[,option][,start]) (ANSI)

With only the first two parameters, it will return the position of the first character in *string* that is not also a character in the string *ref*. If all characters in *string* are also in *ref*, it will return 0.

If *option* is specified, it can be one of:

[N]

(No match) The result will be the position of the first character in *string* that does exist in *ref*, or zero if all exist in *ref*. This is the default option.

[M]

(Match) Reverse the search, and return the position of the first character in *string* that exists in *ref*. If none exists in *ref*, zero is returned.

If *start* (which must be a positive whole number) is specified, the search will start at that position in *string*. The default value for *start* is 1.

VERIFY('foobar', 'barfo')	'2'
VERIFY('foobar', 'barfo', 'M')	'2'
VERIFY('foobar', 'fob', 'N')	'5'
VERIFY('foobar', 'barf', 'N', 3)	'3'
VERIFY('foobar', 'barf', 'N', 4)	'0'

WORD(string,wordno) (ANSI)

Returns the blank delimited word number *wordno* from the string *string*. If *wordno* (which must be a positive whole number) refers to a non-existing word, then a null string is returned. The result will be stripped of any blanks.

WORD('To be or not to be',3)	'or'
WORD('To be or not to be',4)	'not'
WORD('To be or not to be',8)	' '

WORDINDEX(string,wordno) (ANSI)

Returnsthecharacterpositionofthefirstcharacterofblankdelimitedwordnumber *wordno*in *string*, whichisinterpretedasastringofblankdelimitedwords.If *number*(whichmustbeapositivewhole number)referstoawordthatdoesnotexistin *string*,then 0isreturned.

WORDINDEX('To be or not to be',3)	'7'
WORDINDEX('To be or not to be',4)	'10'
WORDINDEX('To be or not to be',8)	'0'

WORDLENGTH(string,wordno) (ANSI)

Returnsthenumberofcharactersinblankdelimitedwordnumber *number*in *string*.If *number*(which mustbeapositivewholenumber)referstoannon-existentword,then 0isreturned.Trailingorleading blanksdonotcountwhencalculatingthelength.

WORDLENGTH('To be or not to be',3)	'2'
WORDLENGTH('To be or not to be',4)	'3'
WORDLENGTH('To be or not to be',0)	'0'

WORDPOS(phrase,string[,start]) (ANSI)

Returnsthewordnumberin *string*whichindicatesatwhich *phrase*begins,providedthat *phrase*isa subphraseof *string*.Ifnot, 0isreturnedtoindicatethatthephrasewasnotfound.Aphrasediffersfrom asubstringinonesignificantway;aphraseisasetofwords,separatedbyanynumberofblanks.

Forinstance," is a"isasubphraseof" This is a phrase".Noticethedifferentamountof whitespacebetween" is"and" a".

If *start*isspecified,itsetsthewordin *string*atwhichthesearchstarts.Thedefaultvaluefor *start*is 1.

WORDPOS('or not','to be or not to be')	'3'
WORDPOS('not to','to be or not to be')	'4'
WORDPOS('to be','to be or not to be')	'1'
WORDPOS('to be','to be or not to be',3)	'6'

WORDS(string) (ANSI)

Returnsthenumberofblankdelimitedwordsinthe *string*.

WORDS('To be or not to be')	'6'
WORDS('Hello world')	'2'
WORDS('')	'0'

WRITECH(file,string) (AREXX)

Writesthestringargumenttothegivenlogicalfile.The returnedvalueistheactualnumberof characterswritten.

WRITECH('outfile','Testing')	'7'
------------------------------	-----

WRITELN(file,string) (AREXX)

Writesthestringargumenttothegivenlogicalfilewitha"newline"appended.The returnedvalueis theactualnumberofcharacterswritten,includingthe“newline”character(s).

WRITELN('outfile','Testing')	'8' /* Unix */
WRITELN('outfile','Testing')	'9' /* DOS */

XRANGE([start][,end]) (ANSI)

Returnsastringthatconsistsofallthecharactersfrom *start*through *end*,inclusive.The defaultvalue forcharacter *start*is '00'x,whilethedefaultvalueforcharacter *end*is 'ff'x.Withoutany parameters,thewholecharactersetin"alphabetic"orderisreturned.Notethattheactualrepresentation oftheoutputfrom XRANGE() dependsonthecharactersetusedbyyourcomputer.

Ifthevalueof *start*islargerthanthevalueof *end*,theoutputwillwraparoundfrom 'ff'xto '00'x. If *start*or *end*isnotastringcontainingexactlyonecharacter,anerrorisreported.

XRANGE('A','J')	'ABCDEFGH IJ'
XRANGE('FC'x)	'FCFD FEFF'x
XRANGE(,'05'x)	'000102030405'x
XRANGE('FD'x,'04'x)	'FD FEFF0001020304'x

X2B(hexstring) (ANSI)

Translate *hexstring*toabinarystring.Eachhexadecimaldigitsin *hexstring*willbetranslatedtofour binarydigitsintheresult.Therewillbenoblanksintheresult.

X2B(' ')	' '
X2B('466f6f 426172')	'010001100110111101101111010000100110000101110010'
X2B('46 6f 6f')	'010001100110111101101111'

X2C(hexstring) (ANSI)

Returns the (packed) string representation of *hexstring*. The *hexstring* will be converted bytewise, and blanks may optionally be inserted into the *hexstring* between pairs or hexadecimal digits, to divide the number into groups and improve readability. All groups must have an even number of hexadecimal digits, except the first group. If the first group has an odd number of hexadecimal digits, it is padded with an extra leading zero before conversion.

X2C(' ')	' '
X2C('466f6f 426172')	'FooBar'
X2C('46 6f 6f')	'Foo'

X2D(hexstring[,length]) (ANSI)

Returns a whole number that is the decimal representation of *hexstring*. If *length* is specified, then *hexstring* is interpreted as a two's complement hexadecimal number consisting of the *number* rightmost hexadecimal numerals in *hexstring*. If *hexstring* is shorter than *number*, it is padded to the left with <NUL> characters (that is: '00'x).

If *length* is not specified, *hexstring* will always be interpreted as an unsigned number. Else, it is interpreted as a signed number, and the leftmost bit in *hexstring* decides the sign.

X2D('03 24')	'792'
X2D('0310')	'784'
X2D('ffff')	'65535'
X2D('ffff', 5)	'65535'
X2D('ffff', 4)	'-1'
X2D('ff80', 3)	'-128'
X2D('12345', 3)	'837'

13 ImplementationspecificdocumentationforRegina

13.1 DeviationsfromtheStandard

- Forthosebuilt-infunctionswherethelastparametercanbeomitted, Reginaallowsthelastcomma tobepreceded,evenwhenthelastparameteritselfhasbeenomitted.
- Theerrormessagesareslightlyredefinedintwoways.Firstly,someofthehaveaslightlymore definitetext,andsecondly,somenewerrormessageshavebeendefined.
- Theenvironmentsavailablearedescribedinchapter[notyetwritten].
- Parametercalling
- StreamI/O
- Conditions
- Nationalcharactersets
- Blanks
- Stackshavethefollowingextrafunctionality: `DROPBUF()`, `DESBUFF()`and `MAKEBUF()`and `BUFTYPE()`.
- `Random()`
- `Sourceline`
- `Time`
- `CharacterSets`

13.2 InterpreterInternalDebuggingFunctions

`ALLOCATED([option])`

Returnstheamountofdynamicstorageallocated,measuredinbytes.Thisisthememoryallocatedby the `malloc()` call,anddoesnotconcernstackspaceorstaticvariables.

Asparameteritmaytakean *option*,whichisoneofthesinglecharacters:

[A]

Itwillreturnastringthatisthenumberofbytesofdynamicmemorycurrentlyallocatedbythe interpreter.

[C]

Returns a number that is the number of bytes of dynamic memory that is currently in use (i.e. not leaked).

[L]

Returns the number of bytes of dynamic memory that is supposed to have been leaked.

[S]

This is the default value if you do not specify an option. Returns a string that is nicely formatted and contains all the other three options, with labels. The format of this string is:

"Memory: Allocated=XXX, Current=YYY, Leaked=ZZZ".

This function will only be available if the interpreter was compiled with the `TRACEMEM` preprocessor macro defined.

DUMPTREE ()

Prints out the internal parse tree for the `REXX` program currently being executed. This output is not very interesting unless you have good knowledge of the interpreter's internal structures.

DUMPVARS ()

This routine dumps a list of all the variables currently defined. It also gives a lot of information which is rather uninteresting for most users.

LISTLEAKED ()

List out all memory that has leaked from the interpreter. As a return value, the total memory that has been listed is returned. There are several options to this function:

[N]

Do not list anything, just calculate the memory.

[A]

List all memory allocations currently in use, not only that which has been marked as leaked.

[L]

Only list the memory that has been marked as leaked. This is the default option.

TRACEBACK ()

Prints out a traceback. This is the same routine which is called when the interpreter encounters an error. Nice for debugging, but not really useful for any other purposes.

13.3 REXXVMSInterfaceFunctions

F\$CVSI

F\$CVTIME

F\$CVUI

F\$DIRECTORY

F\$ELEMENT

F\$EXTRACT

F\$FAO

F\$FILE_ATTRIBUTES

F\$GETDVI

F\$GETJPI

F\$GETQUI

F\$GETSYI

F\$IDENTIFIER

F\$INTEGER

F\$LENGTH

F\$LOCATE

F\$LOGICAL

F\$MESSAGE

F\$MODE

F\$PARSE

F\$PID

F\$PRIVILEGE

F\$PROCESS

F\$SEARCH

F\$SETPRV

F\$STRING

F\$TIME

F\$TRNLNM

F\$TYPE

F\$USER

Conditions

In this chapter, the REXX concept of "conditions" is described. Conditions allow the programmer to handle abnormal control flow, and enable him to assign special pieces of REXX code to be executed in case of certain incidences.

- In the first section the concept of conditions is explained.
- Then, there is a description of how a standard condition in REXX would work, if it existed.
- In the third section, all the existing conditions in REXX are represented, and the differences compared to the standard condition described in the previous section are listed.
- The fourth section contains a collection of random notes on the conditions in REXX.
- The last section describes differences, extensions and peculiarities in Regina on the subject of conditions, and the list specifies behavior.

14 What are Conditions

In this section, the concept of "conditions" are explained: What they are, how they work, and what they mean in programming.

14.1 What Do We Need Conditions for?

14.1.1 Terminology

First, let's look at the terminology used in this chapter. If you don't get a thorough understanding of these terms, you will probably not understand much of what is said in the rest of this chapter.

[Incident :]

A situation, external or internal to the interpreter, which it is required to respond to in certain pre-defined manners. The interpreter recognizes incidents of several different types. The incident will often have a character of "suddenness", and will also be independent of the normal control flow.

[Event :]

Data structure describing one incident, used as a descriptor to the incident itself.

[Condition :]

Name the REXX concept that is equivalent to the incident.

[Raise a Condition :]

The action of transforming the information about an incident into an event. This is done after the interpreter senses the condition. Also includes deciding whether to ignore or produce an event.

[Handle a Condition :]

The act of executing some pre-defined actions as a response to the event generated when a condition was raised.

[(Condition) Trap:]
 DataStructurecontaininginformationabout how to handle a condition.

[(Trap) State:]
 Part of the condition trap.

[(Condition) Handler:]
 Part of the condition trap, which points to a piece of REXX code which is to be used to handle the condition.

[(Trap) Method:]
 Part of the condition trap, which defines how the condition handler is to be invoked to handle the condition.

[Trigger a Trap:]
 The action of invoking a condition handler by the methods specified by the trap method, in order to handle a condition.

[Trap a Condition:]
 Short of triggering a trap for a particular condition.

[Current Trapped Condition:]
 The condition currently being handled. This is the same as the most recent trapped condition on this or higher procedure level.

[(Pending) Event Queue:]
 Data Structure storing zero or more events in a specific order. There are only one event queue. The event queue contains events of all condition types, which have been raised, but not yet handled.

[Default-Action:]
 The pre-defined default way of handling a condition, taken if the trap state for the condition raised is OFF.

[Delay-Action:]
 The pre-defined default action taken when a condition is raised, and the trap state is DELAY.

15 The Mythical Standard Condition

REXX Language Level 4.00 has six different conditions, and REXX Language Level 5.00 has seven. However, each of these is a special case of a mythical, non-existing, standard condition. In order to better understand the real conditions, we start by explaining how a standard condition works.

In the examples below, we will call our non-existing standard condition MYTH. Note that these examples will not be executable on any REXX implementation.

15.1 Information Regarding Conditions (data structures)

There are mainly five conceptual data structures involved in conditions.

[Event queue.]
 There is one interpreter-wide queue of pending conditions. Raising a condition is identical to adding information about the condition to this queue (FIFO). The order of the queue is the same order in which the conditions are to be handled.

Every entry in the queue of pending conditions contains some information about the event: the line number of the REXX script when the condition was raised, a descriptive text and the

conditiontype.

[Default-Action.]

To each, there exists information about the default-action to take if this condition is raised but the trap is in state `OFF`. This is called the "default-action". The standard default-action is to ignore the condition, while some conditions may abort the execution.

[Delay-Action.]

Each condition will also have delay-action, which tells what to do if the condition is raised when condition trap is in state `DELAY`. The standard delay-action is to queue the condition in the queue of pending conditions, while some conditions may ignore it.

[Condition traps.]

For each condition there is a trap which contains three pieces of status information: the state; the handler; and the method. The state can be `ON`, `OFF` or `DELAY`.

The handler names the `REXX` label in the start of the `REXX` code to handle the event. The method can be either `SIGNAL` or `CALL`, and denotes the method in which the condition is to be handled. If the state is `OFF`, then neither handler nor method is defined.

[Current Trapped Condition.]

This is the most recently handled condition, and is set whenever a trap is triggered. It contains information about method, which condition, and a context-dependent description. In fact, the information in the current trapped condition is the same information that was originally put into the pending event queue.

Note that the event queue is a data structure connected to the interpreter itself. You operate on the same event queue, independent of subroutines, even external ones. On the other hand, the condition traps and the current trapped condition are data structures connected to each single routine. When a new routine is called, it will get its own condition traps and a current trapped condition. For internal routines, the initial values will be the same values as those of the caller. For external routines, the values are the defaults.

The initial value for the event queue is to be empty. The default-action and the delay-action are static information, and will always retain their values during execution. The initial values for the condition traps are that they are all in state `OFF`. The initial value for the current trapped condition is that all information is set to the null string to signalize that no condition is currently being trapped.

15.2 How to Setup a Condition Trap

How do you set the information in a condition trap? You do it with a `SIGNAL` or `CALL` clause, with the `ON` or `OFF` subkeyword. Remember that a condition trap contains three pieces of information? Here are the rules for how to set them:

- To set the trap method, use either `SIGNAL` or `CALL` as keyword.
- To set state to `ON` or `OFF`, use the appropriate subkeyword in the clause. Note that there is no clause or function in `REXX`, capable of setting the state of a trap to `DELAY`.
- To set the condition handler, append the term " `NAME handler`" to the command. Note that this term is only legal if you are setting the state to `ON`; you cannot specify a handler when setting the

stateto OFF.

The trap is said to be "enabled" when the state is either ON or DELAY, and "disabled" when the state is OFF. Note that neither the event queue, nor the current trapped condition can be set explicitly by REXX clauses. They can only be set as a result of incidents, when raising and trapping conditions.

It sounds very theoretical, doesn't it? Look at the following examples, which set the trap MYTH:

```
/* 1 */ SIGNAL ON MYTH NAME TRAP_IT
/* 2 */ SIGNAL OFF MYTH
/* 3 */ CALL ON MYTH NAME MYTH_TRAP
/* 4 */ CALL ON MYTH
/* 5 */ CALL OFF MYTH
```

Line 1 sets stateto ON, methodto SIGNAL and handler to TRAP_IT. Line 2 sets stateto OFF, handler and method become undefined. Line 3 sets stateto ON, methodto CALL, and handler to MYTH_TRAP. Line 4 sets stateto ON, methodto CALL and handler to MYTH (the default). Line 5 sets stateto OFF, handler and method become undefined.

Why should method and handler become undefined when the trap in state OFF? For two reasons: firstly, these values are not used when the trap is in state OFF; and secondly, when you set the trap to state ON, they are redefined. So it really does not matter what they are in state OFF.

What happens to this information when you call a subroutine? All information about traps are inherited by the subroutine, provided that it is an internal routine. External routines do not inherit any information about traps, but use the default values. Note that the inheritance is done by copying, so any changes done in the subroutine (internal or external), will only have effect until the routine returns.

15.3 How to Raise a Condition

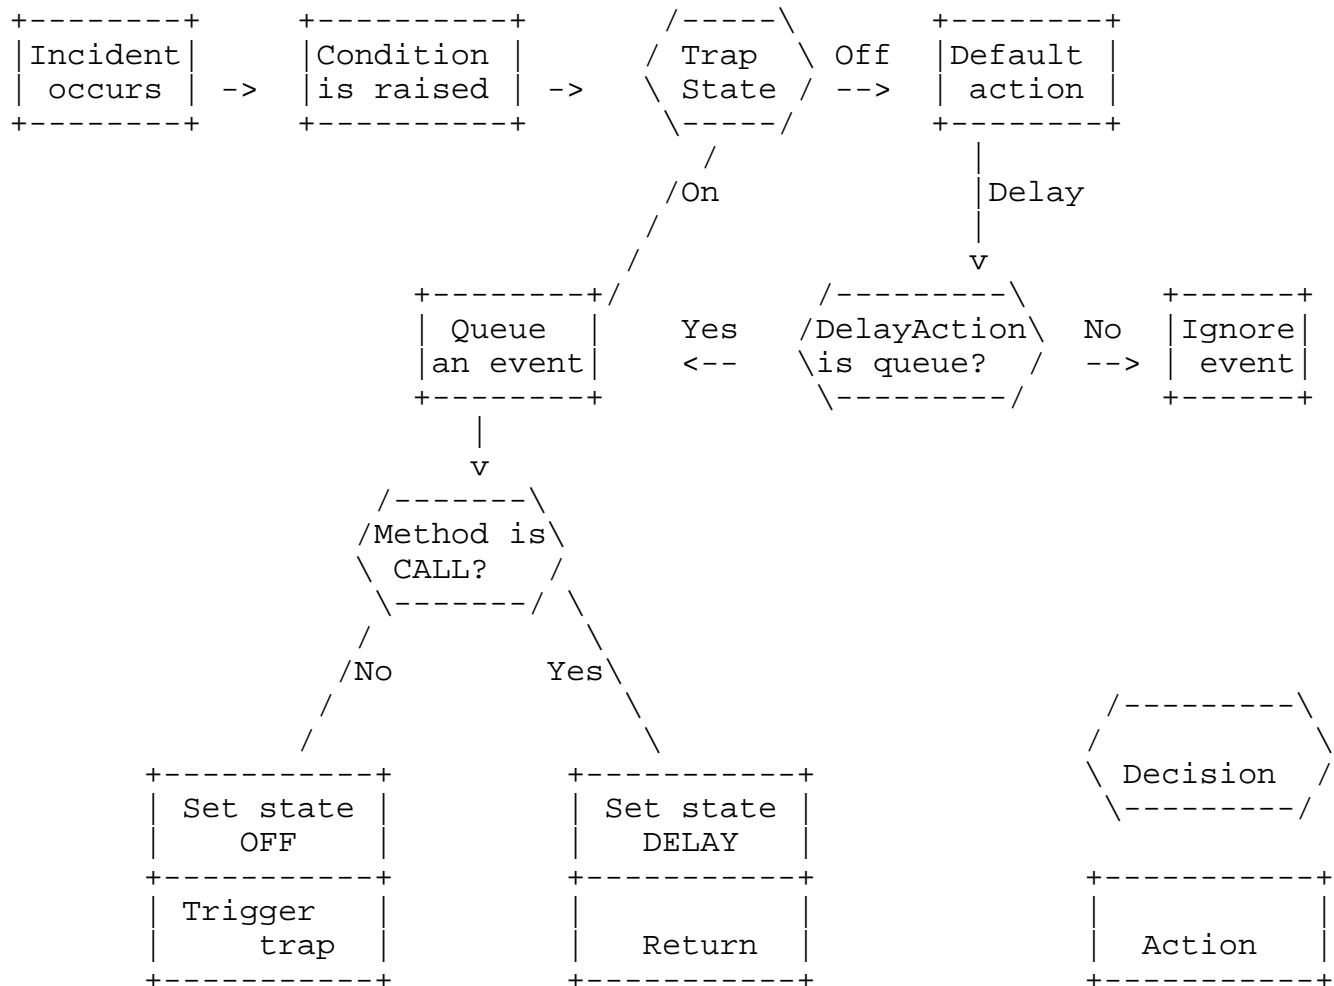
How do you raise a condition? Well, there are really no explicit way in REXX to do that. The conditions are raised when an incident occurs. What sort of situations that is, depends on the context. There are in general three types of incidents, classified by the origin of the event:

- Internal origin. The incident is only dependent on the behavior of the REXX script. The SYNTAX condition is of this type.
- External origin. The REXX script and the interpreter has really no control over when this incident. It happens completely independent of the control of the REXX script or interpreter. The HALT condition is of this type.
- Mixed origin. The incident is of external origin, but the situation that created the incident, was an action by the REXX script or the interpreter. The ERROR condition is of this type: the incident is a command returning error, but it can only occur when the interpreter is executing commands.

For condition trapped by method CALL, standard REXX requires an implementation to at least check for incidents and raise condition at clause boundaries. (But it is allowed to do so elsewhere too;

although the actual triggering must only be performed at clause boundaries.) Consequently, you must be prepared that in some implementations, condition trappable by method CALL might only be raised (and the trap triggered) at clause boundaries, even if they are recurrently trapped by method SIGNAL.

These seven standard conditions will be raised as a result of various situations, read the section describing each one of them for more information.



When an incident occurs and the condition is raised, the interpreter will check the state of the condition trap for that particular condition at the current procedure level.

- If the trap state is OFF, the default-action of the condition is taken immediately. The "standard" default-action is to ignore the condition.
- If the trap state is DELAY, the action will depend on the delay-action of that condition. The standard delay-action is to ignore, then nothing further is done. If the delay-action is to queue, the interpreter

continues as if the state was `ON`.

- If the state of the trap is `ON`, an event is generated which describes the incident, and it is queued in the pending event queue. The further action will depend on the method of trapping.
- If the method is `CALL`, the state of the trap will be set to `DELAY`. Then the normal execution is resumed. The idea is that the interpreter will check the event queue later (at a clause boundary), and trigger the appropriate trap, if it finds any events in the event queue.
- Else, if method of trapping is `SIGNAL`, then the action taken is this: First set the trap to state `OFF`, then terminate clause the interpreter was executing at this procedure level. Then it explicitly triggers the condition trap.

This process has been shown in the figure above. It shows how an incident makes the interpreter raise a condition, and that the state of the condition trap determines what to do next. The possible outcomes of this process are: to take the default-action; to ignore if delay-action is not to queue; to just queue and the continue execution; or to queue and trigger the trap.

15.4 How to Trigger a Condition Trap

What are the situations where a condition trap might be triggered? It depends on the method currently set in the condition trap.

If the method is `SIGNAL`, then the interpreter will explicitly trigger the relevant trap when it has raised the condition after having sensed the incident. Note that only the particular trap in question will be triggered in this case; other traps will not be triggered, even if the pending event queue is non-empty.

In addition, the interpreter will at each clause boundary check for any pending events in the event queue. If the queue is non-empty, the interpreter will not immediately execute the next normal statement, but it will handle the condition(s) first. This procedure is repeated until there are no more events queued. Only then will the interpreter advance to execute the next normal statement.

Note that the `REXX` standard does not require the pending events to be handled in any particular order, although the model shown in this documentation it will be in the order in which the conditions were raised. Consequently, if one clause generates several events that raise conditions before or at the next clause boundary, and these conditions are trapped by method `CALL`. Then, the order on which the various traps are triggered is implementations-dependent. But the order in which the different instances of the same condition is handled, is the same as the order of the condition indicator queue.

15.5 Trapping by Method `SIGNAL`

Assume that a condition is being trapped by method `SIGNAL`, that the state is `ON` and the handler is `MYTH_TRAP`. The following `REXX` clause will set up the trap correctly:

```
SIGNAL ON MYTH NAME MYTH_TRAP
```

Now, suppose the `MYTH` incident occurs. The interpreter will sense it, queue an event, set the trap state to `OFF` and then explicitly trigger the trap, since the method is `SIGNAL`. What happens when the trap is

triggered?

- It collects the first event from the queue of pending events. The information is removed from the queue.
- The current trapped condition is set to the information removed from the pending event queue.
- Then, the interpreter simulates a `SIGNAL` clause to the label named by the trap handler of the trap for the condition in question.
- As all `SIGNAL` clauses, this will have the side-effects of setting the `SIGL` special variable, and terminating all active loops at the current procedure level.

That's it for method `SIGNAL`. If you want to continue trapping condition `MYTH`, you have to execute a new `SIGNAL ON MYTH` clause to set the state of the trap to `ON`. But no matter how quick you reset the trap, you will always have a short period where it is in state `OFF`. This means that you cannot in general use the method `SIGNAL` if you really want to be sure that you don't lose any `MYTH` events, unless you have some control over when `MYTH` condition may arise.

Also note that since the statement being executed is terminated; all active loops on the current procedure level are terminated; and the only indication where the error occurred is the line number (the line may contain several clauses), then it is in general impossible to pick up the normal execution after a condition trapped by `SIGNAL`. Therefore, this method is best suited for a "graceful death" type of traps. If the trap is triggered, you want to terminate what you were doing, and pick up the execution at an earlier stage, e.g. the previous procedure level.

15.6 Trapping by Method `CALL`

Assume that the condition `MYTH` is being trapped by method `CALL`, that the state is `ON` and the handler is `MYTH_HANDLER`.

The following `REXX` clause will set up the trap correctly:

```
CALL ON MYTH NAME MYTH_HANDLER
```

Now, suppose that the `MYTH` incident occurs. When the interpreter senses that, it will raise the `MYTH` condition. Since the trap state is `ON` and the trap method is `CALL`, it will create an event and queue it in the pending event queue and set the trap state to `DELAY`. Then it continues the normal execution. The trap is not triggered before the interpreter encounters the next clause boundary. What happens then?

- At the every clause boundaries, the interpreter checks for any pending events in the event queue. If one is found, it is handled. This action is done repeatedly, until the event queue is empty.
- It will simulate a normal function call to the label named by the trap handler. As with any `CALL` clause, this will set the special variable `SIGL` to the line off from which the call was made. This is done prior to the call. Note that this is the current line at the time when the condition was raised, not when it was triggered. All other actions normally performed when calling a subroutines are done.

Notethatthearguments to the subroutine are set to empty.

- However, just before execution of the routine starts, it will remove the first event in the pending event queue, the information is instead put into the current trapped condition. Notethatthe current trapped condition is information that is saved across subroutine calls. It is set **after** the condition handler is called, and will be local to the condition handler (and functions called by the condition handler). To the "caller" (i.e. the procedure level active when the trap was triggered), it will seem as if the current trapped condition was never changed.
- Then the condition handler finishes execution, and returns by executing the `RETURN` clause. Any expression given as argument to `RETURN` will be ignored, i.e. the special variable `RESULT` will not be set upon return from a condition handler.
- At the return from the condition handler, the current trapped condition and the setup of all traps are restored, as with a normal return from subroutine. As a special case, the state of the trap just triggered, will not be put back into `DELAY` state, but is set to state `ON`.
- Afterwards (and before the next normal clause), the interpreter will again check for more events in the event queue, and it will not continue on the `REXX` script before the queue is empty.

During the triggering of a trap by method `CALL` at a clause boundary, the state of the trap is not normally changed, it will continue to be `DELAY`, as was set when the condition was raised. It will continue to be in state `DELAY` until return from the condition handler, at which the state of the trap in the caller will be changed to `ON`. If, during the execution of the condition trap, the state of the condition being trapped is set, that change will only last until the return from the condition handler.

Since new conditions are generally delayed when a condition handler is executing, new conditions are queued up for execution. If the trap state is changed to `ON`, the pending event queue will be processed as named at the next clause boundary. If the state is changed to `OFF`, the default action of the conditions will be taken at the next clause boundary.

15.7 The Current Trapped Condition

The interpreter maintains a data structure called the current trapped condition. It contains information relating the most recent condition trapped on this or higher procedure level. The current trapped condition is normally inherited by subroutines and functions, and restored after return from these.

- When trapped by method `SIGNAL` the current trapped condition of the current procedure level is set to information describing the condition trapped.
- When trapped by method `CALL`, the current trapped condition at the procedure level which the trap occurred at, is not changed. Instead, the current trapped condition in the condition handler is set to information describing the condition.

The information stored in the current trapped condition can be retrieved by the built-in function `CONDITION ()`. The syntax format of this function is:

`CONDITION(option)`

where *option* is an option string of which only the first character matters. The valid options are: `Condition`, `name`, `Description`, `Instruction` and `State`. These will return: the name of the current trapped condition; the descriptive text; the method; and the current state of the condition, respectively. The default *option* is `Instruction`. See the documentation on the built-in functions. See also the description of each condition below.

Note that the `State` option does not return the state at the time when the condition was raised or the trap was triggered. It returns the current state of the trap, and may change during execution. The other information in the current trapped condition may only change when a new condition is trapped at return from subroutines.

16 The Real Conditions

We have now described how the standard condition and condition trap works in REXX. Let's look at these even conditions defined which do exist. Note that none of these behave exactly as the standard condition.

16.1 The SYNTAX condition

The SYNTAX condition is of internal origin, and is raised when any syntax or runtime error is discovered by the REXX interpreter. It might be any of the situations that would normally lead to the abortion of the program and the report of a REXX error message, except error message number 4 (*Program interrupted*), which is handled by the HALT condition.

There are several differences between this condition and the standard condition:

- It is not possible to trap this condition with the method `CALL`, only method `SIGNAL`. The reason for this is partly that method `CALL` tries to continue execution until next boundary before retriggering the trap. That might not be possible with syntax or runtime errors.
- When this condition is trapped, the special variable `RC` is set to the REXX error number of the syntax or runtime error that caused the condition. This is done just before the setting of the special variable `SIGL`.
- The default action of this condition if the trap state is `OFF`, is to abort the program with a traceback and error message.
- There is no delay-action for condition SYNTAX, since it cannot be trapped by method `CALL`, and consequently never can get into state `DELAY`.

The descriptive text returned by `CONDITION()` when called with the `Description` option for condition SYNTAX, is implementation dependent, and may also be a null string. Consult the implementation-specific documentation for more information.

16.2 The HALT condition

The HALT condition of external origin, which is raised as a result of an action from the user, normally a combination of keys which tries to abort the program. Which combination of keys will vary between operating systems. Some systems might also simulate this event by other means than key combinations. Consult system for more information.

The differences between HALT and the standard condition are:

- The default-action for the HALT condition is to abort execution, although a REXX runtime error number 4 (*Program interrupted*) had been reported. But not that SYNTAX will never be raised if HALT is not trapped.
- The delay-action of this condition is to ignore, not queue.

The standard allows the interpreter to limit the search for situations that would set the HALT condition, to clause boundaries. As a result, the response time from pressing the key combination to actually raising the condition or triggering the trap may vary, even if HALT is trapped by method SIGNAL. If a clause for some reason has blocked execution, and never finish, you may not be able to break the program.

The descriptive text returned by `CONDITION()` when called with the `Description` option for condition HALT, is implementation dependent, and may also be a null string. In general, it will describe the way in which the interpreter was attempted halted, in particular if there are more than one way to do raise a HALT condition. Consult the implementation documentation for more information.

16.3 The ERROR condition

The ERROR is a condition of mixed origin, it is raised when a command returns a return value which indicates error during execution. Often, commands return a numeric value, and a particular value is considered to mean success. Then, other values might raise the ERROR condition.

Differences between ERROR and the standard condition:

- The delay action of ERROR is to ignore, not to queue.
- The special variable `RC` is always set before this condition is raised. So even if it is trapped by method SIGNAL, you can rely on `RC` to be set to the return value of the command.

Unfortunately, there is no universal standard on return values. As stated, they are often numeric, but some operating systems use non-numeric return values. For those which do use numeric values, there are no standard telling which values and ranges are considered errors and which are considered success. In fact, the interpretation of the value might differ between commands within the same operating system.

Therefore, it is up to the REXX implementation to define which values and ranges that are considered errors. You must expect that this information can differ between implementations as well as between different environments within one implementation.

The descriptive text returned by `CONDITION()` when called with the `Description` option for condition `ERROR`, is the command which caused the error. Note that this is the command as the environment saw it, not as it was entered in the `REXX` script source code.

16.4 The `FAILURE` condition

The `FAILURE` is a condition of mixed origin, it is raised when a command returns a return value which indicates failure during execution, abnormal termination, or when it was impossible to execute a command. It is a subset of the `ERROR` condition, and if it is in state `OFF`, then the `ERROR` condition will be raised instead. But note that an implementation is free to consider all return codes from commands as `ERRORS`, and none as `FAILURES`. In that case, the only situation where a `FAILURE` would occur, is when it is impossible to execute a command.

Differences between `FAILURE` and the standard condition:

- The delay action of `FAILURE` is to ignore, not to queue.
- The special variable `RC` is always set before this condition is raised. So even if it is trapped by method `SIGNAL`, you can rely on `RC` to be set to the return value of the command, or the return code that signalizes that the command was impossible to execute.

As for `ERROR`, there is no standard that defines which return values are failures and which are errors. Consult the system and implementation independent documentation for more information.

The descriptive text returned by `CONDITION()` when called with the `Description` option for condition `FAILURE`, is the command which caused the error. Note that this is the command as the environment saw it, not as it was entered in the `REXX` script source code.

16.5 The `NOVALUE` condition

The `NOVALUE` condition is of internal origin. It is raised in some circumstances if the value of an unset symbol (which is not a constant symbol) is requested. Normally, this would return the default value of the symbol. It is considered bad programming practice not to initialize variables, and setting the `NOVALUE` condition is one method of finding the parts of your program that use this programming practice.

Note however, there are only three instances where this condition may be raised: that is when the value of an unset (non-constant) symbol is used requested: in an expression; after the `VAR` subkeyword in a `PARSE` clause; and as an indirect reference in either a template, a `DROP` or a `PROCEDURE` clause. In particular, this condition is not raised if the `VALUE()` or `SYMBOL()` built-in functions refer to an unset symbol.

Differences between `NOVALUE` and the standard condition are:

- It may only be trapped by method `SIGNAL`, never method `CALL`. This requirement might seem somewhat strange, but the idea is that since an implementation is only forced to check for condition trapped by method `CALL` at clause boundaries, incidences that may occur at any point within clauses (like `NOVALUE`) can only be trapped by method `SIGNAL`. (However, condition `NOTREADY` can occur within a clause, and may be trapped by method `CALL` so this does not seem to be absolute consistent.)
- There is not delay-action for condition `NOVALUE`, since it cannot be trapped by method `CALL`, and consequently never can get into state `DELAY`.

The descriptive text returned by calling `CONDITION()` with the `Description` option, is the derived (i.e. tail has been substituted if possible) name of the variable that caused the condition to be raised.

16.6 The `NOTREADY` condition

The condition `NOTREADY` is a condition of mixed origin. It is raised as a result of problems with stream I/O. Exactly what causes it, may vary between implementations, but some of the more probable causes are: waiting for more I/O on transient streams; access to streams not allowed; I/O operation would block if attempted; etc. See the chapter; **Stream Input and Output** for more information.

Differences between `NOTREADY` and the standard condition are:

- It will be ignored rather than queued if condition trap is in state `DELAY`.
- This condition differs from the rest in that it can be raised during execution of a clause, but can still be trapped by method `CALL`.

The descriptive text returned by `CONDITION()` when called with the `Description` option for condition `NOTREADY`, is the name of the stream which caused the problem. This is probably the same string that you used as the first parameter to the function that operates on stream I/O. For the default streams (default input and output stream), the string returned by `CONDITION()` will be null strings.

Note that if the `NOTREADY` trap is in state `DELAY`, then all I/O for files which has tried to raise `NOTREADY` within the current clause will be simulated as if operation had succeeded.

16.7 The `LOSTDIGITS` condition

The condition `LOSTDIGITS` was introduced in Language Level 5.00. It is raised as a result of any arithmetic operation which results in the loss of any digits. i.e. If the number of significant digits in the result of an arithmetic operation would exceed the currently defined number of digits via `NUMERIC DIGITS`, then the `LOSTDIGITS` condition is raised.

Differences between `LOSTDIGITS` and the standard condition are:

- It may only be trapped by method `SIGNAL`, never method `CALL`.

- There is not delay-action for condition `NOVALUE`, since it cannot be trapped by method `CALL`, and consequently never can get into state `DELAY`.

The descriptive text returned by `CONDITION()` when called with the `Description` option for condition `NOTREADY`, is the name of the stream which caused the problem. This is probably the same string that you used as the first parameter to the function that operates on stream I/O. For the default streams (default input and output stream), the string returned by `CONDITION()` will be null strings.

17 Further Notes on Conditions

17.1 Conditions under Language Level 3.50

The concept of conditions was very much expanded from `REXX` language level 3.50 to level 4.00. Many of the central features in conditions are new in level 4.00, these include:

- The `CALL` method is new, previously only the `SIGNAL` method was available, which made it rather difficult to resume execution after a problem. As a part of this, the `DELAY` state has been added too.
- The condition `NOTREADY` has been added, to allow better control over problems involving stream I/O.
- The built-in function `CONDITION()` has been added, to allow extraction of information about the current trapped condition.

17.2 Pitfalls when Using Condition Traps

There are several pitfalls when using conditions:

- Remember that some information is saved across the functions. Both the current trapped condition and the settings of the traps. Consequently, you cannot set a trap in a procedure level from a lower level. (I.e. calling a subroutine to set a trap will not work.)
- Remember that `SIGL` is set when trapped by method `CALL`. This means that whenever a condition might be trapped by `CALL`, the `SIGL` will be set to a new value. Consequently, never trust the contents of the `SIGL` variable for more than one clause at a time. This is very frustrating, but at least it will not happen often. When it does happen, though, you will probably have a hard time debugging it.
- Also remember that if you use the `PROCEDURE` clause in a condition handler called by method `CALL`, remember to `EXPOSE` the special variables `SIGL` if you want to use it inside the condition handler. Else it will be shadowed by the `PROCEDURE`.

17.3 The Correctness of this Description

In this description of conditions in `REXX`, I have gone further in the description of how conditions work, their internal data structures, the order in which things are executed etc., than the standard does. I have tried to interpret these set of distinct statements that is the documentation on condition, and design

a complete and consistent system describing how such conditions work. I have done this to try to clarify an area of REXX which at first glance is very difficult and sometimes non-intuitive.

I hope that the liberties I have taken have helped describe conditions in REXX. I do not feel that the adding of details that I have done in any way change how conditions work, but at least I owe the reader to list which concepts that are genuine REXX, and which have been filled in by me to make the picture more complete. These are not a part of the standard REXX.

- REXX does not have anything called a standard condition. There just "are" a set of conditions having different attributes and values. Sometimes there are default values to some of the attributes, but still there are no default conditions.
- The terms "event" and "incident" are not used. Instead the term "condition" is somewhat overloaded to mean several things, depending on the situation. I have found it advantageous to use different terms for each of these concepts.
- Standard REXX does not have a condition queue, although a structure of such a kind is needed to handle the queuing of pending conditions when the trap state is DELAY.
- The values default-action and delay-action are really non-existing in the Standard REXX documentation. I made them up to make the system more easy to explain.
- The two-step process of first raising the flag, and then (possibly at a later stage) triggering the trap, is not really a REXX concept. Originally, REXX seems to allow implementations to select certain places of the interpreter where events are sought for. All standard conditions that can be called by method CALL, can be implemented by checking only at clause boundaries.
- Consequently, a REXX implementation can choose to trigger the trap immediately after a condition is raised (since conditions are only raised immediately before the trap would trigger anyway). This is also the common way used in language level 3.50, when only method SIGNAL was implemented.
- Unfortunately, the introduction of the state DELAY forces the interpreter to keep a queue of pending conditions, so there is nothing to gain in insisting that raising should happen immediately before triggering. And the picture is even more muddled when the NOTREADY condition is introduced. Since it explicitly allows raising of condition to be done during the clause, even though the triggering of the trap must happen (if method is CALL) at the end of the clause.

I really hope that these changes have made the concept of conditions easier to understand, no harder. Please feel free to flame me for any of these which you don't think is representative for REXX.

18 Conditions in Regina

Here comes documentation that is specific for the Regina implementation of REXX.

18.1 How to Raise the HALT condition

The implementation connects the HALT condition to an external event, which might be the pressing of

certain key combination. The common conventions of the operating system will dictate what that combination of key strokes is.

Below is a list, which describes how to invoke an event that will raise the `HALT` condition under various operating systems which `Regina` runs under.

- Under various variants of the **Unix** operating system, the `HALT` event is connected to the signal "interrupt" (`SIGINT`). Often this signal is bound to special key strokes. Depending on your version of Unix, this might be `<ctrl>-<c>` (mostly BSD-variants) or the `` key (mostly System V). It is also possible to send this signal from the command line, in general using the program `kill(1)`; or from program, in general using the call `signal(3)`. Refer to your Unix documentation for more information.
- Under **VAX/VMS**, the key sequence `<ctrl>-<c>` is used to raise the `HALT` condition in the interpreter.

19 Possible Future extensions

- Here is a list of possible future extensions to `REXX` which has not been implemented into `Regina`. Some of these exist in other implementations of `REXX`, and some of them are just suggestions or ideas thrown around by various people.
- Another extension could have been included, but have been left out so far. It is the delay-action, which in standard `REXX` can be either to ignore or to queue. There is at least one other action that makes sense: to replace. That is, when a trap is in state `DELAY`, and a new condition has been raised, the pending queue is emptied, before the new condition is queued. That way, the new condition will effectively replace any conditions already in the queue.
- If there are several new conditions raised while the condition handler is executing (and the trap state is `DELAY`), only the very last of them is remembered.
- It should be possible to set the state for a trap to `DELAY`, so that any new instances of the condition is handled by the delay-action. As a special case, the `SYNTAX` condition trap might not be set in state `DELAY`.

StreamInputandOutput

And the stream thereof shall be turned into pitch

Isaiah 33:21

*For every one that asketh receivedth;
and he that seeketh findth;
and to him that knocketh it shall be opened.*

Matthew 7:8

This chapter treats the topic of input from and output to streams using the built-in functions. An overview of the other parts of the input/output (I/O) system is also given but not discussed in detail. At the end of the chapter there are sections containing implementation-specific information for this topic.

20 Background and Historical Remarks

Stream I/O is a problem area for languages like REXX. They try to maintain compatibility for all platforms (i.e. to be non-system-specific), but the basic I/O capabilities differ between systems, so the simplest way to achieve compatibility is to include only a minimal, common subset of the functionality of all platforms. With respect to the functionality of the interface to their surrounding environment, non-system-specific script languages like REXX are inherently inferior to system-specific script languages which are hardwired to particular operating systems and can benefit from all their features.

Although REXX formally has its own I/O constructs, it is common for some platforms that most or all of the I/O is performed as operating system commands rather than in REXX. This is how it was originally done under VM/CMS, which was one of the earliest implementations and which did not support REXX's I/O constructs. There, the EXECIO program and the stack (among other methods) are used to transfer data to and from a REXX program.

Later, the built-in functions for stream I/O gained territory, but lots of implementations still rely on special purpose programs for doing I/O. The general recommendation to REXX programmers is to use the built-in functions instead of special purpose programs whenever possible; that is the only way to make compatible programs.

21 REXX's Notion of a Stream

REXX regards a stream as a sequence of characters, conceptually equivalent to what a user might type at the keyboard. Note that a stream is not generally equivalent to a file. [MCGH:DICT] defines a file as "a collection of related records treated as a unit," while [OX:CDICT] defines it as "Information held on backing store [...] in order (a) to enable it to persist beyond the time of execution of a single job and/or (b) to overcome space limitations in main memory." A stream is defined by [OX:CDICT] as "a flow of data characterized by relative long duration and constant rate."

Thus, a file has a flavor of persistency, while a stream has a flavor of sequence and momentariness. For a stream, data read earlier may already have been lost, and the data not yet read may not be currently

defined; for instance the input type data keyboard or the output of a program. Even though much of the REXX literature uses these two terms interchangeably (and after all, there is some overlap), you should bear in mind that there is a difference between them.

In this documentation, the term "file" means "a collection of persistent data on secondary storage, to which random access and multiple retrieval are allowed." The term "stream" means a sequential flow of data from a file or from a sequential device like a terminal, tape, or the output of a program. The term stream is also used in its strict REXX meaning: a handle to/from which a flow of data can be written/read.

22 Short Crash-Course

REXXI/O is very simple, and this short crash course is probably all you need in a first-time reading of this chapter. But note that that, we need to jump a bit ahead in this section.

To read a line from a stream, use the `LINEIN()` built-in function, which returns the data read. To write a stream, use the `LINEOUT()` built-in function, and supply the data to be written as the second parameter. For both operations, give the name of the stream as the first parameter. Some small examples:

```
contents = linein( 'myfile.txt' )
call lineout 'yourfile.txt', 'Data to be written'
```

The first of these reads a line from the stream `myfile.txt`, while the second writes a line to the stream `yourfile.txt`. Both these calls operate on lines and they use as system specific end-of-line marker as a delimiter between lines. The marker is tagged on at the end of any data written out, and stripped off any data read.

Opening a stream in REXX is generally done automatically, so you can generally ignore that in your programs. Another useful method is repositioning to a particular line:

```
call linein 'myfile.txt', 12, 0
call lineout 'yourfile.txt', , 13
```

Where the first of these sets the current read position to the start of line 12 of the stream; the second sets the current write position to the start of line 13. Note that the second parameter is empty, that means no data is to be written. Also note that the current read and write positions are two independent entities; setting one does not affect the other.

The built-in functions `CHARIN()` and `CHAROUT()` are similar to the ones just described, except that they are character-oriented, i.e. the end-of-line delimiter is not treated as a special character.

Examples of use are:

```
say charin( 'myfile.txt', 10 )
call charout 'logfile', 'some data'
```

Here, the first example reads 10 characters, starting at the current input position, while the second writes the eleven characters of "somedata" to the file, without an end-of-file marker afterwards.

It is possible to reposition character-wise too, some examples are:

```
call charin 'myfile', , 8
call charout 'foofile', , 10
```

These two clauses reposition the current read and write positions of the named file to the 8th and 10th characters, respectively.

23 Naming Streams

Unlike most programming languages, REXX does not use file handles; the name of the stream is also in general the handle (although some implementations add an extra level of indirection). You must supply the name to all I/O functions operating on a stream. However, internally, the REXX interpreter is likely to use the native file pointers of the operating system, in order to improve speed. The names specified can generally be the name of an operating system file, a device name, or a special stream name supported by your implementation.

The format of the stream name is very dependent upon your operating system. For portability concerns, you should try not to specify it as a literal string in each I/O call, but set a variable to the stream name, and use that variable when calling I/O functions. This reduces the number of places you need to make changes if you need to port the program to another system. Unfortunately, this approach increases the need for PROCEDURE EXPOSE, since the variable containing the file name must be available to all routines using file I/O for that particular file, and all their non-common ancestors.

Example: Specifying filenames

The following code illustrates a portability problem related to the naming of streams. The variable `filename` is set to the name of the stream operated on in the function call.

```
filename = '/tmp/MyFile.Txt'
say ' first line is' linein( filename )
say 'second line is' linein( filename )
say ' third line is' linein( filename )
```

Suppose this script, which looks like it is written for Unix, is moved to a VMS machine. Then, the stream name might be something like `SYS$TEMP:MYFILE.TXT`, but you only need to change the script at one particular point: the assignment to the variable `filename`; as opposed to three places if the stream name is hard-coded in each of the three calls to `LINEIN()`.

If the stream name is omitted from the built-in I/O functions, a default stream is used: input functions use the default input stream, while output functions use the default output stream. These are implicit references to the default input and output streams, but unfortunately, there is no standard way to explicitly refer to these two streams. And consequently, there is no standard way to refer to the default input or output stream in the built-in function `STREAM()`.

However, most implementations allow you to access the default stream explicitly through a name, maybe the null string or something like `stdin` and `stdout`. However, you must refer to the implementation-specific documentation for information about this.

Also note that standard REXX does not support the concept of a default error stream. On operating systems supporting this, it can probably be accessed through a special name; see system-specific information. The same applies for other special streams.

Sometimes the term "default input stream" is called "standard input stream," "default input devices," "standard input," or just "stdin."

The use of stream names instead of stream descriptors or handles is deeply rooted in the REXX philosophy: Data structures are text strings carrying information, rather than opaque data blocks in internal, binary format. This opens for some intriguing possibilities. Under some operating systems, a file can be referred to by many names. For instance, under Unix, a file can be referred to as `foobar`, `./foobar` and `../foobar`. All which name the same file, although a REXX interpreter may be likely to interpret them as three different streams, because the names themselves differ. On the other hand, nothing prevents an interpreter from discovering that these are names for the same stream, and treat them as equivalent (except concerns for processing time). Under Unix, the problem is not just confined to the use of `./` in filenames, hard-links and soft-links can produce similar effects, too.

Example: Internal file handles

Suppose you start reading from a stream, which is connected to a file called `foo`. You read the first line of `foo`, then you issue a command, in order to rename `foo` to `bar`. Then, you try to read the next line from `foo`. The REXX program for doing this under Unix looks something like:

```
signal on notready
line1 = linein( 'foo' )
'mv foo bar'
line2 = linein( 'foo' )
```

Theoretically, the file `foo` does not exist during the second call, so the second read should raise the NOTREADY condition. However, a REXX interpreter is likely to have opened the stream already, so it is performing the reading on the file descriptor of the open file. It is probably not going to check whether the file exists before each I/O operation (that would require a lot of extra checking). Under most operating systems, renaming a file will not invalidate existing file descriptors. Consequently, the interpreter is likely to continue to read from the original `foo` file, even though it has changed.

Example: Unix temporary files

On some systems, you can delete a file, and still read from and write to the stream connected to that file. This technique is shown in the following Unix-specific code:


```

tmpfile = '/tmp/myfile'
call lineout tmpfile, ''
call lineout tmpfile,, 1
'rm' tmpfile
call lineout tmpfile, 'This is the first line'

```

Under Unix, this technique is often used to create temporary files; you are guaranteed that the file will be deleted on closing, no matter how your program terminates. Unix deletes a file whenever there are no more references to it. Whether the reference is from the file system or from an open descriptor in a user process is irrelevant. After the `rm` command, the only reference to the file is from the REXX interpreter. Whenever it terminates, the file is deleted--since there are no more references to it.

Example: Files in different directories

Here is yet another example of how using the file name directly in the stream I/O functions may give strange effects. Suppose you are using a system that has hierarchical directories, and you have a function `CHDIR ()` which sets a current directory; then consider the following code:

```

call chdir '../dir1'
call lineout 'foobar', 'written to foobar while in dir1'
call chdir '../dir2'
call lineout 'foobar', 'written to foobar while in dir2'

```

Since the file is implicitly opened while you are in the directory `dir1`, the file `foobar` refers to a file located there. However, after changing the directory to `dir2`, it may seem logical that the second call to `LINEOUT ()` operates on a file in `dir2`, but that may not be the case. Considering that these clauses may come a great number of lines apart, that REXX has no standard way of closing files, and that REXX only have one file table (i.e. open files are not local to subroutines); this may open for a significant astonishment in complex REXX scripts.

Whether an implementation treats `../foo` and `./foo` as different streams is system-dependent; that applies to the effects of renaming or deleting the file while reading or writing, too. See your interpreter's system-specific documentation.

Most of the effects shown in the examples above are due to insufficient isolation between the file name of the operating system and the file handle in the REXX program. Whenever a file can be explicitly opened and bound to a file handle, you should do that in order to decrease the possibilities for strange side effects.

Interpreters that allow this method generally have an `OPEN ()` function that takes the name of the files to open as a parameter, and returns a string that uniquely identifies that open file within the current context; e.g. an index into a table of open files. Later, this index can be used instead of the file name.

Some implementations allow only this indirect naming scheme, while others may allow a mix between direct and indirect naming. The latter is likely to create some problems, since some strings are likely to be both valid direct and indirect file IDs.

24 Persistent and Transient Streams

REXX knows two different types of streams: persistent and transient. They differ conceptually in the way they can be operated, which is dictated by the way they are stored. But there is no difference in the data you can read from or write to them (i.e. both can be used for character- or line-wise data), and both are read and written using the same functions.

[Persistent streams]

(often referred to just as "files") are conceptually stored on permanent storage in the computer (e.g. a disk), as an array of characters. Random access to and repeated retrieval of any part of the stream are allowed for persistent streams. Typical examples of persistent streams are normal operating system files.

[Transient streams]

are typically not available for random access or repeated retrieval, either because it is not stored permanently, but read as a sequence of data that is generated on the fly; or because they are available from a sequential storage (e.g. magnetic tape) where random access is difficult or impossible. Typical examples of transient streams are devices like keyboards, printers, communication interfaces, pipelines, etc.

REXX does not allow any repositioning on transient streams; such operations are not conceptually meaningful; a transient stream must be treated sequentially. It is possible to treat a persistent stream as a transient stream, but not vice versa. Thus, some implementations may allow you to open a persistent stream as transient. This may be useful for files to which you have only append access, i.e. writes can only be performed at the end of file. Whether you can open a stream in a particular mode, or change the mode of a stream already open depends on your implementation.

Example: Determining stream type

Unfortunately, there is no standard way to determine whether a given file is persistent or transient. You may try to reposition for the file, and you can assume that the file is persistent if the repositioning succeeded, like in the following code:

```
streamtype: procedure
    signal on notready
    call linein arg(1), 1, 0
    return 'persistent'          /* unless file is empty */
notready:
    return 'transient'
```

Although the idea in this code is correct, there are unfortunately a few problems. First, the `NOTREADY` condition can be raised by other things than trying to reposition a transient stream; e.g. by any repositioning of the current read position in an empty file, if you have write access only, etc. Second, your implementation may not have `NOTREADY`, or it may not use it for this situation.

The best method is to use a `STREAM()` function, if one is available. Unfortunately, that is not very compatible, since no standard stream commands are defined.

25 Opening a Stream

In most programming languages, opening a file is the process of binding a file (given by a filename) to an internal handle. REXX is a bit special, since conceptually, it does not use stream handles, just stream names. Therefore, the stream name is itself also the stream handle, and the process of opening streams becomes apparently redundant. However, not that a number of implementations allow explicit opening, and some even require it.

REXX may open streams "on demand" when they are used for the first time. However, this behavior is not defined in TRL, which says the act of opening the stream is not a part of REXX [TRL2]. This might be interpreted as open-on-demand or that some system-specific program must be executed to open a stream.

Although an open-on-demand feature is very practical, there are situations where you need to open streams in particular modes. Thus, most systems have facilities for explicitly opening a file. Some REXX interpreters may require you to perform some implementation-specific operation before accessing streams, but most are likely to just open them the first time they are referred to in an I/O operation.

There are two main approaches to explicit opening of streams. The first uses a non-standard built-in function normally called `OPEN()`, which generally takes the name of the file to open as the first parameter, and often the mode as the second parameter. The second approach is similar, but uses the standard built-in function `STREAM()` with a `Command` option.

Example: Not closing files

Since there are no open or close operation, a REXX interpreter never knows when to close a stream, unless explicitly told so. It can never predict when a particular stream is to be used next, so it has to keep the current read and write positions in case the stream is to be used again. Therefore, you should always close the streams when you are finished using them. Failure to do so, will fill the interpreter with data about unneeded streams, and more serious, it may fill the file table of your processor system. As a rule, any REXX script that uses more than a couple of streams, should close every stream after use, in order to minimize the number of simultaneously open streams. Thus, the following code might eventually crash for some REXX interpreters:

```
do i=1 to 300
    call lineout 'file.' || i, 'this is file number' i
end
```

A REXX interpreter might try to defend itself against this sort of open-many-close-none programming, using of various programming techniques; this may lead to other strange effects. However, the main responsibility for avoiding this is with you, the REXX script programmer.

Not that if a stream is already open for reading, and you start writing to it, your implementation may have to reopen it in order to open for both reading and writing. There are mainly two strategies for handling this. Either the old file is closed, and then reopened in the new mode, which may leave you

with read and write access to another file. Or a new file handle is opened for the new mode, which may leave you with read and write access to two different files.

These are real-world problems which are not treated by the ideal description of TRL. A good implementation should detect these situations and raise `NOTREADY`.

26 Closing a Stream

As already mentioned, REXX does not have an explicit way of opening a stream. Nor does it have an explicit way of closing a stream. There is one semi-standard method: If you call `LINEOUT ()`, but omit both the data to be written and the new current write position, then the implementation is defined to set the current write position to the end-of-file. Furthermore, it is allowed by TRL to do something "magic" in addition. It is not explicitly defined what this magic is, but TRL suggests that it may be closing the stream, flushing the stream, or committing changes done previously to the stream.

In SAA, the definition is strengthened to state that the "magic" is closing, provided that the environment supports that operation.

A similar operation can be performed by calling `CHAROUT ()` with neither data nor a new position. However, in this case, both TRL and SAA leave it totally up to the implementation whether or not the file is to be closed. One can wonder whether the changes for `LINEOUT ()` in SAA with respect to TRL should also have been done to `CHAROUT ()`, but that this was forgotten.

TRL2 does not indicate that `LINEIN ()` or `CHARIN ()` can be used to close a string. Thus, the closest one gets to a standard way of closing input files is to call, e.g., `LINEOUT ()`; although it is conceptually suspect to call an output routine for an input file. The historical reasons for this omission are perhaps that flushing output files is vital, while the concept of flushing is irrelevant for input files; flushing is an important part of closing a file, and that explains why closing is only indicated for output files.

Thus, the statement:

```
call lineout 'myfile.txt'
```

might be used to close the stream `myfile.txt` in some implementations. However, it is not guaranteed to close the stream, so you cannot depend on this for script portability, but it's better than nothing. However, note that if it closes the stream, then also the current read position is affected. If it merely flushes the stream, then only the current write position is likely to be affected.

27 Character-wise and Line-wise I/O

Basically, the built-in REXX library offers two strategies of reading and writing streams: line-wise and character-wise. When reading line-wise, the underlying storage method of the stream must contain information which describes where each line starts and ends.

Some file systems store this information as one or more special characters; while others structure the file in a number of records; each containing a single line. This introduces a slightly subtle point; even though a stream `foo` returns the same data when read by `LINEIN ()` on two different machines; the

data read from `foo` may differ between the same two machines when the stream is read by `CHARIN()`, and vice versa. This is so because the end-of-line markers can vary between the two operating systems.

Example: Character-wise handling of EOL

Suppose a text file contains the following three lines (ASCII character set is assumed):

```
first
second
third
```

and you first read it line-wise and then character-wise. Assume the following program:

```
file = 'DATAFILE'
foo = ''
do i=1 while chars(file)>0
    foo = foo || c2x(charin(file))' '
end
say foo
```

When the file is read line-wise, the output is identical on all machines, i.e. the three lines shown above. However, the character-wise reading will be dependent on your operating system and its file system, thus, the output might be, e.g., any of:

```
66 69 72 73 74 73 65 6F 63 6E 64 74 68 69 72 64 66 69 72 73 74
0A
```

```
66 69 72 73 74 0A
73 65 6F 63 6E 64 0A
74 68 69 72 64 0A
```

```
66 69 72 73 74 0D 0A
73 65 6F 63 6E 64 0D 0A
74 68 69 72 64 0D 0A
```

If the machine uses records to store the lines, the first one may be the result; here, only the data in the lines of the file is returned. Note that the boxes in the output are put around the data generated by the actual line contents. What is outside the boxes is generated by the end-of-line character sequences.

The second output line is typical for Unix machines. They use the newline ASCII character as line separator, and that character is read immediately after each line. The last line is typical for MS-DOS, where the line separator character sequence is a carriage return following by a newline (ASCII '0D' x and '0A' x).

For maximum portability, the line-wise built-in functions (`LINEIN()`, `LINEOUT()` and `LINES()`) should only be used for line-wise streams. And the character-wise built-in functions (`CHARIN()`,

`CHAROUT ()` and `CHARS ()` should only be used for character-wise data. You should in general be very careful when mixing character- and line-wise data in a single stream; it does work, but may easily lead to portability problems.

The difference between character- and line-wise streams are roughly equivalent to the difference between binary and text streams, but the two concepts are not totally equivalent. In a binary file, the data read is the actual data stored in the file, while in a text file, the character sequences used for denoting end-of-line and end-of-file markers may be translated to actions or other characters during reading.

The end-of-file marker may be differently implemented on different systems. On some systems, this marker is only implicitly present at the end-of-file -- which is calculated from the file size (e.g. Unix). Other systems may put a character signifying end-of-file at the end (or even in the middle) of the file (e.g. <Ctrl-Z> for MS-DOS). These concepts vary between operating systems, interpreters should handle each concept according to the customs of the operating system. Check the implementation-specific documentation for further information. In any case, if the interpreter treats a particular character as end-of-file, then it only gives special treatment to this character during line-wise operations. During character-wise operations, no character has special meanings.

28 Reading and Writing

Four built-in functions provide line- and character-oriented stream reading and writing capabilities: `CHARIN ()`, `CHAROUT ()`, `LINEIN ()`, `LINEOUT ()`.

[`CHARIN ()`]

is a built-in function that takes up to three parameters, which are all optional: the name of the stream to read from, the start point, and the number of characters to read. The stream name defaults to the default input stream, the start point defaults to the current read position, the number of characters to read defaults to one character. Leave out these second parameter in order to avoid all repositioning. During execution, data is read from the streams specified, and returned as the return value.

[`LINEIN ()`]

is a built-in function that takes three parameters too, and they are equivalent to the parameters of `CHARIN ()`. However, if the second parameter is specified, it refers to a line position, rather than a character position; it refers to the character position of the first character of that line. Further, the third parameter can only be 0 or 1, and refers to the number of lines to read; i.e. you cannot read more than one line in each call. The line read is returned by the function, or the null string if no reading was requested.

[`LINEOUT ()`]

is a built-in function that takes three parameters too, the first is the name of the stream to write to, and defaults to the default output stream. The second parameter is the data to be written to the file, and if not specified, now writing occurs. The third parameter is a line-oriented position in the file; if the third parameter is specified, the current position is repositioned at before the data (if any) is written. If data is written, an end-of-line character sequence is appended to the output stream.

[`CHAROUT ()`]

is a built-in function that is used to write characters to a file. It is identical to `LINEOUT ()`,

except that the third parameter refers to a character position, instead of a line position. The second difference is that an end-of-line character sequence is not appended at the end of the data written.

Example: Counting lines, words, and characters

The following REXX program emulates the core functionality of the `wc` program under Unix. It counts the number of lines, words, and characters in a file given as the first argument.

```
file = arg(1)
parse value 0 0 0 with lines words chars
do while lines(file)>0
    line = linein(file)
    lines = lines + 1
    words = words + words(line)
    chars = chars + length(line)
end
say 'lines='lines 'words='words 'chars='chars
```

There are some problems. For instance, the end-of-line characters are not counted, and a last improperly terminated line is not counted either.

29 Determining the Current Position

Standard REXX does not have any seek call that returns the current position in a stream. Instead, it provides two calls that return the amount of data remaining on a stream. These two built-in functions are `LINES()` and `CHARS()`.

- The `LINES()` built-in function returns the number of complete lines left on the stream given as its first parameter. The term "complete lines" does not really matter much, since an implementation can assume the end-of-file to implicitly mean an end-of-line.
- The `CHARS()` built-in function returns the number of character left in the stream given as its first parameter.

This is one of the concepts where REXX/I does not map very well to C/I and vice versa. While REXX reports the amount of data from the current read position to the end of stream, C reports the amount of data from the start of the file to the current position. Further, the REXX method only works for input streams, while the C method works for both input and output files. On the other hand, C has no basic constructs for counting remaining or reposition at lines of a file.

Example: Retrieving current position

So, how does one find the current position in a file, when only allowed to do normal repositioning? The trick is to reposition twice, as shown in the code below.

```
ftell: procedure
    parse arg filename
    now = chars(filename)
    call charin filename, 0, 1
    total = chars(filename)
    call charin filename, 0, total-now
    return total-now
```

Unfortunately, there are many potential problems with this code. First, it only works for input files, since there is no equivalent to `CHARS()` for output files. Second, if the file is empty, none of the repositioning work, since it is illegal to reposition after end-of-file for input files--and the end-of-file is the first position of the file. Third, if the current read position of the file is at the end of file (e.g. all characters have been read) it will not work for similar reasons as for the second case. And fourth, it only works for persistent files, since transient files do not allow repositioning.

Example: Improved `ftell` function

An improved version of the code for the `ftell` routine (given above), which tries to handle these problems is:


```

ftell: procedure
    parse arg filename
    signal on notready name not_persist
    now = chars(filename)
    signal on notready name is_empty
    call charin filename, 0, 1
    total = chars()
    if now>0 then
        call charin filename, 0, total-now+1
    else if total>0 then
        call charin filename, 1, total
    else
        nop /* empty file, should have raised NOTREADY
*/
    return total-now+1

not_persist: say filename 'is not persistent'; return 0

is_empty: say filename 'is empty'; return 0

```

This same method can be used for line-oriented I/O too, in order to return the current line number of an input file. However, a potential problem in that case is that the routine leaves the stream repositioned at the start of the current line, even if it was initially positioned to the middle of a line. In addition, the line-oriented version of this `ftell` routine may prove to be fairly inefficient, since the interpreter may have to scan the whole file twice for end-of-line character sequences.

30 Positioning Within a File

REXX supports two strategies for reading and writing streams: character-wise, and line-wise, this section describes how a program can reposition the current positions for each of these strategies. Note that positioning is only allowed for persistent streams.

For each open file, there is a current read position or a current write position, depending on whether the file is opened for reading or writing. If the file is opened for reading and writing simultaneously, it has both a current read position and a current write position, and the two are independent and in general different. A position within a file is the sequence number of the byte or line that will be read or written in the next such operation.

Note that REXX starts numbering at one, not zero. Therefore, the first character and the first line of a stream are both numbered one. This differs from several other programming languages, which start numbering at zero.

Just after a stream has been opened, the initial value of the current read position is the first character in the stream, while the current write position is the end-of-file, i.e. the position just after the last character in the stream. Then, reading will return the first character (or line) in the stream, and writing will append a new character (or line) to the stream.

These initial values for the current read and write positions are the default values. Depending on your

REXX implementation, other mechanisms for explicitly opening streams (e.g. through the built-in function) may be provided, and may set other initial values for these positions. See the implementation-specific documentation for further information.

STREAM ()

When setting the current read position, it must be set to the position of an existing character in the stream; i.e. a positive value, not greater than the total number of characters in the stream. In particular, it is illegal to set the current read position to the position immediately after the last character in the stream; although this is legal in many other programming languages and operating systems, where it is known as "seeking to the end-of-file".

When setting the current write position, it too must be set to the position of an existing character in the stream. In addition, and unlike the current read position, the current write position may also be set to the position immediately following the last character in the stream. This is known as "positioning at the end-of-file", and it is the initial value for the current write position when a stream is opened. Note that you are not allowed to reposition the current write position further out beyond the end-of-file--which would create a "hole" in the stream--even though this is allowed in many other languages and operating systems.

Depending on your operating system and REXX interpreter, repositioning to after the end-of-file may be allowed as an extension, although it is illegal according to TRL2. You should avoid this technique if you wish to write portable programs.

REXX only keeps one current read position and one current write position for each stream. So both line-wise and character-wise reading as well as positioning of the current read position will operate on the same current read position, and similarly for the current write position.

When repositioning line-wise, the current write position is set to the first character of the line positioned at. However, if positioning character-wise so that the current read position is in the middle of a line in the file, a subsequent call to `LINEIN ()` will read from (and including) the current position until the next end-of-line marker. Thus, `LINEIN ()` might under some circumstances return only the last part of a line. Similarly, if the current write position has been positioned in the middle of an existing line by character-wise positioning, and `LINEOUT ()` is called, then the line written out becomes the last part of the line stored in the stream.

Note that if you want to reposition the current write position using a line count, the stream may have to be open for read, too. This is because the interpreter may have to read the contents of the stream in order to find where the lines start and end. Depending on your operating system, this may even apply if you reposition using character count.

Example: Repositioning in empty files

Since the current read position must be at an existing character in the stream, it is impossible to reposition in or read from an empty stream. Consider the following code:

```
filename = '/tmp/testing'
call lineout filename,, 1 /* assuming truncation */
call linein filename, 1, 0
```

One might believe that this would set the current read and write position to the start of the stream. However, assume that the `LINEOUT()` call truncates the file, so that it is zero bytes long. Then, the last call can never be legal, since there is no byte in the file at which it is possible to position the current read position. Therefore, a `NOTREADY` condition is probably raised.

Example: Relative repositioning

It is rather difficult to reposition a current read or write position relative to the current position. The only way to do this within the definition of the standard is to keep a counter which tells you the current position. That is, if you want to move the current read position five lines backwards, you must do it like this:

```
filename = '/tmp/data'
linenum = 0 ;
say linein(filename,10); linenum = 10
do while random(100)>3
    say linein(filename); linenum = linenum+1
end
call linein(filename,linenum-5,0); linenum = linenum-5
```

Here, the variable `linenum` is updated for each time the current read position is altered. This may not seem too difficult, and it is not in most cases. However, it is nearly impossible to do this in the general case, since you must keep an account of both line numbers and character numbers. Setting one may invalidate the other: consider the situation where you want to reposition the current read position to the 10th character before the 100th line in the stream. Except from mixing line-wise and character-wise I/O (which can have strange effects), this is nearly impossible. When repositioning character-wise, the line number count is invalidated, and vice versa.

The "only" proper way of handling this is to allow one or more (non-standard) `STREAM()` built-in function operations that return the current character and line count of the stream in the interpreter.

Example: Destroying line count

This example shows how overwriting text to the middle of a file can destroy the line count. In the following code, we assume that the file `foobare` exists, and contains ten lines which are "first line", "second line", etc. up to "tenth line". Then consider the following code:

```

filename = 'foobar'
say linein(filename, 5)    /* says 'fifth line' */
say linein(filename)      /* says 'sixth line' */
say linein(filename)      /* says 'seventh line' */
call lineout filename, 'This is a very long line', 5
say linein(filename, 5)    /* says 'This is a very long line' */
say linein(filename)      /* says 'venth line' */
say linein(filename)      /* says 'eight line' */

```

As you can see from the output of this example, the call to `LINEOUT()` inserts a long line and overwrites the fifth and sixth lines completely, and the seventh line partially. Afterwards, the sixth line is the remaining part of the old seventh line, and the new seventh line is the old eighth line, etc.

31 Errors: Discovery, Handling, and Recovery

TRL2 contains two important improvements over TRL1 in the area of handling errors in stream I/O: the `NOTREADY` condition and the `STREAM()` built-in function. The `NOTREADY` condition is raised whenever a stream I/O operation did not succeed. The `STREAM()` function is used to retrieve status information about a particular stream or to execute a particular operation for a stream.

You can discover that an error occurred during an I/O operation in one of the following ways: a) it may trigger a `SYNTAX` condition; b) it may trigger a `NOTREADY` condition; or c) it may just not return that data it was supposed to. There is no clear border between which situations should trigger `SYNTAX` and which should trigger `NOTREADY`. Errors in parameters to the I/O functions, like a negative start position, is clearly a `SYNTAX` condition, while reading off the end-of-file is equally clearly a `NOTREADY` condition. In between lay more uncertain situations like trying to position the current write position after the end-of-file, or trying to read a non-existent file, or using an illegal filename.

Some situations are likely to be differently handled in various implementations, but you can assume that they are handled as either `SYNTAX` or `NOTREADY`. Defensive, portable programming requires you to check for both. Unfortunately, `NOTREADY` is not allowed in TRL1, so you have to avoid that condition if you want maximum compatibility. And due to the very lax restrictions on implementations, you should always perform very strict verification on all data returned from any file I/O built-in function.

If neither are trapped, `SYNTAX` will terminate the program while `NOTREADY` will be ignored, so the implementor's decision about which of these to use may even depend on the severity of the problem (i.e. if the problem is small, raising `SYNTAX` may be a little too strict). Personally, I think `SYNTAX` should be raised in this context only if the value of a parameter is outside its valid range for all contexts in which the function might be called.

Example: General `NOTREADY` condition handler

Under TRL2 the "correct" way to handle `NOTREADY` conditions and errors from I/O operations is unfortunately very complex. It is shown in this example, in order to demonstrate the procedure:

```

myfile = 'MYFILE.DAT'
signal on syntax name syn_handler
call on notready name IO_handler
do i=1 to 10 until res=0
    res = lineout(myfile, 'line #'i)
    if (res=0) then
        say 'Call to LINEOUT() didn't manage to write
out data'
end
exit

IO_handler:
syn_handler:
    file = condition('D')
    say condition('C') 'raised for file' file 'at line'
sigl': '
    say ' ' sourceline(sigl)
    say ' State='stream(file,'S') 'reason:'
stream(file,'D')
    call lineout( condition( 'D' )) /* try to close */
    if condition('C')=='SYNTAX' then
        exit 1
    else
        return

```

Notethedoublecheckinginthisexample: firsttheconditionhandlerissetuptotrapany NOTREADY conditions, andthenthe returncode from LINEOUT() is checked for each call.

As you can see, there is not really that much information that you can retrieve about what went wrong. Some systems may have additional sources from which you can get information, e.g. special commands for the STREAM() built-in function, but these are non-standard and should be avoided when writing compatible programs.

32 Common Differences and Problems with Stream I/O

This section describes some of the common traps and pitfalls of REXX I/O.

32.1 Where Implementations are Allowed to Differ

TRL is rather relaxed in its specifications of what an interpreter must implement of the I/O system. It recognizes that operating systems differ, and that some details must be left to the implementor to decide, if REXX is to be effectively implemented. The parts of the I/O subsystem of REXX where implementations are allowed to differ, are:

- The functions LINES() and CHARS() are not required to return the number of lines or characters left in a stream. TRL says that if it is impossible or difficult to calculate the numbers, these functions may return 1 unless it is absolutely certain that there are no more data left. This leads to some rather kludgy programming techniques.

- Implementations are allowed to ignore closing streams, since TRL does not specify a way to do this. Often, the closing of streams is implemented as a command, which only makes it more incompatible.
- Check the implementation-specific documentation before using the function `LINEOUT (file)` for closing files.
- The difference in the action of closing and flushing a file, can make a `REXX` script that works under one implementation crash under another, so this feature is of very limited value if you are trying to write portable programs.

TRL says that because the operating system environments will differ a lot, and an efficient and useful interpreter is the most important goal, implementations are allowed to deviate from the standard in any respect necessary in the domain of I/O [TRL2]. Thus, you should never assume anything about the I/O system, as the "rules" listed in TRL are only advisory.

32.2 Where Implementations might Differ anyway

In the section above, some areas where the standard allows implementations to differ are listed. In an ideal world, that ought to be the only trap that you should need to look out for, but unfortunately, the world is not ideal. There are several areas where the requirements set up by the standard is quite high, and where implementations are likely to differ from the standard.

These areas are:

- Repositioning at (for the current write position) or beyond the end-of-file may be allowed. On some systems, to prohibit that would require a lot of checking, so some systems will probably skip that check. At least for some operating systems, the act of repositioning after end-of-file is a useful feature.
- Under Unix, it can be used for creating a dynamically sized random access file; do not bother about how much space is allocated for the file, just position to the correct "slot" and write the data there. If the data file is sparse, holes might occur in the file; that is parts of the file which has not been written, and which is all zeros (and which are therefore not stored on disk).
- Some implementations will use the same position for both the current read position and the current write position to overcome these implementations. Whenever you are doing a read, and the previous operation was a write (or vice versa), it may prove useful to reposition the current read (or write) position.
- There might be a maximum line size for your `REXX` interpreter. At least the 50 Kb limit on string length may apply.
- Handling the situation where another program writes data to a file which is used by the `REXX` interpreter for reading.

32.3 LINES () and CHARS () are inaccurate

Because of the large differences between various operating systems, REXX allows some fuzz in the implementation of the LINES () and CHARS () built-in functions. Sometimes, it is difficult to calculate the number of lines or characters in a stream; generally because the storage format of the file often requires a linear search through the whole stream to determine that number. Thus, REXX allows an implementation to return the value 1 for any situation where the real number is difficult or impossible to determine. Effectively, an implementation can restrict the domain of return values for these two functions only 1 and 0 from these two functions.

Many operating systems store lines using a special end-of-line character sequence. For these systems, it is very time-consuming to count the number of lines in a file, as the file must be scanned for such character sequences. Thus, it is very tempting for an implementor to return the value 1 for any situation where there are more than zero lines left.

A similar situation arises for the number of characters left, although it is more common to know this number, thus it is generally a better chance of CHARS () returning the true number of characters left than LINES () returning the true number of lines left.

However, you can be fairly sure that if an implementation returns a number greater than 1, then that number is the real number of lines (or characters) left in the stream. And simultaneously, if the number returned is 0, then there is no lines (or characters) left to be read in the stream. But if the number is 1, then you will never know until you have tried.

Example: File reading idiom

This example shows a common idiom for reading all contents of a file into REXX variables using the LINES () and LINEIN () built-in functions.

```
i = 1
signal on notready
lleft = lines(file)
do while lleft > 0
    do i = i to i + lleft
        line.i = linein(file)
    end
    lleft = lines(file)
end
notready:
lines.0 = i - 1
```

Here, the two nested loops iterate over all the data to be read. The innermost loop reads all data currently available, while the outermost loop checks for more available data. Implementations having a LINES () that return only 0 and 1 will generally iterate the outermost loop many times; while implementations that return the "true" number from LINES () generally only iterate the outermost loop once.

There is only one place in this code that `LINEIN()` is called. The `I` variable is incremented at only one place, and the variable `LINES . 0` is set in one clause, too. Some redundancy can be removed by setting the `WHILE` expression to:

```
do while word(value('lleft',lines(file)) lleft,2)>0
```

The two assignments to the `LEFT` variable must be removed. This may look more complicated, but it decreases the number of clauses having a call to `LINES()` from two to one. However, it is less certain that this second solution is more efficient, since using `VALUE()` built-in function can be inefficient over "normal" variable references.

32.4 The Last Line of a Stream

How to handle the last line in a stream is sometimes a problem. If you use a system that stores end-of-lines as special character sequences, and the last part of the data of a stream is an unterminated line, then what is returned when you try to read that part of data?

There are three possible solutions: First, it may interpret the end-of-file itself as an implicit end-of-line, in this case, the partial part of the line is returned, as if it was properly terminated. Second, it may raise the `NOTREADY` condition, since the end-of-file was encountered during reading. Third, if there is any chance of additional data being appended, it may wait until such data are available. The second and third approaches are suitable for persistent and transient files, respectively.

The first approach is sometimes encountered. It has some problems though. If the end of a stream contains the data `ABC<NL>XYZ`, then it might return the string `XYZ` as the last line of the stream. However, suppose the last line was an empty line, then the last part of the stream would be: `ABC<NL>`. Few would argue that there is any line in this stream after the line `ABC`. Thus, the decision whether the end-of-file is an implicit end-of-line depends on whether there would be a last line has zero length or not.

A pragmatic solution is to let the end-of-file only be an implicit end-of-file if the characters immediately in front of it are not an explicit end-of-line character sequence.

However, TRL gives some indication that an end-of-file is not an implicit end-of-line. It says that `LINES()` returns the number of complete lines left, and that `LINEIN()` returns a complete line. On the other hand, the end-of-line sequence is not rigidly defined by TRL, so an implementor is almost free to define end-of-line in just about any terms that are comfortable. Thus, the last line of a stream may be a source of problem if it is not explicitly terminated by an end-of-line.

32.5 Other Parts of the I/O System

This section lists some of the other parts of `REXX` and the environments around `REXX` that may be considered a part of the I/O system.

[Stack.]

The stack be used to communicate with external environments. At the `REXX` side, the interface to the stack is the instructions `PUSH`, `PULL`, `PARSE PULL`, and `QUEUE`; and the built-in function `QUEUED()`. These can be used to communicate with external programs by storing data to be transferred on the stack.

[The `STREAM()` built-in function.]

This function is used to control various aspects about the files manipulated with the other standard I/O functions. The standard says very little about this function, and leaves it up to the implementor to specify the rest. Operations like opening, closing, truncating, and changing modes

[The `SAY` instruction.]

The `SAY` instruction can be used to write data to the default output stream. If you use redirection, you can indirectly use it to write data to a file.

[The `ADDRESS` instruction.]

The `ADDRESS` instruction and commands can be used to operate on files, depending on the power of your host environments and operating system.

[The `VALUE()` built-in function.]

The function `VALUE()`, when used with three parameters, can be used to communicate with external host environments and the operating system. However, this depends on the implementation of your interpreter.

[SAA API.]

The SAA API provides several operations that can be used to communicate between processes. In general, SAA API allows you to perform the operations listed above from a binary program written in a language other than `REXX`.

And of course, I/O is performed whenever a `REXX` program or external function is started.

32.6 Implementation-Specific Information

This section describes some implementations of stream I/O in `REXX`. Unfortunately, this has become a very large section, reflecting the fact that stream I/O is an area of many system-specific solutions.

In addition, the variations within this topic are rather large. `Regina` implements a set of functions that are very close to that of `TRL2`. The other extremes are `ARexx` and `BRexx`, which contain a set of functions which is very close to the standard I/O library of the C programming language.

32.7 Stream I/O in Regina 0.07a

`Regina` implements stream I/O in a fashion that closely resembles how it is described in `TRL2`. The following list gives the relevant system-specific information.

[Names for standard streams.]

`Regina` uses `<stdout>` and `<stdin>` as names for the standard output and input streams. Note that the angle brackets are part of the names. You may also access the standard error stream (on systems supporting this stream) under the name `<stderr>`. In addition, the null string is taken to be equivalent to an empty first parameter in the I/O-related built-in functions.

[Implicit opening.]

`Regina` implicitly opens any file whenever it is first used.

If the first operation is a read, it will be opened in read-only mode. If the first operation is a write, it is opened in read-write mode. In this case if the read-write opening does not succeed, the file is opened in write-only mode. If the file exists, the opening is non-destructive, i.e. that

the file is not truncated or overwritten when opened, else it is created if opened in read-write mode.

If you name a file currently open in read-only mode in a write operation, `Regina` closes the file, and reopens it in read-write mode. The only exception is when you call `LINEOUT ()` with both second and third arguments unspecified, which always closes a file, both for reading and writing. Similarly, if the file was opened in write-only mode, and you use it in a read operation, `Regina` closes and reopens in read-write mode.

This implicit reopening is enabled by default. You can turn it off by unsetting the extension `ExplicitOpen`.

[Separate current positions.]

The environment in which `Regina` operates (ANSI C and POSIX) does not allow separate read and write positions, but only supplies one position for both operations. `Regina` handles this by maintaining the two positions internally, and move the "real" current position back and forth depending on whether a read or write operation is next.

[Swapping out file descriptors.]

In order to defend itself against "open-many-close-none" programming, `Regina` tries to "swap out" files that have been unused for some time. Assuming that your operating system limits `Regina` to 100 simultaneously open files; when you try to open your 101st file, `Regina` closes the least recently used stream, and recycles its descriptor for the new file. You can enable or disable this recycling with the `SwapFileP` extension.

During this recycling, `Regina` only closes the file in the operating system, but retains all vital information about the file itself. If you re-access the file later, `Regina` reopens it, and positions the current read and write positions at the correct (i.e. previous) positions. This introduces some uncertainties into stream processing. Renaming a file affects it only if it gets swapped out. Since the swap operation is something the users don't see, it can cause some strange effects.

`Regina` will not allow a transient stream to be swapped out, since they often are connected to some sort of active partner in the other end, and closing the file might kill the partner or make it impossible to reestablish the stream. So only persistent files are swapped out. Thus, you can still fill the file table in `Regina`.

[Explicit opening and closing.]

`Regina` allows streams to be explicitly opened or closed through the use of the built-in function `STREAM ()`. The exact syntax of this function is described in section `stream`. Old versions of `Regina` supported two non-standard built-in functions `OPEN ()` and `CLOSE ()` for these operations. These functions are still supported for compatibility reasons, but might be removed in future releases. Their availability is controlled by the `OpenBif` and `CloseBif` extensions.

[Truncation after writing lines.]

If you reposition line-wisely the current write position to the middle of a file, `Regina` truncates the file at the new position. This happens whether the data is written during the `LINEOUT ()` or not. If not, the file might contain half a line, some lines might disappear, and the line count would in general be disrupted. The availability of this behavior is controlled by `LineOutTrunc`, which is turned on by default.

Unfortunately, the operation of truncating a file is not part of POSIX, and it might not exist on

all systems, soon some are systems, this truncating will not occur. In order to be able to truncate a file, your machine must have the `ftruncate()` system call in C. If you don't have this, the truncating functionality is not available.

[Caching info on lines left.]

When Regina executes the built-in function `LINES()` for a persistent stream, it caches the number of lines left as an attribute to the stream. In subsequent calls to `LINEIN()`, this number is updated, so that subsequent calls to `LINES()` can retrieve the cached number instead of having to re-scan the rest of the stream, provided that the number is still valid. Some operations will invalidate the count: repositioning the current read position; reading using the character oriented I/O, i.e. `CHARIN()`; and any write operation by the same interpreter on the stream. Ideally, any write operations should invalidate the count, but that might require a large overhead before any operation, in order to check whether the file has been written to by other programs.

This functionality can be controlled by the extension called `CacheLineNo`, which is turned on by default. Note that if you turn that off, you can experience a serious decrease in performance.

The following extra built-in functions relating to stream I/O are defined in Regina. They are provided for extra support and compatibility with other systems. Their support may be discontinued in later versions, and they are likely to be moved to a library of extra support.

CLOSE(*streamid*)

Closes the stream named by *streamid*. This stream must have been opened by implicit open or by the `OPEN` function call earlier. The function returns 1 if there was any file to close, and 0 if the file was not opened. Note that the return value does not indicate whether the closing was successful. You can use the extension named `CloseBif` with the `OPTIONS` instruction to select or remove this function. This function is now obsolete, instead you should use:

```
STREAM( streamid, 'Command', 'CLOSE' )
```

CLOSE(<i>myfile</i>)	1	if stream was open
CLOSE('NOSUCHFILE')	0	if stream didn't exist

OPEN(*streamid*, *access*)

Open the stream named *streamid* with the access *access*. If *access* is not specified, the access `R` will be used. *access* may be the following characters. Only the first character of the *access* is needed.

[R]

(Read)Open for read access. The file pointer will be positioned at the start of the file, and only read operations are allowed.

[W]

(Write)Open for write access and position the current write position at the end of the file. An error is returned if it was not possible to get appropriate access.

The return value from this function is either 1 or 0, depending on whether the named stream is in opened state after the operation has been performed.

Note that if you open the files " foobar" and " ./foobar" they will point to the same physical file, but **Regina** interprets them as two different streams, and will open an internal file descriptor for each one. If you try to open an already open stream, using the same name, it will have no effect.

You can use the extension `OpenBif` with the `OPTIONS` instruction to control the availability of this function. This function is now obsolete, but is still kept for compatibility with other interpreters and old versions of **Regina**. Instead, with **Regina** you should use:

```
STREAM( streamid, 'C', 'READ' | 'WRITE' | 'APPEND' | 'UPDATE' )
```

OPEN(myfile,'write')	1	maybe, if successful
OPEN(passwd,'Write')	0	maybe, if no write access
OPEN('DATA','READ')	0	maybe, if successful

The return value from this function is either 1 or 0, depending on whether the named stream is in opened state after the operation has been performed.

32.8 Functionality to be Implemented Later

This section lists the functionality not yet in **Regina**, but which is intended to be added later. Most of these are fixest problems, compatibility modes, etc.

[Indirect naming of streams.]

Currently, streams are named directly, which is a convenient. However, there are a few problems: for instance, it is difficult to write to a file which name is `<stdout>`, simply because that is a reserved name. To fix this, an indirect naming scheme will be provided through the `STREAM()` built-in function. The functionality will resemble the `OPEN()` built-in function of **ARexx**.

[Consistence in file handle swapping.]

When a file handle is currently swapped out in order to avoid filling the system file table, very little checking of consistency is currently performed. At least, vital information about the file should be retained, such as the inode and file system for Unix machines retrieval by the `fstat()` call. When the file is swapped in again, this information must be checked against the file which is reopened. If there is a mismatch, `NOTREADY` should be raised. Similarly, when reopening a file because of a new access mode is requested, the same checkings should be performed.

[Files with holes.]

Regina will be changed to allow it to generate files with holes for system where this is relevant.

Although standard REXX does not allow this, it is a very common programming idiom for certain systems, and should be allowed. It will, however, be controllable through an extension called `SparseFiles`.

32.9 Stream I/O in ARexx 1.15

ARexx differs considerably from standard REXX with respect to stream I/O. In fact, none of the standard stream functionality of REXX is available in ARexx. Instead, a completely distinct set of functions are used. The difference is so big, that it is useless to describe ARexx stream I/O in terms of standard REXX stream I/O, and everything said so far in this chapter is irrelevant for ARexx. Therefore, we explain the ARexx functionality from scratch.

All in all, the ARexx file I/O interface resembles the functions of the Standard C I/O library, probably because ARexx is written in C, and the ARexx I/O functions are "just" interfaces to the underlying C functions. You may want to check up the documentation for the ANSI C I/O library as described in [ANSIC], [KR], and [PJPlauger].

ARexx uses a two-level naming scheme for streams. The file names are bound to a stream name using the `OPEN()` built-in function. In all other I/O functions, only the stream name is used.

`OPEN(name, filename[, mode])`

You use the `OPEN()` built-in function to open a stream connected to a file called *filename* in AmigaDOS. In subsequent I/O calls, you refer to the stream as *name*. These two names can be different.

The *name* parameter cannot already be in use by another stream. If so, the `OPEN()` function fails. Note that the *name* parameter is case-sensitive. The *filename* parameter is not strictly case-sensitive: the case used when creating a new file is preserved, but when referring to an existing file, the name is case-insensitive. This is the usual behavior of AmigaDOS.

If any of the other I/O operations uses a stream name that has not been properly opened using `OPEN()`, that operation fails, because ARexx has no auto-open-on-demand feature.

The optional parameter *mode* can be any of `Read`, `Write`, or `Append`. The mode `Read` opens an existing file and sets the current position to the start of the file. The mode `Append` is identical to `Read`, but sets the current position to the end-of-file. The mode `Write` creates a new file, i.e. if a file with that name already exists, it is deleted and a new file is created. Thus, with `Write` you always start with an empty file. Note that the terms "read," "write," and "append" are only remotely connected to the mode in which the file is opened. Both reading and writing are allowed for all of these three modes; the mode names only reflect the typical operations of these modes.

The result from `OPEN()` is a boolean value, which is 1 if a file by the specified *name* was successfully opened during the `OPEN()` call, and 0 otherwise.

The number of simultaneously open files is no problem because AmigaDOS allocates file handles dynamically, and thus only limited by the available memory. One system managed 2000 simultaneously open files during a test.

OPEN('infile','work:DataFile')	1	if successful
OPEN('work','RAM:FooBar','Read')	0	if didn't exist
OPEN('output','TmpFile','W')	1	(re)creates file

CLOSE(*name*)

You use the `CLOSE()` built-in function to close a stream. The parameter *name* must match the first parameter in a call to `OPEN()` earlier in the same program, and must refer to an open stream. The return value is a boolean value that reflects whether there was a file to close (but not whether it was successfully closed).

CLOSE('infile')	1	if stream was previously open
CLOSE('outfile')	0	if stream wasn't previously open

WRITELN(*name*,*string*)

The `WRITELN()` function writes the contents of *string* as a line to the stream *name*. The *name* parameter must match the value of the first parameter in an earlier call to `OPEN()`, and must refer to an open stream. The data written is all the characters in *string* immediately followed by the newline character (ASCII <Ctrl-J> for AmigaDOS).

The return value is the number of characters written, including the terminating newline. Thus, a return value of 0 indicates that nothing was written, while a value which is one more than the number of characters in *string* indicates that all data was successfully written to the stream.

When writing a line to the middle of a stream, the old contents is written over, but the stream is not truncated; there is no way to truncate a stream with the **ARexx** built-in functions. This overwriting can leave partial lines in the stream.

WRITELN('tmp','Hello,world!')	14	if successful
WRITELN('work','Hithere')	0	nothing was written
WRITELN('tmp','Hithere')	5	partially successful

WRITECH(*name*,*string*)

The WRITECH () function is identical to WRITELN (), except that the terminating newline character is not added to the data written out. Thus, WRITELN () is suitable for line-wise output, while WRITECH () is useful for character-wise output.

WRITECH('tmp','Hello,world!')	13	if successful
WRITECH('work','Hithere')	0	nothing was written
WRITECH('tmp','Hithere')	5	partially successful

READLN(*name*)

The READLN () function reads a line of data from the stream referred to by *name*. The parameter *name* must match the first parameter of an earlier call to OPEN (), i.e. it must be an open stream.

The return value is a string of characters which corresponds to the characters in the stream from and including the current position forward to the first subsequent newline character found. If no newline character is found, the end-of-file is implicitly interpreted as a newline and the end-of-file state is set. However, the data returned to the user never contains the terminating end-of-line.

To differ between the situation where the last line of the stream was implicitly terminated by the end-of-file and where it was explicitly terminated by an end-of-line character sequence, use the EOF () built-in function. The EOF () returns 1 in the former case and 0 in the latter case.

There is a limit in AREXX on the length of lines that you can read in one call to READLN (). If the length of the line in the stream is more than 1000 characters, then only the first 1000 characters are returned. The rest of the line can be read by additional READLN () and READCH () calls. Note that whenever READLN () returns a string of exactly 1000 characters, then no terminating end-of-line was found, and a new call to READLN () must be executed in order to read the rest of the line.

READLN('tmp')	Hello world!	maybe
READLN('work')		maybe, if unsuccessful

READCH(*name*[,*length*])

The READCH () built-in function reads characters from the stream named by the parameter *name*, which must correspond to the first parameter in a previous call to OPEN (). The number of characters

read is given by *length*, which must be a non-negative integer. The default value of *length* is 1.

The value returned is the data read, which has the length corresponding to the *length* parameter if no errors occurred.

There is a limit in **ARexx** for the length of strings that can be read in one call to **READCH ()**. The limit is 65535 bytes, and is a limitation in the maximum size of an **ARexx** string.

READCH('tmp',3)	Hel	maybe
READCH('tmp')	l	maybe
READCH('tmp',6)	oworl	maybe

EOF (*name*)

The **EOF ()** built-in function tests to see whether the end-of-file has been seen on the stream specified by *name*, which must be an open stream, i.e. the first parameter in a previous call to **OPEN ()**.

The return value is 1 if the stream is in end-of-file mode, i.e. if a read operation (either **READLN ()** or **READCH ()**) has seen the end-of-file during its operation. However, reading the last character of the stream does not put the stream in end-of-file mode; you must try to read at least one character past the last character. If the stream is not in end-of-file mode, the return value is 0.

Whenever the stream is in end-of-file mode, it stays there until a call to **SEEK ()** is made. No read or write operation can remove the end-of-file mode, only **SEEK ()** (and closing followed by reopening).

EOF('tmp')	0	maybe
EOF('work')	1	maybe

SEEK (*name* , *offset* [, *mode*])

The **SEEK ()** built-in function repositions the current position of the file specified by the parameter *name*, which must correspond to an open file, i.e. to the first parameter of a previous call to **OPEN ()**. The current position in the file is set to the byte referred to by the parameter *offset*. Note that *offset* is zero-based, so the first byte in the file is numbered 0. The value returned is the current position in the file after the seek operation has been carried through, using **Beginning** mode.

If the current position is attempted set past the end-of-file or before the beginning of the file, then the current position is not moved, and the old current position is returned. Note that it is legal to position at the end-of-file, i.e. the position immediately after the last character of the file. If a file contains 12

characters, the valid range for the resulting new current position is 0-12.

The last parameter, *mode*, can take any of the following values:

Beginning, Current, or End. It specifies the base of the seeking, i.e. whether it is relative to the first byte, the end-of-file position, or the old current position. For instance: for a 20-byte file with current position 3, then offset 7 for base Beginning is equivalent to offset -13 for base End and offset 4 for Current. Note that only the first character of the *mode* parameter is required, the rest of that parameter is ignored.

SEEK('tmp',12,'B')	12	if successful
SEEK('tmp',-4,'Begin')	12	if previously at 12
SEEK('tmp',-10,'E')	20	if length is 30
SEEK('tmp',5)	17	if previously at 12
SEEK('tmp',5,'Celcius')	17	only first character in mode matters
SEEK('tmp',0,'B')	0	always to start of file

32.10 Main Differences from Standard REXX

Now, as the functionality has been explained, let me point out the main conceptual differences from standard REXX; they are:

[Current position.]

ARexx does not differ between a current read and write position, but uses a common current position for both reading and writing. Further, this current position (which it is called in this documentation) can be set to any byte within the file, and to the end-of-file position. Note that the current position is zero-based.

[Indirect naming.]

The stream I/O operations in ARexx do not get a parameter which is the name of the file. Instead, ARexx uses an indirect naming scheme. The `OPEN ()` built-in function binds a REXX stream name for a file to a named file in the AmigaDOS operating system; and later, only the REXX stream name is used in other stream I/O functions operating on that file.

[Special stream names.]

There are two special filenames in ARexx: `STDOUT` and `STDIN`, which refer to the standard input file and standard output file. With respect to the indirect naming scheme, these are not file names, but names for open streams; i.e. they can be used in stream I/O operations other than `OPEN ()`. For some reason, it is possible to close `STDIN` but not `STDOUT`.

[NOTREADY not supported.]

ARexx has no `NOTREADY` condition. Instead, you must detect errors by calling `EOF ()` and checking the return codes from each I/O operation.

[Other things missing.]

In ARexx, all files must be explicitly opened. There is no way to reposition line-wise, except for reading lines and keeping a count yourself.

Of course, ARexx also has a lot of functionality which is not part of standard REXX, like relative

repositioning, explicit opening, an end-of-file indicator, etc. But this functionality is descriptive above in the descriptions of extended built-in functions, and it is of less interest there.

When an **ARexx** script has opened a file in **Write** mode, other **ARexx** scripts are not allowed to access that file. However, if the file is opened in **Read** or **Append** mode, then other **ARexx** scripts can open the file too, and the same state of the contents of the file is seen by all scripts.

Note that it is difficult to translate between using standard **REXX** **stream I/O** and **ARexx** **stream I/O**. In particular, the main problem (other than missing functionality in one of the systems) is the processing of end-of-lines. In standard **REXX**, the end-of-file is detected by checking whether there is more data left, while in **ARexx** one checks whether the end-of-file has been read. The following is a common standard **REXX** idiom:

```
while lines('file') > 0 /* for each line available */
    say linein('file') /* process it */
end
```

In **ARexx** this becomes:

```
tmp = readln('file') /* attempt to read first line */
do until eof('file') /* if EOF was not seen */
    say tmp /* process line */
    tmp = readln('file') /* attempt to read next line */
end
```

It is hard to mechanically translate between them,

because of the lack of an **EOF()** built-in function in standard **REXX**, and the lack of a **LINES()** built-in function in **ARexx**.

Note that in the **ARexx** example, an improperly terminated last line is not read as an independent line, since **READLN()** searches for an end-of-line character sequence. Thus, in the last invocation **tmp** is set to the last unterminated line, but **EOF()** returns true too. To make this different, make the **UNTIL** subterm of the **DO** loop check for the expression **EOF('file') && TMP <> ''**.

The limit of 1000 characters for **READLN()** means that a generic line reading routine in **ARexx** must be similar to this:

```
readline: procedure
    parse arg filename
    line = ''
    do until length(tmp) < 1000
        tmp = readln(filename)
        line = line || tmp
    end
    return line
```

This routine calls **READLN()** until it returns a line that is shorter than 1000 characters. Note that end-

of-file checking is ignored, since `READLN()` returns an empty string at the end-of-stream.

32.11 Stream I/O in BRexx1.0b

BRexx contains a set of I/O which shows very close relations with the C programming language I/O library. In fact, you should consider consulting the C library documentation for in-depth documentation on this functionality.

BRexx contains a two-level naming scheme: in REXX, streams are referred to by a stream handle, which is an integer; in the operating system files are referred to by a filename, which is a normal string. The function `OPEN()` is used to bind a filename to a stream handle. However, BRexx I/O functions generally have the ability to get a reference either as a filename and a stream handle, and open the file if appropriate. However, if the name of a file is an integer which can be interpreted as a file descriptor number, it is interpreted as a descriptor rather than a name. Whenever you use BRexx and want to program robust code, always use `OPEN()` and the descriptor.

If a file is opened by specifying the name in a I/O operation other than `OPEN()`, and the name is an integer and only one or two higher than the highest current file descriptor, strange things may happen.

Five special streams are defined, having the pseudofilenames: `<STDIN>`, `<STDOUT>`, `<STDERR>`, `<STDAUX>`, and `<STDPRN>`; and are assigned pre-defined stream handles from 0 to 4, respectively. These refer to the default input, default output, and default error output, default auxiliary output, and printer output. The two last generally refer to the `COM1:` and `LPT1:` devices under MS-DOS. Either upper or lowercase letter can be used when referring to these four special names.

However, note that if any of these five special files are closed, they cannot be reopened again. The reopened file will be just a normal file, having the name e.g. `<STDOUT>`.

There is a few things you should watch out for with the special files. I/O involving the `<STDAUX>` and `<STDPRN>` can cause the `Abort`, `Retry`, `Ignore` message to be shown once for each character that was attempted read or written. It can be boring and tedious to answer `Ror` if the text string is long. If `A` is answered, BRexx terminates.

You should never write data to file descriptor 0 (`<STDIN>`), apparently, it will only disappear. Likewise, never read data to file descriptors 1 and 2 (`<STDOUT>` and `<STDERR>`), the former seem to terminate the program while the latter apparently just return the null string. Also be careful with reading from file descriptors 3 and 4, since your program may hang if no data is available.

`OPEN(file, mode)`

The `OPEN()` built-in function opens a file named by *file*, in mode *mode*, and returns an integer which is the number of the stream handle assigned to the file. In general, the stream handle is a non-negative integer, where 0 to 4 are pre-defined for the default streams. If an error occurred during the open operation, the value `-1` is returned.

The *mode* parameter specifies the mode in which the file is opened. It consists of two parts: the access mode, and the file mode. The access mode part consists of one single character, which can be `r` for read, `w` for write, and `a` for append. In addition, the `+` character can be appended to open a file in both read and write mode. The file mode part can also have one additional character which can be `t` for text files and `b` for binary files. The `t` mode is default.

The following combination of `+` and access mode are possible:

`r` is non-destructive open for reading; `w` is destructive open for write-only mode; `a` is non-destructive open for in-append-only mode, i.e. only write operations are allowed, and all write operations must be performed at the end-of-file; `r+` is non-destructive open for reading and writing; `w+` is destructive open for reading and writing; and `a+` is non-destructive open in append update, i.e. reading is allowed anywhere, but writing is allowed only at end-of-file. Destructive mode means that the file is truncated to zero length when opened.

In addition, the `b` and `t` characters can be appended in order to open the file in binary or text mode.

These modes are the same as under C, although the `t` mode character is strictly not in ANSI C. Also note that `r`, `w`, and `a` are mutually exclusive, but one of them must always be present. The mode `+` is optional, but if present, it must always come immediately after `r`, `w`, or `a`. The `t` and `b` modes are optional and mutually exclusive; the default is `t`. If present, `t` or `b` must be the last character in the mode string.

<code>open('myfile','w')</code>	7	perhaps
<code>open('no.such.file','r')</code>	-1	if non-existent
<code>open('c:tmp','r+b')</code>	6	perhaps

If two file descriptors are opened to the same file, only the most recently of them works. However, if the most recently descriptor is closed, the least recently starts working again. There may be other strange effects too, so try to avoid reopening a file that is already open.

CLOSE(*file*)

The `CLOSE()` built-in function closes a file that is already open. The parameter *file* can be either a stream handler returned from `OPEN()` or a filename which has been opened (but for which you do not know the correct stream handle).

The return value of this function seems to be the null string in all cases.

<code>close(6)</code>		if open
<code>close(7)</code>		if not open

close('foobar')		perhaps
-----------------	--	---------

EOF(*file*)

The EOF () built-in function checks the end-of-file state for the stream given by *file*, which can be either a stream descriptor or a filename. The value returned is 1 if the end-of-file status is set for the stream, and 0 if it is cleared. In addition, the value -1 is returned if an error occurred, for instance if the file is not open.

The end-of-file indicator is set whenever an attempt was made to read at least one character past the last character of the file. Note that reading the last character itself will not set the end-of-file condition.

eof(foo)	0	if not at eof
eof('8')	1	if at eof
eof('no.such.file')	-1	if file is not open

READ([*file*][, *length*])

The READ () built-in function reads data from the file referred to by the *file* parameter, which can be either a filename or a stream descriptor. If it is a filename, and that file is not currently open, then BRExx opens the file in mode `rt`. The default value of the first parameter is the default input stream. The data is read from and including the current position.

If the *length* parameter is not specified, a whole line is read, i.e. reading forward to and including the first end-of-line sequence. However, the end-of-line sequence itself is not returned. If the *length* parameter is specified, it must be a non-negative integer, and specifies the number of characters to read.

The data returned is the data read, except that if *length* is not specified, the terminating end-of-line sequence is stripped off. If the last line of a file contains a string untruncated by the end-of-string character sequence, then the end-of-file is implicitly interpreted as an end-of-line. However, in this case the end-of-file state is entered, since the end-of-stream was found while looking for an end-of-line.

read('foo')	one line	reads a complete line
read('foo',5)	another	reads part of a line
read(6)	error line	using a file descriptor
read()	hello there	perhaps, reads line from default input stream

```
WRITE([file][,[string][,dummy]])
```

The `WRITE ()` built-in function writes a string of data to the stream specified by the `file` parameter, or by default the default output stream. If specified, `file` can be either a filename or a stream descriptor. If it is a filename, and that file is not already open, it is opened using `wt` mode.

The data written is specified by the `string` parameter.

The return value is an integer, which is the number of bytes written during the operation. If the file is opened in text mode, all ASCII new line characters are translated into ASCII `CRLF` character sequences. However, the number returned is not affected by this translation; it remains independent of any text of binary mode. Unfortunately, errors while writing is seldom trapped, so the number returned is generally the number of character that was supposed to be written, independent of whether they were actually written or not.

If a third parameter is specified, the data is written as a line, i.e. including the end-of-line sequence. Else, the data is written as-is, without any end-of-line sequence. Note that with `BR`exx, the third parameter is considered present if at least the comma in front of it--the second comma--is present. This is a bit inconsistent with the standard operation of the `ARG ()` built-in function. The value of the third parameter is always ignored, only its presence is considered.

If the second parameter is omitted, only an end-of-line action is written, independent of whether the third parameter is present or not.

<code>write('bar','data')</code>	4	writes four bytes
<code>write('bar','data','nl')</code>	4+??	write a line
<code>write('bar','data',)</code>	4+??	same as previous

```
SEEK(file[,[offset][,origin]])
```

The `SEEK ()` built-in function moves the current position to a location in the file referred to by `file`. The parameter `file` can be either a filename (which must already be open) or a stream descriptor. This function does not implicitly open files that are not currently open.

The parameter `offset` determines the location of the stream and must be an integer. It defaults to zero. Note that the addressing of bytes within the stream is zero-based.

The third parameter can be any of `TOF`, `CUR`, or `EOF`, in order to set the reference point in which to recompute the `offset` location. The three strings refer to top-of-file, current position, and end-of-file, and either upper or lower case can be used. The default value is ???.

The return value of this function is the absolute position of the position in the file after the seek

operation has been performed.

The `SEEK ()` function provides a very important additional feature. Whenever a file opened for both reading and writing has been used in a read operation and is to be used in a write operation next (or vice versa), then a call to `SEEK ()` must be performed between the two I/O calls. In other words, after a read only a seek and reading may occur; after a write, only seeking and writing may occur; and after a seek, reading, writing, and seeking may occur.

32.12 Problems with Binary and Text Modes

Under the MS-DOS operating system, the end-of-line character sequence is `<CR><LF>`, while in C, the end-of-line sequence is only `<LF>`. This opens for some very strange effects.

When an MS-DOS file is opened for read in text mode by `BRexx`, all `<CR><LF>` character sequences in file data are translated to `<LF>` when transferred into the C program. Further, `BRexx`, which is a C program, interprets `<LF>` as an end-of-line character sequence. However, if the file is opened in binary mode, then the first translation from `<CR><LF>` in the file to `<LF>` into the C program is not performed. Consequently, if a file that really is a text file is opened as a binary file and read line-wise, all lines would appear to have a trailing `<CR>` character.

Similarly, `<LF>` written by the C program is translated to `<CR><LF>` in the file. This is always done when the file is opened in text mode. When the file is opened in binary mode, all data is transferred without any alterations. Thus, when writing lines to a file which is opened for write in binary mode, the lines appear to have only `<LF>`, not `<CR><LF>`. If later opened as a text file, this is not recognized as an end-of-line sequence.

Example: Differing end-of-lines

Here is an example of how an incorrect choice of file type can corrupt data. Assume `BRexx` running under MS-DOS, using `<CR><LF>` as an end-of-line sequence in text files, but the system calls translating this to `<LF>` in the file I/O interface. Consider the following code.

```
file = open('testfile.dat', 'wt')      /* text mode */
call write file, '45464748'x, 'dummy'  /* i.e. 'abcd' */
call write file, '65666768'x, 'dummy'  /* i.e. 'ABCD' */
call close file
file = open('testfile.dat', 'rb')      /* binary mode */
say c2x(read(file))                   /* says '454647480D' */
say c2x(read(file))                   /* says '656667680D' */
call close file
```

Here, two lines of four characters each are written to the file, while when reading, two lines of five characters are read. The reason is simply that the writing was in text mode, so the end-of-line character sequence was `<CR><LF>`; while the reading was in binary mode, so the end-of-line character sequence was just `<LF>`. Thus, the `<CR>` preceding the `<LF>` is taken to be part of the line during the read.

To avoid this, be very careful about using the correct mode when opening files. Failure to do so will

almost certainly give strange effects.

Extensions

This chapter describes how extensions to Regina are implemented. The whole contents of this chapter is specific for Regina.

33 Why Have Extensions

Why do we need extensions? Well, there are a number of reasons, although not all of these are very good reasons:

- Adaptation to new environments may require new functionality in order to easily interface to the operating system.
- Extending the language with more power, to facilitate programming.
- Sometimes, a lot of time can be saved if certain assumptions are met, so an extension might be implemented to allow programmers to take shortcuts.
- When a program is ported from one platform to another, part of the code may depend on non-standard features not available on the platform being ported to. In this situation, the availability of extensions that implement the feature may be of great help to the programmer.
- The implementor had some good idea during development.
- Backwards compatibility.

Extensions arise from holes in the functionality. Whether they will survive or not depends on how they are perceived by programmers; if perceived as useful, they will probably be used and thus supported in more interpreters.

34 Extensions and Standard REXX

In standard REXX, the `OPTIONS` instruction provides a "hook" for extensions. It takes any type of parameters, and interprets them in a system-dependent manner.

The format and legal values of the parameters for the `OPTIONS` instruction is clearly implementation dependent [TRL2, p62].

35 Specifying Extensions in Regina

In Regina there are three levels of extensions. Each independent extension has its own name. Exactly what an independent extension is, will depend on the viewer, but a classification has been done, and is listed at the end of this chapter.

At the lowest level are these "atomic" extensions. Then there are some "meta-extensions". These are collections of other extensions which belong together in some manner. If you need the extension for

creating "buffers" on the stack, it would be logical to use the extension to remove buffers from the stack too. Therefore, all the individual extensions for operations that handle buffers in the stack can be named by such a "meta-extension". At the end of this chapter, there is a list of all the meta-extensions, and which extension they include.

At the top is "standards". These are a set of extensions that make the interpreter behave in a fashion compatible with some standard. Note that "standard" is used very liberally, since it may refer to other implementations of REXX. However, this description of how the extensions are structured is only followed to some extent. Where practical, the structure has been deviated.

36 The Trouble Begins

There is one very big problem with extensions. If you want to be able to turn them on and off during execution, then your program has to be a bit careful.

More and more REXX interpreters (including Regina) seem to do parsing when the interpreter is started. The "old" way was to postpone the parsing of each clause until it was actually executed. This leads to the problem mentioned.

Suppose you want to use an extension that allows a slightly different syntax, for the sake of the argument, let us assume that you allow an expression after the `SELECT` keyword. Also assume that this extension is only allowed in extended mode, not in "standard mode". However, since Regina parses the source code only once (typically at the start of the program), the problem is a catch-22: the extension can only be turned on after parsing the program, but it is needed before parsing. This also applies to a lot of other REXX interpreters, and all REXX compilers and preprocessors.

If the extension is not turned on during parsing, it will generate a syntax error, but the parsing is all done before the first clause is executed. Consequently, this extension cannot be turned on during execution, it has to be set before the parsing starts.

Therefore, there are two alternative ways to invoke a set of extensions; neither of which is implemented in Regina.

- It can be invoked by using the `-e` option to the interpreter. The word following the option is the extension or standard to invoke. Multiple `-e` options can be specified.
- It can be invoked by setting the environment variable `REXXEXTS`, which must be a string of the same format as the parameters to the `OPTIONS` clause.

37 The Format of the `OPTIONS` clause

The format of the `OPTIONS` clause is very simple, it is followed by any REXX string expression, which is interpreted as a set of space-separated words. The words are treated strictly in order from left to right, and each word can change zero or more extension settings.

Each extension has a name. If the word being treated matches that name, that extension will be turned on. However, if the word being treated matches the name of an extension but has the prefix `NO`, then

that extension is turned off. If the word does not match any extensions, then it is simply ignored, without creating any errors or raising any conditions.

Example: Extensions changing parsing

An example of this is the `LINES` BIF. In the following piece of code the same BIF returns different data:

```
/* file 'aa' contains 5 lines */
options FAST_LINES_BIF_DEFAULT
do i=1 to 2
    if i=2 then OPTIONS NOFAST_LINES_BIF_DEFAULT
    say lines('aa')
end
```

In the first iteration of the loop, `LINES('aa')` returns 1, indicating that there is at least 1 line remaining in the stream 'aa'. However, in the second iteration of the loop, `LINES('aa')` will return 5, indicating that there are 5 lines remaining in the stream.

Regina's frequent usage of extensions may slow down execution. To illustrate how this can happen, consider the `OPEN()` extra built-in function. As this is an extension, it might be dynamically included and excluded from the scope of currently defined function. Thus, if the function is used in a loop, it might be in the scope during the first iteration, but not the second. Thus, Regina cannot cache anything relating to this function, since the cached information may be outdated later. As a consequence, Regina must look up the function in the table of functions for each invocation. To avoid this, you can set the extension `CACHEEXT`, which tells Regina to cache in if whenever possible, without regard to whether this may render useless later execution of `OPTIONS`.

38 The Fundamental Extensions

Here is a description of all "atomic" extensions in Regina:

[AREXX_BIFS]

This option allows the user to enable or disable the `AREXX` BIFs introduced into Regina 3.1. The default is `AREXX_BIF` on Amiga and AROS, but `NOAREXX_BIF` on all other platforms.

[AREXX_SEMANTICS]

With the introduction of `AREXX` BIFs into Regina 3.1, differences in the semantics of a number of BIFs resulted. These BIFs that differ between *Standard Regina* and *AREXX* are `OPEN()`, `CLOSE()` and `EOF()`. This `OPTION` specifies that the `AREXX` semantics be used for these BIFs. The default is to use Regina semantics for these BIFs.

[BUFTYPE_BIF]

Allows calling the built-in function `BUFTYPE()`, which will write out all the contents of the stack, indicating the buffers, if there are any. The idea is taken from VM/CMS, and its command named `BUFTYPE`.

[CACHEEXT]

Tells Regina that information should be cached whenever possible, even when this will render future execution of the `OPTIONS` instruction useless. Thus, if you use e.g. the `OPEN ()` extra built-in function, and you set `CACHEEXT`, then you may experience that the `OPEN ()` function does not disappear from the current scope when you set the `NOOPEN_BIF` extension.

Whether or not removal of an extension really does happen is unspecified when `CACHEEXT` has been called at least once. Effectively, info is cached during the period when `CACHEEXT` was in effect might not be "uncached". The advantage of `CACHEEXT` is efficiency when you do not need to do a lot of toggling of some extension.

[DESBUF_BIF]

Allows calling the built-in function `DESBUF ()`, to remove all contents and all buffers from the stack. This function is an idea taken from the program by the same name under VM/CMS.

[DROPBUF_BIF]

Allows calling the built-in function `DROPBUF ()`, to remove one of more buffers from the stack. This function is an idea taken from the program by the same name under VM/CMS.

[EXT_COMMANDS_AS_FUNCS]

When Regina resolves an expression to a function, and that function is not a built-in or a registered external function, Regina attempts to execute the function as an operating system command. With `NOEXT_COMMANDS_AS_FUNCS` set, Regina will return error 43; "Routine not found". `EXT_COMMANDS_AS_FUNCS` is the default.

[FAST_LINES_BIF_DEFAULT]

The `LINES () BIF` in versions of Regina prior to 0.08 returned the actual number of lines available in a stream. Since then, the `LINES () BIF` has been changed to only return 0 or 1. This was done for two reasons. First, it is faster, and secondly, the `ANSI` standard allows for an option to return the actual number of lines. This `OPTION` is for backwards compatibility with programs written assuming the prior behaviour of the `LINES () BIF`. `FAST_LINES_BIF_DEFAULT` is the default.

[FLUSHSTACK]

Tells the interpreter that whenever a command clause instructs the interpreter to flush the command's output on the stack, and simultaneously take the input from the stack, then the interpreter will not buffer the output but flush it to the real stack before the command has terminated. That way, the command may read its own output. The default setting for Regina is not to flush, i.e. `NOFLUSHSTACK`, which tells the interpreter to temporarily buffer all output lines, and flush them to the stack when the command has finished.

[INTERNAL_QUEUES]

Regina implements multiple named queues both as part of the interpreter, and as an external resource. The use of the `RXQUEUE () BIF`, will make Regina use the external queueing mechanism. This `OPTION` allows the exclusive use of Regina's internal queueing mechanism. `NOINTERNAL_QUEUES` is the default.

[LINEOUTTRUNC]

This option tells the interpreter that whenever the `LINEOUT ()` built-in function is executed for a persistent file, the file will be truncated after the newly written line, if necessary. This is the default setting of Regina, unless your system does not have the `ftruncate ()` system call. The complement option is `NOLINEOUTTRUNC`.

[MAKEBUF_BIF]

Allows calling the built-in function `MAKEBUF ()`, to create a buffer on the stack. This function is an idea taken from a program by the same name under VM/CMS.

[PRUNE_TRACE]

Makes deeply nested routines be displayed on one line. Instead of indenting the trace output a very long line (possibly wrapping over several lines on the screen). It displays `[. . .]` at the start of the line, indicating that part of the whitespace of the line has been removed.

[REGINA_BIFS]

This OPTION allows the user to turn off any non-`ANSI` extension BIFs. The default is `REGINA_BIFS`.

[STDOUT_FOR_STDERR]

All output that `Regina` would normally write to `stderr`, such as `TRACE` output and errors, are written to `stdout` instead. This is useful if you need to capture `TRACE` output and normal output from `SAY` to a file in the order in which the lines were generated. The default is `NOSTDOUT_FOR_STDERR`.

[STRICT_ANSI]

This OPTION results in interpretation of a program to strict `ANSI` standards, and will reject any `Regina` extensions. `NOSTRICT_ANSI` is the default.

[STRICT_WHITE_SPACE_COMPARISON]

This OPTION specifies if `ANSI` rules for non-strict comparisons are applied. Under `ANSI`, when doing non-strict comparisons, only the space character is stripped from the two comparitors. Under `Regina`'s default behaviour, all whitespace characters are stripped. `NOSTRICT_WHITE_SPACE_COMPARISON` is the default.

[TRACE_HTML]

This OPTION generates HTML `<PRE>` and `</PRE>` tags around `TRACE` output, to enable tracing from within CGI scripts. The default is `NOTRACE_HTML`.

Note: `OPEN_BIF`, `FIND_BIF`, `CLOSE_BIF` and `FILEIOOPTION` have been removed in `Regina 3.1`

39 Meta-extensions

[BUFFERS]

Combination of `BUFTYPE_BIF`, `DESBUF_BIF`, `DROPBUF_BIF` and `MAKEBUF_BIF`.

40 Semi-standards

[CMS]

A set of extensions that stems from the VM/CMS operating system. Basically, this includes the most common extensions in the VM/CMS version of `REXX`, in addition of some functions that perform tasks normally done with commands under VM/CMS.

[VMS]

A set of interface functions to the VMS operating system. Basically, this makes the `REXX` programming under VMS as powerful as programming directly in DCL.

[UNIX]

A set of interface functionality to the Unix operating system. Basically, this includes some

functions that are normally called as commands when programming Unix shell scripts. Although it is possible to call these as commands in **Regina**, there are considerable speed improvements in implementing them as built-in functions.

41 Standards

The following table shows which options are available in different **REXXL** language levels, and the default settings applicable for **Regina**.

[ANSI]

REXX Language level 5.0, as described in [ANSI]. This can be set with the **STRICT_ANSI** **OPTION** in **Regina**.

[REGINA]

REXX Language level 5.0, plus extensions, as implemented by **Regina 3.1** and above.

[SAA]

REXX Language level ??, as defined by IBM's System Application Architecture [SAA].

[TRL1]

REXX Language level 3.50, as described in [TRL1].

[TRL2]

REXX Language level 4.00, as described in [TRL2].

Option	ANSI	REGINA	SAA	TRL1	TRL2
AREXX_BIFS	no	yes	no	no	no
AREXX_SEMANTICS	no	no	no	no	no
BUFTYPE_BIF	no	yes	??	no	no
CACHEEXT	no	no	no	no	no
DESBUF_BIF	no	yes	??	no	no
DROPBUF_BIF	no	yes	??	no	no
EXT_COMMANDS_AS_FUNCS	no	yes	??	no	no
FAST_LINES_BIF_DEFAULT	no	yes	??	no	no
FLUSHSTACK	no	no	??	no	no
INTERNAL_QUEUES	no	no	??	no	no
LINEOUTTRUNC	no	yes	??	no	no
MAKEBUF_BIF	no	yes	??	no	no
PRUNE_TRACE	no	yes	no	no	no
REGINA_BIFS	no	yes	no	no	no
STDOUT_FOR_STDERR	no	no	no	no	no
STRICT_ANSI	no	no	no	no	no
STRICT_WHITE_SPACE_COMPARISON	no	no	no	no	no
TRACE_HTML	no	no	no	no	no

The Stack

In this chapter, the stack and operations manipulating the stack are discussed. Since the stack is external to the REXX language, there are large differences between implementations with respect to the stack. These differences are attempted described in the latter part of this chapter.

Another goal of this chapter is to try to describe both the "real" standards and some of the most commonly used de facto standards related to stack operation. Where something is not a part of any defined standard, this is clearly labeled. Also, some liberties have been taken in order to create a coherent vocabulary on a field where very little standardization has taken place.

42 Background and history

In the various definitions of REXX, there are numerous references to the "stack" (often called the "external data queue", or just the "queue"). It is a structure capable of storing information, but it is not a part of the REXX language itself. Rather, it is a part of the external environments supporting a REXX implementation.

Originally, the reference to the stack was introduced into REXX because of the strong binding between REXX and IBM mainframes in the early history of REXX [BMARKS]. Most (all?) of the operating systems for these machines support a stack, and many of their script programming idioms involve the stack. Therefore, it was quite natural to introduce an interface to the stack into REXX, and consequently today many of the programming paradigms of REXX involve a stack.

Unfortunately, this introduced an element of incompatibility into REXX, as the stack is not in general supported for other operating systems. Consequently, REXX implementors often must implement a stack as well as the core REXX interpreter. Since no authoritative definition of the stack exists, considerable differences between various implementations. Ironically, although the stack was introduced to help communication between separate programs, the interpreter-specific implementations of stacks may actually be a hindrance against compatibility between different interpreters.

The stack may have "seemed like a good idea at the time", but in hindsight, it was probably a bad move, since it made REXX more dependent on the host operating system and its interfaces.

43 General functionality of the stack

This section describes the functionality generally available in implementations of stacks. The basic functionality described here will be complemented with information on specific implementations later. Unless explicitly labeled otherwise, this functionality is available in all standards treated in this documentation.

43.1 Basic functionality

Below is listed the general functionality of the stack, in order of decreasing compatibility. I.e. the functionality listed first is more likely to be a part of all implementations than the ones listed at the end of the list.

- The stack is a data structure, which strings can either be inserted into or extracted from. The strings in the stack are stored in a linear order. Extraction and insertion works at a granularity of a complete string, i.e. it is not possible to insert or extract parts of a string.
- The stack has two ends: a top and a bottom. New strings can be inserted into the stack in both ends, but strings can only be extracted from the top of the stack.
- There exists a way of counting the number of strings currently stored in the stack.

A stack is often compared with the pile of plates you often find in a canteen. It allows you to either add new plates at the top of the pile or take old plates from the top. When a plate is taken from the pile, it will be the most recently added plate (that is still present). Stack operating in REXX work the same way, although there also allow "plates" to be added to the bottom of the pile.

- There might be an implementation-specific limit on the length and number of strings stored in the stack. Ideally, the maximum length will be fairly large, at least 2^{16} , although some implementations are likely to enforce shorter limits. Similarly, there might be a limit on the number of strings that can be simultaneously stored in the stack. Ideally, there should be no such limit.
- It is natural that there are limits imposed on the amount of memory occupied by the strings in the stack. Some implementations are likely to reserve a fixed (but perhaps configurable) amount of memory for this purpose while others can dynamically re-size the stack as long as enough memory is available.
- Some implementations might restrict the set of characters allowed in strings in the stack, although ideally, all characters should be allowed, even characters normally used for end-of-line or end-of-string.

This documentation uses the term "string", while "line" is in common use elsewhere. The term is used because the strings in the stack are not inherently interpreted as lines (having an implied end-of-line), only as a string.

Note that the stack itself is not a part of REXX, only the parts which interface to the stack.

Example: Using the stack to transfer parameters

This is a common REXX idiom used in several situations for special parameter passing. The following code illustrates its use:


```

do i=1 to 10          /* for each parameter string      */
    queue string.1    /* put the string on the stack */
*/
end
call subrout 10        /* call the subroutine        */
exit

subrout: procedure    /* the definition of the subroutine */
    do j=1 to arg(1)  /* for each parameter passed */
        parse pull line.j /* retrieve the parameter
    */
    end
    ...                /*do something with the
parameters*/
    return

```

In this example, ten parameter strings are transferred to the subroutine SUBROUT. The parameters are stored in the stack, and only the number of parameters are transferred as a "real" argument.

There are several advantages: first, one avoids problems related to exposing variable names. Since the data is stored on the stack, there is no need to refer to the variable names and bind the variables in the subroutine to variables in the caller routine. In [TRL1], indirect reference to variables in PROCEDURE EXPOSE is illegal, and this method circumvents the problem.

Two other ways around this problemistouse INTERPRET for the PROCEDURE EXPOSE instruction in order to dynamically determine which variables to expose; or to use the VALUE () built-in function (with its two first parameters). The former is incompatible with TRL2, while the latter is incompatible with TRL1. Using the stack can solve the problem in a fashion compatible with both standards. Anyway, if the called routine is an external routine, then exposing does not work, so using the stack to transfer values may be the only solution.

Another advantage of this idiom; TRL only requires implementation to support 10 parameters for subroutines. Although there are no reasons why an implementation should set a limit for the number of parameters a routine can get, you should use another mechanism than arguments when the number of strings is greater than 10. Using the stack fixes this.

43.2 LIFO and FIFO stack operations

As already mentioned, the stack is a linear list of strings. Obviously, this list has two ends. Strings can only be extracted from one end, while strings can be added to both ends.

If a set of new strings are added to the same end as they are later extracted from, the strings will be extracted in the reversed order with respect to the order in which they were added. This is called stacking "LIFO", which means "last-in-first-out", meaning that the last string stacked, will be the first string extracted, i.e. reversal of the order.

Similarly, when a set of strings are stacked in the end opposite to the end which they are later extracted from, they will be extracted in the same order in which they were stacked. This is referred to as "FIFO" stacking, meaning "first-in-first-out".

The FIFO method of stacking is also sometimes referred to as "queueing", while the LIFO method is sometimes referred to as "stacking" or "pushing".

43.3 Using multiple buffers in the stack

The concept of buffers and everything directly related to buffers lay without the domain of standard REXX. Thus, this section describes a de facto standard.

Not that Regina supports multiple buffers only in internal stacks.

Some implementations support "buffers", which are a means of focusing on a part of the stack. When creating a new buffer, the old contents of the stack is somewhat insulated from the effects of stack operations. When the buffer is removed, the state of the old buffer is restored, to some extent: Whenever a string is read from the stack, and the topmost buffer on the stack is empty, then that buffer will be destroyed. Consequently, if this situation has arisen, dropping buffers will not restore the state of the stack before the buffer was created.

The functionality of buffers, and their effect on other stack operations may differ considerably between implementations.

Whenever a queueing operation is performed (e.g. by the `QUEUE` instruction), then the new string is inserted into the bottom of the topmost buffer, not the bottom of the stack. This is the same if the stack has no buffers, but else, the outcome of the queueing operation can be very different.

With IBM mainframe operating systems like CMS, buffers can be inserted on the top of the stack. To perform buffer operations, operating system commands are used. It may be instructive to list the buffer operations of CMS:

[DESBUF]

Removes all strings and buffers from the stack, and leaves the stack clean and empty. It is often used instead of repeated calls to `DROPBUF`. It always returns the value zero.

[DROPBUF]

Removes zero or more buffers from the stack. It takes one parameter which can be omitted, and which must be an integer position if specified, and is the assigned number of the bottom-most buffer to be removed, i.e. that buffer and all buffers above it (and of course, all the strings in these buffers) are to be removed. If the parameter is not specified, only the topmost buffer is removed. The return value is always zero, unless an error occurred.

[MAKEBUF]

Makes a new buffer on the stack, starting at the current top of the stack. The return code (as stored in the special variable `RC`) is the number of buffers currently on the stack after the new buffer has been added. Obviously, this will be a positive integer. This program takes no parameters.

One might regard a buffer as a sort of bookmark, which is inserted into the stack, so that a subsequent `DROPBUF` command can remove the stack down to a particular such bookmark.

When such a mark is located on the top of the stack, and a `PULL` instruction is executed, the buffer

mark is implicitly destroyed when the `PULL` instruction reads the string below the buffer mark. This is to say that a buffer can be destroyed by either a `DESBUFF` command, a `DROPBUF` command, or aread from the stack (by either the `PULL` or `PARSE PULL` instructions).

43.4 The zeroth buffer

Normally, data pushed on the stack is added to the top of the stack. When a stack contains only one buffer, the strings in that buffer are the strings stored above that buffer-mark. The strings below it are not part of the first buffer; instead, they are said to belong to the zeroth buffer.

Thus, all strings from the bottom of the stack, up till the first buffer mark (or the top of the stack if no buffer exists) is said to be the strings in the zeroth buffer. However, note that the zeroth buffer is only defined implicitly. Thus, it cannot really be removed by calling `DROP`; only the strings in the zeroth buffer are removed. Afterwards, the zeroth buffer will still contain all strings at the bottom of the stack, up till the first buffer mark (if existing).

Example: Process all strings in the stack

This is a common `REXX` idiom, where a loop iterates over all the strings currently in the stack, but otherwise leaves the stack untouched. Supposing the routine `PROCESS()` exists, and dotoprocessing with its parameter and return the processed string:

```
do i=1 to 5                                /* just to fill the stack */
    push 'line #' i
end

do queued()                                /* foreach line in the stack */
    parse pull line                        /* fetch the line */
    /*
        queue process(line)                /* put back the processed line */
    */
end
```

Here, it is important to use `QUEUE` to put the strings back into the stack, not `PUSH`, else the loop will iterate the correct number of times, but only operate on the same data string. It is also important that the stack does not contain any buffers. Since `QUEUE` will insert into the bottom of the topmost buffer, the loop would iterate the correct number of times, but only on a part of the stack. Thus, the topmost part of the strings in the stack would be processed multiple times.

Example: How to empty the stack

The following short examples show how you can most easily empty the stack:

```

do i=1 to 5                /* Just to fill the stack */
    push 'line #' i
end

do queued()                /* For each line in the stack */
    pull                    /* Remove the line from the
stack */
end

```

This is trivially simple, but there are several interesting and subtle notes to make about this example. First, if the number of strings in the stack is likely to change, due to some external process, then the clauses should perhaps better be written as:

DO

```

do i=1 to 5                /* Just to fill the stack */
    push 'line #' i
end

do while queued()>0        /* While the stack is not empty */
    pull                    /* Remove a line from the stack
*/
end

```

This will in general mean more work for the interpreter, as it is now required to check the number of strings in the stack for each iteration, while for the previous code fragment, the number of strings is only checked once. Another point is that this might not remove all buffers from the stack. Suppose the zeroth buffer is empty, i.e. there exists an buffer which was put on the stack when the stack was empty. This buffer is removed in any of the following situations: calling `DESBUF`, calling `DROPBUF` (sometimes), or reading a string below the buffer mark. Since there are no strings below the buffer mark, pulling a string from the stack would make the interpreter read from the keyboard, and hang the interpreter.

Thus, the only "safe" way to remove the string and buffers from the stack, without side effects, is to call `DESBUF` or `DROPBUF`. On the other hand, if you only want to make sure that there are no strings in the buffer, the method described here is more suitable, since it is far more compatible (although possibly not so efficient). But anyway, buffers are not a compatible construct, so it does not matter so much.

43.5 Creating new stacks

The description of multiple stack operations in this section, is not part of standard REXX, nor is it implemented in Regina. Thus, this section describes a de facto standard and you may find that few implementations support these operations.

Just as the operations described above let the REXX programmer use multiple buffers within one stack, there exists another set of operations which let the programmer create multiple stacks. There is really nothing fancy about this, except that a command will swap the stack the interpreter correctly uses with another stack.

To the interpreter this is really equivalent to a situation where a command empties the current stack, and sets up a new stack. When one stack is empty, and the REXX program tries to read from the stack,

therequestwillnot"overflow"tothepreviousstack(asrequeststoanemptybuffer"overflows"tothe previousbuffer). Thus, the use of multiple stacks has even less direct impact on REXX interpreter than multiple buffers.

Here, it is instructive to list the commands operating multiple stacks that exist. This list has been taken from the MVS environment, according to [REXXSAA].

[DELSTACK]

Is used to remove the most current stack, and make the most recent of the saved stacks the current stack. When there are no saved stacks, the current stack is emptied.

[NEWSTACK]

Creates a new stack, which becomes the current stack. The old current stack is put on the top of the list of saved stacks, and can be retrieved as the current stack by a subsequent DELSTACK.

[QBUF]

Counts the number of buffers in the current stack, and return that number as the return value. A REXX program starting this command can retrieve this value as the special variable RC.

[QLELEM]

Counts the number of strings (i.e. elements) in the current stack, and return that value as the return value of the command. This value can be retrieved in REXX as the special variable RC. This operation is equivalent to the QUEUED() built-in function in REXX; it has been probably included for the benefit of other script languages that have less functionality than REXX.

[QSTACK]

Counts the number of stacks (including the current stack) and return the value as the return value from the command. This number can be retrieved in REXX as the special variable RC.

One can regard multiple buffers and stacks as two ways of insulating the stack; where multiple stacks are a deeper and more insulating method than buffers. Note that each stack can contain multiple buffers, while a buffer cannot contain any stacks. The term "hard buffers" has been used about multiple stacks, as opposed to normal buffers, which are sometimes called "soft buffers".

Also note that neither multiple stacks nor buffers are part of standard REXX, so you might come across implementations that support only multiple stacks, only buffers, or even none of them.

Example: Counting the number of buffers

In order to count the number of buffers on the stack, the following method can be used (Regina syntax has been used for buffer handling). This method is equivalent to the QBUF command described above.

```
buffers = makebuf() - 1
call dropbuf
```

This will store the number of buffers in the stack in the variable buffers. However, just as for the other examples using buffers, this example also suffers from the fact that buffer handling is fairly non-standard. Thus, you will have to adapt the code to whatever system you want to use.

44 The interface between REXX and the stack

As defined in TRL, the interface to the stack consists of the `PARSE PULL, PULL, PUSH, and QUEUE` instructions; and the `QUEUED()` built-in function.

There exists a binary interface to the stack in SAA, see the chapter on the SAA API interface. This interface consists of the `RXMSQex` handler and the `QUENAME` value of the `RXSHV_PRIV` request of the `RexxVariablePool()` function of the variable pool interface.

45 Strategies for implementing stacks

As mentioned, stacks are rarely a part of the operating system. Therefore, under most operating systems, REXX interpreters have to implement their own stacks. There are several strategies for doing this, some of which are listed below.

[In the operating system.]

This is of course "the right way" to do it. However, it requires that the definition of the operating system is such that stacks are supported. Currently, only IBM mainframe-based systems support stack, together with a few others systemsthathaveincludedstacksas a consequence of making REXX a main scripting language (Amiga and OS/2 come to mind).

[As a device driver.]

This is really just a variation of making the stack a part of the operating system. However, in some systems, drivers can be added very easily to the system. Drivers are often file system-based, in which case driver-based stack operations must operate on a file or pseudo-file. But for some systems, adding a driver requires much more profound changes, reconfiguration, and often system privileges. In all cases, drivers are likely to be very system specific.

[As a daemon.]

A "daemon" is a background process that does some housekeeping service, e.g. handling mail from remote systems. Implementing a stack as a daemon is only slightly simpler than using a driver, but the main idea is the same for both approaches.

[In the interpreter.]

Using this approach, the stack is built into the interpreter as a sort of extension. This is often the simplest way, since it requires very little coordination with other programs during run-time. The main problem is that the stack becomes private to the interpreter, so two interpreters cannot use the same stack; not even if they are two invocations of the same interpreter.

These items are listed in the order of how closely they are coupled to the operating system: the first items are very closely, while the last items are loosely coupled. The more closely coupled the implementation of a stack is coupled to the operating system, the better is the chance that several interpreters on the same system can communicate in a compatible way, using the stack.

There is room for several hybrid solutions, based on the four fundamental approaches. For instance, a built-in stack can also act as a daemon.

Regina supports the stack as both a daemon and internal to the interpreter.

Example: Command takes input from the stack

In the example above, the routine that is called takes its arguments from the stack. Similarly, command to an external environment can get their arguments in the same way. Here is an example of how to do it:

```
queue 'anonymous'           /* the username */
queue 'user@node'           /* the password */
queue 'dir'                 /* first command */
queue 'exit'                /* second command */
address command 'FTP flipper.pvv.unit.no'
```

Although this is very convenient in some situations, there is also considerable disadvantages with this method: There is no real interactive communication between the interpreter and the command; i.e. all input meant for the command must be set up before the command itself is invoked. Consequently, if one of the input lines to the command provokes an error, there is very little error handling facility. Commonly, such an error might start a cascade of errors, as the remaining input lines are likely to be invalid, or even be interpreted in a context different from what they were intended.

As with all commands involving the stack, it is important to push or queue the correct order.

Using this technique, a program can "fool" a command to do almost anything, by storing the correct input on the stack. However, there is a big disadvantage: Since the stack is implementation-dependent, it is not certain that a command will take its input from the stack. For some systems, this is the default, while for others systems, this is only possible through some explicit action. Some systems might not even allow commands to take their input from the stack at all.

Example: "Execing" commands

Many script programming languages can only execute commands while still running, or at most start a new command immediately after the termination (like the `exec ()` system call in Unix). However, the stack can be used on some systems to set up the system to execute one or more commands after the current script terminates. Here is an example:

```
push 'ls'                   /* finally execute 'ls' */
push 'who'                  /* then execute 'who' */
push 'pwd'                  /* first execute 'pwd' */
exit 0
```

Supposing that the system reads its commands from the stack if the stack is not empty, then this script will terminate after having set up the stack so that the three commands `pwd`, `who` and `ls` will be run in that sequence. Note the order, if `QUEUE` had been used, the order would be the opposite, which is perhaps more intuitive (assuming the topmost buffer is empty).

As with the example above, this too is only relevant for some systems, thus is not very compatible, and you should be careful when using it. It also suffers from the lack of interactivity, error handling, and the

importance of the order in which the strings are pushed or queued. For all practical reasons, this is just a special case.

Using the stack to "leave behind" command names and input only works for systems where command interpreters and commands read their input from the stack. This is in general true for IBM mainframe systems, but very few other systems.

46 Implementation of the stack in Regina

In Regina, the stack is implemented as both an integral, private part of the interpreter and as a cross-platform external stackable to be used by multiple clients on multiple machines. Internal stacks provide the obvious advantage of speed at the expense of data sharing. External stacks are considerably slower, but do enable data sharing between instances of Regina and/or other programs.

Regina supports the standard TRL (and ANSI) REXX stack interface functionality, like `PARSE PULL`, `PULL`, `QUEUE`, `PUSH`, the `QUEUED()` built-in function, and in future versions, support the SAAPI stack interface. These commands and functions operate on both the internal and external stacks.

46.1 Implementation of the internal stack in Regina 2.2

Whenever the REXX programmer wants to execute a command and let that command either flush the output to the internal stack, or read its input from the internal stack, this has to be arranged by the interpreter itself. In Regina this is normally done by prepending or appending certain terms to the command to be executed.

Consider the following command clauses for Regina:

```
'ls >LIFO'
'who >FIFO'
'LIFO> wc'
'LIFO> sort >FIFO'
```

For all these commands, the "piping" terms are stripped off the command string before the command is sent to the command interpreter of the operating system. Thus, the command interpreter only sees the commands `ls`, `who`, `wc`, and `sort`. The terms stripped off, are used as indicators of how the input and output is to be coupled with the stack. The use of input/output redirection as above is only available with the internal stack.

Note that it is important not to confuse the redirection of output to the stack and input from the stack in Regina with the redirection of the Unix shells. The two can be mixed in command lines, but are still two different concepts.

The first command will execute the `ls` command, and redirect the output from it to the stack in a LIFO fashion. The second executes the command `who` and redirects the output to the stack to, but in a FIFO fashion. The third command executes the `wc`, but lets the standard input of that command come from the stack. Actually, it is irrelevant whether `FIFO>` or `LIFO>` is used for input; the strings are read from the top of the stack in both cases. The fourth command is a plain `ps` command without any

redirection to or from the stack. The last command executes the
from the stack, and redirect the output to the stack.

`sort` program and lets it read its input

Regina allows a command to take both an input and an output "redirection" to a stack, as showed in the last example above. However, it also guarantees that the output is not available in the stack before the command has terminated. The output from the command is stored in a temporary stack, and flushed to the ordinary stack after the command is terminated. Thus, the command will not start to read its own output.

Note that this temporary buffering of command output is the default behavior, which might be set up to something different at your site.

In addition, you can change it through the `OPTIONS` instruction, by using either `FLUSHSTACK` or `BUFFERSTACK` as "parameters".

Note the difference between **Regina's** redirection and Unix redirection. In **Regina**, only the term `LIFO>` (when first in the command string), and the terms `>LIFO` and `>FIFO` (when last in the command string), will be interpreted as redirection directives. These terms will be stripped off the command string. All other redirection directives will be left untouched. If you should happen to need to redirect output from a Unix command to the file `FIFO` or `LIFO`, then you can append a space at the end or specify the file as `/FIFO` or `/LIFO`. That will make **Regina** ignore the redirection term.

Note that this particular form of redirection of command input and output will most probably disappear in future versions of **Regina**, where it will probably be replaced by an extended `ADDRESS` instruction.

In addition to the ANSI standard, there are a few extra built-in functions, which are supposed to provide compatibility with other **REXX** implementations, principally **CMSREXX**. These are `BUFTYPE`, `DESBUF`, `DROPBUF` and `MAKEBUF`. See the descriptions of these functions in the built-in functions section above.

46.2 Implementation of the external stack in Regina 2.2

The implementation of the external stack follows the model used by OS/2 **REXX**, but is implemented as an operating system daemon. This daemon is `rxstack`.

rxstack

Under most operating systems, `rxstack` is started from the operating system's startup process and terminates when the machine is shutdown. Under Windows NT/2000, it runs as a Service.

Communication between `rxstack` and **Regina** is done via TCP/IP sockets. Using sockets as the IPC mechanism on a local machine is somewhat slow compared to other mechanisms such as shared memory or named pipes. It does however enable operation between machines on different operating systems to function seamlessly.

The full syntax of the `rxstack` command is:

`rxstack[switch]`

switch	is one of the following switches	
-install	install the NT Service; RexxStack- Windows NT/2000 only	
-remove	remove the NT Service; RexxStack- Windows NT/2000 only	
-run	runs rxstack in a command prompt- Windows NT/2000 only	
-d	run rxstack as a daemon- Unix only	
-k	kills (stops) rxstack- subject to being a valid killer- see	Security of External Queues

To stop rxstack, the process can be killed with a SIGINT or SIGTERM or by running rxstack with the **-k** switch.

rxqueue

To allow non- REXX programs to interface to the rxstack daemon, a companion program; rxqueue, is provided. rxqueue communicates with non- REXX programs via its stdin and stdout.

Consider the following equivalents for Regina's internal and external stack

'ls >LIFO'	'ls rxqueue /lifo'
'who >FIFO'	'who rxqueue /fifo'
'LIFO> wc'	rxqueue /pull wc'
'LIFO> sort >FIFO'	rxqueue /pull sort rxqueue /fifo'

The full syntax of the rxqueue command is:

rxqueue[queue][switch]

queue	is a Regina external queue name—see the next section for structure. If no queue is specified, rxqueue uses the queue name; SESSION	
switch	is one of the following switches—see OS/2 REXX	
	/fifo	queue lines from stdin LIFO onto the
queue	/lifo	queue lines from stdin FIFO onto the
queue	/clear	remove all lines from the queue
	the following switches are Regina extensions	
	/queued	return the number of lines on the queue
	/pull	pull all lines from the queue and display on stdout

rxqueue Built-in Function

REXX programs communicate with rxstack via the normal queueing mechanisms of QUEUE, PUSH, PULL and QUEUED(). These commands operate on the current queue and have no mechanism for changing the queue to use. This is where RXQUEUE() is used. Its primary purpose is to control the queue that the remainder of the REXX program operates on.

QueueNames

To enable the use of the REXX stack as a cross-platform, multi-machine IPC, the naming conventions adopted by OS/2 REXX have been modified. As OS/2 REXX queues are local to a single machine, queue names have no structure. To enable identification of queues on different machines, some structure must be built into external queue names on Regina. An external queue name on Regina has the following format:

[queue][@machine[:port]]

The components of the queue name are:

queue	the name of the queue. The only criteria for the name is that it contains none of the following characters: @, ., or :. The queue component can be blank, when specifying the default queue on a specified machine.
machine	the machine that hosts the specified queue. This can either be a standard IPv4 IP address or a machine name that can be resolved to a standard IPv4 IP address. The machine name is optional, and defaults to 127.0.0.1
port	The port number that rxstack on machine is listening to. The default port number for rxstack is 5757.

When referring to queues on the local machine, the machine and port components need not be specified. The behaviour of the external stack is then the same as for OS/2 REXX, with the exception that the queues on the local machine can still be manipulated by Regina on another machine.

Some examples may make this clearer. TBD

Security of External Queues

(Not implemented yet)

A daemon process like rxstack, waiting on a TCP/IP socket for anyone to connect to and use is open to abuse. To reduce the openness of rxstack, it uses a security mechanism much like the Unix hosts.allow and hosts.deny files is used to control access to rxstack.

Environment Variables

RXQUEUE
RXSTACK

Interfacing Rexx to other programs

This chapter describes an interface between a REXX interpreter and another program, typically written in C or another high level, compiled language. It is intended for application programmers who are implementing REXX support in their programs. It describes the interface known as the REXXSAA API.

47 Overview of functions in SAA

The functionality of the interface is divided into some main areas:

- Subcommand handlers
which trap and handle a command to an external environment.
- External function handlers
extend the REXX language with external functions
- Interpreting
REXX scripts, either from a disk file, or from memory.
- Variable interface
which makes it possible to access the variables in the interpreter, and allows operations like setting, fetching and dropping variables.
- System exits
which are used to hook into certain key points in the interpreter while it executes a script.
- External Queue interface
which allows access to Regina's external queuing mechanism.
- Macro space functions
which are used to load and save external macros into Regina's macro space for faster execution.
- Memory Allocation functions
which provide for platform-independent memory allocating/deallocation functions.
- Callback functions
which are used to allow the API program to execute a procedure within the running script.

In the following section each of these areas are described in detail, and a number of brief but complete examples are given at the end of the chapter.

The description is of a highly technical nature, since it is assumed that the reader will be an application programmer seeking information about the interface. Therefore, much of the content is given as prototypes and C styled data type definitions. Although this format is cryptic for non-C programmers, it will convey exact, compact, and complete information to the intended readers. Also, the problems with ambiguity and incompleteness that often accompany a descriptive prose text are avoided.

47.1 Include Files and Libraries

All the C code that uses the REXX application interface, must include a special header file that contains the necessary definitions. This file is called `rexxsaa.h`. Where you will find this file, will depend on your system and which compiler you use.

Also, the interface part between the application and the REXX interpreter may be implemented as a library, which you link with the application using the functions described in this chapter. The name of this library, and its location might differ from system to system. Under Unix, this library can be implemented as a static (`libregina.a`) or dynamic library (`libregina.[so|sl]`). Under other platforms Regina is also implemented as a static or dynamic library.

47.2 Preprocessor Symbols

Including a header file ought to be enough; unfortunately, that is not so. Each of the domains of functionality listed above are defined in separate *sections* in the `rexxsaa.h` header file. In order for these to be made available, certain preprocessor symbols have to be set. For instance, you have to include the following definition:

```
#define INCL_RXSHV
```

in order to make available the definitions and data types concerning the variable pool interface. The various definitions that can be set are:

- **INCL_RXSUBCOM**
Must be defined in order to get the prototypes, data types and symbols needed for the subcommand interface of the API.
- **INCL_RXFUNC**
Must be defined in order to get the prototypes, data types and symbols needed for the external function interface of the API.
- **INCL_RXSYSEXIT**
Must be defined in order to get the prototypes, data types, and symbols needed for the system exit functions.
- **INCL_RXSHV**
Must be set in order to get the prototypes, symbols and data type definitions necessary to use the REXX variable pool.
- **INCL_RXQUEUE**
Must be set in order to get the prototypes, symbols and data type definitions necessary to use the REXX external queues.
- **INCL_RXMACRO**
Must be set in order to get the prototypes, symbols and data type definitions necessary to use the REXX macro space interface of the API.

47.3 Data structures and data types

In this section, some data structures and data types relevant to the application interface to REXX are defined and described. The data types defined are:

- **RXSTRING**
Holds a REXX string.
- **RXSYSEXIT**
Holds a definition of a system exit handler. Used when starting a REXX script with `RexxStart()`, and when defining the system exit handlers.

The data types used in the SAA API are defined in `rexxsaa.h`. They are:

```

typedefcharCHAR;
typedefshortSHORT;
typedeflongLONG;
typedefchar*PSZ;
typedefCHAR*PCHAR;
typedefSHORT*PSHORT;
typedefLONG*PLONG;
typedefunsignedcharUCHAR;
typedefunsignedshortUSHORT;
typedefunsignedlongULONG;
typedefUSHORT*PUSHORT;
typedefchar*PCH;
typedefunsignedchar*PUCHAR;
typedefvoidVOID;
typedefvoid*PVOID;
typedefULONGAPIRET;
typedefAPIRET(APIENTRY*PFN)();

```

One other item needs mentioning; `APIENTRY`. This value is used to specify the linkage type on OS/2 and Win32 platforms. It is assumed that this value `#defined` by inclusion of compiler-specific header files in `rexxsaa.h`. Under Unix, this is `#defined` to nothing.

47.3.1 The `RXSTRING` structure

The SAA API interface uses `Rexxstring` which are stored in the structure `RXSTRING`. There is also a datatype `PRXSTRING`, which is a pointer to `RXSTRING`. Their definitions are:

```

typedefstruct{
    unsignedchar*strptr; /*Pointer to string contents*/
    unsignedlongstrlength; /*Length of string*/
}RXSTRING;

typedefRXSTRING*PRXSTRING;

```

The `strptr` field is a pointer to an array of characters making up the contents of the `Rexxstring`, while `strlength` holds the number of characters in that array.

Unfortunately, there are some inconsistencies in naming of various special kinds of strings. In `REXX` (TRL), a `'nullstring'` is a string that has zero length. On the other hand, the SAA API operates with two kinds of special strings: *nullstrings* and *zerolengthstrings*. The latter is a string with zero length (equals nullstrings in `REXX`), while the former is a sort of *undefined* or *empty* string, which denotes a string without a value. The *nullstrings* of SAA API are used to denote unspecified values (e.g. a parameter left out in a subroutine call). In this chapter, when the terms *nullstrings* and *zerolengthstrings* are italicized, they refer to the SAA API style meaning.

A number of macros are defined, which simplifies operation on `RXSTRING`s for the programmer. In the list below, all parameters called `x` are of type `RXSTRING`.

- **MAKERXSTRING(x,content,length)**
The parameter **content** must be a pointer to **char**, while **length** is integer. The **x** parameter will be set to the contents and length supplied. The only operations are assignments; no new space is allocated and the contents of the string is not copied.
- **RXNULLSTRING(x)**
Return true only if **x** is a *null string*.
i.e. **x.strptr** is **NULL**.
- **RXSTRLEN(x)**
Return the length of the string **x** as an unsigned long. Zero is returned both when **x** is a *null string* or a *zero length string*.
- **RXSTRPTR(x)**
Return a pointer to the first character in the string **x**, or **NULL** if **x** is a *null string*. If **x** is a *zero length string*, and non-**NULL** pointer is returned.
- **RXVALIDSTRING(x)**
Return true only if **x** is neither a *null string* nor a *zero length string*.
i.e. **x** must have non-empty contents.
- **RXZEROLENSTRING(x)**
Return true only if **x** is a *zero length string*.
i.e. **x.strptr** is non-**NULL**, and **x.strlength** is zero.

These definitions are most likely to be defined as preprocessors or macros, so you should never *call* them with *parameters* having any side effects. Also note that at least **MAKERXSTRING()** is likely to be implemented as two statements, and might not work properly if following e.g. an **if** statement. Check the actual definitions in the **rexxsaa.h** header file before using them in a fancy context.

One definition of these might be (don't rely on this to be the case with your implementation):

```
#define MAKERXSTRING(x,c,l)((x).strptr=(c),(x).strlength=(l))
#define RXNULLSTRING(x) (!(x).strptr)
#define RXSTRLEN(x)((x).strptr?(x).strlength:0UL)
#define RXSTRPTR(x)((x).strptr)
#define RXVALIDSTRING(x)((x).strptr&&(x).strlength)
#define RXZEROLENSTRING(x)((x).strptr&&!(x).strlength)
```

Note that these definitions of strings differ from the normal definition in C programs; where a string is an array of characters, and its length is implicitly given by a terminating **ASCII NUL** character. In the **RXSTRING** definition, a string can contain any character, including an **ASCII NUL**, and the length is explicitly given.

47.3.2 The **RXSYSEXIT** structure

This structure is used for defining which system exit handlers are to handle which system exits. The two relevant data types are defined as:

```
typedef struct {
    unsigned char* sysexit_name;
    short sysexit_code;
```

```
}RXSYSEXIT;
```

```
typedefRXSYSEXIT*PRXSYSEXIT;
```

In this structure, `sysexit_name` is a pointer to the ASCII NUL terminated string containing the name of a previously registered (and currently active) system exit handler. The `sysexit_code` field is main function code of a system exit.

The system exits are divided into main functions and sub-functions. An exit is defined to handle a main function, and must thus handle all the sub-functions for that main function. All the functions and sub-functions are listed in the description of the `EXIT` structure.

48 TheSubcommandHandlerInterface

This section describes the subcommand handler interface, which enables the application to trap commands in a REXX script being executed and handle this command itself.

48.1 What is a Subcommand Handler

A subcommand handler is a piece of code, that is called to handle a command to an external environment in REXX. It must be either a subroutine in the application that started the interpreter, or a subroutine in a dynamic link library. In any case, when the interpreter needs to execute a command to an external environment, it will call the subcommand handler, passing the command as a parameter. Typically, an application will set up a subcommand handler before starting a REXX script. That way, it can trap and handle any command being executed during the course of the script.

Each subcommand handler handles one environment, which is referred to by a name. It seems to be undefined whether upper and lowercase letters differ in the environment name, so you should assume they differ. Also, there might be an upper limit for the length of an environment name, and some letters may be illegal as part of an environment name.

Regina allows any letter in the environment name, except ASCII NUL; and sets no upper limit for the length of an environment name. However, for compatibility reasons, you should avoid *uncommon* letters and keep the length of the name fairly short.

The prototype of a subcommand handler function is:

```
APIRET APIENTRY handler(  
    PRXSTRING command,  
    USHORT flags,  
    PRXSTRING returnstring  
);
```

After registration, this function is called whenever the application is to handle a subcommand for a given environment. The value of the parameters are:

[command]

The command string that is to be executed. This is the resulting string after the command expression has been evaluated in the REXX interpreter. It cannot be empty, although it can be a *zero-length-string*.

[flags]

Points to an unsigned short which is to receive the status of the completion of the handler. This can be one of the following: RXSUBCOM_OK, RXSUBCOM_ERROR, or RXSUBCOM_FAILURE. The contents will be used to determine whether to raise any condition at return of the subcommand. Do not confuse it with the return value.

[returnstring]

Points to a RXSTRING which is to receive the return value from the subcommand. Passing the return value as a string makes it possible to return non-numeric return codes. As a special case, you might set `returnstring.strptr` to NULL, instead of specifying a return string of the ASCII representation of zero.

Notethatit is not possible to return *nothing* in a subcommand, since this is interpreted as zero. Nor is it possible to return a numeric return code as such; you must convert it to ASCII representation before you return.

The `returnstring` will provide a 256 byte array which the programmer might use if the return data is not longer than that. If that space is not sufficient, the handler can provide another area itself. In that case, the handler should not de-allocate the default area, and the new area should be allocated in a standard fashion.

48.2 The `RexxRegisterSubcomExe()` function

This function is used to register a subcommand handler with the interface. The subcommand handler must be a procedure located within the code of the application. After registration, the `REXX` interpreter can execute subcommands by calling the subcommand handler with parameters describing the subcommand.

The prototype for `RexxRegisterSubcomExe()` is:

```
APIRETAPIENTRYRexxRegisterSubcomExe(
    PSZEnvName,
    PFNEntryPoint,
    PCHARUserArea
);
```

All the parameters are input, and their significance are:

[EnvName]

Points to an ASCII NUL terminated character string which defines the name of the environment to be registered. This is the same name as the `REXX` interpreter uses with the `ADDRESS` clause in order to select an external environment.

[EntryPoint]

Points to the entry point of the subcommand handler routine for the environment to be registered. See the section on Subcommand Handlers for more information. There is an upper limit for the length of this name.

[UserArea]

Points to an 8 byte area of information that is to be associated with this environment. This pointer can be `NULL` if no such area is necessary.

The areas pointed to by `EnvName` and `UserArea` are copied to a private area in the interface, so the programmer may de-allocate or reuse the area used for these parameters after the call has returned.

The `RexxRegisterSubcom()` returns an `unsigned long`, which carries status information describing the outcome of the operation. The status will be one of the `RXSUBCOM` values:

[RXSUBCOM_OK]

The subcommand handler was successfully registered.

[RXSUBCOM_DUP]

The subcommand handler was successfully registered. There already existed another subcommand handler which was registered with `RexxRegisterSubcomDll()`, but this will be shadowed by the newly registered handler.

[RXSUBCOM_NOTREG]

Due to some error, the handler was not registered. Probably because a handler for `EnvName` was already defined at a previous call to `RexxRegisterSubcomExe()`.

[RXSUBCOM_NOEMEM]

The handler was not registered, due to lack of memory.

[RXSUBCOM_BADTYPE]

Indicates that the handler was not registered, due to one or more of the parameters having invalid values.

48.3 The `RexxRegisterSubcomDll()` function

This function is used to set up a routine that is located in a module in a dynamic link library, as a subcommand handler. Some operating systems don't have dynamic linking, and thus cannot make use of this facility. The prototype of this function is:

```
APIRET APIENTRY RexxRegisterSubcomDll(  
    PSZ EnvName,  
    PSZ ModuleName,  
    PFN EntryPoint,  
    PCHAR UserArea,  
    ULONG DropAuth  
);
```

This function is not yet supported by `Regina`.

48.4 The `RexxDeregisterSubcom()` function

This function is used to remove a particular environment from the list of registered environments. The prototype of the function is:

```
APIRET APIENTRY RexxDeregisterSubcom(  
    PSZ EnvName,  
    PSZ ModuleName  
);
```

Both parameters are input values:

[EnvName]

Pointer to an ASCII NUL terminated string, which represents the name of the environment to be removed.

[ModuleName]

Also an ASCII NUL terminated string, which points to the name of the module containing the subcommand handler of the environment to be deleted.

The list of defined environments is searched, and if an environment matching the one named by the

first parameter are found, it is deleted.

The returned value from `RexxDeregisterSubcom()` can be one of:

- [RXSUBCOM_OK]
The subcommand handler was successfully deleted.
- [RXSUBCOM_NOTREG]
The subcommand handler was not found.
- [RXSUBCOM_BADTYPE]
One or more of the parameters had illegal values, and the operation was not carried through.

Most systems that do have dynamic linking have no method for reclaiming the space used by dynamically linked routines. So, even if you were able to load a *dll*, there are no guarantees that you will be able to unload it.

48.5 The `RexxQuerySubcom()` function

This function retrieves information about a previously registered subcommand handler. The prototype of the function is:

```
APIRETAPIENTRYRexxQuerySubcom(  
    PSZEnvName,  
    PSZModuleName,  
    PUSHORTFlag,  
    PCHARUserWord  
);
```

The significance of the parameters are:

- [EnvName]
Points to an ASCII NUL terminated character string, which names the subcommand handler about which information is to be returned.
- [ModuleName]
Points to an ASCII NUL terminated character string, which names a dynamic link library. Only the named library will be searched for the subcommand handler named by `EnvName`. This parameter must be `NULL` if all subcommand handlers are to be searched.
- [Flag]
Points to a short which is to receive the value `RXSUBCOM_OK` or `RXSUBCOM_NOTREG`. In fact, this is the same as the return value from the function.
- [UserWord]
Points to an area of 8 bytes. The *user area* of the subcommand handler is copied to the area pointed to by `UserWord`. This parameter might be `NULL` if the data of the *user area* is not needed.

The returned value from `RexxQuerySubcom()` can be one of:

[RXSUBCOM_OK]

The subcommand handler was found, and the required information has been returned in the `Flag` and `UserWord` variables.

[RXSUBCOM_NOTREG]

The subcommand handler was not found. The `Flag` variable will also be set to this value, and the `UserWord` variable is not changed.

[RXSUBCOM_BADTYPE]

One or more of the parameters had illegal values, and the operation was not carried through.

49 The External Function Handler Interface

This section describes the external function handler interface, which extends the language by enabling external functions to be written in a language other than REXX.

49.1 What is an External Function Handler

An external function handler is a piece of code, that is called to handle external functions and subroutine calls in REXX. It must be either a subroutine in the application that started the interpreter, or a subroutine in a dynamic link library. In any case, when the interpreter needs to execute a function registered as an external function, it will call the external function handler, passing the function name as a parameter.

All external functions written in a language other than REXX must be registered with the interpreter before starting a REXX script.

An external function handler can handle one or more functions. The handler can determine the function actually called by examining one of the parameters passed to the handler and act accordingly.

The prototype of a subcommand handler function is:

```
APIRET APIENTRY handler(  
    PSZ name,  
    ULONG argc,  
    PRXSTRING argv,  
    PSZ queue name,  
    PRXSTRING return string  
);
```

After a function is registered with this function defined as the handler, this function is called whenever the application calls the function. The value of the parameters are:

[name]	The function called.
[argc]	The number of parameters passed to the function. Argv will contain argc RXSTRINGS.
[queue name]	The name of the currently defined data queue.
[return string]	Points to a RXSTRING which is to receive the return value from the function. Passing the return value as a string makes it possible to return non-numeric return codes. As a special case, you might set return string.strptr to NULL, instead of specifying a return string of the ASCII representation of zero.

The return string will provide a 256 byte array which the programmer might use if the return data is not longer than that. If that space is not sufficient, the handler can provide another area itself. In that case, the handler should not de-allocate the default area, and the new area should be allocated in a standard fashion. If the external function does not return a value, it should set return string to an empty RXSTRING. This will enable the interpreter to raise error 44; *Function did not return data* , if the

external function is called as a function. If the external function is invoked via a `CALL` command, the interpreter drops the special variable `RESULT`.

The handler returns zero if the function completed successfully. When the handler returns a non-zero value, the interpreter will raise error 40; *Invalid call to routine*.

49.2 The `RexxRegisterFunctionExe()` function

This function is used to register an external function handler with the interface. The external function handler must be a procedure located within the code of the application. After registration, the REXX interpreter can execute external functions as if they were built-ins.

The prototype for `RexxRegisterFunctionExe()` is:

```
APIRETAPIENTRY RexxRegisterFunctionExe(
    PSZFuncName,
    PFNEntryPoint
);
```

All the parameters are input, and their significance are:

- [FuncName]
Points to an ASCII NUL terminated character string which defines the name of the external function to be registered. This is the same name as the REXX interpreter uses with a function call or via the `CALL` command.
- [EntryPoint]
Points to the entry point of the external function handler routine for the function to be registered. See this section on External Function Handlers for more information.

The area pointed to by `FuncName` is copied to a private area in the interface, so the programmer may de-allocate or reuse the area used for this parameter after the call has returned.

The `RexxRegisterFunctionExe()` returns an `unsigned long`, which carries status information describing the outcome of the operation. The status will be one of the `RXFUNC` values:

- [RXFUNC_OK]
The handler was successfully registered.
- [RXFUNC_DUP]
The handler was successfully registered. There already existed another external function handler which was registered with `RexxRegisterFunctionExe()`, but this will be shadowed by the newly registered handler.
- [RXFUNC_NOEMEM]
The handler was not registered, due to lack of memory.

49.3 The `RexxRegisterFunctionDll()` function

This function is used to set up an external function handler that is located in a module in a dynamic link library. Some operating systems don't have dynamic linking, and thus cannot make use of this facility. The prototype of this function is:

```
APIRETAPIENTRYRexxRegisterFunctionDll(
    PSZExternalName,
    PSZLibraryName,
    PSZInternalName
);
```

All the parameters are input, and their significance are:

- [ExternalName]
Point to an ASCII NUL terminated character string which defines the name of the external function to be registered. This is the same name as the REXX interpreter uses with a function call or via the CALL command.
- [LibraryName]
Point to an ASCII NUL terminated character string which defines the name of the dynamic library. This string may require a directory specification.
- [InternalName]
Point to an ASCII NUL terminated character string which defines the name of the entry point within the dynamic library. On systems where the case of function names in dynamic libraries is relevant, this name **must** be specified in the same case as the function name within the dynamic library.

The areas pointed to by all parameters are copied to a private area in the interface, so the programmer may de-allocate or reuse the areas used for these parameters after the call has returned.

The `RexxRegisterFunctionDll()` returns an `unsigned long`, which carries status information describing the outcome of the operation. The status will be one of the `RXFUNC` values:

- [RXFUNC_OK]
The handler was successfully registered.
- [RXFUNC_DUP]
The handler was successfully registered. There already existed another external function handler which was registered with `RexxRegisterFunctionDll()`, but this will be shadowed by the newly registered handler.
- [RXFUNC_NOEMEM]
The handler was not registered, due to lack of memory.

49.4 The `RexxDeregisterFunction()` function

This function is used to remove a particular external function handler from the list of registered external function handlers. The prototype of the function is:

```
APIRETAPIENTRYRexxDeregisterFunction(
    PSZFuncName
);
```

The parameter is an input value:

[FuncName]

Points to an ASCII NUL terminated character string which defines the name of the external function to be registered. This is the same name as the REXX interpreter uses with a function call or via the CALL command.

The list of defined function handlers is searched, and if an environment matching the one named by the parameter are found, it is deleted. This call is used to de-register function handlers registered with either RexxRegisterFunctionExe() or RexxRegisterFunctionDll().

The returned value from RexxDeregisterFunction() can be one of:

[RXFUNC_OK]

The handler was successfully deleted.

[RXFUNC_NOTREG]

The handler was not found.

Most systems that do have dynamic linking have no method for reclaiming the space used by dynamically linked routines. So, even if you were able to load a *dll*, there are no guarantees that you will be able to unload it.

49.5 The RexxQueryFunction() function

This function retrieves the status of an external function handler. The prototype of the function is:

```
APIRET APIENTRY RexxQueryFunction(  
    PSZ FuncName  
);
```

The significance of the parameters is:

[FuncName]

Points to an ASCII NUL terminated character string which defines the name of the external function to be registered. This is the same name as the REXX interpreter uses with a function call or via the CALL command.

The returned value from RexxQueryFunction() can be one of:

[RXFUNC_OK]

The external function handler was found.

[RXFUNC_NOTREG]

The handler was not found.

50 Executing REXX Code

This section describes the `RexxStart()` function, which allows the application to startup the interpreter and make it interpret pieces of REXX code.

50.1 The `RexxStart()` function

This function is used to invoke the REXX interpreter in order to execute a piece of REXX code, which may be located on disk, as a pre-tokenized macro, or as ASCII source code in memory.

```
APIRETAPIENTRY RexxStart(  
    LONG ArgCount,  
    PRXSTRING ArgList,  
    PSZ ProgramName,  
    PRXSTRING Instore,  
    PSZ EnvName,  
    LONG CallType,  
    PRXSYSEXIT Exits,  
    PUSHORT ReturnCode,  
    PRXSTRING Result  
);
```

Of these parameters, `ReturnCode` and `Result` are output-only, while `Instore` is both input and output. The rest of the parameters are input-only. The significance of the parameters are:

[ArgCount]

The number of parameter strings given to the procedure. This is the number of defined REXX-strings pointed to by the `ArgList` parameter. The default maximum number of arguments that can be passed is 32, but this can be changed by the `MAX_ARGS_TO_REXXSTART` macro in `rexx.h`.

[ArgList]

Points to an array of REXX-strings, constituting the parameters to this call to REXX. The size of this array is given by the parameter `ArgCount`. If `ArgCount` is greater than one, the first and last parameters are `ArgList[0]` and `ArgList[ArgCount-1]`. If `ArgCount` is 0, the value of `ArgList` is irrelevant.

If the `strptr` of one of the elements in the array pointed to by `ArgList` is `NULL`, that means that this parameter is empty (i.e. unspecified, as opposed to a string of zero size).

[ProgName]

An ASCII NUL terminated string, specifying the name of the REXX script to be executed. The value of `Instore` will determine whether this value is interpreted as the name of a (on-disk) script, or a pre-tokenized macro. If it refers to a filename, the syntax of the contents of this parameter depends on the operating system.

[Instore]

Parameter used for storing tokenized REXX scripts. This parameter might either be `NULL`, else it will be a pointer to two `PRXSTRING` structures, the first holding the ASCII version of a REXX program, the other holding the tokenized version of that program. See below for more information about how to use `Instore`.

[EnvName]

PointertoASCII NULterminatedstring namingtheenvironmentwhichistobethe initialcurrentenvironmentwhenthe scriptisstarted.Ifthisparameterissetto NULL, thefiletypeisusedastheinitialenvironmentname.Whatthefiletypeis,maydependon youroperatingsystem,butingeneralitiseverythingafterthelastperiod'.'inthe filename.

[CallType]

Avaluedescribingwhetherthe REXXinterpreteristobeinvokedincommand, functionorsubroutinemode.Actually,thishaslittlesignificance.Themaindifferenceis thatincommandmode,onlyoneparameterstringcanbepassed,andinfunctionmode, avaluemustbereturned.Inaddition,themodechosenwillaffecttheoutputofthe PARSESOURCE instructionin REXX.

Threesymbolicvaluesofintegraltypearedefined,whichcanbeusedforthisparameter: RXCOMMAND, RXFUNCTIONand RXSUBROUTINE.

Avalueof RXRESTRICTEDcanbeOR'edwithoneoftheabovetypestospecifythat Reginawillrunin *restricted* mode.Thisisparticularlyusefulwhen Reginaisusedas anembeddedintepreterinapplicationssuchasadatabaseprocedurallanguageoraweb-browserscriptinglanguage.

[SysExists]

Apointertoanarrayofexithandlerstobeused.Ifnoexithandlersaretobefined, NULLmaybespecified.Eachelementinthearraydefinesoneexithandler,andthe elementimmediatelyfollowingthelastdefinitionmusthavea `sysexit_codes`setto RXENDLST.

[ReturnCode]

Pointertoa SHORTintegerwherethereturncodeisstored,providedthatthereturned valueisnumeric,andwithinthelange- $(2^{**}15)$ to $2^{**}15-1$.Idon'tknowwhathappens to ReturnCodeifeitheroftheseconditionsisnotsatisfied.Itprobablybecomes undefined,whichmeansthatitistotallyuselessincetheprogramhastoinspectthe returnstringinordertodeterminewhether ReturnCodeisvalid.

[Result]

Pointsto a REXXstringintowhichtheresultstringiswritten.Thecallermayormay notletthe `strptr`fieldbesupplied.If supplied(i.e.itisnon-NULL),that areawillbe used,elseanewareawillbeallocated.Ifthesuppliedareaisused,itssizeissupposed tobegivenbythe `strlen`field.Ifthesizeifnotsufficient,anewareawillbe allocated,by `RexxAllocateMemory()`,andthecallermustsetothatit isproperlyde-allocatedusing `RexxFreeMemory()`.

Notethatthe `ArgCount`parameterneednotbethesameasthe `ARG()`built-infunctionwouldreturn. Differenceswilloccurifthelastentriesin `ArgList`are *nullstrings*.

The `Instore`parameterneedssomespecialattention.Itisusedtodirectlyorindirectlyspecifywheretofetchthecodetoexecute.Thefollowing *algorithm*isusedtodeterminewhattotoexecute:

If `Instore`is NULL,then `ProgName`namesthefilenameofanon-disk REXXscriptwhichit tobereadandexecuted.

Else, if `Instore` is not `NULL`, the script is somewhere in memory, and no reading from disk is performed. If both `Instore[0].strptr` and `Instore[1].strptr` are `NULL`, then the script to execute is a pre-loaded macro which must have been loaded with a call to either `RexxAddMacro()` or `RexxLoadMacroSpace()`; and `ProgName` is the name of the macro to execute.

Else, if `Instore[1].strptr` is non-`NULL`, then `Instore[1]` contains the pre-tokenized image of a REXX script, and it is used for the execution.

Else, if `Instore[0].strptr` is non-`NULL`, then `Instore[0]` contains the ASCII image of a REXX script, just as if the script had been read directly from the disk (i.e. including line feeds and such). This image is passed to the interpreter, which tokenizes it, and stores the tokenized script in the `Instore[1]` string, and then proceeds to execute that script. Upon return, the `Instore[1]` will be set, and can later be used to re-execute the script within the same process, without the overhead of tokenizing.

The user is responsible for de-allocating any storage used by `Instore[1]`. Note that after tokenizing, the source code in `Instore[0]` is strictly speaking not needed anymore. It will only be consulted if the user calls the `SOURCELINE()` built-in function. It is not an error to use `SOURCELINE()` if the source is not present, but null strings and zero will be returned.

To tokenize a REXX script and save it for execution by a later execution by `RexxStart()` either in the currently running process or outside the current process, you need to call `RexxStart()` with the following arguments:

<i>Parameter</i>	<i>Value</i>	<i>Notes</i>
<code>ArgCount</code>	1	
<code>ArgList.strlength</code>	3	
<code>ArgList.strptr</code>	//T	
<code>ProgName</code>		Ignored
<code>Instore[0].strptr</code>	ASCII image of REXX script	
<code>Instore[0].strlength</code>	Length of <code>Instore[0].strptr</code>	
<code>Instore[1].strptr</code>	Ignored	This will be populated with the tokenised code.
<code>Instore[1].strlength</code>	Ignored	This will be set to the length of <code>Instore[1].strptr</code>
<code>EnvName</code>	SYSTEM	
<code>CallType</code>	RXCOMMAND	
<code>SysExits</code>	NULL	
<code>ReturnCode</code>	Ignored	
<code>Result.strptr</code>	Ignored	
<code>Result.strlength</code>	Ignored	

The valid return values from `RexxStart()` are:

[Negative]

indicates that a syntax error occurred during interpretation. In general, you can expect the error value to have the same absolute value as the `REXX` syntax error (but opposite signs, of course).

[Zero]

indicates that the interpreter finished executing the script without errors.

[Positive]

indicates probably that some problem occurred, that made it impossible to execute the script, e.g. a bad parameter value. However, I can't find any references in the documentation which states which values it is supposed to return.

During the course of an execution of `RexxStart()`, subcommand handlers and exit handlers might be called. These may call any function in the application interface, including another invocation of `RexxStart()`.

Often, the application programmer is interested in providing support simplifying the specification of filenames, like an environment variable search path or a default filetype. The `REXX` interface does support a default filetype: `.CMD`, but the user may not set this to anything else. Therefore, it is generally up to the application programmer to handle search paths, and also default filetypes (unless `.CMD` is OK).

If the initial environment name (`EnvName`) is `NULL`, then the initial environment during interpretation will be set equal to the filetype of the script to execute. If the script does not have a filetype, it is probably set to some interpreter specific value.

51 VariablePoolInterface

This section describes the variable pool part of the application interface, which allows the application programmer to set, retrieve and drop variables in the REXX interpreter from the application program. It also allows access to other information.

The C preprocessor symbol `INCL_RXSHV` must be defined if the definitions for the variable pool interface are to be made available when `rexxsaa.h` is included.

51.1 Symbolic or Direct

First, let us define two terms, *symbolic variable name* and *direct variable name*, which are used in connection with the variable pool.

A symbolic variable name is the name of a variable, but it needs normalization and tails substitution before it names the real variable. The name `foo.bar` is a symbolic variable name, and it is transformed by normalization, to `FOO.BAR`, and then by tails substitution to `FOO.42` (assuming that the current value of `BAR` is 42).

Normalization is the process of uppercasing all characters in the symbolic name; and tails substitution is the process of substituting each distinct simple symbol in the tail for its value.

On the other hand, a direct variable refers directly to the name of the variable. In a sense, it is a symbolic variable that has already been normalized and tails substituted. For instance, `foo.bar` is not a valid direct variable name, since lowercase letters are not allowed in the variable stem. The direct variable `FOO.42` is the same as the variable above. For simple variables, the only difference between direct and symbolic variable names is that lowercase letters are allowed in symbolic names.

Note that the two direct variable names `FOO.bar` and `FOO.BAR` refer to different variables, since upper and lowercase letters differ in the tail. In fact, the tail of a compound direct variable may contain any character, including ASCII NUL. The stem part of a variable, and all simple variables cannot contain any lowercase letters.

As a remark, what would the direct variable `FOO.` refer to: the stem `FOO.` or the compound variable having stem `FOO.` and an null string as tail? Well, I suppose the former, since it is the more useful. Thus, the latter is inaccessible as a direct variable.

51.2 The SHVBLOCK structure

All requests to manipulate the REXX variable pool are controlled by a structure which is called SHVBLOCK, having the definition:

```
typedef struct shvnode {
    struct shvnode *shvnext; /* ptr to next in blk in chain */
    RXSTRING shvname; /* name of variable */
    RXSTRING shvvalue; /* value of variable */
    ULONG shvnamelen; /* length of shvname.strptr */
    ULONG shvvaluelen; /* length of shvvalue.strptr */
    UCHAR shvcode; /* operation code */
}
```

```

        UCHARshvret; /*returncode*/
    }SHVBLOCK;

typedefSHVBLOCK*PSHVBLOCK;

```

The fields `shvnext` and `shvcode` are purely input, while `shvret` is purely output. The rest of the fields might be input or output, depending on the requested operation, and the value of the fields. The significance of each field is:

- [`shvnext`]

One call to `RexxVariablePool()` may sequentially process many requests. The `shvnext` field links on one request to the next inline. The last request must have set `shvnext` to `NULL`. The requests are handled individually and thus, calling `RexxVariablePool()` with several requests is equivalent to making one call to `RexxVariablePool()` for each request.
- [`shvname`]

Contains the name of the variable to operate on, as a `RXSTRING`. This field is only relevant for some requests, and its use may differ.
- [`shvvalue`]

Contains the value of the variable to operate on as a `RXSTRING`. Like `shvname`, this might not be relevant for all types of requests.
- [`shvnamelen`]

The length of the array that `shvname.strptr` points to. This field holds the maximum possible number of characters in `shvname.strptr`. While `shvname.strlength` holds the number of characters that are actually in use (i.e. defined).
- [`shvvaluelen`]

The length of the array that `shvvalue.strptr` points to. Related to `shvvalue`, like `shvnamelen` related to `shvname`.
- [`shvcode`]

The code of operation; decides what type of request to perform. A list of all the available requests is given below.
- [`shvret`]

A return code describing the outcome of the request. This code is a bit special. The lower seven bits are flags which are set depending on whether some condition is met or not. Values above 127 are not used in this field.

There is a difference between `shvnamelen` and `shvname.strlength`. The former is the total length of the array of characters pointed to by `shvname.strptr` (if set). While the latter is the number of these characters that are actually in use. When a `SHVBLOCK` is used to return data from `RexxVariablePool()`, and a pre-allocated string space has been supplied, both these will be used; `shvname.strlength` will be set to the length of the data returned, while `shvnamelen` is never changed, only read to find the maximum number of characters that `shvname` can hold.

Even though `shvnamelen` is not really needed when `shvname` is used for input, it is wise to set it to its proper value (or at least set it to the same as `shvname.strlength`). The same applies for `shvvalue` and `shvvaluelen`.

The field `shvcode` can take one of the following symbolic values:

[RXSHV_DROPV]

The variable named by the direct variable name `shvname` is dropped (i.e. becomes undefined). The fields `shvvalue` and `shvvalueLen` do not matter.

[RXSHV_EXIT]

This is used to set the return value for an external function or exit handler.

[RXSHV_FETCH]

The value of the variable named by the direct variable name `shvname` is retrieved and stored in `shvvalue`. If `shvvalue.strptr` is `NULL`, the interpreter will allocate sufficient space to store the value (but it is the responsibility of the application programmer to release that space). Else, the value will be stored in the area allocated for `shvvalue`, and `shvvalueLen` is taken to be the maximum size of that area.

[RXSHV_NEXTV]

This code is used to retrieve the names and values of all variables at the current procedure level; i.e. excluding variables shadowed by `PROCEDURE`. The name and value of each variable are retrieved

simultaneously into `shvname` and `shvvalue`, respectively.

Successive requests for `RXSHV_NEXTV` will traverse the interpreter's internal data structure for storing variables, and return a new pair of variable name and value for each request. Each variable that is visible in the current scope, is returned once and only once, but the order is non-deterministic.

When all available variables in the `REXX` interpreter have already been retrieved, subsequent `RXSHV_NEXTV` requests will set the flag `RXSHV_LVAR` in the `shvret` field. There are a few restrictions. The traversal will be reset whenever the interpreter resumes execution, so an incomplete traversal cannot be continued in a later external function, exit handler, or subcommand handler. Also, any set, fetch or drop operation will reset the traversal. These restrictions have been added to ensure that the variable pool is static throughout one traversal.

[RXSHV_PRIV]

Retrieve some piece of information from the interpreter, other than a variable value, based on the value of the `shvname` field. The value is stored in `shvvalue` as for a *normal* fetch. A list of possible names is shown below.

[RXSHV_SET]

The variable named by the direct variable name `shvname` is set to the value given by `shvvalue`.

[RXSHV_SYFET]

Like `RXSHV_FETCH`, except that `shvname` is a symbolic variable name.

[RXSHV_SYDRO]

Like `RXSHV_DROPV`, except that `shvname` is a symbolic variable name.

[RXSHV_SYSET]

Like `RXSHV_SET`, except that `shvname` is a symbolic variable name.

One type of request that needs some special attention is the `RXSHV_PRIV`, which retrieves a kind of *meta-variable*. Depending on the value of `shvname`, it returns a value in `shvvalue` describing some aspect of the interpreter. For `RXSHV_PRIV` the possible values for `shvname` are:

[PARM]

Returns the ASCII representation of the number of parameters to the currently active REXX procedure. This may not be the same value as the built-in function ARG() returns, but is the number ArgCount in RexxStart(). The two might differ if a routine was called with trailing omitted parameters.

[PARM.n]

The n must be a positive integer; and the value returned will be the n'th parameter at the current procedure level. This is not completely equivalent to the information that the built-in function ARG() returns. For parameters where ARG() would return the state omitted, the returned value is a *nullstring*, while for parameters where ARG() would return the state *existing*, the return value will be the parameter string (which may be a *zero length string*).

[QUENAME]

The name of the currently active external data queue. This feature has not yet been implemented in Regina, which always returns *default*.

[SOURCE]

Returns the same string that is used in the PARSE SOURCE clause in REXX, at the current procedure level of interpretation.

[VERSION]

Returns the same string that is used in the PARSE VERSION clause in REXX.

The value returned by a variable pool request is a bit uncommon. A return value is computed for each request, and stored in the shvret field. This is a one-byte field, of which the most significant bit is never set. A symbolic value RXSHV_OK is defined as the value zero, and the shvret field will be equal to this name if none of the flags listed below is set. The symbolic values for these flags are:

[RXSHV_BADF]

The shvcode of this request contained a bad function code.

[RXSHV_BADN]

The shvname field contained a string that is not valid in this context. What exactly is a valid value depends on whether the operation is a private, a symbolic variable, or a direct variable operation.

[RXSHV_LVAR]

Set if and only if the request was RXSHV_NETXV, and all available variables have already been retrieved by earlier requests.

[RXSHV_MEMFL]

There was not enough memory to complete this request.

[RXSHV_NEWV]

Set if and only if the referenced variable did not previously have a value. It can be returned for any set, fetch or drop operation.

[RXSHV_TRUNC]

Set if the retrieved value was truncated when it was copied into either the shvname or shvvalue fields. See below.

These flags are directly suitable for logical OR, without shifting, e.g. to check for truncation and no variables left, you can do something like:

```
if(req->shvret&(RXSHV_TRUNC|RXSHV_LVAR))
    printf("Truncation or no vars left\n");
```

`RXSHV_TRUNC` can only occur when the interface is storing a retrieved value in a `SHVBLOCK`, and the pre-allocated space is present, but not sufficiently large. As described for `RXSHV_FETCH`, the interpreter will allocate enough space if `shvvalue.strptr` is `NULL`, and then `RXSHV_TRUNC` will never be set. Else the spaces supplied by `shvvalue.strptr` is used, and `shvvalue.len` is taken as the maximum length of `shvvalue`, and truncation will occur if the supplied space is too small.

Some implementations will consider `SHV_MEMFL` to be so severe as to skip the rest of the operations in a chain of requests. In order to write compatible software, you should never assume that requests following in a chain after a request that returned `SHV_MEMFL` have been performed.

`RXSHV_BADN` is returned if the supplied `shvname` contains a value that is not legal in this context. For the symbolic set, fetch and drop operations, that means a symbol that is a legal variable name; both upper and lower case letters are allowed. For the direct set, fetch and drop operations, that means a variable name after normalization and tail substitution is not a legal variable name. For `RXSHV_PRIV`, it must be one of the values listed above.

There is a small subtlety in the above description. TRL states that when a `REXX` assignment assigns a value to a stem variable, all possible variables having that stem are assigned a new value (independent of whether they had an explicit value before). So, strictly speaking, if a stem is set, then a `RXSHV_NETV` sequence should return an (almost) infinite sequence of compound variables for that stem. Of course, that is completely useless, so you can assume that only compound variables of that stem given an explicit value after the stem was assigned a value will be returned by `RXSHV_NEXTV`. However, because of that subtlety, the variables returned by `RXSHV_NEXTV` for compound variables might not be representative for the state of the variables.

e.g. what would a sequence of `RXSHV_NEXT` requests return after the following `REXX` code?:

```
foo.='bar'
drop foo.bar
```

The second statement there, might not change the returned values! After the first statement, only the stem `foo.` would probably have been returned, and so also if all variables were fetched after the second statement.

51.3 Regina Notes for the Variable Pool

Due to the subtleties described at the end of the previous subsection, some notes on how Regina handles `RXSHV_NEXTV` requests for compound variables are in order. The following rules apply:

- Both the stem variable `FOO.` and the compound variable having `FOO.` as stem and a null string as tail, are returned with the name of `FOO.`. In this situation, a sequence of `RXSHV_NEXTV` requests may seem to return values for the same variable twice. This is unfortunate, but it seems to be the only way. In any case, you'll have to perform the `RXSHV_SYFET` in order to determine which is which.
- If a stem variable has not been assigned a value, its compound variables are only returned if they have been assigned an explicit value. i.e. compound variables for that stem that have either never

been assigned a value, or have been dropped, will not be reported by `RXSHV_NEXTV`. There is nothing strange about this.

- If a stem variable has been assigned a value, then its compound variables will be reported in two cases: Firstly, the compound variables having explicitly been assigned a value afterwards. Secondly, the compound variables which have been dropped afterwards, which are reported to have their initial value, and the flag `RXSHV_NEWV` is set in `shvret`.

It may sound a bit stupid that unset variables are listed when the request is to list all variables which have been set, but that is about the best I can do, if I am to stay within the standard definition and return a complete and exact status of the variable pool.

If the return code from `RexxVariablePool()` is less than 128, Regina is guaranteed to have tried to process all requests in the chain. If the return code is above 127, some requests may not have been processed. Actually, the number 127 (or 128) is a bit inconvenient, since it will be a problem for later expansion of the standard. A much better approach would be to have a preprocessor symbol (say, `RXSHV_FATAL`, and if the return code from the `RexxVariablePool()` function was larger than that, it would be a *direct* error code, and not a *composite* error code built from the `shvret` fields of the requests. The `RXSHV_FATAL` would then have to be the addition of all the atomic composite error codes.

(Warning: author mounting the soapbox.)

The *right* way to fix this, is to let the function `RexxVariablePool()` set another flag in `shvret` (e.g. named `RXSHV_STEM`) during `RXSHV_NEXTV` if and only if the value returned is a stem variable. That way, the application programmer would be able to differ between stem variables and compound variables with a null string tail.

To handle the other problem with compound variables and `RXSHV_NEXTV`, I would have liked to return a *null string* in `shvvalue` if and only if the variable is a compound variable having its initial value, and the stem of that compound variable has been assigned a value. Then, the value of the compound variable is equal to its name, and is available in the `shvname` field.

I'd also like to see that the `shvret` value contained other information concerning the variables, e.g. whether the variable was exposed at the current procedure level. Of course, Regina does not contain any of these extra, non-standard features.

(Author is dismounting the soapbox.)

When Regina is returning variables with `RXSHV_NEXTV`, the variables are returned in the order in which they occur in the open hash table in the interpreter, i.e. the order in which variables belonging to different bins are returned is inconsistent, but the order in which variables hashed to the same bin are returned, is non-deterministic. Note that all compound variables belonging to the same stem are returned in one sequence.

51.4 The `RexxVariablePool()` function

This function is used to process a sequence of variable requests, and process them sequentially. The prototype of this function is:

```
APIRETAPIENTRYULONGRexxVariablePool(
```

SHVBLOCK*Request
);

Its only parameter is a pointer to a SHVBLOCK structure, which may be the first of the linked list. The function performs the operations specified in each block. If an error should occur, the current request is terminated, and the function moves on to the next request in the chain.

The result value is a bit peculiar. If the returned value is less than 128, it is calculated by logically OR'ing the returned shvret field of all the requests in the chain. That way, you can easily check whether any of the requests was e.g. skipped because of lack of memory. To determine which request, you have to iterate through the list.

If the result value is higher than 127, it signifies an error. If any of these values are set, you cannot assume that any of the requests have been processed. The following symbolic name gives its meaning.

[RXSHV_NOAVL]

Means that the interface is not available for this request. This might occur if the interface was not able to start the interpreter, or if an operation requested a variable when the interpreter is not currently executing any script (i.e. idle and waiting for a script to execute).

52 TheSystemExitHandlerInterface

The exit handlers provide a mechanism for governing important aspects of the REXX interpreter from the application. It can trap situations like the interpreter writing out text, and then handle them itself, e.g. by displaying the text in a special window. You can regard system exits as a sort of *hooks*.

52.1 TheSystemExitHandler

Just like the subcommand handler, the system exit handler is a routine supplied by the application, and is called by the interpreter when certain situations occur. These situations are described in detail later. For the examples below, we will use the output from `SAY` as an example.

If a system exit handler is enabled for the `SAY` instruction, it will be called with a parameter describing the text that is to be written out. The system exit handler can choose to handle the situation (e.g. by writing the text itself), or it can ignore it and let the interpreter perform the output. The return code from the system exit tells the interpreter whether a system exit handled the situation or not.

A system exit handler must be a routine defined according to the prototype:

```
LONGAPIENTRY my_exit_handler(  
    LONG ExitNumber,  
    LONG Subfunction,  
    PEXIT ParmBlock  
);
```

In this prototype, the type `PEXIT` is a pointer to a parameter block containing all the parameters necessary to handle the situation. The actual definition of this parameter block will vary, and is described in detail in the list of each system exit.

The exits are defined in a two-level hierarchy. The `ExitNumber` defines the main function for a system exit, while the `Subfunction` defines the subfunction within that main function. e.g. for `SAY`, the main function will be `RXSIO` (the system exit for standard I/O) and the subfunction will be `RXSIO SAY`. The `RXSIO` main function has other sub-functions for handling trace output, interactive trace input, and `PULL` input from standard input.

The value returned from the system exit handler must be one of the following symbolic values:

[RXEXIT_HANDLED]

Signal that the system exit handler took care of the situation, and that the interpreter should not proceed to do the default action. For the `SAY` instruction, this means that the interpreter will not print out anything.

[RXEXIT_NOT_HANDLED]

Signal that the system exit handler did not take care of the situation, and the interpreter will proceed to perform the default action. For the `SAY` instruction, this means that it must print out the argument to `SAY`.

[RXEXIT_RAISE_ERROR]

Signal that the interpreter's default action for this situation should not be performed, but instead a `SYNTAX` condition should be raised. Don't get confused by the name, it is not the `ERROR` condition, but the `SYNTAX` condition is raised, using the syntax error

Failure in system service (normally numbered 48).

In addition to returning information as the numeric return value, information may also be returned by setting variables in the parameter block. For instance, if the system exit is to handle interactive trace input, that is how it will supply the interpreter with the input string.

It is a good and disciplined practice to let your exit handlers start by verifying the `ExitNumber` and `Subfunction` codes, and immediately return `RXEXIT_NOT_HANDLED` if it does not recognize both of them. That way, your application will be upwards compatible with future interpreters which might have more sub-functions for any given main function.

52.2 List of System Exit Handlers

52.2.1 RXFNC-The External Function Exit Handler

The `RXFNC` system exit handler provides hooks for external functions. It has only one subfunction; `RXFNCAL`, which allows an application program to intervene and handle any external function or subroutine.

Do not confuse this exit handler with the external function routines which allow you to define new `REXX`, semi-built-in functions. The exit handler is called for all invocations of external routines, and can be called for function names which you were unaware of.

The parameter `ParmBlock` for `RXFNCAL` is defined as:

```
typedef struct {
    typedef struct {
        unsigned int rxferr: 1;
        unsigned int rxfnfd: 1;
        unsigned int rxffsub: 1;
    } rxfnc_flags;
    unsigned char *rxfnc_address;
    unsigned short rxfnc_addressl;
    unsigned char *rxfnc_que;
    unsigned short rxfnc_quel;
    unsigned short rxfnc_argc;
    RXSTRING *rxfnc_argv;
    RXSTRING *rxfnc_retc;
} RXFNCCAL_PARM;
```

The significance of each variable is:

[`rxfnc_flags.rxferr`]

Is an output parameter that is set on return in order to inform the interpreter that the function or subroutine was incorrectly called, and thus the `SYNTAX` conditions should be raised.

[`rxfnc_flags.rxfnfd`]

Is an output parameter that tells the interpreter that the function was not found. Note the inconsistency: it is only effective if the exit handler returns `RXEXIT_HANDLED`, which looks like a logic contradiction to setting the not-found flag.

[*rxfunc_flags*.*rxffsub*]

Is an input parameter that tells the exit handler whether it was called for a function or subroutine call. If set, the call being handled is a subroutine call and returning a value is optional; else it was called for a function, and must return a value in *rxfunc_ret* if *RXEXIT_HANDLED* is to be returned.

[*rxfunc_name*]

Is a pointer to the name of the function or subroutine to be handled, stored as a character array. This is an input parameter, and its length is given by the *rxfunc_name1* parameter.

[*rxfunc_name1*]

Holds the length of *rxfunc_name*. Note that the last character is the letter *ell*, not the number one.

[*rxfunc_que*]

Points to a character array holding the name of the currently active queue. This is an input parameter. The length of this name is given by the *rxfunc_que1* field.

[*rxfunc_que1*]

Holds the length of *rxfunc_que*. Note that the last character is the letter *ell*, not the number one.

[*rxfunc_argc*]

Is the number of arguments passed to the function or subroutine. It defines the size of the array pointed to by the *rxfunc_argv* field.

[*rxfunc_argv*]

Points to an array holding the parameters for the routines. The size of this array is given by the *rxfunc_argc* field. If *rxfunc_argc* is zero, the value of *rxfunc_argv* is undefined.

[*rxfunc_ret*]

Holds an *RXSTRING* structure suitable for storing the return value of the handler. It is the responsibility of the handler to allocate space for the contents of this string (i.e. the array pointed to by the *rxfunc_ret.strptr*).

52.2.2 *RXCMD*-The Subcommand Exit Handler

The main function code for this exit handler is given by the symbolic name *RXCMD*. It is called whenever the interpreter is about to call a subcommand, i.e. a command to an external environment. It has only one subfunction: *RXCMDHST*.

The *ParmBlock* parameter for this subfunction has the following definition:

```
typedef struct {
    typedef struct {
        unsigned int rxfcfail:1;
        unsigned int rxfccerr:1;
    } rxcmd_flags;
    unsigned char *rxcmd_address;
    unsigned short rxcmd_address1;
    unsigned char *rxcmd_dll;
    unsigned short rxcmd_dll_len;
    RXSTRING rxcmd_command;
    RXSTRING rxcmd_ret;
} RXCMDHST_PARM;
```

The significance of each variable is:

- [rxcmd_flags.rxfcfail] If this flag is set, the interpreter will raise a **FAILURE** condition at the return of the exit handler.
- [rxcmd_flags.rxfcerr] Like the previous, but the **ERROR** condition is raised instead.
- [rxcmd_address] Points to a character array containing the name of the environment to which the command normally would be sent.
- [rxcmd_addressl] Holds the length of `rxcmd_address`. Note that the last character is the letter *ell*, not the number one.
- [rxcmd_dll] Defines the name for the DLL which is to handle the command. I'm not sure what this entry is used for. It is not currently in use for **Regina**.
- [rxcmd_dll_len] Holds the length of `rxcmd_dll`. If this length is set to zero, the subcommand handler for this environment is not a DLL, but an EXE handler.
- [rxcmd_command] Holds the command string to be executed, including command name and parameters.
- [rxcmd_retc] Set by the exit handler to the string which is to be considered the return code from the command. It is assigned to the special variable **RC** at return from the exit handler. The user is responsible for allocating space for this variable. I have no clear idea what happens if `rxcmd_retc.strptr` is set to **NULL**; it might set **RC** to zero, to the null string, or even drop it.

It seems that this exit handler is capable of raising both the **ERROR** and the **FAILURE** conditions simultaneously. I don't know whether that is legal, or whether only the **FAILURE** condition is raised, since the **ERROR** condition is a sort of *subset* of **FAILURE**.

Note that the return fields of the parameter block are only relevant if the value **RXEXIT_HANDLED** was returned. This applies to the `rxcmd_flags` and `rxcmd_retc` fields of the structure.

52.2.3 RXMSQ-The External Data Queue Exit Handler

The external data queue exit handler is used as a hook for operations manipulating the external data queue (or the stack). Unfortunately, the stack is a borderline case of what is relevant to the **REXX SAA** API. Operations like putting something on, retrieving a string from, obtaining the size, etc. of the stack is not part of the **SAA** API.

However, some of this functionality is seemingly here; but not all. For instance for the **RXMSQPLL** subfunction, **SAA** API is called by the interpreter before the interpreter calls whatever system-specific call is available for retrieving a string from the stack.

Thus the **SAA** API can be used by an application to provide the interpreter with a fake stack, but it is not a suitable means for the application itself to manipulate the *real* stack.

The **RXMSG** exit has four subfunctions:

[RXMSQPLL]

This is called before a line is retrieved from the stack and the application may itself provide the interpreter with an alternative line. On entry, the third parameter points to a structure having the following definition:

```
typedef struct {
    RXSTRING rxmsq_ret;
} RXMSQPLL_PARM;
```

The `rxmsq_ret` field holds the string to be retrieved from the stack. Note that it is an output parameter, so its value on entry is undefined.

[RXMSQPSH]

This is called before the interpreter puts a line on the stack, and it may grab the line itself, and thus prevent the interpreter from putting the line on the stack. Note that this exit handles both pushing and queuing. The third parameter is:

```
typedef struct {
    struct {
        unsigned rxmflifo: 1;
    } rxmsq_flags;
    RXSTRING rxmsq_value;
} RXMSQPSH_PARM;
```

Here the field `rxmsq_value` holds the string to be put on the stack. Whether the string is to be pushed or queued is determined by the boolean value `rxmsq_flags.rxmflifo`, which is `TRUE` if the string is to be pushed.

All values are input values. What happens if you change them is not defined in the SAA API. Some implementations may let you modify the contents of `rxmsq_value` and return `RXEXIT_NOT_HANDLED` and the string pushed by the interpreter contains the modified string. However, you should not rely on this since it is highly incompatible. You may not de-allocate `rxmsq_value`.

[RXMSQSIZ]

this is called before the interpreter tries to determine the size of the stack, and it may present an alternative size to the interpreter. The third parameter is:

```
typedef struct {
    ULONG rxmsq_size;
} RXMSQSIZ_PARM;
```

The field `rxmsq_size` can be set to the number the application wants the `QUEUED()` function to return. Note that this parameter is undefined on entry, so it cannot be used to retrieve the number of lines on the stack.

[RXSQNAM]

This is called before the interpreter tries to retrieve the name of the current stack, and it may present the interpreter with an alternative name. Note that this functionality is part of SAA but not TRL; it supports the `Get` option of the `RXQUEUE()` built-in function. Note that there are no other exits supporting the other options of `RXQUEUE()`. The

third parameter for this exit is:

```
typedef struct {
    RXSTRING rxmsq_name;
} RXMSQNAM_PARM;
```

As with `RXSQMSIZ`, the field `rxmsq_name` can be set to the name which the application wants to return to the interpreter as the name of the current stack. Note that this is an output-only parameter; its value on input is undefined, and in particular is not the name of the real stack.

Note that this area is troublesome. In `TRL`, external data queues are not defined as part of the language, while in `SAA` it is. Thus, `TRL`-compliant interpreters are likely to implement stacks in various ways that may not be compatible with the `SAA`.

52.2.4 RXSIO-The Standard I/O Exit Handler

The main code for this exit handler has the symbolic value `RXSIO`. There are four sub-functions:

[RXSIODTR]

Called whenever the interpreter needs to read a line from the user during interactive tracing. Note the difference between this subfunction and `RXSIOTRD`.

[RXSIOSAY]

Called whenever the interpreter tries to write something to standard output in a `SAY` instruction, even a `SAY` instruction without a parameter.

[RXSIOTRC]

Called whenever the interpreter tries to write out debugging information, e.g. during tracing, as a traceback, or as a syntax error message.

[RXSIOTRD]

Called whenever the interpreter needs to read from the standard input stream during a `PULL` or `PARSEPULL` instruction. Note that it will not be called if there is sufficient data on the stack to satisfy the operation.

Note that these functions are only called for the exact situations that are listed above. e.g. the `RXSIOSAY` is not called during a call to the `REXX` built-in function `LINEOUT()` that writes to the default output stream. `TRL` says that `SAY` is identical to calling `LINEOUT()` for the standard output stream, but `SAA` API still manages to see the difference between stem variables and compound variables with a ``zero-length-string" tail. Please bear with this inconsistency.

Depending on the subfunction, the `ParmBlock` parameter will have four only slightly different definitions. It is kind of frustrating that the `ParmBlock` takes so many different data types, but it can be handled easily using unions, see a later section. The definitions are:

```
typedef struct {
    RXSTRING rxsiodtr_ret; /* the interactive trace input */
} RXSIODTR_PARM;
```

```
typedef struct {
    RXSTRING rxsio_string; /* the SAY line to write out */
}
```

```

}RXSIOSAY_PARM;

typedef struct {
    RXSTRINGrxsio_string; /* the debug line to write out */
}RXSIOTRC_PARM;

typedef struct {
    RXSTRINGrxsiotrd_retc; /* the line to read in */
}RXSIOTRD_PARM;

```

In all of these, the `RXSTRING` structure either holds the value to be written out (for `RXSIOSAY` and `RXSIOTRC`), or the value to be used instead of reading standard input stream (for `RXSIOTRD` and `RXSIODTR`). Note that the values set by `RXSIOTRD` and `RXSIODTR` are ignored if the exithandler does not return the value `RXEXIT_HANDLED`.

Any end-of-line markers are stripped off the strings in this context. If the exithandler writes out the string during `RXSIOSAY` or `RXSIOTRC`, it must supply any end-of-line action itself. Similarly, the interpreter does not expect an end-of-line marker in the data returned from `RXSIODTR` and `RXSIOTRD`.

The space used to store the return data for the `RXSIODTR` and `RXSIOTRD` sub-functions, must be provided by the exithandler itself, and the space is not de-allocated by the interpreter. The space can be reused by the application at any later time. The space allocated to hold the data given by the `RXSIOSAY` and `RXSIOTRC` sub-functions, will be allocated by the interpreter, and must neither be de-allocated by the exithandler, nor used after the exithandler has terminated.

52.2.5 RXHLT-The Halt Condition Exit Handler

Note: Because the `RXHLT` exithandler is called after every `REXX` instruction, enabling this exits slows `REXX` program execution.

The main code for this exithandler has the symbolic value `RXHLT`. There are two sub-functions:

[`RXHLTTST`]

Called whenever the interpreter polls externally raised `HALT` conditions; i.e. after every `REXX` instruction.

The definition of the `ParmBlock` is:

```

typedef struct {
    unsigned rxfhlt: 1;
}RXHLTTST_PARM;

```

The `rxfhlt` parameter is set to the state of the `HALT` condition in the interpreter; either `TRUE` or `FALSE`.

[`RXHLTCLR`]

Called to acknowledge processing of the `HALT` condition when the interpreter has recognized and raised a `HALT` condition.

52.2.6 RXTRC-The TraceStatusExitHandler

52.2.6.1 RXINI-The InitializationExitHandler

RXTER and this exit handler are a bit different from the others. RXINI provides the application programmer with a method of getting control before the execution of the script starts. Its main purpose is to enable manipulation of the variable pool in order to set up certain variables before the script starts, or set the trace mode.

It has only one subfunction, RXINIEXT, called once during each call to RexxStart(): just before the first REXX statement is interpreted. Variable manipulations performed during this exit will have effect when the script starts.

As there is no information to be communicated during this exit, the value of ParmBlock is undefined. It makes no difference whether you return RXEXIT_HANDLED or RXEXIT_NOT_HANDLED, since there is no situation to handle.

52.2.7 RXTER-The TerminationExitHandler

This exit resembles RXINI. Its sole subfunction is RXTEREXT, which is called once, just after the last statement of the REXX script has been interpreted. The state of all variables are intact during this call; so it can be used to retrieve the values of the variables at the exit of a script. (In fact, that is the whole purpose of this exit handler.)

Like RXINI, there is no information to be communicated during the exit, so ParamBlock is undefined in this call. And also like RXINI, it is more of a hook than an exit handler, so it does not matter whether you return RXEXIT_HANDLED or RXEXIT_NOT_HANDLED.

52.2.8 RXENV-The External EnvironmentExitHandler

This SystemExit is specific to **Regina**, so cautions should be exercised if you plan on making your code portable to other **Rexx** interpreters.

The main code for this exit handler has the symbolic value RXENV. There are four sub-functions:

[RXGETENV]

Called whenever the BIF; VALUE() is called to obtain a value from the external environment. i.e. the call to VALUE() is of the form:
VALUE(' VARNAME ' , , ' ENVIRONMENT ').

[RXSETENV]

Called whenever the BIF; VALUE() is called to set a value in the external environment. i.e. the call to VALUE() is of the form:
VALUE(' VARNAME ' , newvalue , ' ENVIRONMENT ').

[RXGETCWD]

Called whenever the current working directory is needed to be obtained from the environment. The DIRECTORY() BIF respects this system exit.

[RXSETCWD]

Called whenever the current working directory is changed by a call to DIRECTORY() or CHDIR() BIFs.

The ParmBlock parameter has the following definitions for each sub-function type:

```

typedef struct{
    RXSTRINGrxenv_name;/*thenameoftheexternalenvironmentvariable*/
    RXSTRINGrxenv_value;/*thereturnedvalueoftheexternalenvironment
variable*/
}RXGETENV_PARM;

typedef struct{
    RXSTRINGrxenv_name;/*thenameoftheexternalenvironmentvariable*/
    RXSTRINGrxenv_value;/*thevalueoftheexternalenvironmentvariable*/
}RXSETENV_PARM;

```

In both of these, the `RXSTRING;rxenv_name` structure holds the name of the environment variable as known by the external environment. Note that the values set by `RXSIOTRD` and `RXSIODTR` are ignored if the exithandler does not return the value `RXEXIT_HANDLED`.

The space used to store the return data for the `RXSIODTR` and `RXSIOTRD` sub-functions, must be provided by the exithandler itself, and the space is not de-allocated by the interpreter. The space can be reused by the application at any later time. The space allocated to hold the data given by the `RXSIO SAY` and `RXSIOTRC` sub-functions, will be allocated by the interpreter, and must neither be de-allocated by the exithandler, nor used after the exithandler has terminated.

53 TheExternalQueueInterface

The external queue interface provides a mechanism for interacting with the interpreter's external queues. This interface is analogous to a Rexx program's use of PUSH, QUEUE, PULL, and RXQUEUE(). Note that this interface only works with the external queues; it cannot interface to the internal named queues that exist within the interpreter.

53.1 TheRexxCreatQueue()function

This function is used to create a new, named, external queue.

The prototype for `RexxCreateQueue()` is:

```
APIRETAPIENTRY REXXCreateQueue(
    PSZBuffer,
    ULONGBuffLen,
    PSZRequestedName,
    ULONG*DupFlag
);
```

The following parameters are input, and their significance are:

[RequestedName]

Pointsto an ASCII NUL terminated character string which specifies the name of the queue to be created. See `QueueNames` for the structure of a queue name. If the user wishes to have the interpreter create a unique queue name on the local queues server at the default port number, then this value should be set to NULL. To request an interpreter-generated queue name, on the machine *fred* listening on port 5858, then specify *@fred:5858*. I leave the queue name portion blank.

The following parameters are output, and their significance are:

[Buffer]

Points to an ASCII NUL terminated character string allocated by the user. The name of the queue that is created will be copied into this area.

[BufLen]

Specifies the size of the memory area pointed to by *Buffer*.

[**DupFlag**]

Indicates if the queue name that was requested, already existed. If a queue name was specified, and the queue already existed, DupFlag is set to `RXQUEUE_DUP`, otherwise it is set to 0.

The `RexxCreateQueue()` returns an `unsigned long`, which carries status information describing the outcome of the operation. The status will be one of the `RXQUEUE` values:

[RXQUEUE OK]

The queue was successfully created.

[RXQUEUE_NOEMEM]

The queue was not created, due to lack of memory.

53.2 The REXXDeleteQueue() function

This function is used to delete a named, external queue.

The prototype for REXXDeleteQueue() is:

```
APIRETAPIENTRY REXXDeleteQueue(  
    PSZQueueName  
);
```

The only parameter is an input, and its significance is:

[QueueName]

Points to an ASCII NUL terminated character string which specifies the name of the queue to be deleted. See QueueNames for the structure of a queue name.

The REXXDeleteQueue() returns an unsigned long, which carries status information describing the outcome of the operation. The status will be one of the RXQUEUE values:

[RXQUEUE_OK]

The queue was successfully deleted.

[RXQUEUE_NOTREG]

The queue name specified does not exist.

[RXQUEUE_BADQNAME]

The queue name was not specified.

53.3 The REXXQueryQueue() function

This function is used to determine the number of items that are available on the named, external queue.

The prototype for REXXQueryQueue() is:

```
APIRETAPIENTRY REXXQueryQueue(  
    PSZQueueName,  
    ULONG*Count  
);
```

One parameter is an input, and its significance is:

[QueueName]

Points to an ASCII NUL terminated character string which specifies the name of the queue to be queried. See QueueNames for the structure of a queue name.

The following parameter is output, and its significance is:

[Count]

Points to an unsigned long which indicates the number of items on the specified queue.

The `RexxQueryQueue()` returns an `unsigned long`, which carries status information describing the outcome of the operation. The status will be one of the `RXQUEUE` values:

[`RXQUEUE_OK`]
The queue was successfully queried, and *Count* contains the number of items on the queue.
[`RXQUEUE_NOTREG`]
The queue name specified does not exist.
[`RXQUEUE_BADQNAME`]
The queue name was not specified

53.4 The `RexxAddQueue()` function

This function is used to determine add an item to a named, external queue.

The prototype for `RexxAddQueue()` is:

```
APIRETAPIENTRYRexxAddQueue(  
    PSZQueueName,  
    PRXSTRINGEntryData,  
    ULONGAddFlag  
);
```

All parameters are input, and their significance are:

[*QueueName*]
Points to an ASCII NUL terminated character string which specifies the name of the queue on which the data is to be added. See *QueueNames* for the structure of a queue name.
[*EntryData*]
Points to a `RXSTRING` structure containing the data to be added to the queue.
[*AddFlag*]
Indicates how the data is to be added. Can be one of:
`RXQUEUE_FIFO`, to indicate that the data is to be added in a first-in-first-out order. This is equivalent to the `QUEUE` keyword.
`RXQUEUE_LIFO`, to indicate that the data is to be added in a last-in-first-out order. This is equivalent to the `PUSH` keyword.

The `RexxAddQueue()` returns an `unsigned long`, which carries status information describing the outcome of the operation. The status will be one of the `RXQUEUE` values:

[`RXQUEUE_OK`]
The data was successfully added to the specified queue.
[`RXQUEUE_NOTREG`]
The queue name specified does not exist.
[`RXQUEUE_BADQNAME`]
The queue name was not specified

53.5 The REXXPullQueue() function

This function is used to extract an item from the specified named, external queue. When successful, the item from the queue is returned, and that item deleted from the queue.

The prototype for REXXPullQueue() is:

```
APIRET APIENTRY REXXPullQueue(  
    PSZQueueName,  
    PRXSTRINGDataBuf,  
    PDATETIMETimeStamp,  
    ULONGWaitFlag  
);
```

The following parameters are input, and their significance are:

[QueueName]

Points to an ASCII NUL terminated character string which specifies the name of the queue from which the data is to be extracted. See QueueNames for the structure of a queue name.

[WaitFlag]

Indicates if the process should wait until there is data in the specified queue before returning. This could cause the process to block forever, if no data is due in the queue. Regin does not support this option at this stage; RXQUEUE_NOWAIT is assumed. Value can be one of:

RXQUEUE_WAIT, the process is to block and wait for data if the queue is currently empty.

RXQUEUE_NOWAIT, the process does not wait for data in the queue if it is currently empty. REXXPullQueue() will return RXQUEUE_EMPTY if there is no data in the queue.

The following parameters are output, and their significance are:

[DataBuf]

Points to a RXSTRING structure into which the contents of the extracted item are placed. The memory associated with the RXSTRING strptr, should be deallocated using REXXFreeMemory().

[TimeStamp]

Points to a PDATETIME structure, which on return, contains the time detail of when the item was added to the external queue. Regin does not support this option at this stage.

The REXXPullQueue() returns an unsigned long, which carries status information describing the outcome of the operation. The status will be one of the RXQUEUE values:

[RXQUEUE_OK]

The data was successfully added to the specified queue.

[RXQUEUE_NOTREG]

The queue name specified does not exist.

[RXQUEUE_BADQNAME]

The queue name was not specified

[RXQUEUE_EMPTY]

The queue was empty and RXQUEUE_NOWAIT was specified.

[RXQUEUE_BADWAITFLAG]

The value of the *WaitFlag* parameter was not RXQUEUE_WAIT or RXQUEUE_NOWAIT.

54 TheMacroSpaceInterface

The macrospace interface provides a mechanism for pre-loading external Rexx programs into the current interpreter's macrospace, so that the macros can be executed faster than reading them from disk each time they are called. This interface is not available in Regina at this stage.

54.1 TheRexxAddMacro()function

54.2 TheRexxDropMacro()function

54.3 TheRexxSaveMacroSpace()function

54.4 TheRexxLoadMacroSpace()function

54.5 TheRexxQueryMacro()function

54.6 TheRexxReorderMacro()function

54.7 TheRexxClearMacroSpace()function

55 Allocating and De-allocating Space

For several of the functions described in this chapter, the application calling them must allocate or de-allocate dynamic memory. Depending on the operating system, compiler and REXX interpreter, the method for these allocations and de-allocations vary. Because of this, Reginas supplies the API function calls `RexxAllocateMemory()` and `RexxFreeMemory()`. These functions are wrappers for the appropriate compiler or operating system memory functions.

55.1 The `RexxAllocateMemory()` function

The prototype for `RexxAllocateMemory()` is:

```
PVOID APIENTRY RexxAllocateMemory(  
    ULONG size  
);
```

The parameter is an input, and its significance is:

[size]
The number of bytes of dynamic memory requested.

`RexxAllocateMemory()` returns a pointer to the newly allocated block of memory, or `NULL` if no memory could be allocated.

55.2 The `RexxFreeMemory()` function

The prototype for `RexxFreeMemory()` is:

```
APIRET APIENTRY RexxFreeMemory(  
    PVOID MemoryBlock  
);
```

The parameter is an input, and its significance is:

[MemoryBlock]
A void pointer to the block of memory allocated by the interpreter, or allocated by a previous call to `RexxAllocateMemory()`.

`RexxFreeMemory()` always returns 0.

56 Calling back into running REXX Code

This section describes the `RexxCallBack()` function, which allows the application to execute a procedure within the running REXX program. This function is particularly useful for building a Rexx interface to those library applications that operate using a callback mechanism.

A callback mechanism is one where certain events within a particular application can be *connected* to a particular function, so that when a particular event occurs, the *connected* function is executed. Many C library applications use a callback mechanism.

This function is specific to Regina, so caution should be exercised if you plan on making your code portable to other REXX interpreters.

56.1 The `RexxCallBack()` function

This function is used to invoke the REXX interpreter in order to execute a piece of REXX code, which may be located on disk, as a pre-tokenized macro, or as ASCII source code in memory.

```
APIRETAPIENTRYRexxCallBack(  
    PSZProcedureName,  
    LONGArgCount,  
    PRXSTRINGArgList,  
    PUSHORTReturnCode,  
    PRXSTRINGResult  
);
```

Of these parameters, `ReturnCode` and `Result` are output-only. The rest of the parameters are input-only. The significance of the parameters are:

[ProcedureName]

An ASCII NUL terminated string, specifying the name of the internal procedure of the running REXX script to be executed. This internal procedure name must exist or this function will return with `RX_CB_BADN`.

[ArgCount]

The number of parameter strings given to the procedure. This is the number of defined REXX-strings pointed to by the `ArgList` parameter. The default maximum number of arguments that can be passed is 32, but this can be changed by the `MAX_ARGS_TO_REXXSTART` macro in `rexx.h`.

[ArgList]

Point to an array of REXX-strings, constituting the parameters to this call to REXX. The size of this array is given by the parameter `ArgCount`. If `ArgCount` is greater than one, the first and last parameters are `ArgList[0]` and `ArgList[ArgCount-1]`. If `ArgCount` is 0, the value of `ArgList` is irrelevant.

If the `strptr` of one of the elements in the array pointed to by `ArgList` is `NULL`, that means that this parameter is empty (i.e. unspecified, as opposed to a string of zero size).

[ReturnCode]

Point to a `SHORT` integer where the return code from the called `ProcedureName` is stored, provided that the returned value is numeric, and within the range $-(2^{*}15)$ to

2**15-1. I don't know what happens to `ReturnCode` if either of these conditions is not satisfied. It probably becomes undefined, which means that it is totally useless since the program has to inspect the return string in order to determine whether `ReturnCode` is valid.

[Result]

Points to a REXX string into which the result string from the called `ProcedureName` is written. The caller may or may not let the `strptr` field be supplied. If supplied (i.e. it is non-NULL), that area will be used, else a new area will be allocated. If the supplied area is used, its size is supposed to be given by the `strlength` field. If the size is not sufficient, a new area will be allocated, by `RexxAllocateMemory()`, and the caller must see to that it is properly de-allocated using `RexxFreeMemory()`.

Note that the `ArgCount` parameter need not be the same as the `ARG()` built-in function would return. Differences will occur if the last entries in `ArgList` are *null strings*.

The arguments passed to the `ProcedureName` will be passed individually, i.e. the `PARSE ARG` command in the `ProcedureName` must use commas to separate the arguments.

The valid return values from `RexxCallBack()` are:

[Negative]

indicates that a syntax error occurred during interpretation. In general, you can expect the error value to have the same absolute value as the REXX syntax error (but opposite signs, of course).

[Zero, or `RX_CB_OK`]

indicates that the interpreter finished executing the procedure without errors.

[Positive]

indicates probably that some problem occurred, that made it impossible to execute the procedure, e.g. a bad parameter value. The values that can be returned are:

<code>RX_CB_BADP</code>	bad parameters
<code>RX_CB_NOTSTARTED</code>	there is no running REXX program
<code>RX_CB_TOOMANYP</code>	too many parameters supplied
<code>RX_CB_BADN</code>	the <code>ProcedureName</code> does not exist

During the course of an execution of `RexxCallBack()`, subcommand handlers and exit handlers might be called. These may call any function in the application interface, including another invocation of `RexxCallBack()`.

Implementation Limits

This chapter lists the implementation limits required by the REXX standard. All implementations are supposed to support at least these limits.

57 Why Use Limits?

Why use implementation limits at all? Often, a program (ab) uses a feature in a language to an extent that the implementor did not foresee. Suppose an implementor decides that variable names cannot be longer than 64 bytes. Sooner or later, a programmer gets the idea of using very long variable names to encode special information in the name; maybe the output of a machine-generated program. The result will be a program that works only for some interpreters or only for some problems.

By introducing implementation limits, REXX tells the implementor to what extent an implementation is required to support certain features, and simultaneously it tells the programmers how much functionality they can assume is present.

Note that these limited are required minimums for what an implementation must allow. An interpreter is not supposed to enforce these limits unless there is a good reason to.

58 What Limits to Choose?

A limit must not be perceived as an absolute limit, the implementor is free to increase the limit. To some extent, the implementor may also decrease the limit, in which case this must be properly documented as a non-standard feature. Also, the reason for this should be noted in the documentation.

Many interpreters are likely to have "memory" as an implementation limit, meaning that they will allow any size as long as there is enough memory left. Actually, this is equivalent to no limit, since running out of memory is an error with limits forcing interpreters as well. Some interpreters let the users set the limits, often controlled through the `OPTIONS` instruction.

For computers, limit choices are likely to be powers of two, like 256, 1024, 8192, etc. However, the REXX language takes the side of the user, and defines the limits in units which look a more "sensible" to computer non-experts: most of the limits in REXX are numbers like 250, 500, 1000, etc.

59 Required Limits

These are the implementation minimums defined by REXX:

[Binary strings]

Must be able to hold at least 50 characters after packing. That means that the unpacked size might be at least 400 characters, plus embedded whitespace.

[Elapse time clock]

Must be able to run for at least 10^{10} seconds, which is approximately 31.6 years. In general, this is really a big overkill, since virtually no program will run for such a period. Actually, few computers will be operational for such a period.

[Hexadecimal strings]

Must be able to hold at least 50 characters after packing. This means that the unpacked size might be at least 100 characters, plus embedded whitespace.

[Literal strings]

Must be able to hold at least 100 characters. Note that a double occurrence of the quote character (the same character used to delimit the string) in a literal string counts as a single character. In particular, it does not count as two, nor does it start a new string.

[Nesting of comments]

Must be possible to nest at least 10 levels. What happens then is not really defined. Maybe one of the syntax errors is issued, but none is obvious for this use. Another, more dangerous way of handling this situation would be to ignore new start-of-comments designators when on level 10. This could, under certain circumstances, lead to running of code that is actually commented out. However, most interpreters are likely to support nesting of comments to an arbitrary level.

[The Number of Parameters]

In calls must be supported up to at least 10 parameters. Most implementations support somewhat more than that, but quite a few enforce some sort of upper limit. For the built-in function, this may be a problem only for `MIN ()` and `MAX ()`.

[Significant digits]

Must be supported to at least 9 decimal digits. Also, if an implementation supports floating point numbers, it should allow exponents up to 9 decimal digits. An implementation is allowed to operate with different limits for the number of significant digits and the number of digits in exponents.

[Subroutine levels]

May be nested to a total of 100 levels, which counts both internal and external functions, but probably not built-in functions. You may actually trip in this limit if you are using recursive solution for large problems. Also, some tail-recursive approaches may crash in this limit.

[Symbol (name) length]

Can be at least 50 characters. This is the name of the symbol, not the length of the value if it names a variable. Nor is it the name of the variable after tail substitution. In other words, it is the symbol as it occurs in the source code. Note that this applies not only to simple symbols, but also to compound symbols and constant symbols. Consequently, you cannot write numbers of more than 50 digits in the source code, even if `NUMERIC DIGITS` is set high.

[Variable name length]

Of at least 50 characters. This is the name of a variable (which may or may not be set) after tail substitution.

60 Older(Obsolete)Limits

First edition of TRL1 contained some additional limits, which have been relaxed in the second edition in order to make implementation possible for a large set of computers. These limits are:

[Clock granularity]

Was defined to be at least of a millisecond.

Far from all computers provide this granularity, so the requirements have been relaxed. The current requirement is a granularity of at least one second, although a millisecond granularity is advised.

61 What the Standard does not Say

An implementation might enforce a certain limit even though one is not specified in the standard. This section tries to list most of the places where this might be the case:

[The stack]

(Also called: the external data queue) is not formally defined as a concept of the language itself, but a concept to which the REXX language has an interface. Several limits might apply to the stack, in particular the maximum length of a line in the stack and the maximum number of lines the stack can hold at once.

There might also be other limits related to the stack, like a maximum number of buffers or a maximum number of different stacks. These concepts are not referred to by REXX, but the programmer ought to be aware of them.

[Files]

May have several limits not specified by the definition of REXX, e.g. the number of files simultaneously open, the maximum size of a file, and the length and syntax of filenames. Some of these limits are enforced by the operating system rather than an implementation. The programmers should be particularly aware of the maximum number of simultaneously open files, since REXX does not have a standard construct for closing files.

[Expression nesting]

Can in some interpreters only be performed to a certain level. No explicit minimum limit has been put forth, so take care in complex expressions, in particular machine-generated expressions.

[Environment name length]

May have some restrictions, depending on your operating system. There is not defined any limit, but there exists an error message for use with too long environment names.

[Clause length]

May have an upper limit. There is defined an error message "Clause too long" which is supposed to be issued if a clause exceeds a particular implementation-dependent size. Note that a "clause" does not mean a "line" in this context; a line can contain multiple clauses.

[Source line length]

Might have an upper limit. This is not the same as a "clause" (see above). Typically, the source line limit will be much larger than the clause limit. The source line limit ought to be as large as the string limit.

[Stack operations]

Might be limited by several limits; first there is the number of strings in the stack, then there is the maximum length of each string, and at last there might be restrictions on the character set allowed in strings in the stack. Typically, the stack will be able to hold any character. It will either have "memory" as the limit for the number of strings and the length of each string, or it might have a fixed amount of memory set aside for stack strings. Some implementations also set a maximum length of stack strings, often 2^8 or 2^{16} .

62 What an Implementation is Allowed to "Ignore"

In order to make the REXX language implementable on as many machines as possible, the REXX standard allows implementation to ignore certain features. The existence of these features are

recommended, but not required. These features are:

[Floating point numbers]

Are not required; integers will suffice. If floating points are not supported, numbers can have not fractional or exponential part. And then normal division will not be available, i.e. the operator `"/"` will not be present. Use integer division instead.

[File operations]

Are defined in REXX, but an implementation seem to be allowed to differ in just about any file operation feature.

63 Limits in Regina

Regina tries not to enforce any limits. Wherever possible, "memory" is the limit, at the cost of some CPU whenever internal data structures must be expanded if their initial size were too small. Note that Regina will only increase the internal areas, not decrease them afterwards. The rationale is that if you happen to need a large internal area once, you may need it later in the same program too.

In particular, Regina has the following limits:

Binary strings	source linesize
Clock granularity	0.001-1 second (note 3)
Elapsed time clock	until ca. 2038 (note 1)
Named Queues	100
Hexadecimal strings	source linesize
Interpretable string	source linesize
Literal string length	source linesize
Nesting of comments	memory
Parameters	memory
Significant digits	memory (note 2)
Subroutine levels	memory
Symbol length	source linesize
Variable name length	memory (note 2)

Notes:

1) Regina uses the Unix-derived call `time()` for the elapsed time (and time in general). This is a function which returns the number of seconds since January 1st 1970. According to the ANSI C standard, in which Regina is written, this is a number which will at least hold the number $2^{31}-1$. Therefore, these machines will be able to work until about 2038, and Regina will satisfy the requirement of the elapsed time clock until 2006. By then, computers will hopefully be 64 bit.

Unfortunately, the `time()` C function call only returns whole seconds, so Regina is forced to use other (less standardized) calls to get a finer granularity. However, most of what is said about `time()` applies for these too.

2) The actual upper limit for these are the maximum length of a string, which is at least 2^{32} . So for

all practical purposes, the limit is "memory".

3) The clock granularity is a bit of a problem to define. All systems can be trusted to have a granularity of about 1 second. Except from that, it's very difficult to say anything more specific for certain. Most systems allow an alternative way to retrieve the time, giving a more accurate result. Wherever these alternatives are available, **Regina** will try to use them. If everything else fails, **Regina** will use 1 second granularity.

For most machines, the granularities are in the range of a few milliseconds. Some typical examples are: 20ms for Sun3, 4ms for Decstations 3100, and 10ms for SGI Indigo. Since this is a hardware restriction, this is the best measure anyone can get for these machines.

Appendixes

64 Definitions

In order to make the definitions more readable, but still have a rigid definition of the terms, some extra comments have been added to some of the definitions. These comments are enclosed in square brackets.

Argument is an *expression* supplied to a *function* or *subroutine*, and it provides data on which the call can work on.

Assignment is a *clause* in which second *token* is the equals sign. [Not that the statements "a=b" and "3=4" are an (invalid) assignment, not an expression. The type of the first token is irrelevant; if the second token is the equals sign, then the clause is assumed to be an assignment.]

Blanks are characters which *glyphs* are empty space, either vertically or horizontally. A blank is not a *token* (but may sometimes be embedded into tokens), but acts as *token separators*. [Exactly which characters are considered blank will differ between operating systems and implementations, but the <space> character is always a blank. The <tab> character is also often considered a blank. Other characters considered blank might be the end-of-line (<eol>), vertical tab (<vt>), and form feed (<ff>). See specific documentation for each interpreter for more information.]

Buffer

Caller routine

Character is a piece of information about a mapping from a storage unit (normally a byte) and a *glyph*. Often used as "the meaning of the glyph mapped to a particular storage unit". [The glyph "A" is the same in EBCDIC and ASCII, but the character "A" (i.e. the mapping from glyph to storage unit) differs.]

Character string is an finite, ordered, and possibly empty set of *characters*.

Clause is a non-empty collection of *tokens* in a REXX script. The tokens making up a clause are all the consecutive tokens delimited by two consecutive *clause delimiters*. [Clauses are further divided into *null clauses*, *instructions*, *assignments*, and *commands*.]

Clause delimiter is a non-empty sequence of elements of a subset of *tokens*, normally the line feed and the semicolon. Also the start and end of a *REXX script* are considered clause delimiters. Also colon is a clause separator, but it is only valid after a label.

Command

Compound variable is a *variable* which name has at least one " ." character that is not positioned at the end of the name.

Current environment is a particular *environment* to which *commands* is routed if no explicit environment is specified for their routing.

Current procedure level is the *procedure level* in effect at a certain point during execution.

Daemon

Decimal digit

Device driver

Digit is a single character having a numeric value associated with its glyph.

Empty string

Environment is an interface to which REXX can route *commands* and afterwards retrieve status information like *return values*.

Evaluation is the process applied to an *expression* in order to derive a *character string*.

Exposing is the binding of a *variable* in the *current procedure level* to the variable having the same name in the *caller routine*. This binding will be in effect for as long as the current procedure level is active.

Exponential form is a way of writing particularly large or small *numbers* in a fashion that makes them more readable. The number is divided into a mantissa and an exponent of base 10.

Expression is a non-empty sequence of *tokens*, for which there exist syntactic restrictions on which tokens can be members, and the order in which the tokens can occur. [Typically, an expression may consist of literal strings or symbols, connected by concatenation and operators.]

External data queue see "stack".

External subroutine is a *script* of REXX code, which is executed as a response to a *subroutine* or *function call* that is neither internal nor built-in.

FIFO

Glyph is an atomic element of text, having a meaning and an appearance; like a letter, a digit, a punctuation mark, etc.

Hex is used as a general abbreviation for term *hexadecimal* when used in compound words like hex digit and hex string.

Hexadecimal digit is a *digit* in the number system having a base of 16. The first ten digits are identical with the *decimal digits* (0-9), while for the last six digits, the first six letters of the Latin alphabet (A-F) are used.

Hexadecimalstring is a *characterstring* that consists only of the *hexadecimaldigits*, and with optional *whitespace* to divide the hexadecimal digits into groups. Leading or trailing whitespace is illegal. All groups except the first must consist of an even number of digits. If the first group has an odd number of digits, an extra leading zero is implied under some circumstances.

Instruction is a *clause* that is recognized by the fact that the first *token* is a special *keyword*, and that the clause is not an *assignment* or label. Instructions typically are well-defined REXX language components, such as loops and function calls.

Interactivetrace is a *trace* mode, where the *interpreter* halts execution between each *clause*, and offer the user the possibility to specify arbitrary REXX *statements* to be executed before the execution continues.

Label

LIFO

Literalname is a name which will always be interpreted as a constant, i.e. that no variable substitution will take place.

Literalstring is a *token* in a REXX *script*, that basically is surrounded by quotation marks, in order to make a *characterstring* containing the same *characters* as the literal string.

Keyword is an element from a finite set of symbols.

Mainlevel

Mainprogram

Namespace is a collection of named *variables*. In general, the expression is used when referring to the set of variables available to the *program* at some point during interpretation.

Nullstring is a *characterstring* having the length zero, i.e. an empty character string. [Not the difference from the undefined string.]

Operating system

Parameters

Parsing

Procedure level

Program is a collection of REXX code, which may be zero or more *scripts*, or other repositories of REXX code. However, a program must contain all the code to be executed.

Queue see "external data queue" or "stack".

Routine is a unit during run-time, which is a procedural level. Certain settings are saved across *routines*. One *routine* (the caller *routine*) can be temporarily suspended while another *routine* is executed (the called *routine*). With such nesting, the called *routine* must be terminated before execution of the caller *routine* can be resumed. Normally, the CALL instruction or a function call is used to do this. Note that the main level of a REXX script is also a *routine*.

Script is a single file containing REXX code.

Space separated

Stack

Statement is a *clause* having in general some action, i.e. a clause other than a *null clause*.
[Assignments, commands and instructions are statements.]

Stem collection

Stem variable

Strictly order

Subkeyword is a *keyword*, but the prefix "sub" stresses the fact that a *symbol* is a keyword only in certain contexts [e.g. inside a particular instruction].

Subroutine is a *routine* which has been invoked from another REXX *routine*; i.e. it cannot be the "main" program of a REXX script.

Symbol

Symbol table

Tail substitution

Term

Token

Token separator

Uninitialized

Variable name

Variable symbol

Whitespace One or several consecutive *blank* characters.

hexliteral

norm.hexstring

bin{digit,string,literal}

norm.binstring

packedcharstring

Character strings is the only type of data available in Rexx, but to some extent there are 'subtypes' of character strings; character strings which contents has certain format. These special formats is discussed below.

65 Bibliography

[KIESEL]

Peter C. Kiesel, *REXX-Advanced Techniques for Programmers* .McGraw-Hill, 1993, ISBN0-07-034600-3

[CALLAWAY]

Merill Callaway, *The ARexx Cookbook* .511-A Girard Blvd. SE, Albuquerque, NM 87106: Whitestone, 1992, ISBN0-9632773-0-8

[TRL2]

M.F. Cowlishaw, *The REXX Language-Second Edition* .Englewood Cliffs: Prentice-Hall, 1990, ISBN0-13-780651-5

[TRL1]

M.F. Cowlishaw, *The REXX Language-First Edition* .Englewood Cliffs: Prentice-Hall, 1985, ISBN0-13-780735-X

[SYMPOS3]

Proceedings of the REXX Symposium for Developers and Users .Stanford: Stanford Linear Accelerator Center, 1992

[TRH:PRICE]

Stephen G. Price, *SAAPortability* ,chapter 37, pp477-498. In Goldberg and Smith III [TRH], 1992

[TRH]

Gabriel Goldberg and Smith III, Philip H., *The REXX Handbook* .McGraw-Hill, 1992, ISBN0-07-023682-8

[DANEY]

Charles Daney, *Programming in REXX*. McGraw-Hill, 1992, ISBN 0-07-015305-1

[BMARKS]

Brian Marks, *Advanced REXX programming*. McGraw-Hill, 1992

[ZAMARA]

Chris Zamara and Nick Sullivan, *Using A Rexx on the Amiga*. Abacus, 1991, ISBN 1-55755-114-6

[REXXSAA]

W. David Ashley, *SAA Procedure Language REXX Reference*. 5 Timberline Dr., Trophy Club, TX 76262: Pedagogic Software, 1991

[MCGH:DICTIONARY]

Sybil P. Parker, *McGraw-Hill Dictionary of Computers*. McGraw-Hill, 1984, ISBN 0-07-045415-9

[PJPLAUGER]

P. J. Plauger, *The Standard C Library*. Englewood Cliffs: Prentice Hall, 1992, ISBN 0-13-131509-9

[KR]

Brian W. Kernighan and Dennis M. Ritchie, *The C Programming Language - Second Edition*. Englewood Cliffs: Prentice Hall, 1988, ISBN 0-13-110362-8

[ANSI C]

Programming languages - C, Technical Report ISO/IEC 9899:1990, ISO, Case postale 56, CH-1211 Geneva 20, Switzerland, 1990

[OX:CDICTIONARY]

Edward L. Glaser and I. C. Pyle and Valerie Illingsworth, *Oxford Reference Dictionary of Computing - Third Edition*. Oxford University Press, 1990, ISBN 0-19-286131-X

[ANSI C]

Programming Languages - REXX, ANSI X3.274-1996, 11 West 42nd Street, New York, New York 10036

66 GNUFreeDocumentationLicense

GNUFreeDocumentationLicense

Version1.1, March2000

Copyright(C)2000FreeSoftwareFoundation, Inc.

59TemplePlace, Suite330, Boston, MA02111-1307USA

Everyoneispermittedtocopyanddistributeverbatimcopiesofthislicenseddocument, butchangingit
isnotallowed.

0. PREAMBLE

ThepurposeofthisLicenseistomakeamanual, textbook, orotherwrittendocument"free" inthe
senseoffreedom: toassureeveryonetheeffectivefreedomtocopyandredistributeit, withorwithout
modifyingit, eithercommerciallyornoncommercially. Secondly, thisLicensepreservesforthe
authorandpublisherawaytogetcreditfortheirwork, whilenotbeingconsideredresponsiblefor
modificationsmadebyothers.

ThisLicenseisakindof"copyleft", whichmeansthatderivativeworksofthedocumentmust
themselvesbefreeinthesamesense. ItcomplementstheGNUGeneralPublicLicense, whichisa
copyleftlicensedesignedforfreesoftware.

WehavedesignedthisLicenseinordertouseitformanualsforfreesoftware, becausefreesoftware
needsfreedocumentation: afreeprogramshouldcomewithmanualsprovidingthesamefreedomthat
thesoftwaredoes. ButthisLicenseisnotlimitedtosoftwaremanuals; itcanbeusedforanytextual
work, regardless ofsubjectmatterorwhetheritispublishedasaprintedbook. Werecommendthis
Licenseprincipallyforworkswhosepurposeisinstructionorreference.

1. APPLICABILITY AND DEFINITIONS

ThisLicenseappliestoanymanualorotherworkthatcontainsanoticeplacedbythecopyrightholder
sayingitcanbedistributedunderthetermsofthisLicense. The"Document", below, referstoanysuch
manualorwork. Anymemberofthepublicisalicensee, andisaddressedas"you".

A"ModifiedVersion"oftheDocumentmeansanyworkcontainingtheDocumentoraportionofit,
eithercopiedverbatim, orwithmodificationsand/ortranslatedintoanotherlanguage.

A"SecondarySection"isanamedappendixorafront-mattersectionoftheDocumentthatdeals
exclusivelywiththerelationshipofthepublishersorauthorsoftheDocumenttotheDocument's
overallsubject(ortorelatedmatters)andcontainsnothingthatcouldfalldirectlywithinthatoverall
subject. (Forexample, iftheDocumentisinpartatextbookofmathematics, aSecondarySectionmay
notexplainanymathematics.) Therelationshipcouldbeamatterofhistoricalconnectionwiththe
subjectorwithrelatedmatters, oroflegal, commercial, philosophical, ethicalorpoliticalposition
regardingthem.

The"InvariantSections"arecertainSecondarySectionswhosetitlesaredesignated, asbeingthoseof

Invariant Sections, in the notice that says that the Document is released under this License.

The "Cover Texts" are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License.

A "Transparent" copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, whose contents can be viewed and edited directly and straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup has been designed to thwart or discourage subsequent modification by readers is not Transparent. A copy that is not "Transparent" is called "Opaque".

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML designed for human modification. Opaque formats include PostScript, PDF, proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML produced by some word processors for output purposes only.

The "Title Page" means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, "Title Page" means the text near the most prominent appearance of the work's title, preceding the beginning of the body of the text.

2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notices saying this License apply to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

3. COPYING IN QUANTITY

If you publish printed copies of the Document numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition.

Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If there are required texts for either cover that are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a publicly-accessible computer-network location containing a complete Transparent copy of the Document, free of added material, which the general network-using public has access to download anonymously at no charge using public-standard network protocols. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the author of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do the following in the Modified Version:

- A. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any, be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.
- B. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has less than five).
- C. State on the Title page the name of the publisher of the Modified Version, as the publisher.
- D. Preserve all the copyright notices of the Document.
- E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
- F. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
- G. Preserve in that license notice the full list of Invariant Sections and required Cover Texts given in the Document's license notice.
- H. Include an unaltered copy of this License.
- I. Preserve the section entitled "History", and its title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
- J. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a

work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.

K. In any section entitled "Acknowledgements" or "Dedications", preserve the section's title, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.

L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.

M. Delete any section entitled "Endorsements". Such a section may not be included in the Modified Version.

N. Do not retitله any existing section as "Endorsements" or to conflict in title with any Invariant Section.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their title to the list of Invariant Sections in the Modified Version's license notice.

These titles must be distinct from any other section titles.

You may add a section entitled "Endorsements", provided it contains nothing but endorsements of your Modified Version by various parties--for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number.

Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections entitled "History" in the various original documents, forming one section entitled "History"; likewise combine any sections entitled "Acknowledgements", and any sections entitled "Dedications". You must delete all sections entitled "Endorsements."

6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, does not as a whole count as a Modified Version of the Document, provided no compilation copyright is claimed for the compilation. Such a compilation is called an "aggregate", and this License does not apply to the other self-contained works thus compiled with the Document, on account of their being thus compiled, if they are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one quarter of the entire aggregate, the Document's Cover Texts may be placed on covers that surround only the Document within the aggregate.

Otherwise they must appear on covers around the whole aggregate.

8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4.

Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License provided that you also include the original English version of this License. In case of a disagreement between the translation and the original English version of this License, the original English version will prevail.

9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided for under this License. Any other attempt to copy, modify, sublicense or distribute the Document is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See <http://www.gnu.org/copyleft/>.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License "or any later version" applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation.

ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

Copyright (c) YEAR YOURNAME.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.1 or any later version published by the Free Software Foundation; with the Invariant Sections being LIST THEIR TITLES, with the Front-Cover Texts being LIST, and with the Back-Cover Texts being LIST.

A copy of the license is included in this section entitled "GNU Free Documentation License".

If you have no Invariant Sections, write "with no Invariant Sections" instead of saying which ones are invariant. If you have no Front-Cover Texts, write "no Front-Cover Texts" instead of "Front-Cover Texts being LIST"; likewise for Back-Cover Texts.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.

Table of Contents

1 Purpose of this document	3
2 Implementation.....	3
3 Ports of Regina	4
4 Executing Rexx programs with Regina	4
4.1 Switches.....	5
4.2 External Rexx programs	6
5 Definitions.....	7
6 Null clauses	8
7 Commands.....	10
7.1 Assignments.....	10
8 Instructions.....	11
8.1 The ADDRESS Instruction	14
8.2 The ARG Instruction	20
8.3 The CALL Instruction	20
8.4 The DO/END Instruction	23
8.5 The DROP Instruction	27
8.6 The EXIT Instruction	28
8.7 The IF/THEN/ELSE Instruction	29
8.8 The INTERPRET Instruction	30
8.9 The ITERATE Instruction	32
8.10 The LEAVE Instruction	32
8.11 The NOP Instruction	33
8.12 The NUMERIC Instruction	33
8.13 The OPTIONS Instruction	35
8.14 The PARSE Instruction	36
8.15 The PROCEDURE Instruction	38
8.16 The PULL Instruction	41
8.17 The PUSH Instruction	42
8.18 The QUEUE Instruction	42
8.19 The RETURN Instruction	42
8.20 The SELECT/WHEN/OTHERWISE Instruction	44
8.21 The SIGNAL Instruction	45
8.22 The TRACE Instruction	48
8.23 The UPPER Instruction	50
9 Operators.....	50
9.1 Arithmetic Operators	50
9.2 Assignment Operators	50
9.3 Comparative Operators	50
9.4 Concatenation Operators	51
9.5 Logical Operators	51
10 Implementation-Specific Information	52
10.1 Miscellaneous.....	52
10.2 Implementation of the ADDRESS environment	52
10.2.1 SYSTEM aka ENVIRONMENT aka OS2 ENVIRONMENT	53
10.2.2 PATH.....	53

10.2.3CMDakaCOMMAND	53
10.2.4REXXorREGINA	54
10.3ADDRESSWITHonWindows	54
10.4ReginaRestrictedMode	55
10.5NativeLanguageSupport	55
10.5.1ErrorMessages	55
10.5.2Implementation.....	56
11GeneralInformation	57
11.1TheSyntaxFormat	57
11.2PrecisionandNormalization	58
11.3StandardParameterNames	58
11.4ErrorMessages	59
11.5PossibleSystemDependencies	59
11.6Blanksvs.Spaces	61
12ReginaBuilt-inFunctions	62
ABBREV(long,short[,length]).....	(ANSI)
.....	62
ABS(number).....	(ANSI)
.....	62
ADDRESS().....	(ANSI)
.....	62
ARG([argno[,option]]).....	(ANSI)
.....	62
B2C(binstring).....	(AREXX)
.....	64
B2X(binstring).....	(ANSI)
.....	64
BEEP(frequency[,duration]).....	(OS/2)
.....	64
BITAND(string1[, [string2][,padchar]]).....	(ANSI)
.....	64
BITCHG(string,bit).....	(AREXX)
.....	65
BITCLR(string,bit).....	(AREXX)
.....	65
BITCOMP(string1,string2,bit[,pad]).....	(AREXX)
.....	65
BITOR(string1[, [string2][,padchar]]).....	(ANSI)
.....	65
BITSET(string,bit).....	(AREXX)
.....	66
BITTST(string,bit).....	(AREXX)
.....	66
BITXOR(string1[, [string2][,padchar]]).....	(ANSI)
.....	66
BUFTYPE().....	(CMS)
.....	66

C2B(string).....	(AREXX)
.....	67
C2D(string[,length]).....	(ANSI)
.....	67
C2X(string).....	(ANSI)
.....	68
CD(directory).....	(REGINA)
.....	68
CHDIR(directory).....	(REGINA)
.....	68
CENTER(string,length[,padchar])	(ANSI)
.....	68
CENTRE(string,length[,padchar])	(ANSI)
.....	68
CHANGESTR(needle,haystack,newneedle)	(ANSI)
.....	69
CHARIN([streamid][,[start][,length]]).....	(ANSI)
.....	69
CHAROUT([streamid][,[string][,start]]).....	(ANSI)
.....	70
CHARS([streamid]).....	(ANSI)
.....	71
CLOSE(file).....	(AREXX)
.....	71
COMPARE(string1,string2[,padchar]).....	(ANSI)
.....	71
COMPRESS(string[,list]).....	(AREXX)
.....	72
CONDITION([option]).....	(ANSI)
.....	72
COPIES(string,copies).....	(ANSI)
.....	72
COUNTSTR(needle,haystack).....	(ANSI)
.....	73
CRYPT(string,salt).....	(REGINA)
.....	73
DATATYPE(string[,option]).....	(ANSI)
.....	73
DATE([option_out[,date[,option_in]])	(ANSI)
.....	75
DELSTR(string,start[,length]).....	(ANSI)
.....	77
DELWORD(string,start[,length]).....	(ANSI)
.....	77
DESBUF().....	(CMS)
.....	78
DIGITS().....	(ANSI)

LASTPOS(needle,haystack[,start]).....	(ANSI)
.....	86
LEFT(string,length[,padchar]).....	(ANSI)
.....	86
LENGTH(string).....	(ANSI)
.....	86
LINEIN([streamid][,[line][,count]]).....	(ANSI)
.....	87
LINEOUT([streamid][,[string][,line]]).....	(ANSI)
.....	88
LINES([streamid][,option]).....	(ANSI)
.....	89
MAKEBUF().....	(CMS)
.....	89
MAX(number1[,number2]...)	(ANSI)
.....	90
MIN(number[,number]...)	(ANSI)
.....	90
OPEN(file,filename,['Append' 'Read' 'Write']).....	(AREXX)
.....	90
OVERLAY(string1,string2[,start][,[length][,padchar]]).....	(ANSI)
.....	91
POPEN(command[,stem.].....	(REGINA)
.....	91
POS(needle,haystack[,start]).....	(ANSI)
.....	92
QUALIFY([streamid]).....	(ANSI)
.....	92
QUEUED().....	(ANSI)
.....	92
RANDOM(max).....	(ANSI)
.....	92
RANDOM([min][,[max][,seed]]).....	(ANSI)
.....	92
RANDU([seed]).....	(AREXX)
.....	93
READCH(file,length).....	(AREXX)
.....	94
READLN(file).....	(AREXX)
.....	94
REVERSE(string).....	(ANSI)
.....	94
RIGHT(string,length[,padchar]).....	(ANSI)
.....	94
RXFUNCADD(externalname,library,internalname).....	(SAA)
.....	94
RXFUNCDROP(externalname).....	(SAA)

.....	95
RXFUNCERRMSG().....	(REGINA)
.....	95
RXFUNCQUERY(externalname).....	(SAA)
.....	95
RXQUEUE(command[,queue]).....	(OS/2)
.....	95
SEEK(file,offset,['Begin' 'Current' 'End']).....	(AREXX)
.....	96
SHOW(option,[name],[pad]).....	(AREXX)
.....	96
SIGN(number).....	(ANSI)
.....	97
SLEEP(seconds).....	(CMS)
.....	97
SOURCELINE([lineno]).....	(ANSI)
.....	97
SPACE(string[, [length] [, padchar]]).....	(ANSI)
.....	97
STATE(streamid).....	(CMS)
.....	98
STORAGE([address],[string],[length],[pad]).....	(AREXX)
.....	98
STREAM(streamid[,option[,command]]).....	(ANSI)
.....	98
STRIP(string[, [option] [, char]]).....	(ANSI)
.....	103
SUBSTR(string,start[, [length] [, padchar]]).....	(ANSI)
.....	104
SUBWORD(string,start[,length]).....	(ANSI)
.....	104
SYMBOL(name).....	(ANSI)
.....	104
TIME([option_out[,time[option_in]])	(ANSI)
.....	105
TRACE([setting]).....	(ANSI)
.....	106
TRANSLATE(string[, [tableout] [, [tablein] [, padchar]]).....	(ANSI)
.....	107
TRIM(string).....	(AREXX)
.....	107
TRUNC(number[,length]).....	(ANSI)
.....	107
UNAME([option]).....	(REGINA)
.....	108
UNIXERROR(errorno).....	(REGINA)
.....	108

UPPER(string).....	(AREXX)	109
USERID().....	(REGINA)	109
VALUE(symbol[, [value], [pool]]).....	(ANSI)	109
VERIFY(string, ref[, [option] [, start]]).....	(ANSI)	110
WORD(string, wordno).....	(ANSI)	110
WORDINDEX(string, wordno).....	(ANSI)	111
WORDLENGTH(string, wordno).....	(ANSI)	111
WORDPOS(phrase, string[, start]).....	(ANSI)	111
WORDS(string).....	(ANSI)	112
WRITECH(file, string).....	(AREXX)	112
WRITELN(file, string).....	(AREXX)	112
XRANGE([start] [, end]).....	(ANSI)	112
X2B(hexstring).....	(ANSI)	112
X2C(hexstring).....	(ANSI)	113
X2D(hexstring[, length]).....	(ANSI)	113
13 Implementations specific documentation for Regina		114
13.1 Deviations from the Standard		114
13.2 Interpreter Internal Debugging Functions		114
13.3 REXX VMS Interface Functions		116
14 What are Conditions		118
14.1 What Do We Need Conditions for?		118
14.1.1 Terminology.....		118
15 The Mythical Standard Condition		119
15.1 Information Regarding Conditions (data structures)		119
15.2 How to Setup a Condition Trap		120
15.3 How to Raise a Condition		121
15.4 How to Trigger a Condition Trap		123
15.5 Trapping by Method SIGNAL		123
15.6 Trapping by Method CALL		124
15.7 The Current Trapped Condition		125
16 The Real Conditions		126
16.1 The SYNTAX condition		126

16.2TheHALTcondition	127
16.3TheERRORcondition	127
16.4TheFAILUREcondition	128
16.5TheNOVALUEcondition	128
16.6TheNOTREADYcondition	129
16.7TheLOSTDIGITScondition	129
17FurtherNotesonConditions	130
17.1ConditionsunderLanguageLevel3.50	130
17.2PitfallswhenUsingConditionTraps	130
17.3TheCorrectnessofthisDescription	130
18ConditionsinRegina	131
18.1HowtoRaisetheHALTcondition	131
19PossibleFutureextensions	132
20BackgroundandHistoricalRemarks	133
21REXX'sNotionofaStream	133
22ShortCrash-Course	134
23NamingStreams	135
24PersistentandTransientStreams	138
25OpeningaStream	139
26ClosingaStream	140
27Character-wiseandLine-wiseI/O	140
28ReadingandWriting	142
29DeterminingtheCurrentPosition	144
30PositioningWithin aFile	145
31Errors:Discovery,Handling,andRecovery	148
32CommonDifferencesandProblemswithStreamI/O	149
32.1WhereImplementationsareAllowedtoDiffer	149
32.2WhereImplementationsmightDifferanyway	150
32.3LINES()andCHARS()areInaccurate	151
32.4TheLastLineofaStream	152
32.5OtherPartsoftheI/OSystem	152
32.6Implementation-SpecificInformation	153
32.7StreamI/OinRegina0.07a	153
32.8FunctionalitytobeImplementedLater	156
32.9StreamI/OinARexx1.15	157
32.10MainDifferencesfromStandardREXX	161
32.11StreamI/OinBRexx1.0b	163
32.12ProblemswithBinaryandTextModes	167
33WhyHaveExtensions	169
34ExtensionsandStandardREXX	169
35SpecifyingExtensionsinRegina	169
36TheTroubleBegins	170
37TheFormatoftheOPTIONSclosure	170
38TheFundamentalExtensions	171
39Meta-extensions.....	173
40Semi-standards.....	173
41Standards.....	174

42	Background and history	175
43	General functionality of the stack	175
43.1	Basic functionality	175
43.2	LIFO and FIFO stack operations	177
43.3	Using multiple buffers in the stack	178
43.4	The zeroth buffer	179
43.5	Creating new stacks	180
44	The interface between REXX and the stack	182
45	Strategies for implementing stacks	182
46	Implementations of the stack in Regina	184
46.1	Implementation of the internal stack in Regina 2.2	184
46.2	Implementation of the external stack in Regina 2.2	185
	rxstack	185
	rxqueue	186
	rxqueueBuilt-inFunction	186
	QueueNames	187
	Security of External Queues	187
	Environment Variables	187
47	Overview of functions in SAA	188
47.1	Include Files and Libraries	188
47.2	Preprocessor Symbols	189
47.3	Data structures and data types	189
47.3.1	The RXSTRING structure	190
47.3.2	The RXXSYSEXIT structure	191
48	The Subcommand Handler Interface	193
48.1	What is a Subcommand Handler	193
48.2	The RexxRegisterSubcomExe() function	194
48.3	The RexxRegisterSubcomDll() function	195
48.4	The RexxDeregisterSubcom() function	195
48.5	The RexxQuerySubcom() function	196
49	The External Function Handler Interface	198
49.1	What is an External Function Handler	198
49.2	The RexxRegisterFunctionExe() function	199
49.3	The RexxRegisterFunctionDll() function	199
49.4	The RexxDeregisterFunction() function	200
49.5	The RexxQueryFunction() function	201
50	Executing REXX Code	202
50.1	The RexxStart() function	202
51	Variable Pool Interface	206
51.1	Symbolic or Direct	206
51.2	The SHVBLOCK structure	206
51.3	Regina Notes for the Variable Pool	210
51.4	The RexxVariablePool() function	211
52	The System Exit Handler Interface	213
52.1	The SystemExitHandler	213
52.2	List of SystemExitHandlers	214
52.2.1	RXFNC-The External FunctionExitHandler	214

52.2.2RXCMD-TheSubcommandExitHandler	215
52.2.3RXMSQ-TheExternalDataQueueExitHandler	216
52.2.4RXSIO-TheStandardI/OExitHandler	218
52.2.5RXHLT-TheHaltConditionExitHandler	219
52.2.6RXTRC-TheTraceStatusExitHandler	220
52.2.6.1.1RXINI-TheInitializationExitHandler	220
52.2.7RXTER-TheTerminationExitHandler	220
52.2.8RXENV-TheExternalEnvironmentExitHandler	220
53TheExternalQueueInterface	222
53.1TheRexxCreatQueue()function	222
53.2TheRexxDeleteQueue()function	223
53.3TheRexxQueryQueue()function	223
53.4TheRexxAddQueue()function	224
53.5TheRexxPullQueue()function	225
54TheMacroSpaceInterface	227
54.1TheRexxAddMacro()function	227
54.2TheRexxDropMacro()function	227
54.3TheRexxSaveMacroSpace()function	227
54.4TheRexxLoadMacroSpace()function	227
54.5TheRexxQueryMacro()function	227
54.6TheRexxReorderMacro()function	227
54.7TheRexxClearMacroSpace()function	227
55AllocatingandDe-allocatingSpace	228
55.1TheRexxAllocateMemory()function	228
55.2TheRexxFreeMemory()function	228
56CallingbackintorunningREXXCode	229
56.1TheRexxCallBack()function	229
57WhyUseLimits?	231
58WhatLimitstoChoose?	231
59RequiredLimits	231
60Older(Obsolete)Limits	232
61WhattheStandarddoesnotSay	233
62WhatanImplementationisAllowedto"Ignore"	233
63LimitsinRegina	234
64Definitions.....	236
65Bibliography.....	240
66GNUFreeDocumentationLicense	242